

A screenshot from the video game Super Mario Bros. The screen displays a level where every surface is covered in large, hand-drawn style numbers. The numbers form a repeating pattern of 6s, 4s, 2s, 9s, 7s, and 1s. In the bottom right corner, there is a blue rectangular area containing a white cloud icon and the text "MARIO" above "00000".



Generative Models II

Phillip Isola, MIT, OpenAI

DLSS

7/27/18

What's a generative model?

For this talk: models that output high-dimensional data

(Or, anything involving a GAN, VAE, PixelCNN, etc)

Useful for lots of problems beyond density estimation and sampling random images!

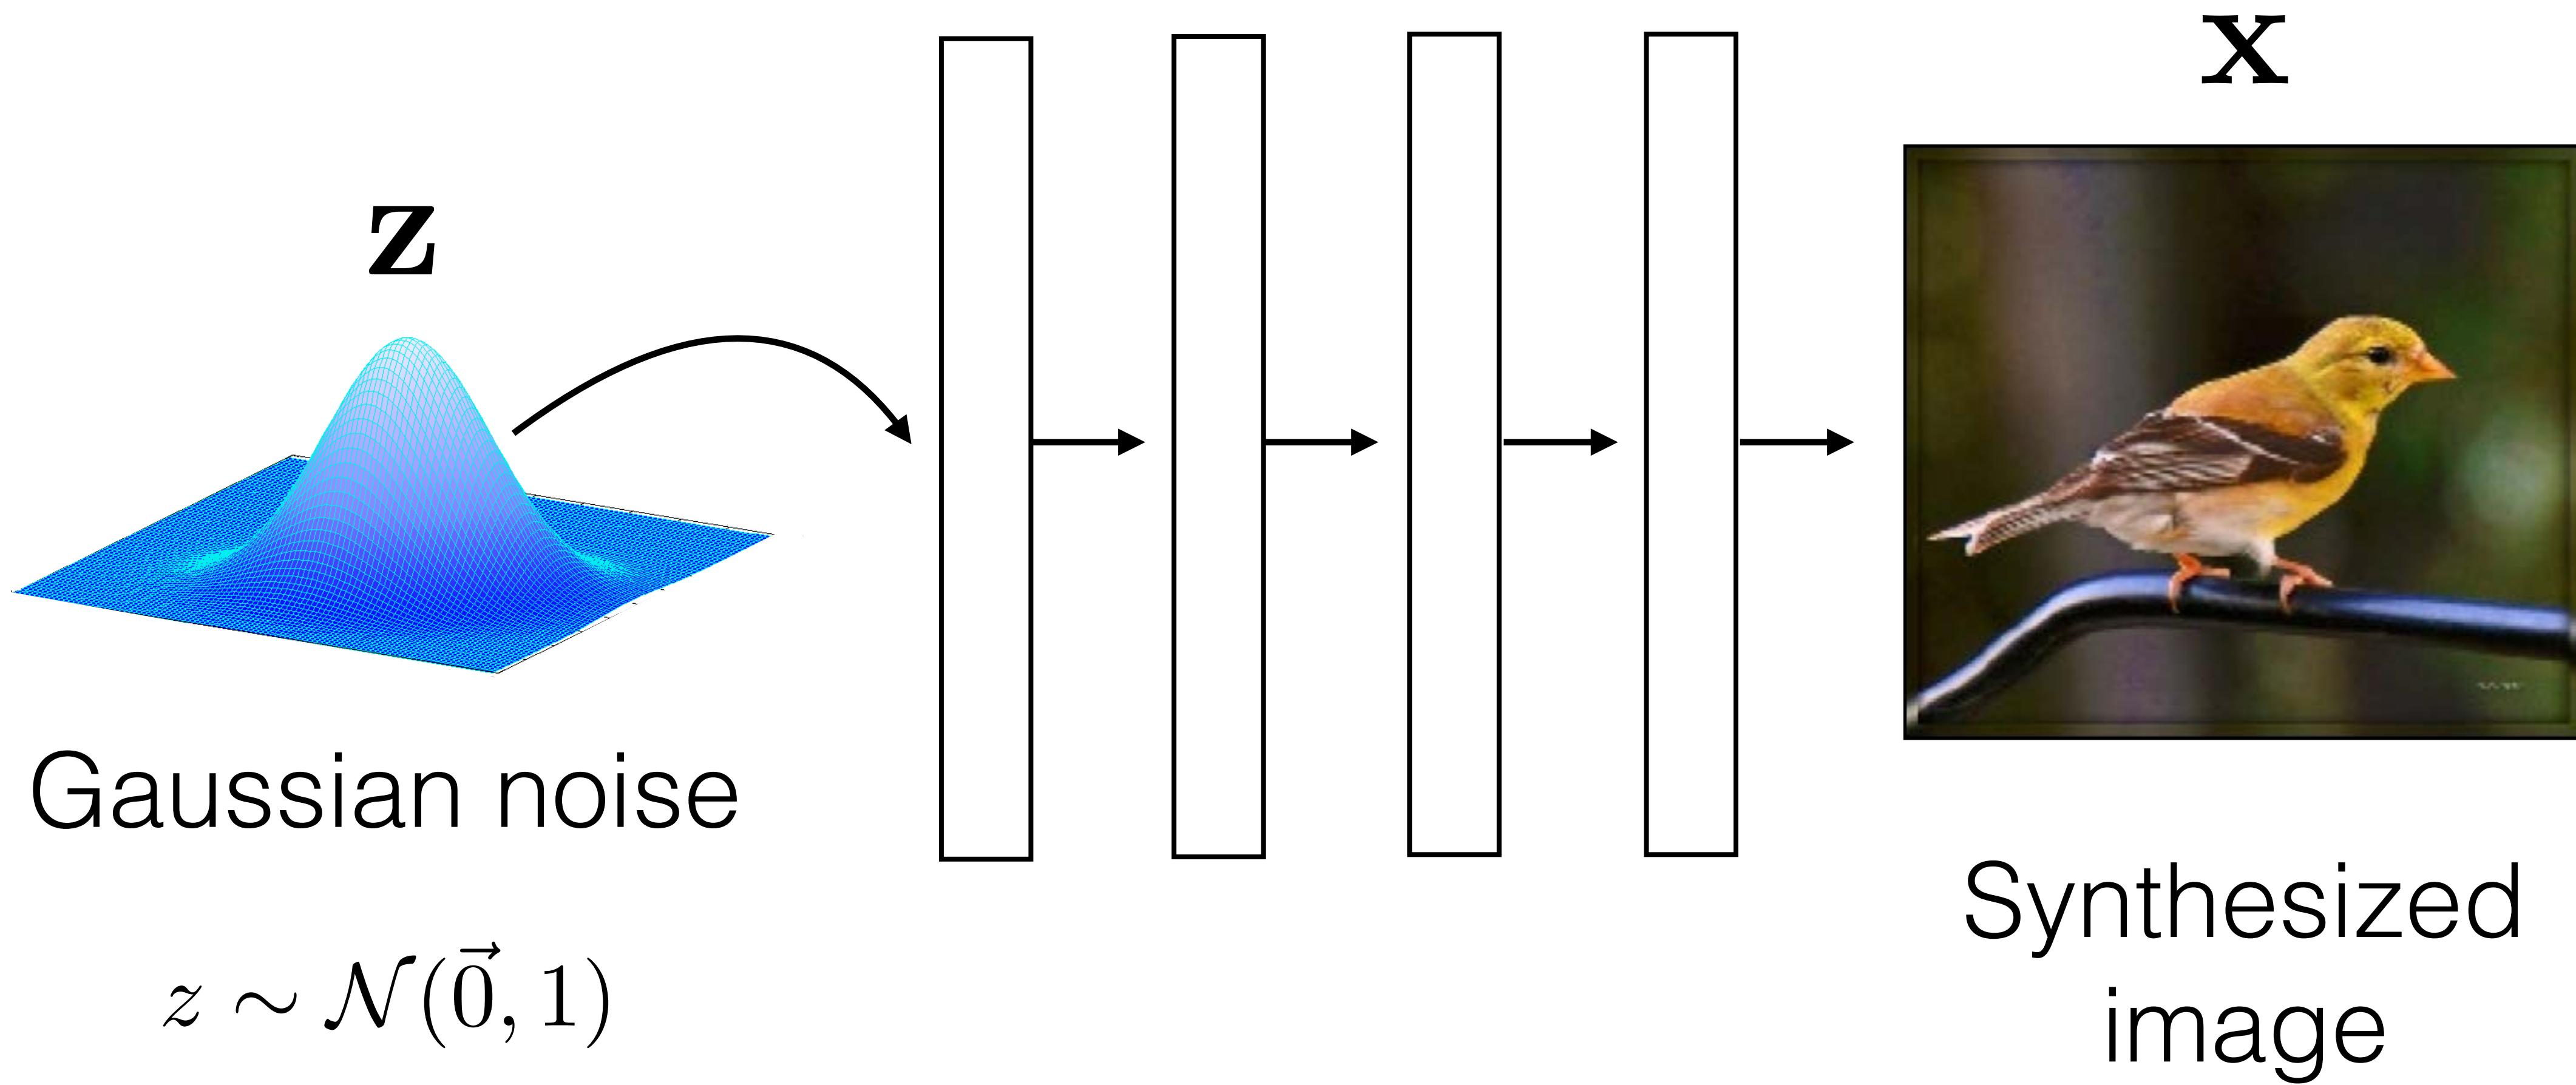
What can you do with generative models?

1. Data prediction
2. Domain mapping
3. Representation learning
4. Model-based intelligence

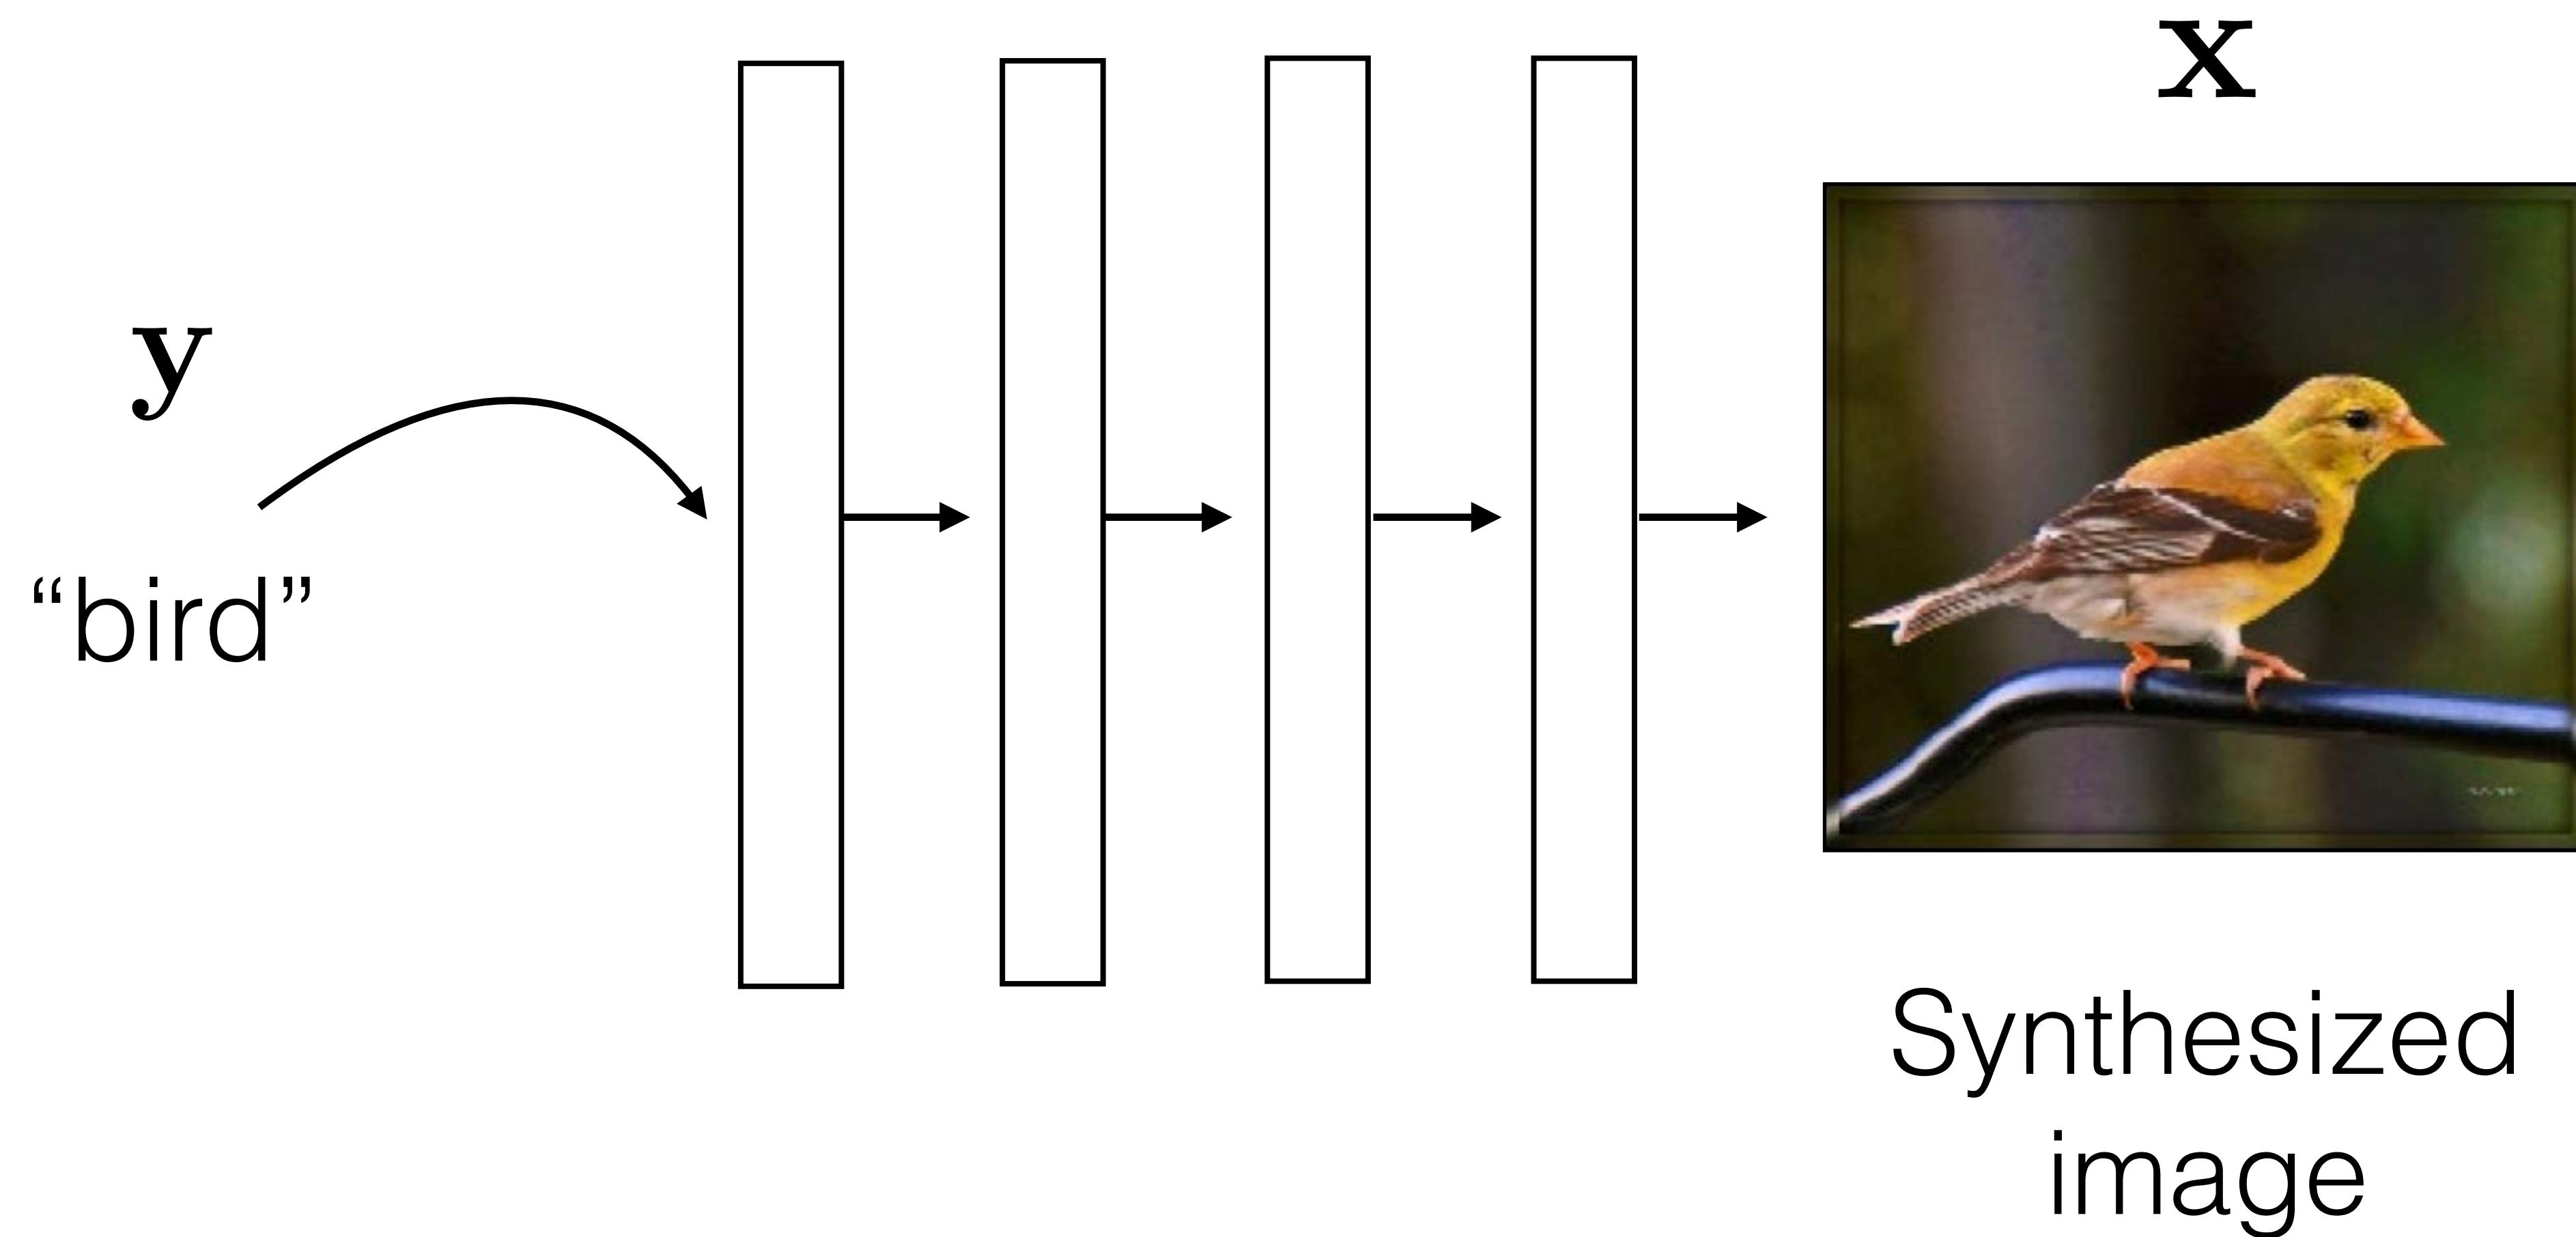
What can you do with generative models?

- 1. Data prediction**
2. Domain mapping
3. Representation learning
4. Model-based intelligence

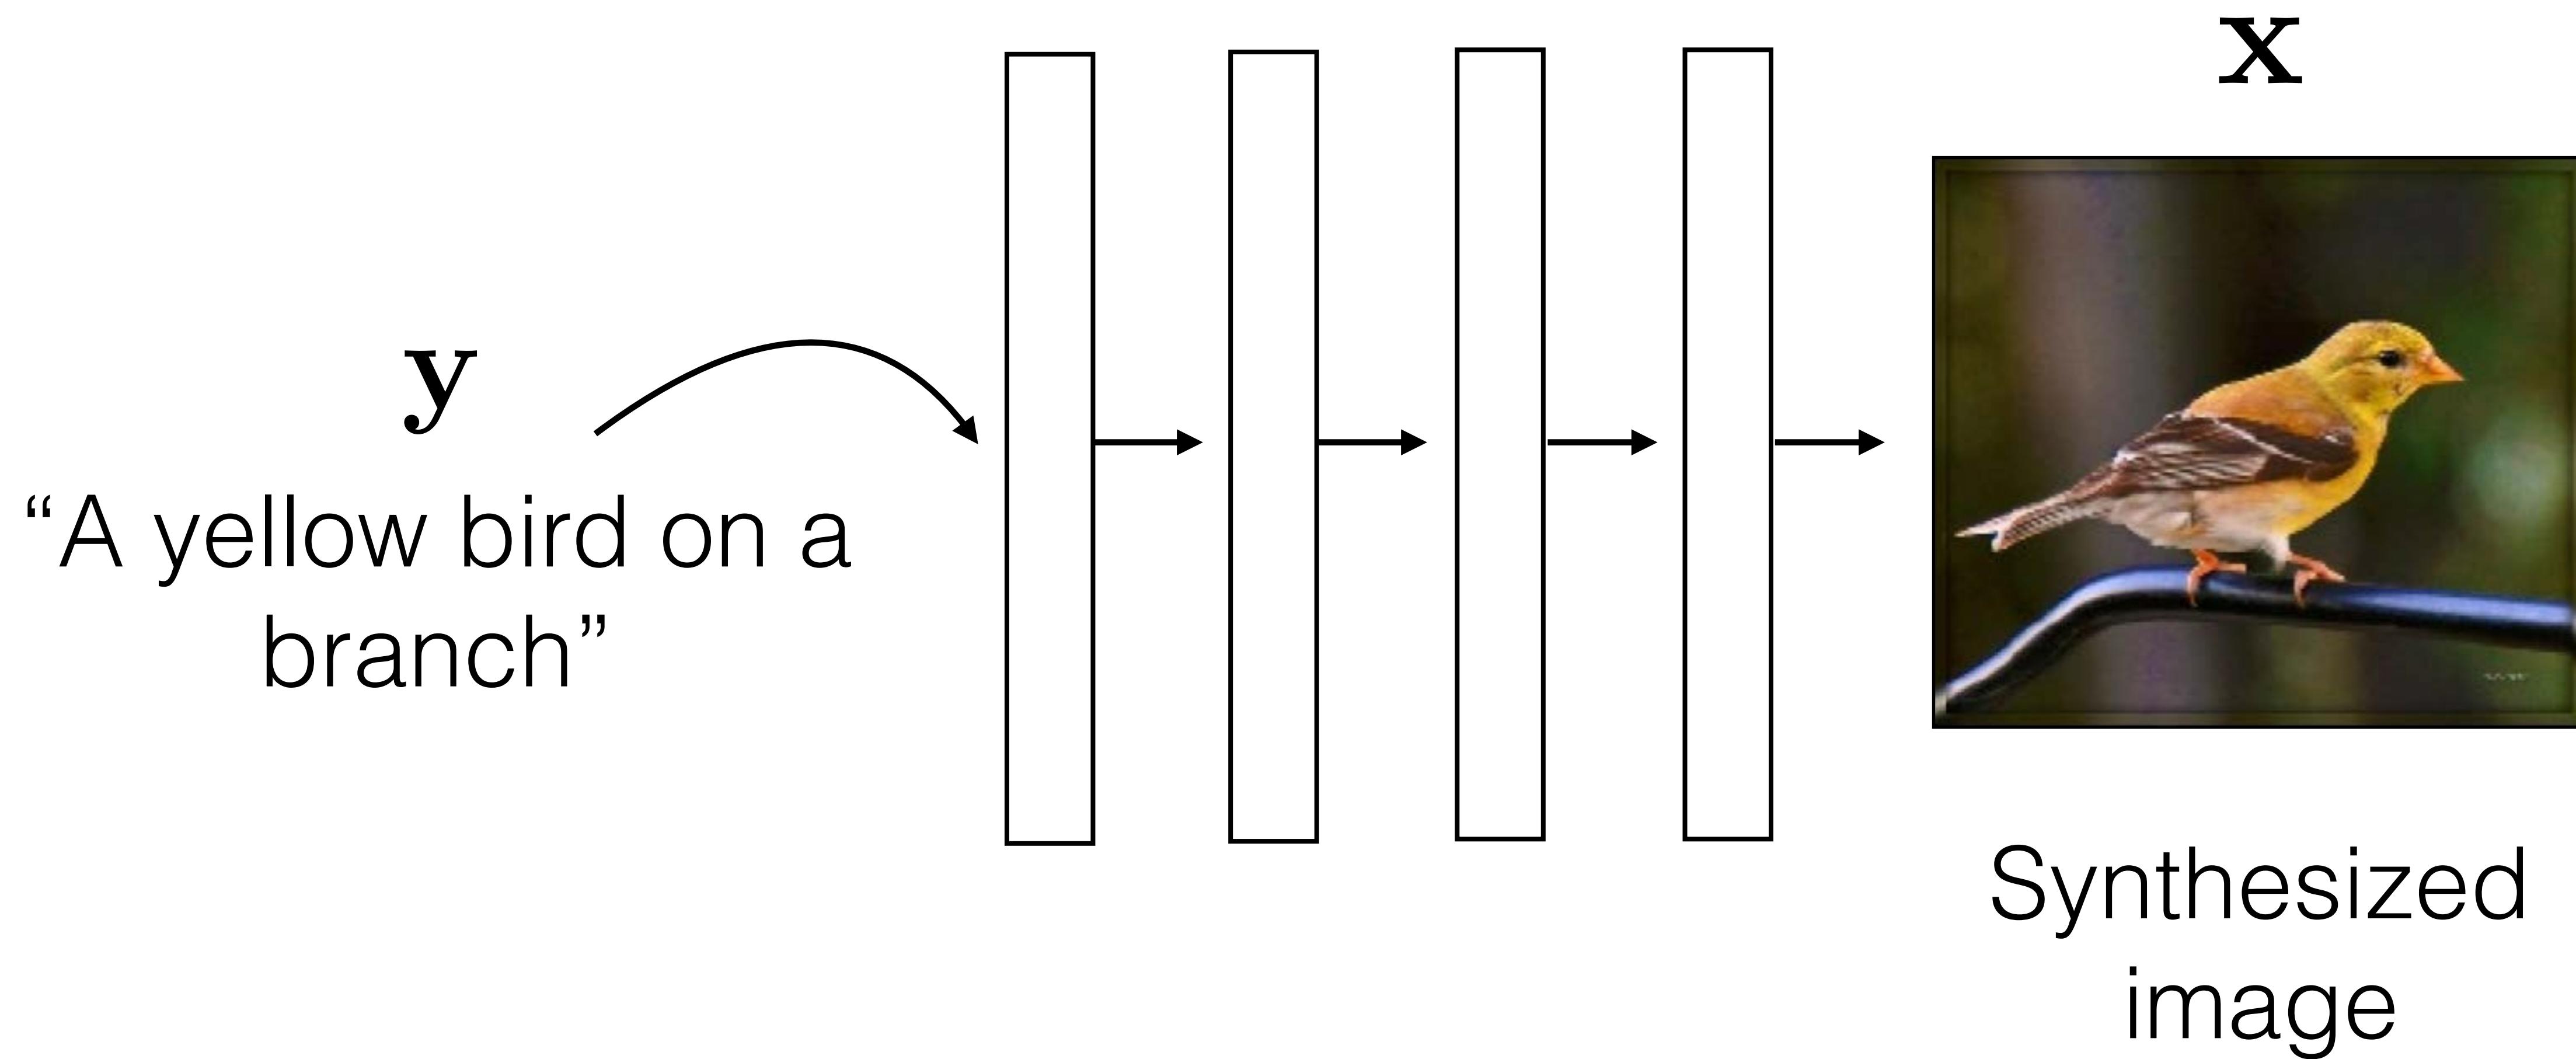
Generative Model



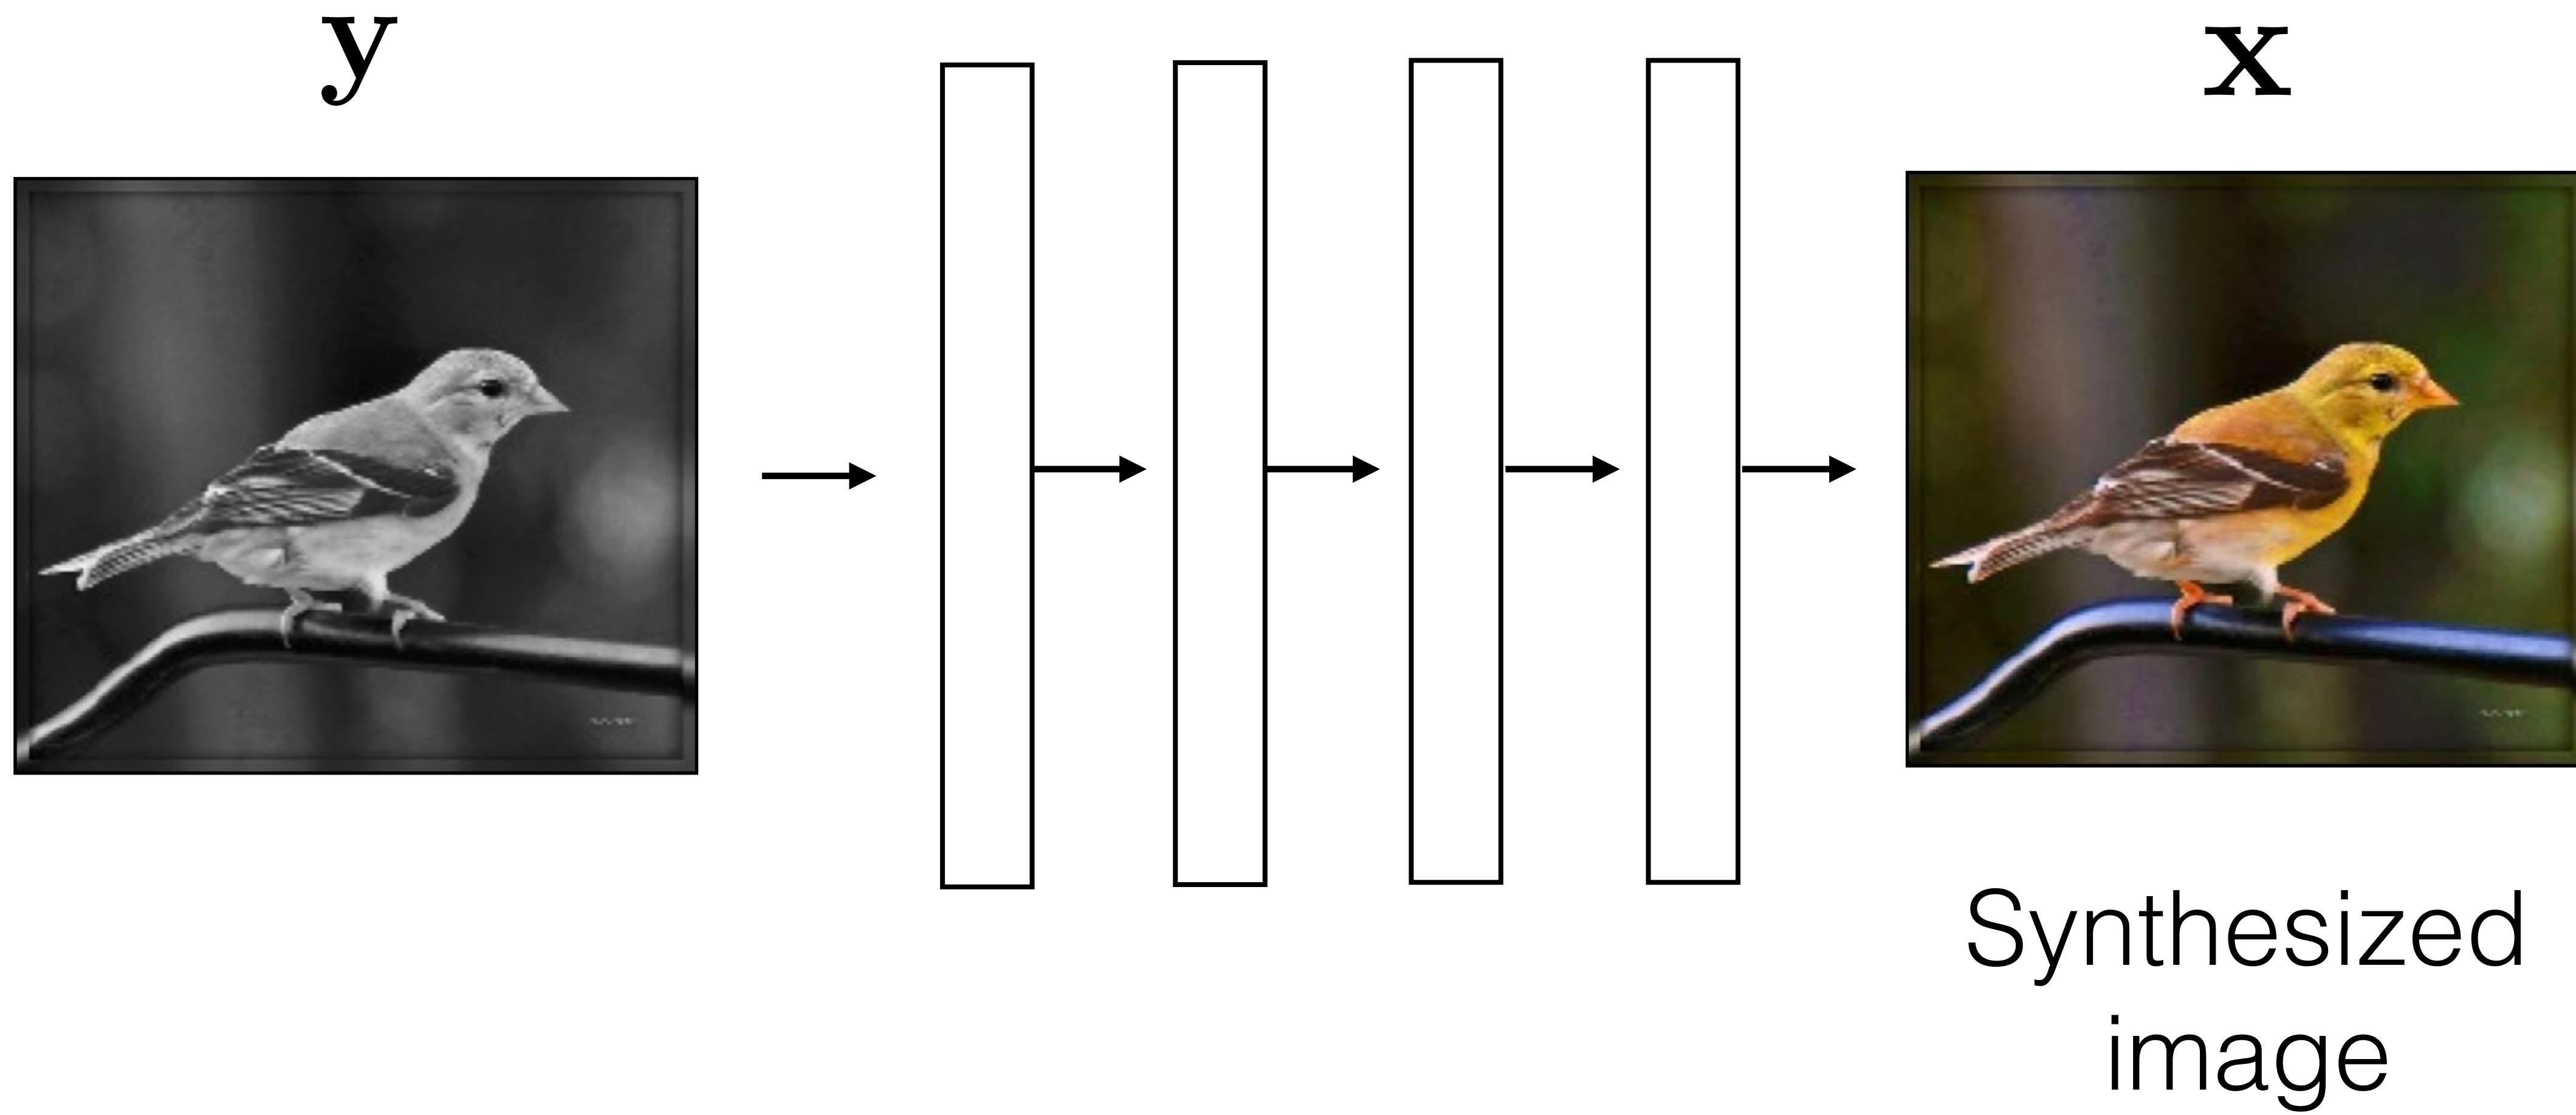
Conditional Generative Model



Conditional Generative Model

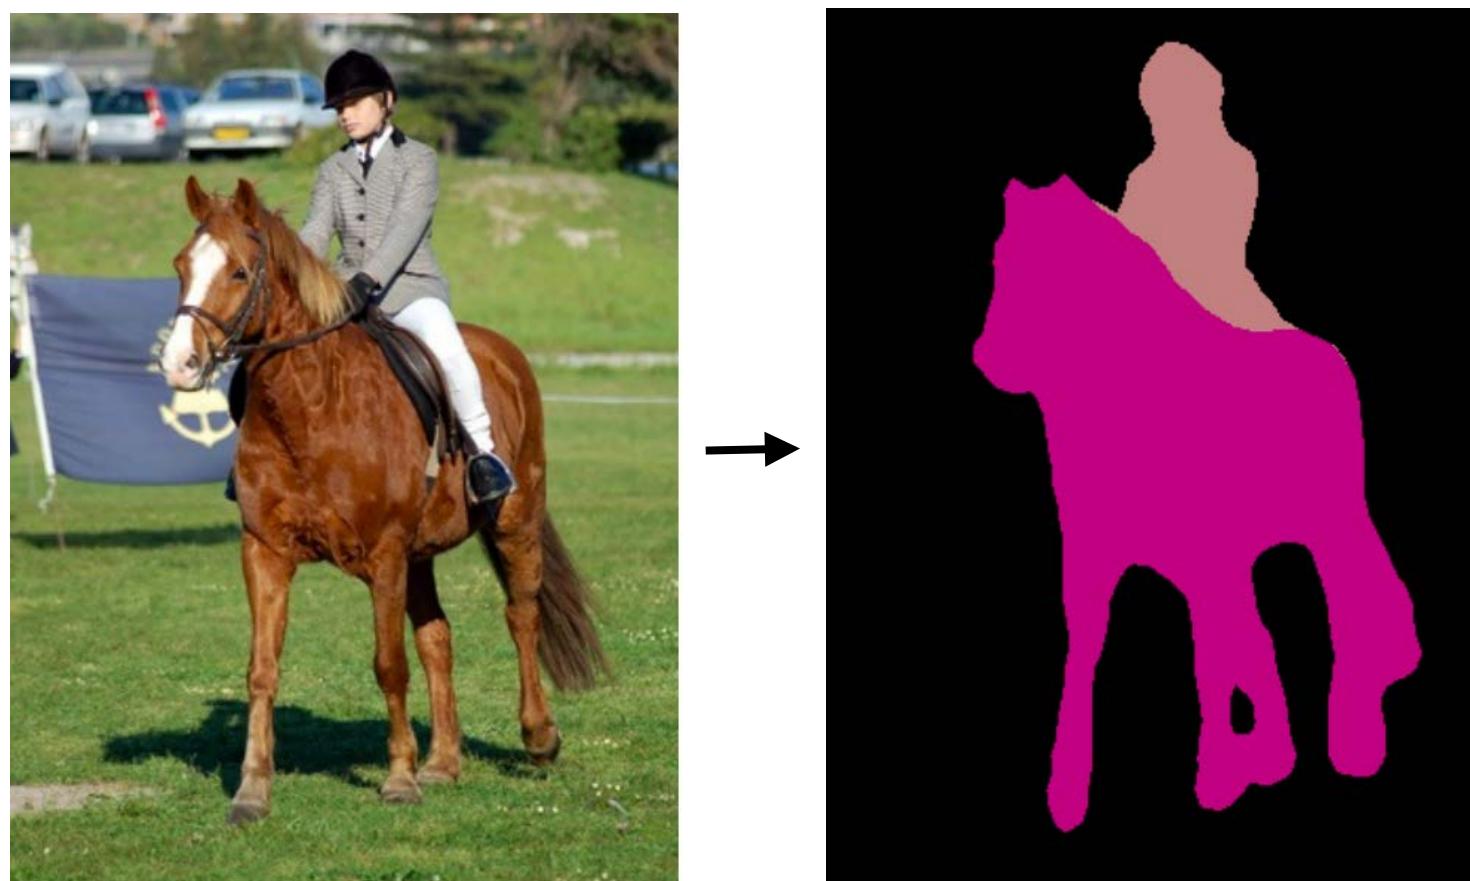


Conditional Generative Model



Data prediction problems (“structured prediction”)

Object labeling



[Long et al. 2015, ...]

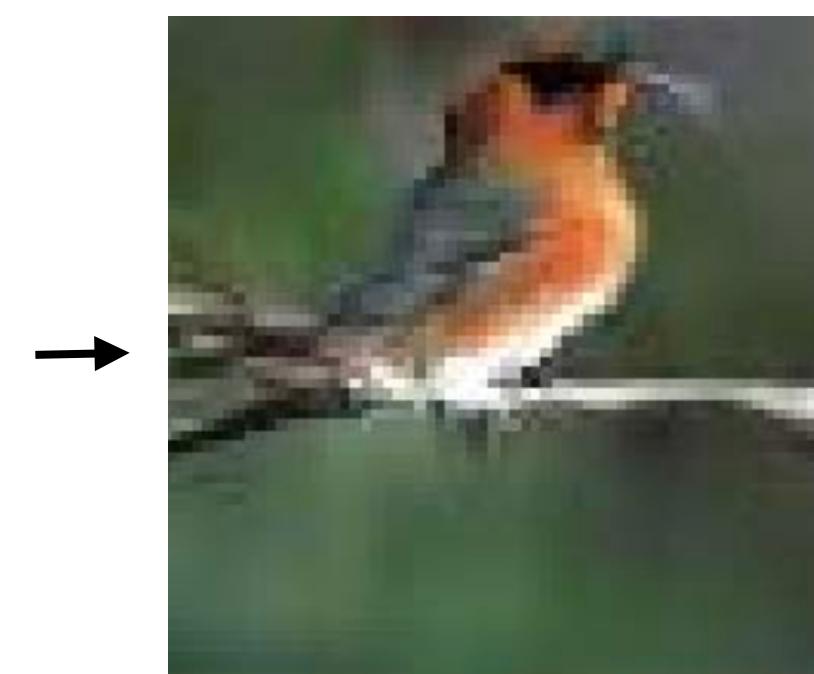
Edge Detection



[Xie et al. 2015, ...]

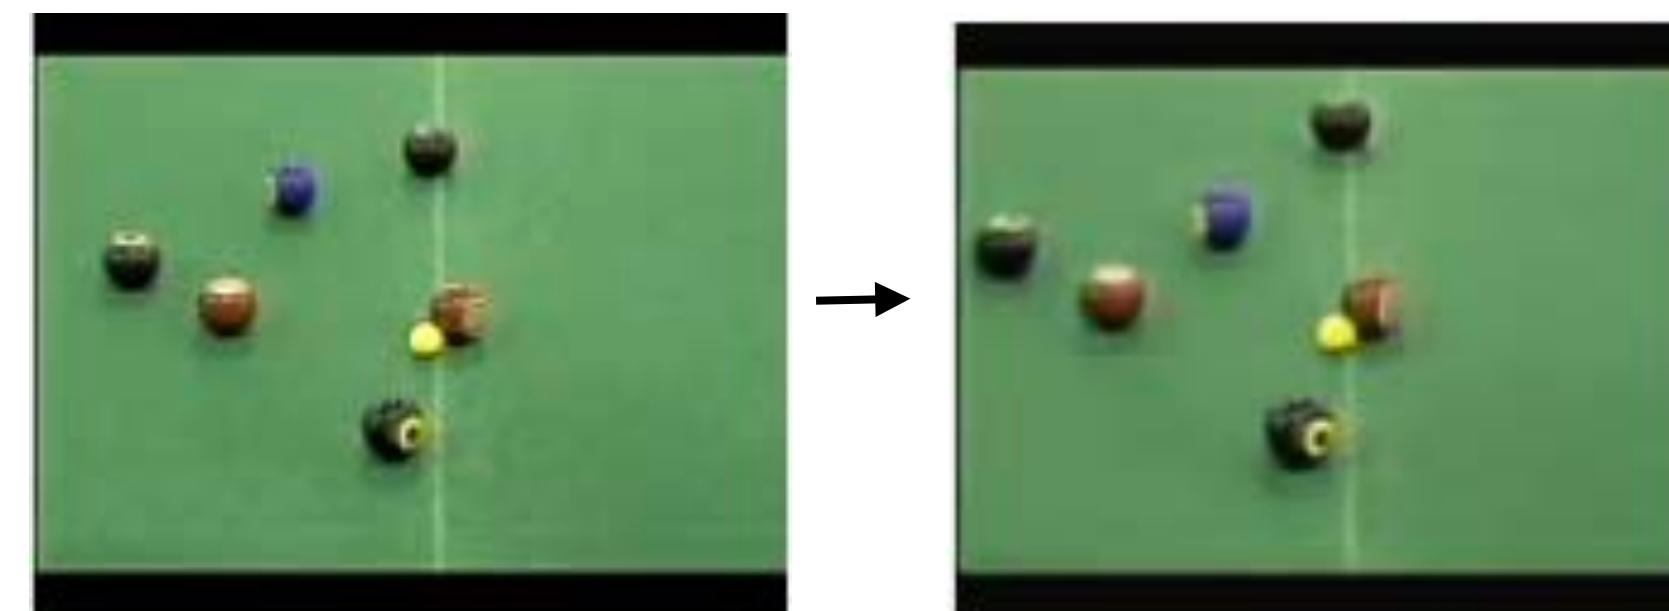
Text-to-photo

“this small bird has a pink
breast and crown...”



[Reed et al. 2014, ...]

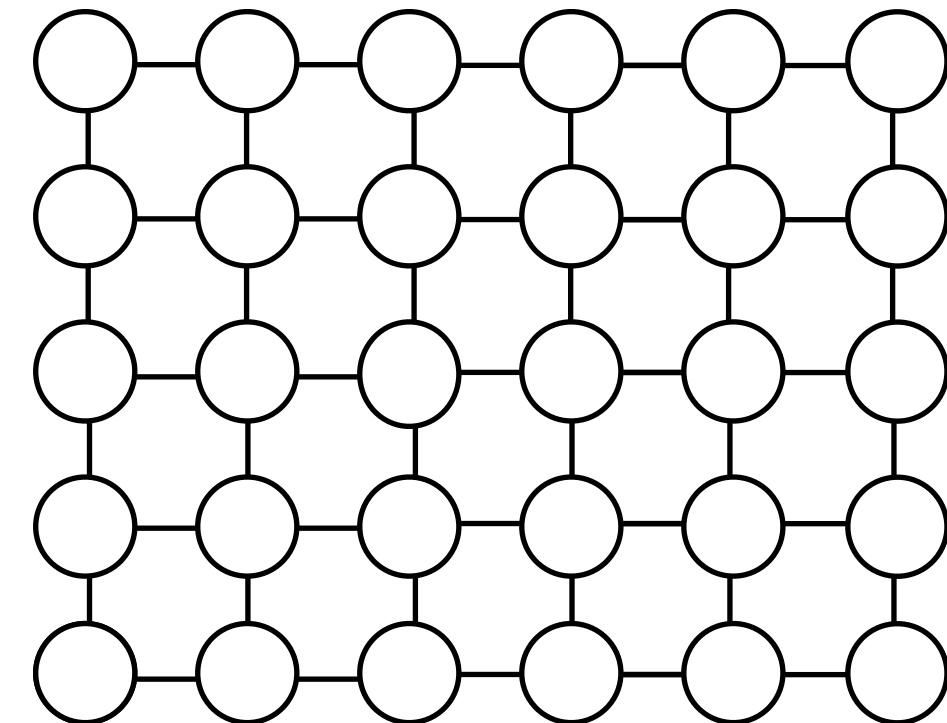
Future frame prediction



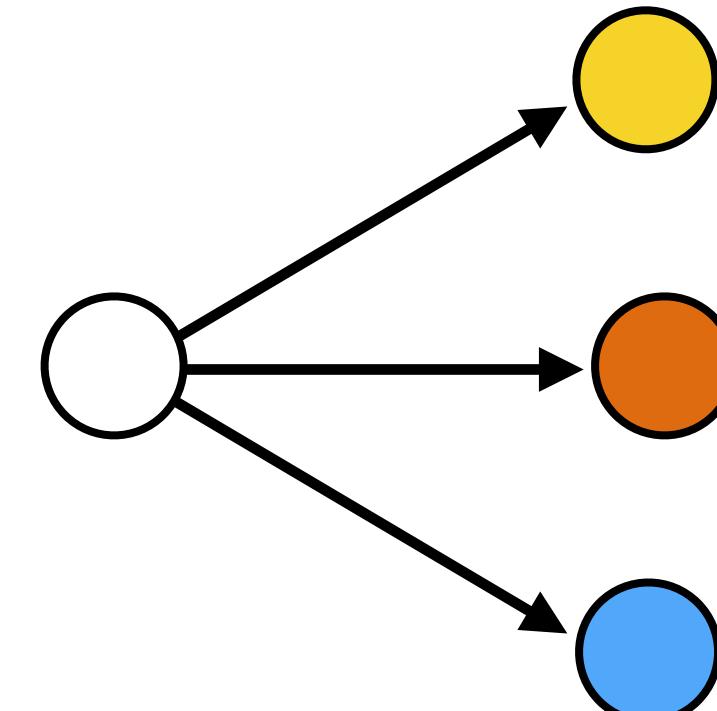
[Mathieu et al. 2016, ...]

Challenges in data prediction

1. Output is a high-dimensional, structured objects

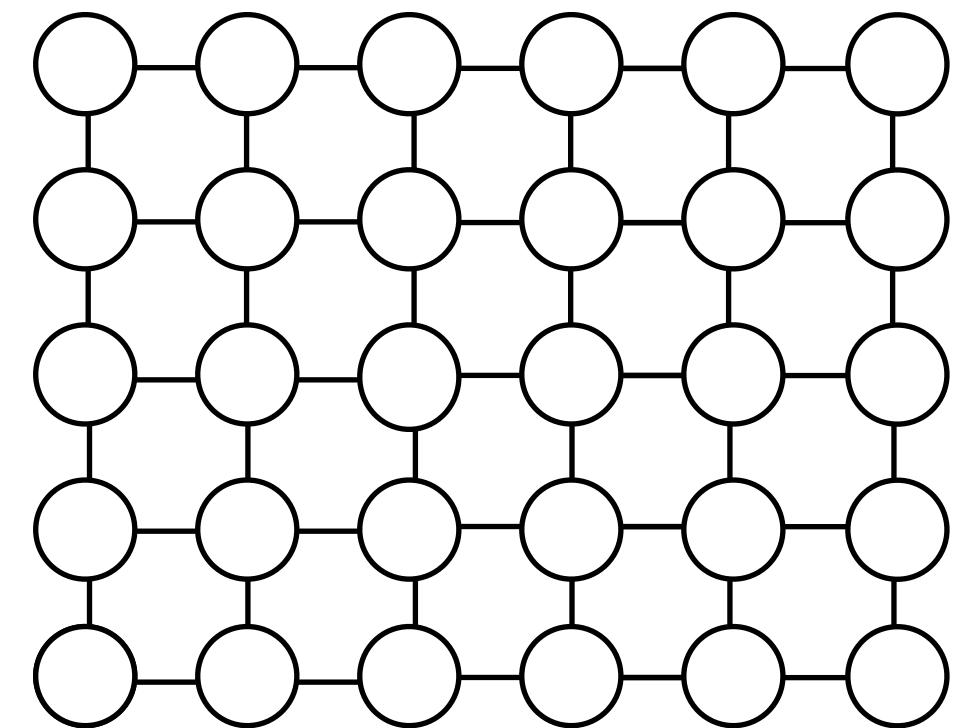


2. Uncertainty in the mapping, many plausible outputs

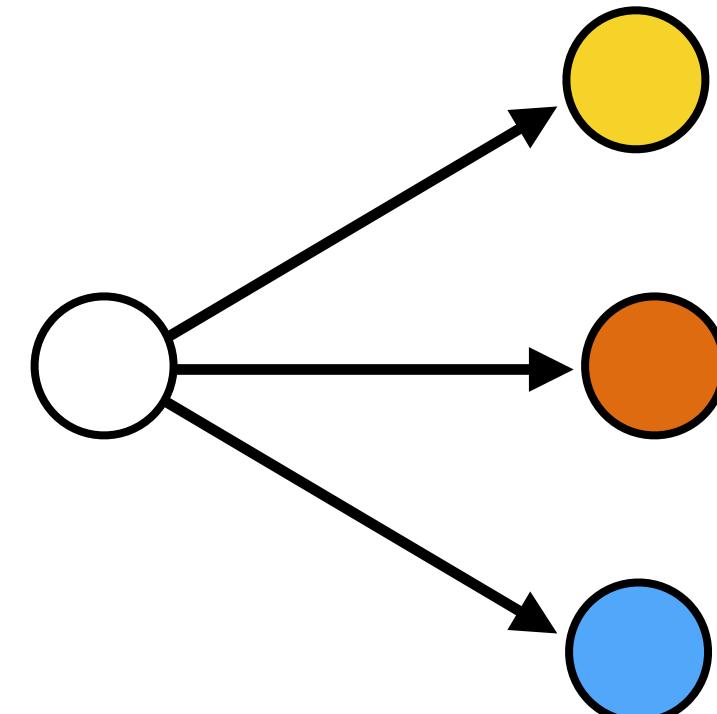


Properties of generative models

1. Model high-dimensional, structured objects



2. Model uncertainty; a whole distribution of possible outputs



Conditional Generative Models

Two ways to do it:

1. Model $p(x,y)$, then do inference to get conditional $p(y|x)$

$$\arg \max_y p(x, y) = \arg \max_y p(y|x)$$

2. Directly model $p(y|x)$

Image-to-Image Translation

Input \mathbf{x}

<i>Training data</i>	
\mathbf{x}	\mathbf{y}
{  ,  }	
{  ,  }	
{  ,  }	
:	

Output \mathbf{y}



$$\arg \min_{\mathcal{F}} \mathbb{E}_{\mathbf{x}, \mathbf{y}} [L(\mathcal{F}(\mathbf{x}), \mathbf{y})]$$

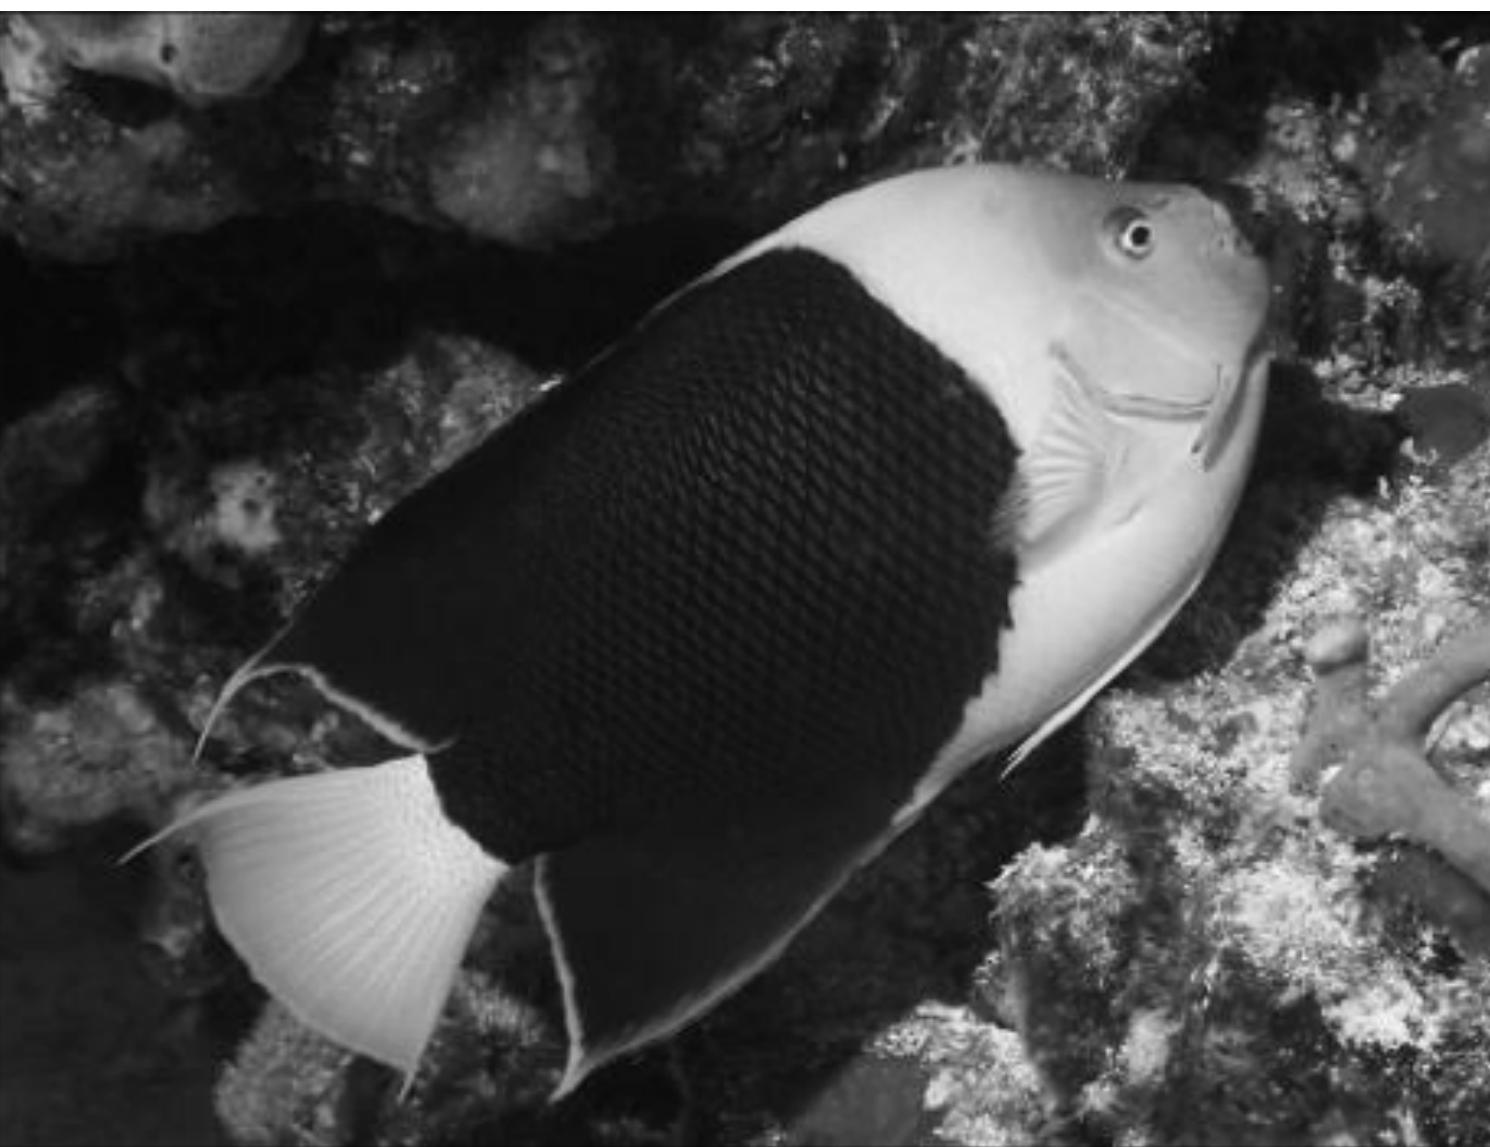
Objective function
(loss)

Neural Network

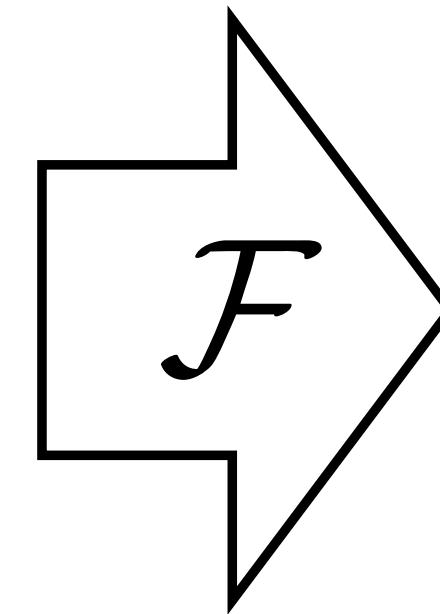
[Zhang et al., ECCV 2016]

Image-to-Image Translation

Input \mathbf{x}



Output \mathbf{y}



$$\arg \min_{\mathcal{F}} \mathbb{E}_{\mathbf{x}, \mathbf{y}} [L(\mathcal{F}(\mathbf{x}), \mathbf{y})]$$

“What should I do”

“How should I do it?”

Designing loss functions

Input



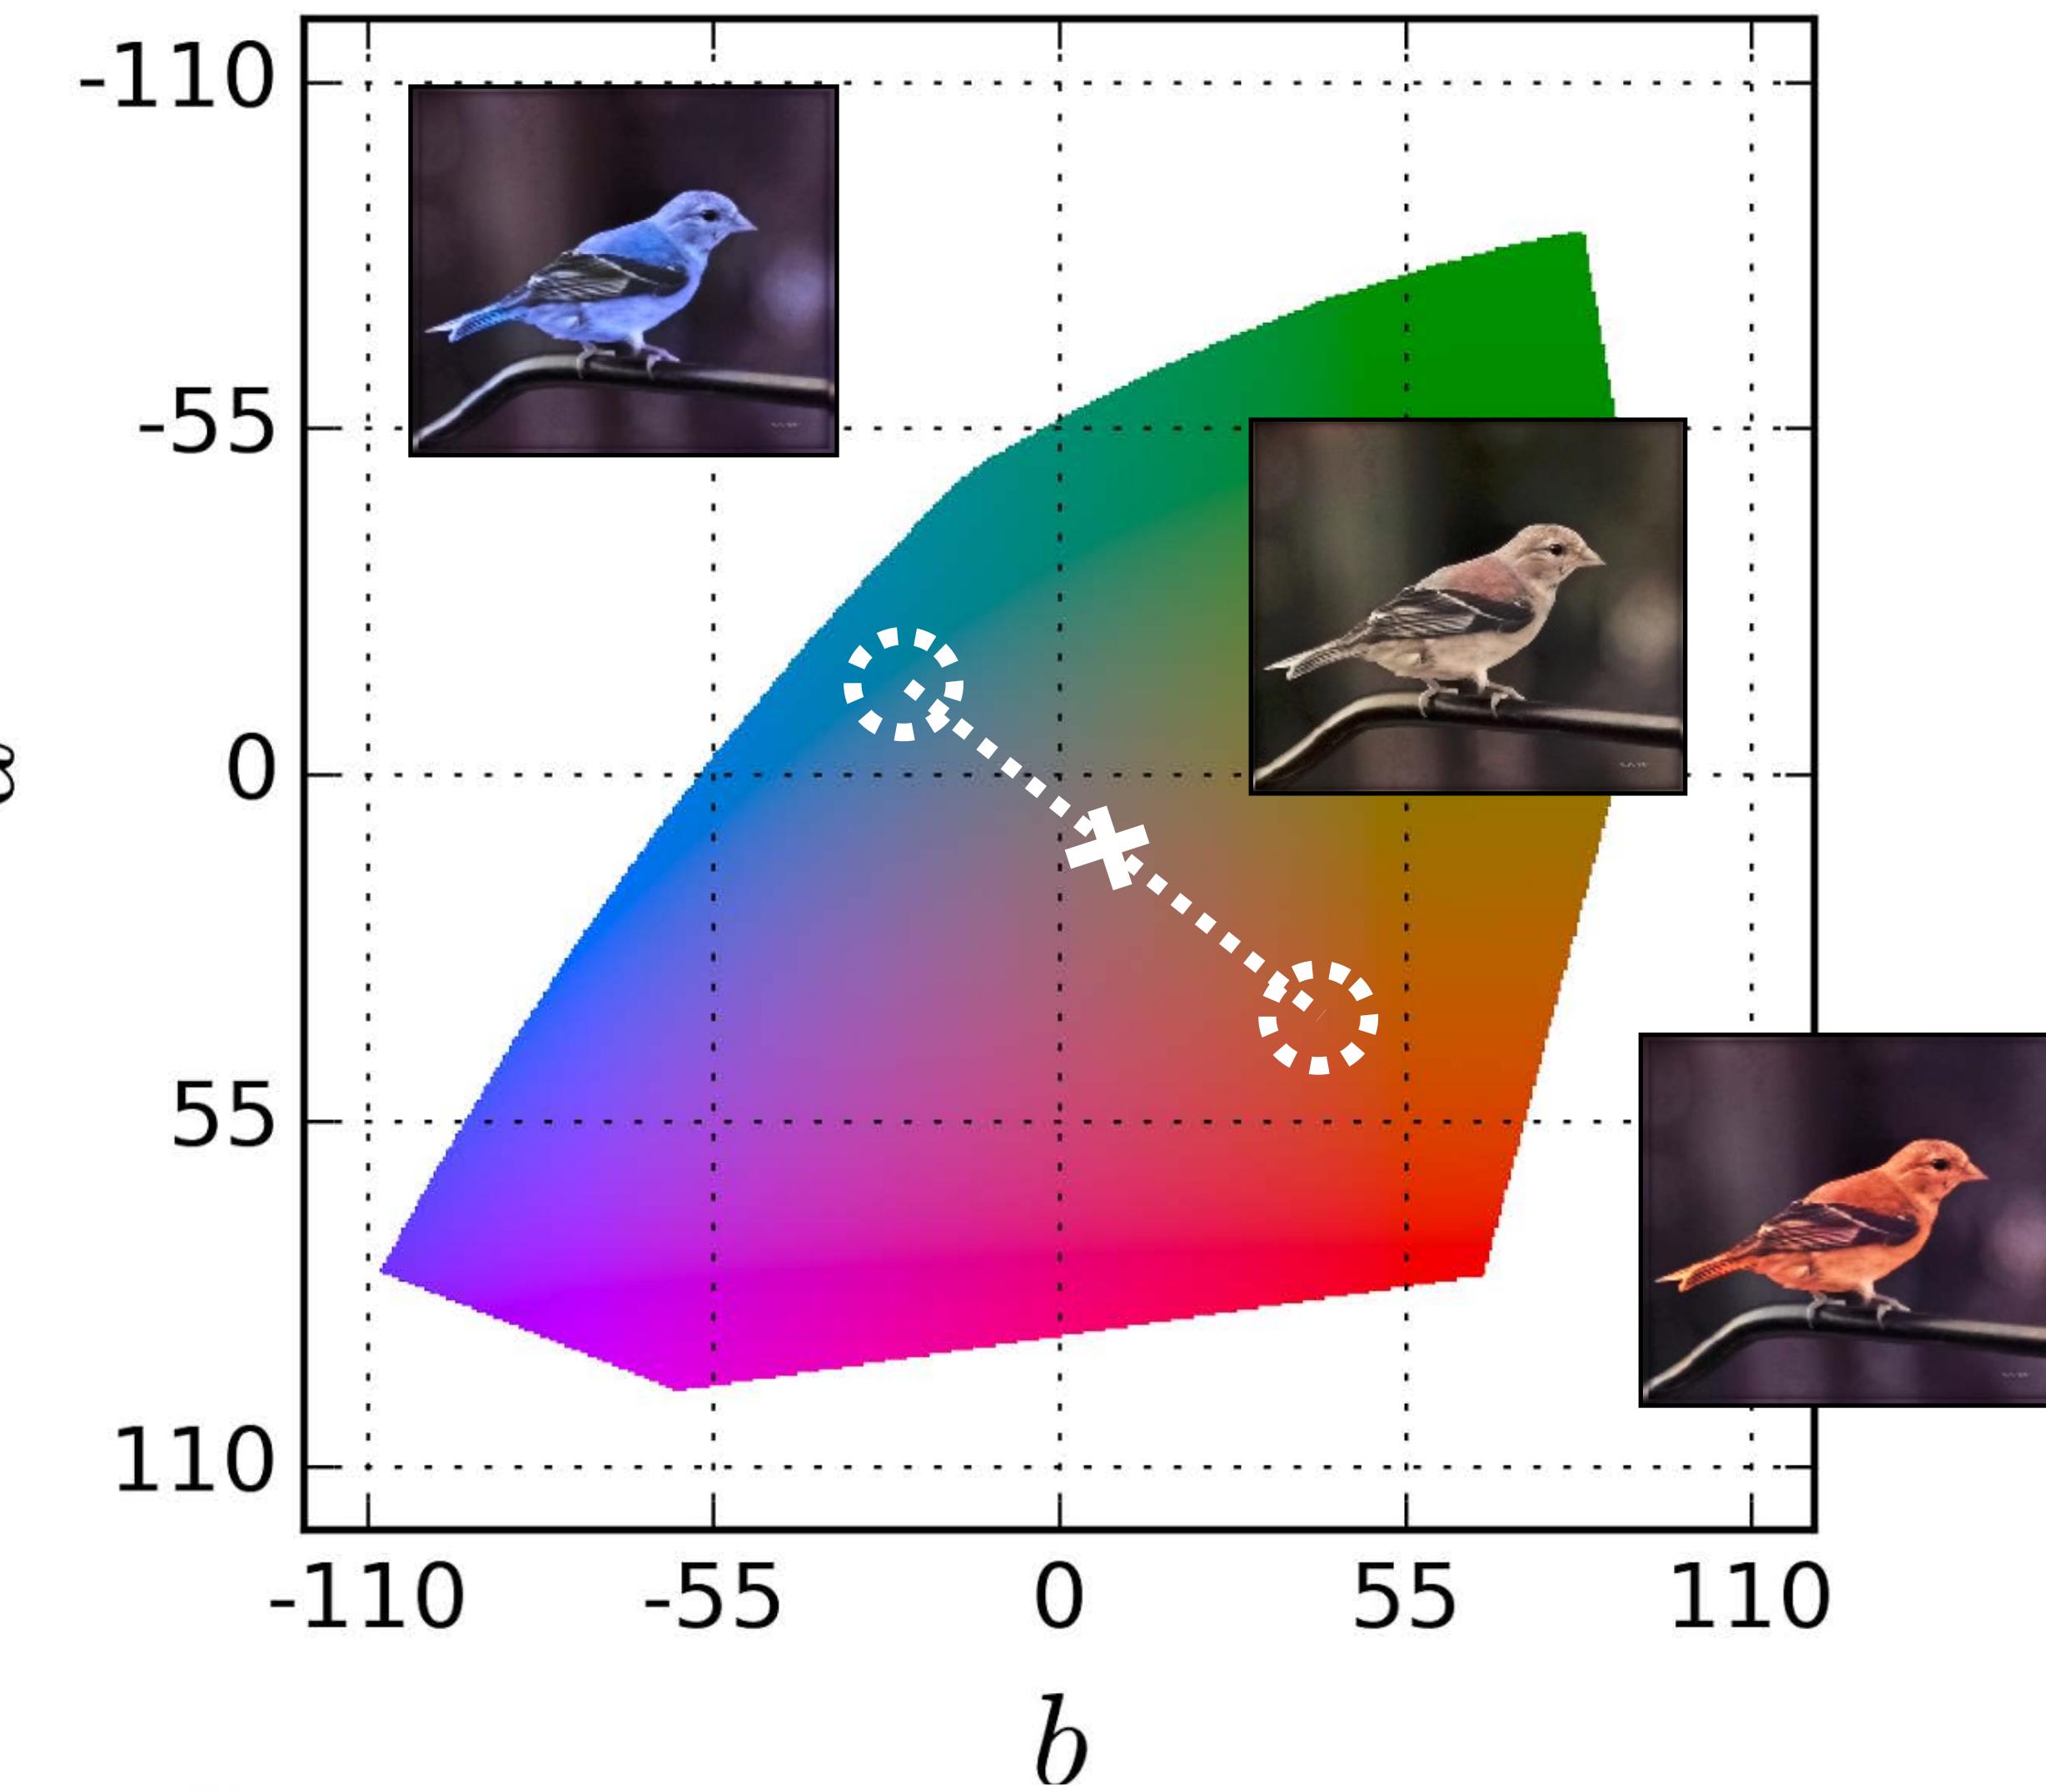
Output



Ground truth



$$L_2(\hat{\mathbf{Y}}, \mathbf{Y}) = \frac{1}{2} \sum_{h,w} \|\mathbf{Y}_{h,w} - \hat{\mathbf{Y}}_{h,w}\|_2^2$$



$$L_2(\hat{\mathbf{Y}}, \mathbf{Y}) = \frac{1}{2} \sum_{h,w} \|\mathbf{Y}_{h,w} - \hat{\mathbf{Y}}_{h,w}\|_2^2$$

Designing loss functions

Input



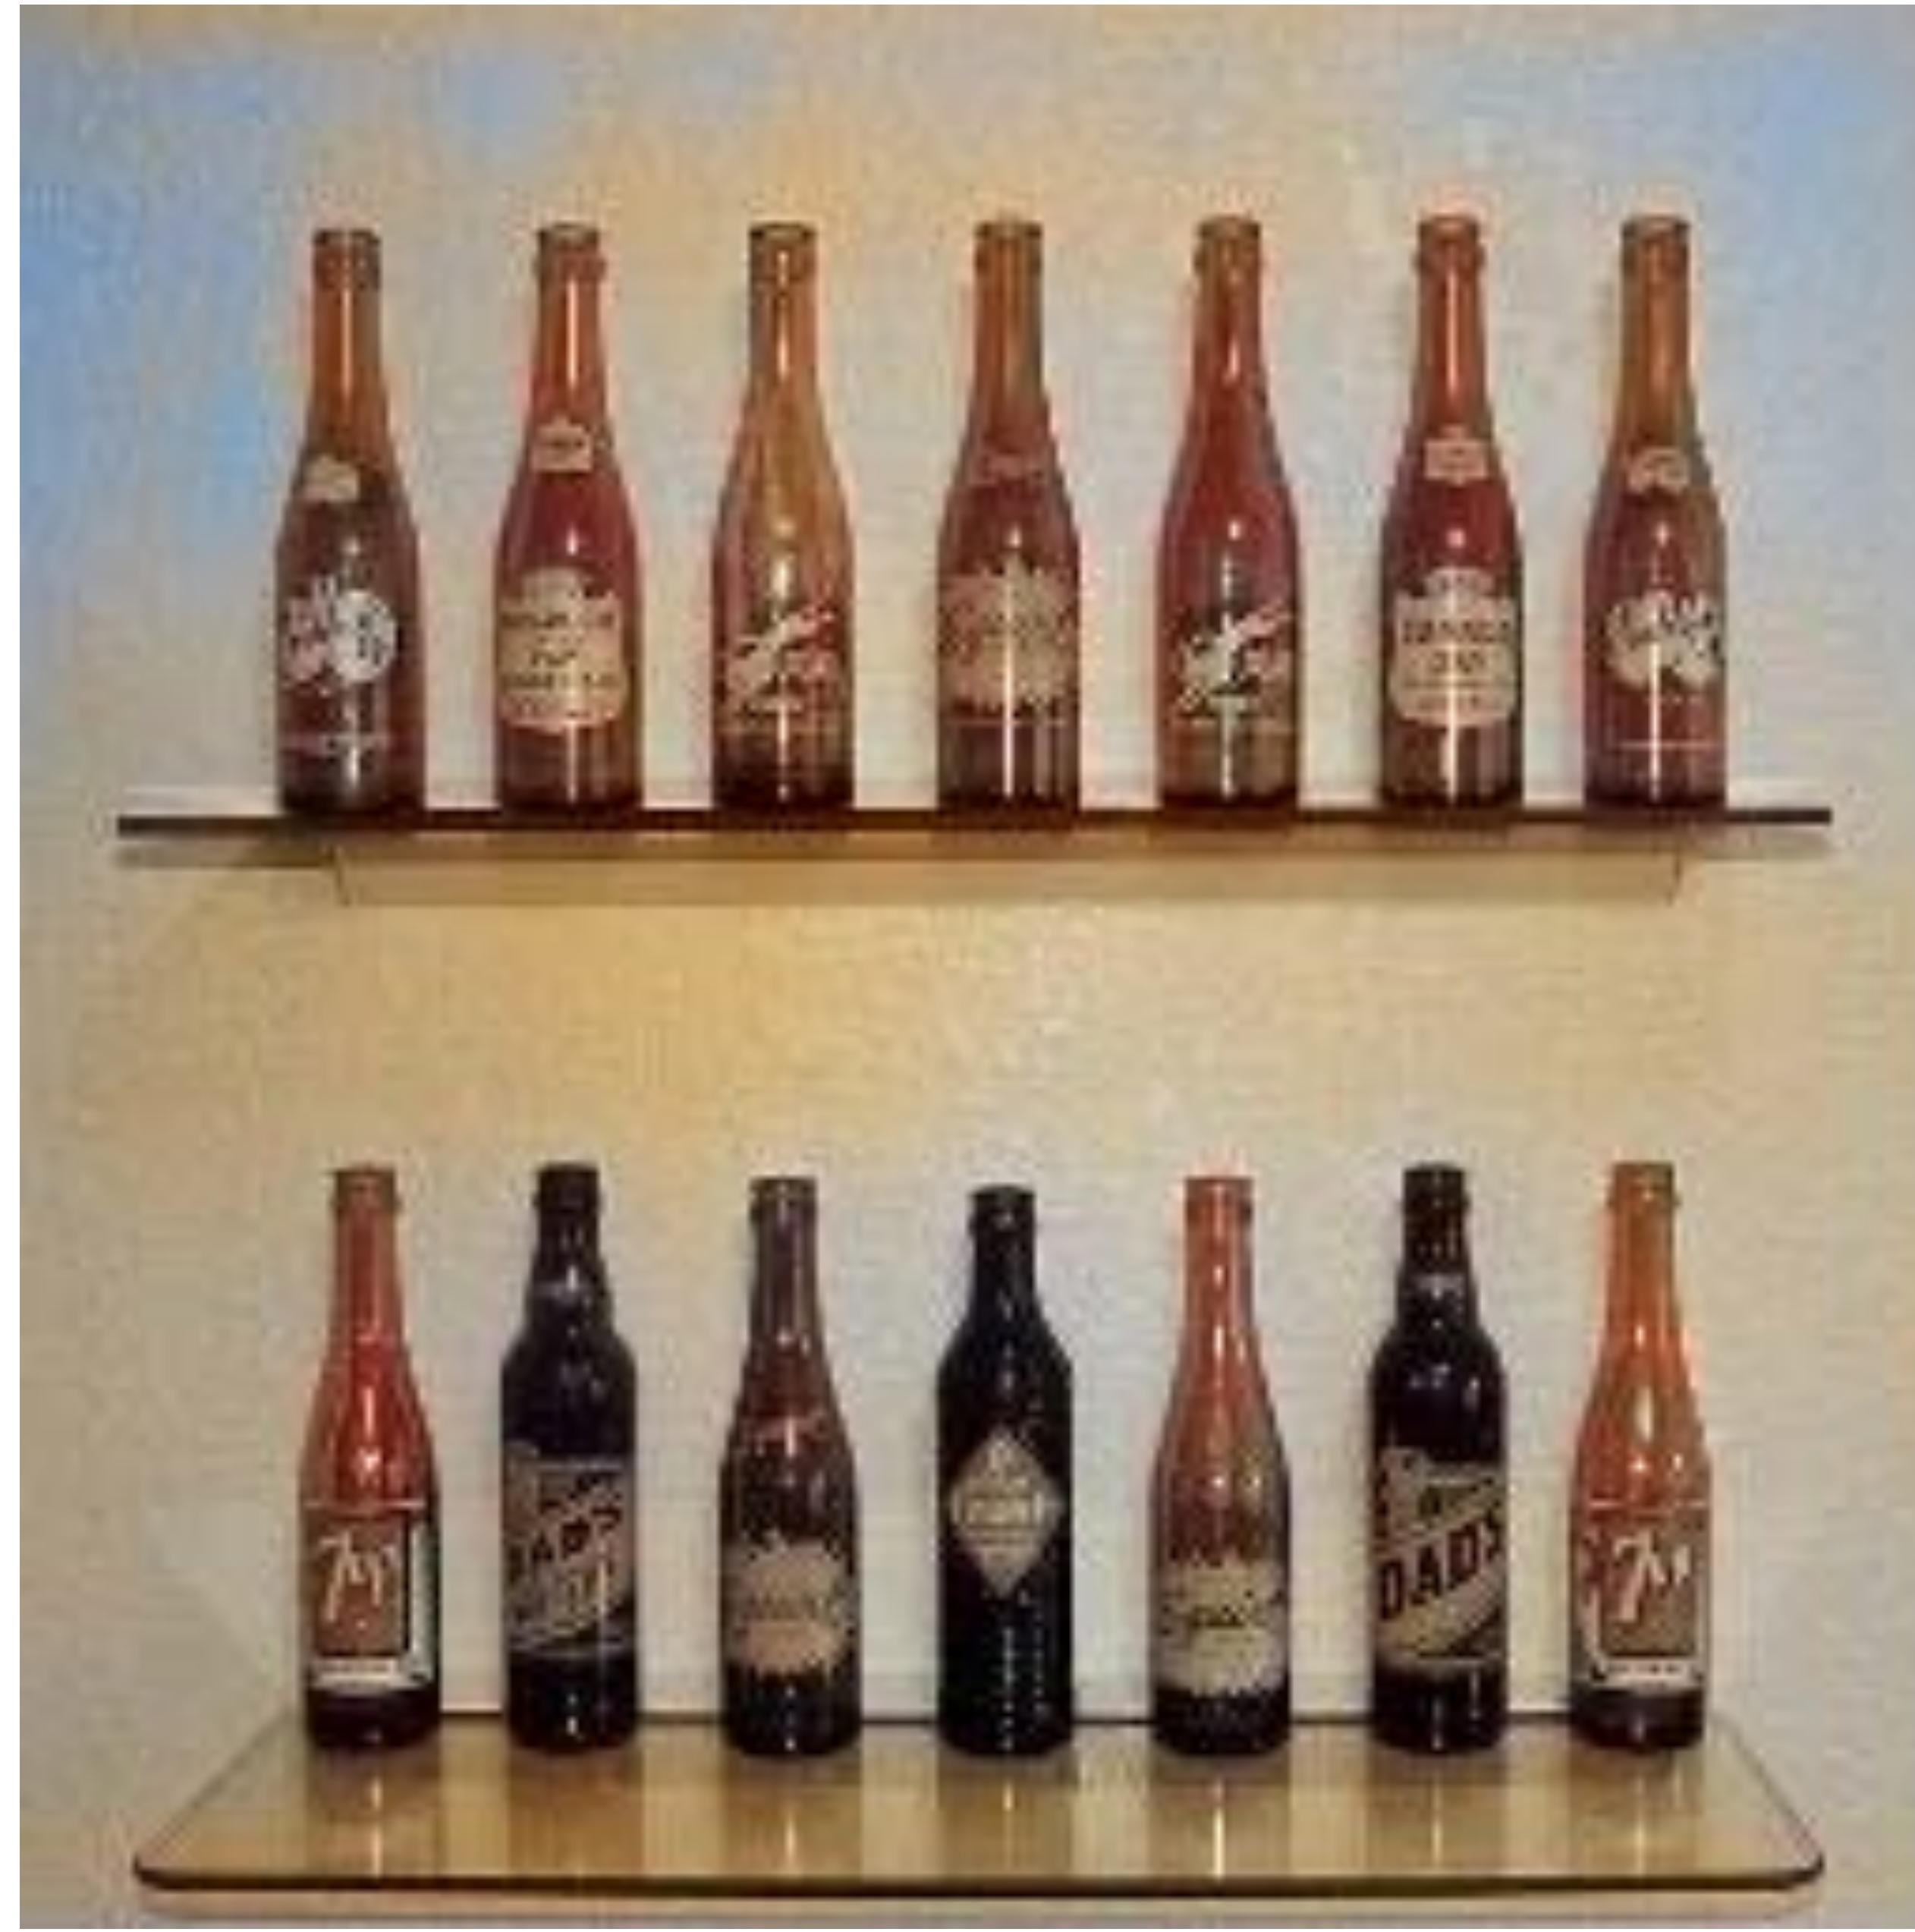
Zhang et al. 2016



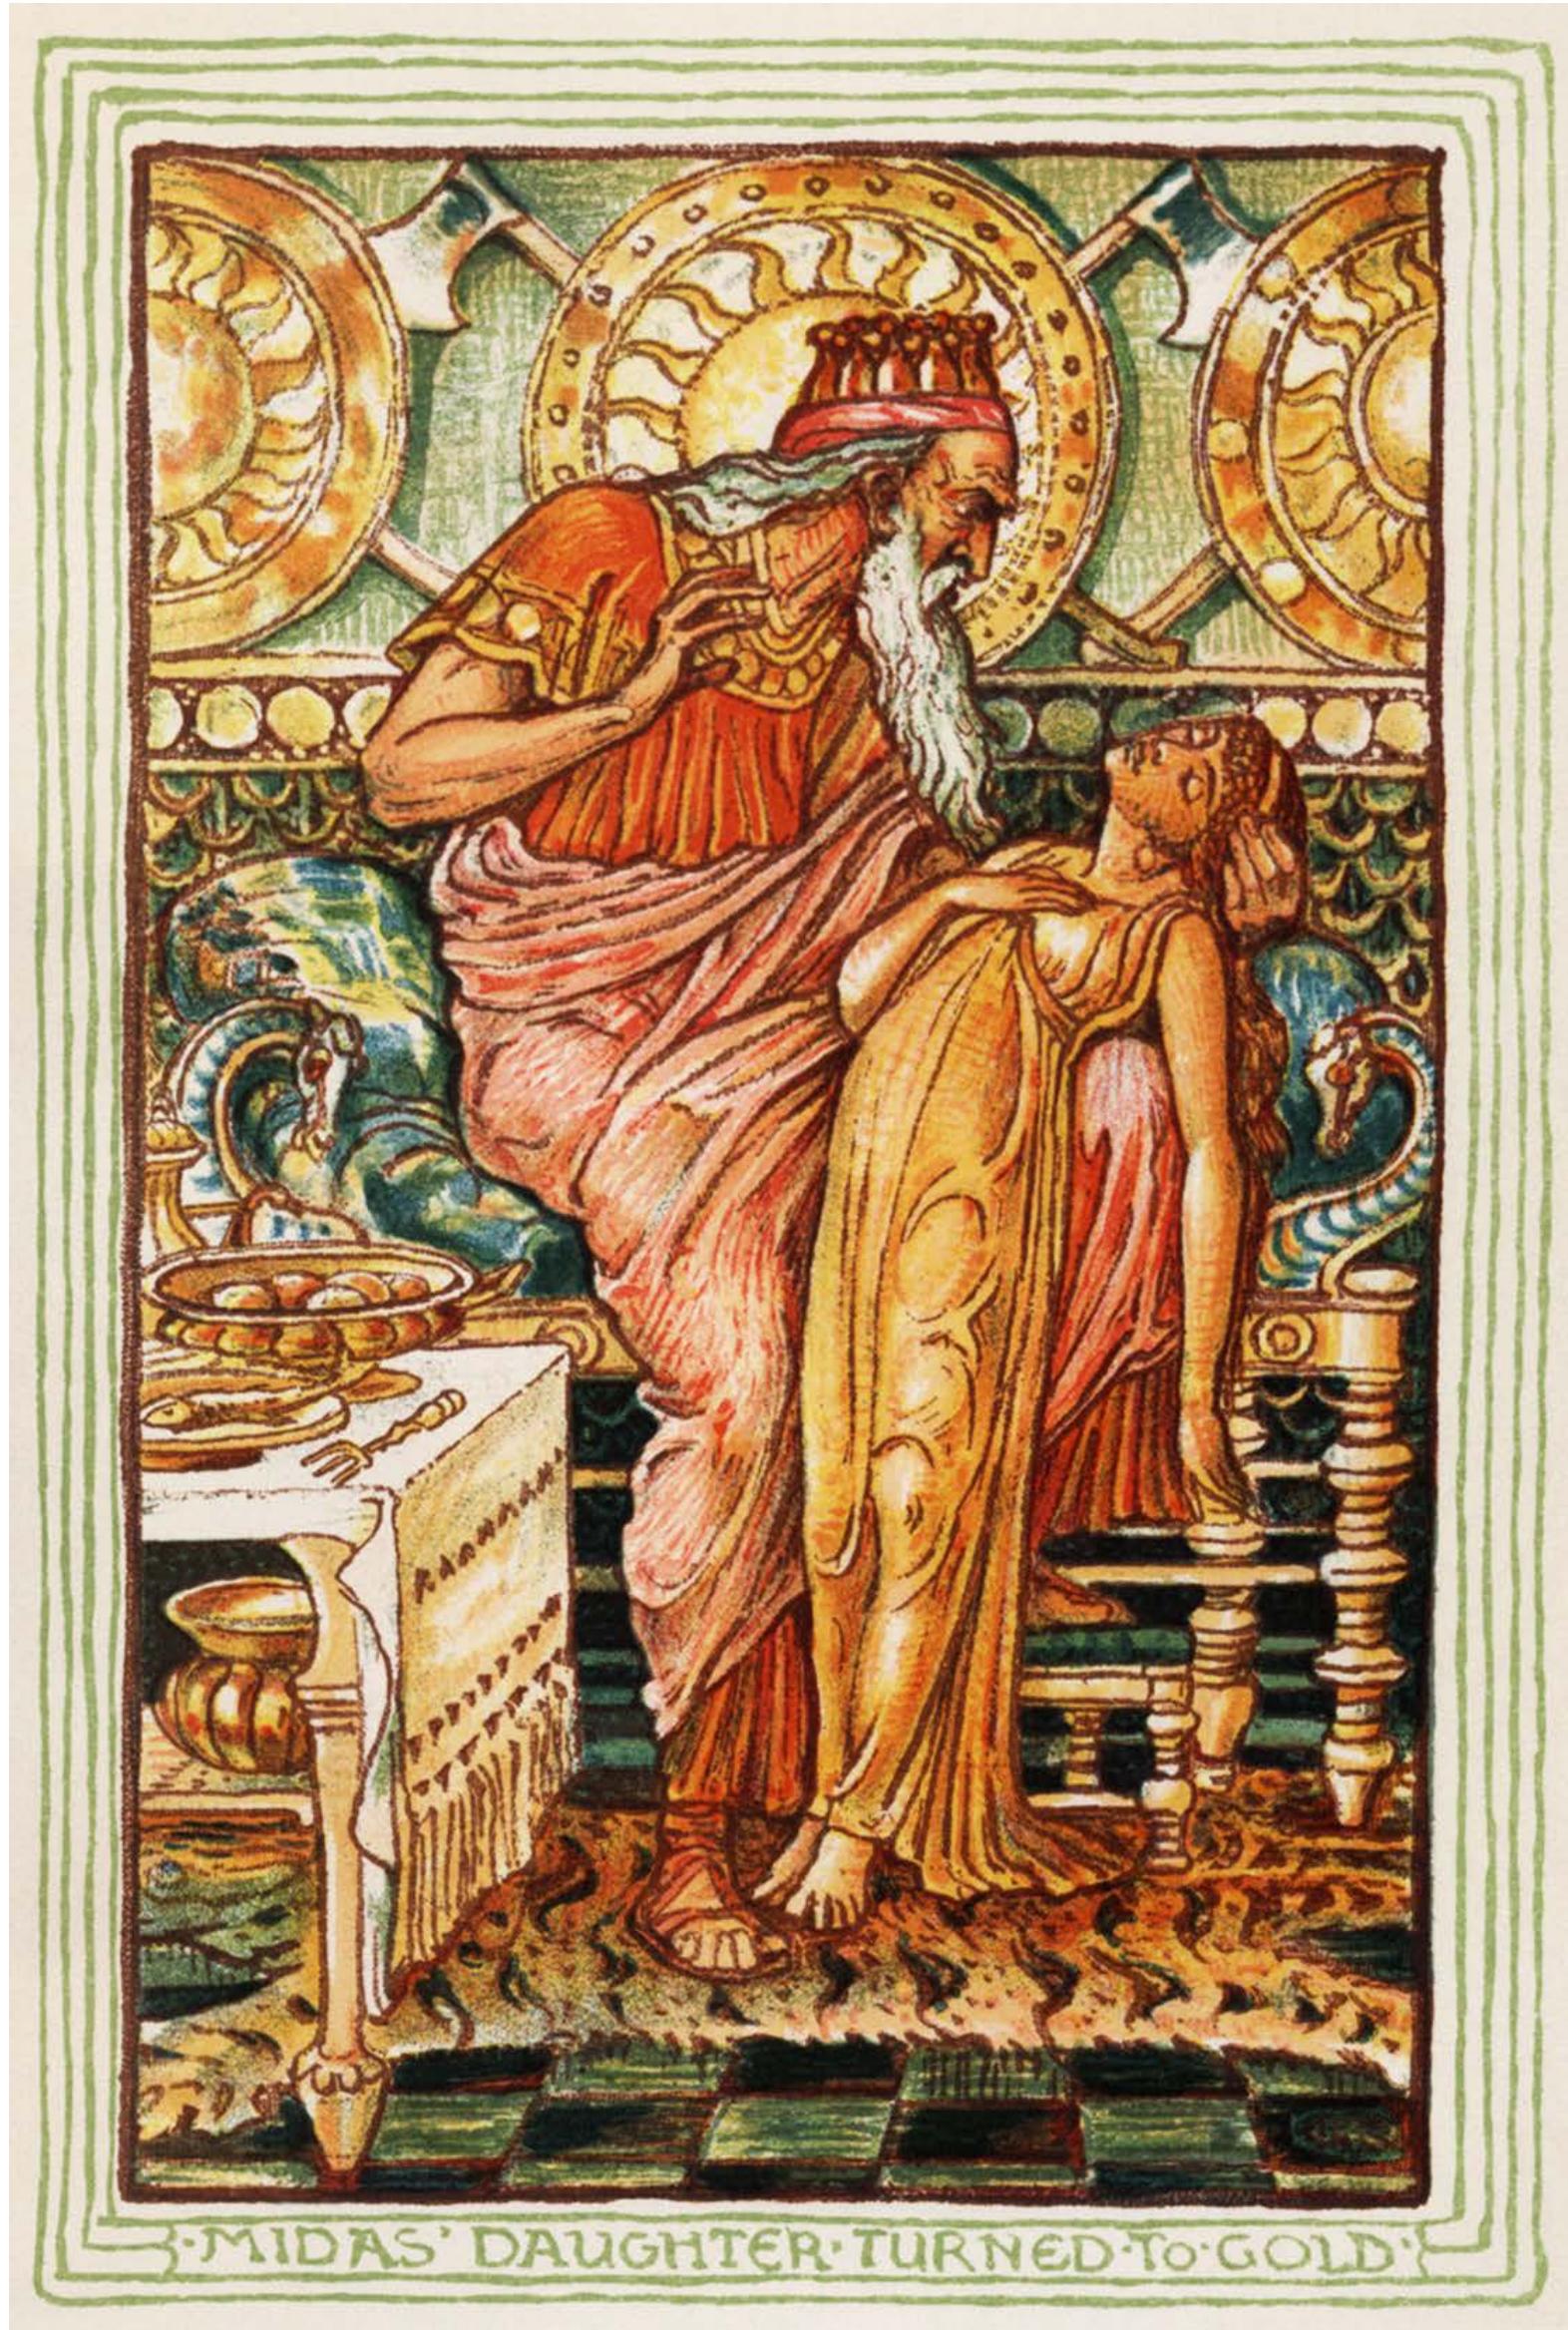
Ground truth



Color distribution cross-entropy loss with colorfulness enhancing term.



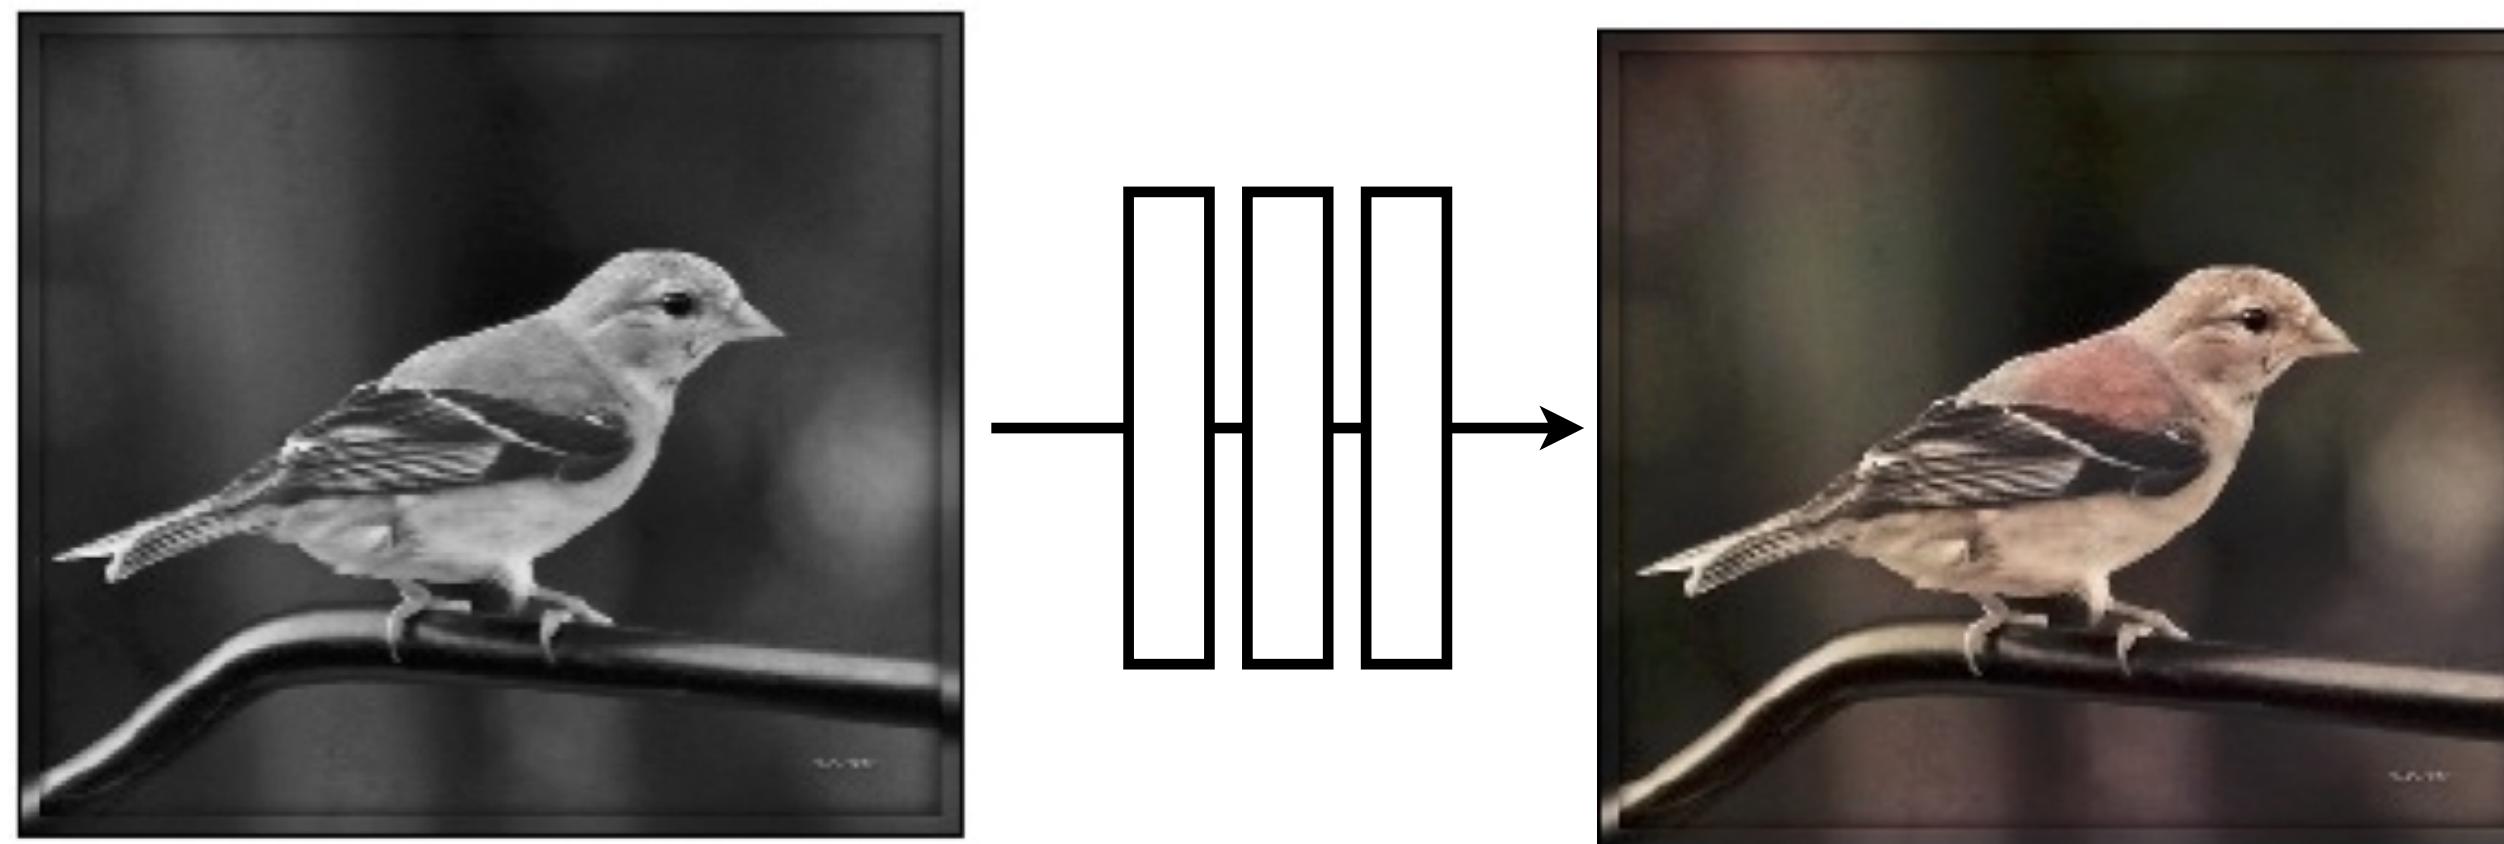
Designing loss functions



Be careful what you wish for!

Designing loss functions

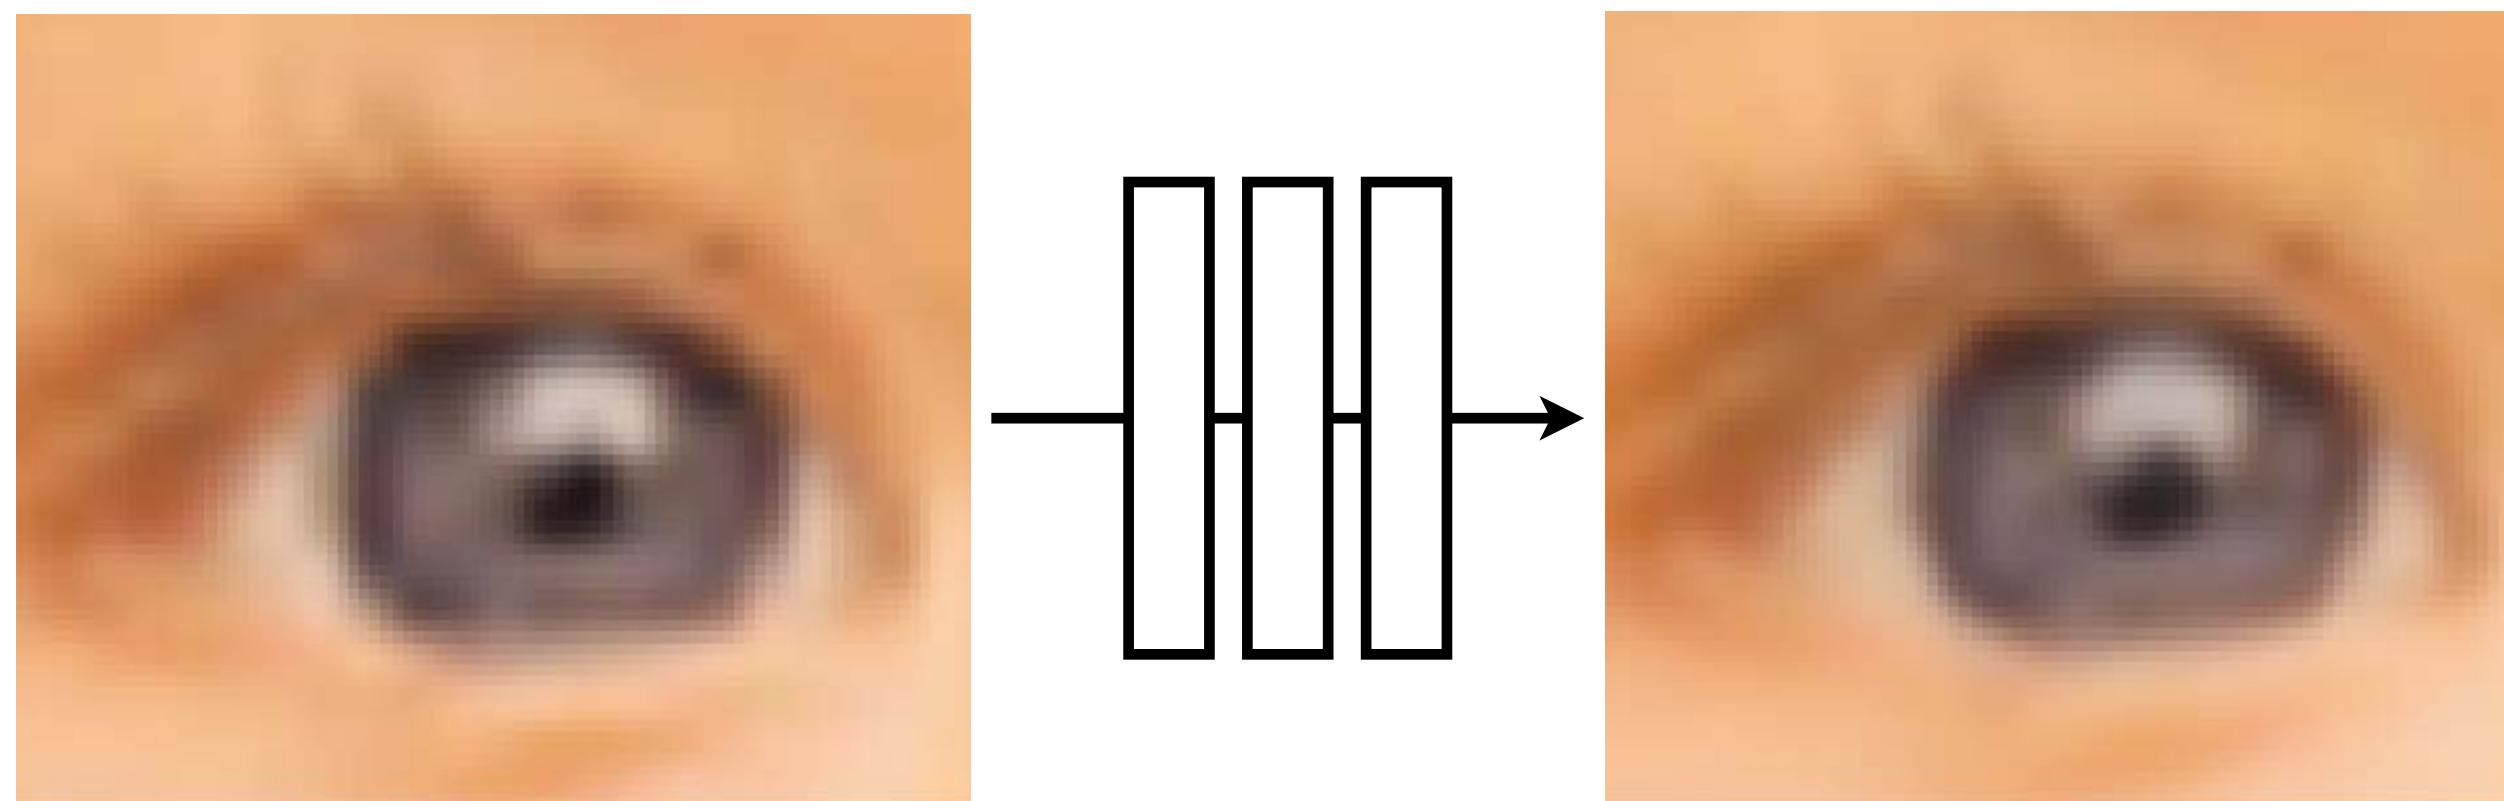
Image colorization



L2 regression

[Zhang, Isola, Efros, ECCV 2016]

Super-resolution

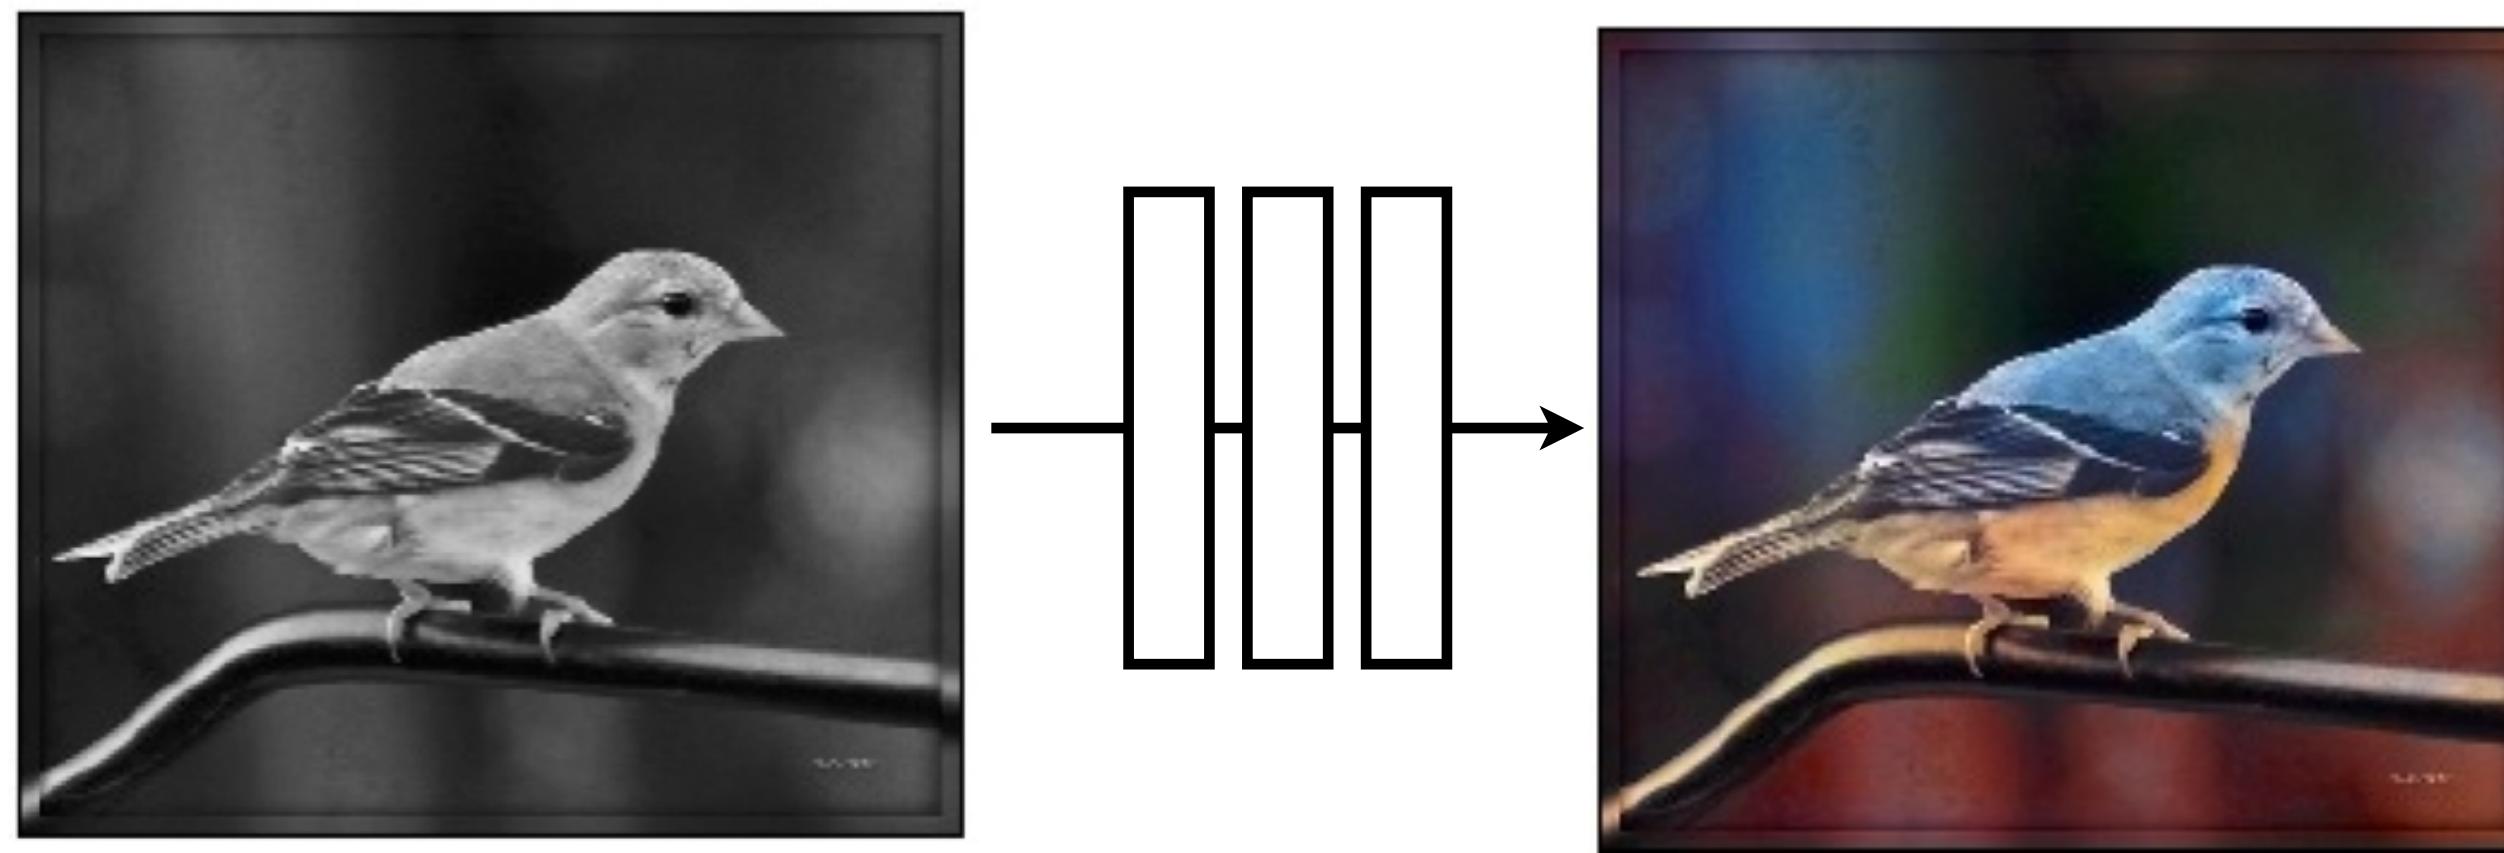


L2 regression

[Johnson, Alahi, Li, ECCV 2016]

Designing loss functions

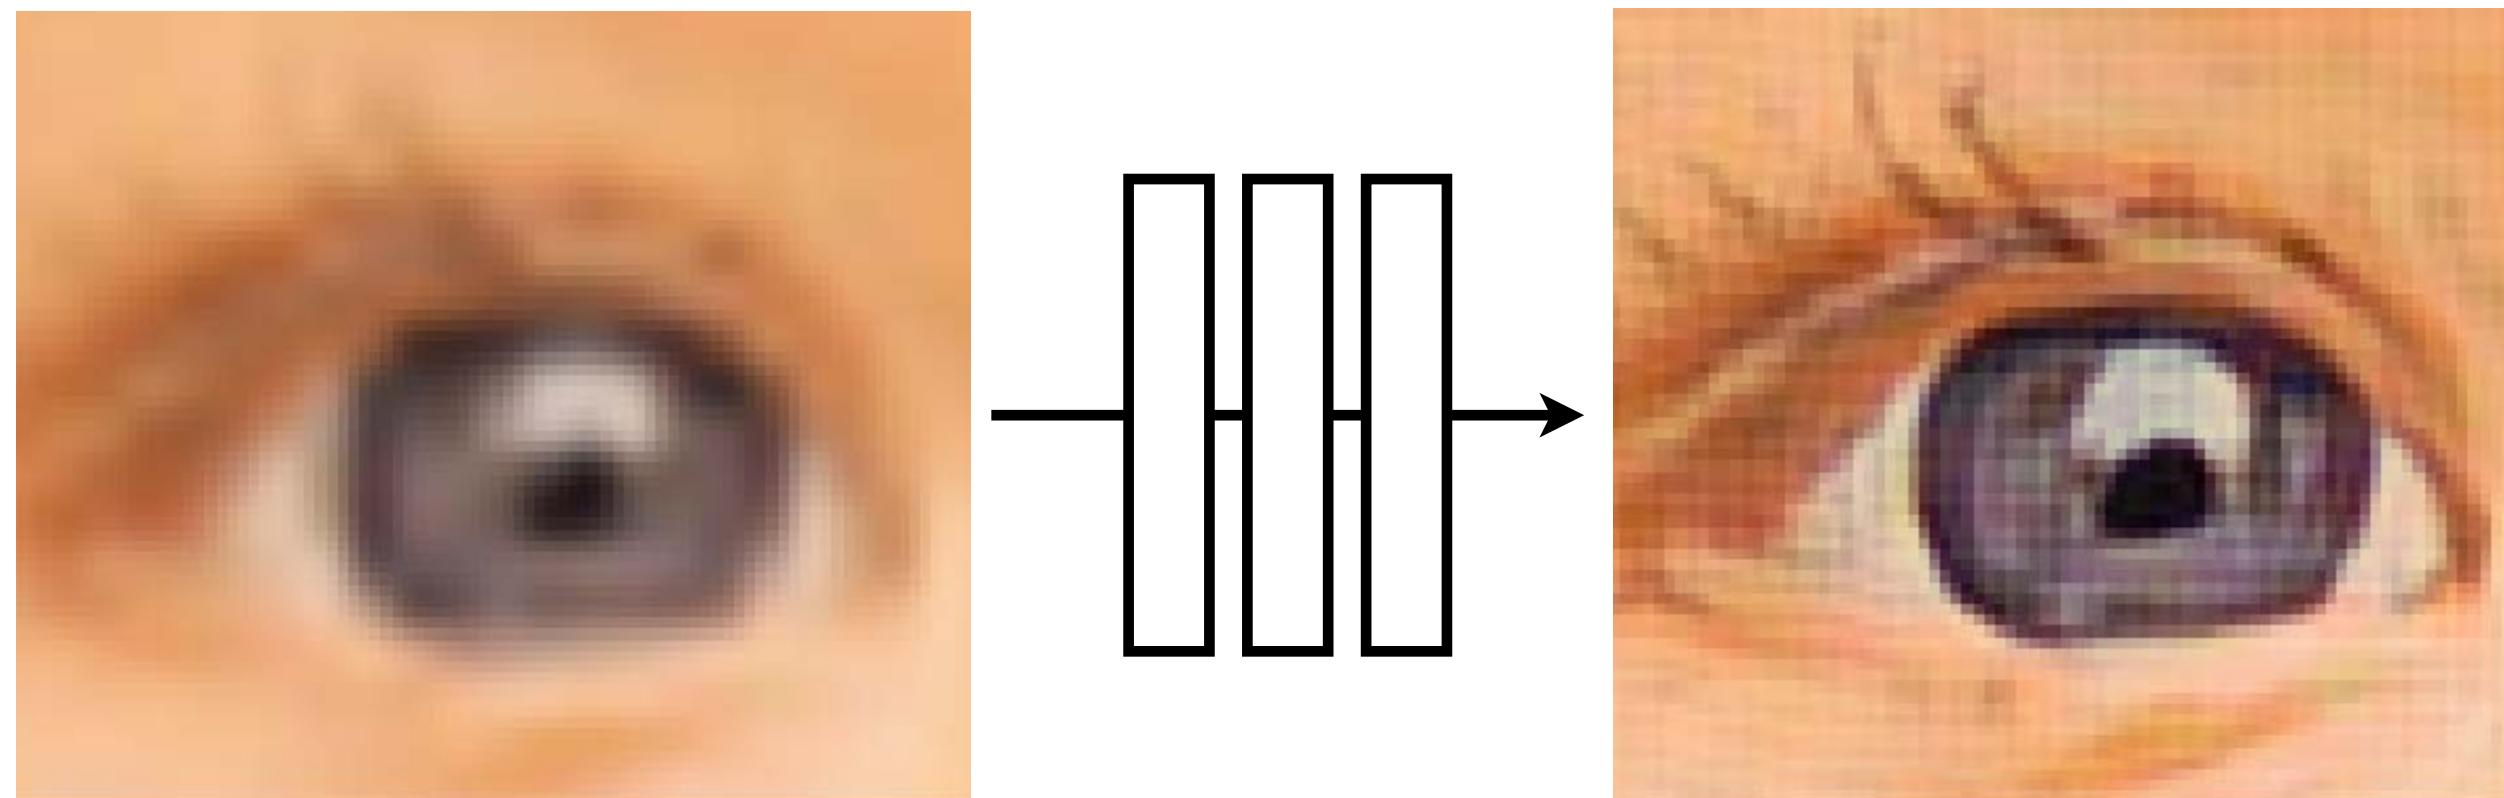
Image colorization



[Zhang, Isola, Efros, ECCV 2016]

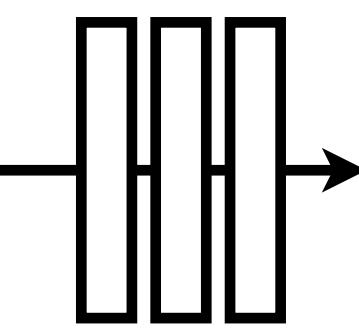
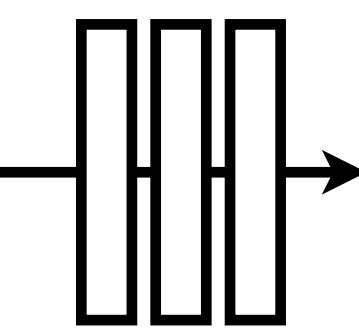
Cross entropy objective,
with colorfulness term

Super-resolution



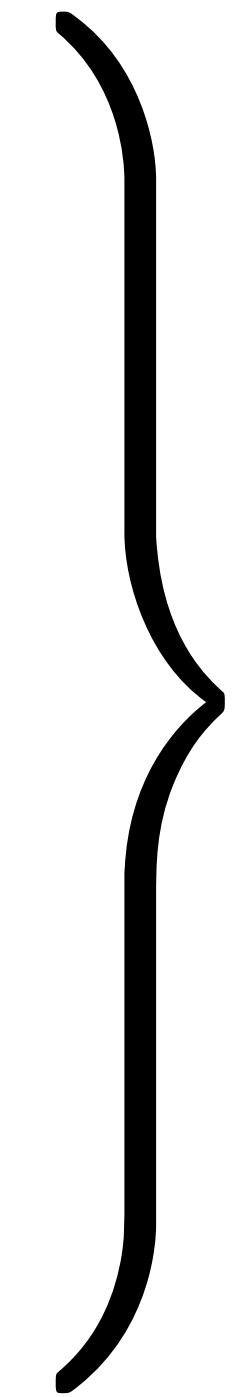
[Johnson, Alahi, Li, ECCV 2016]

Deep feature covariance
matching objective



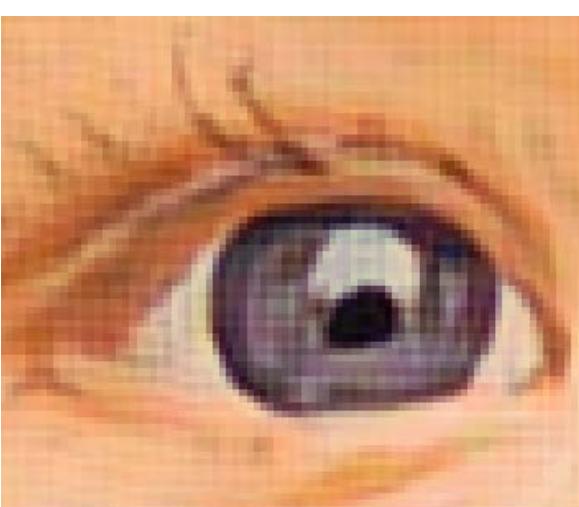
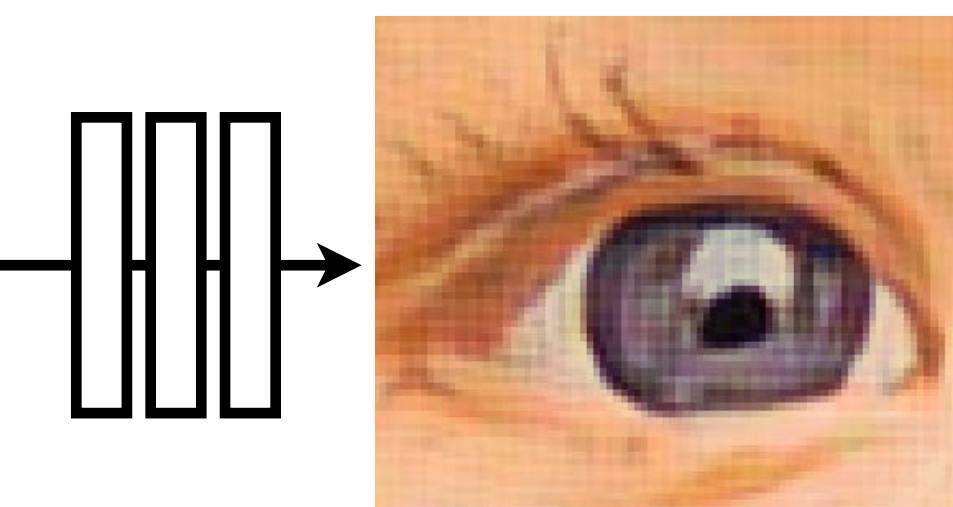
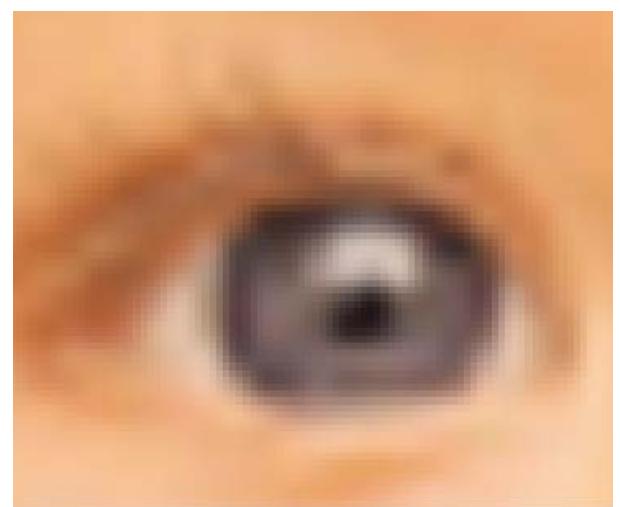
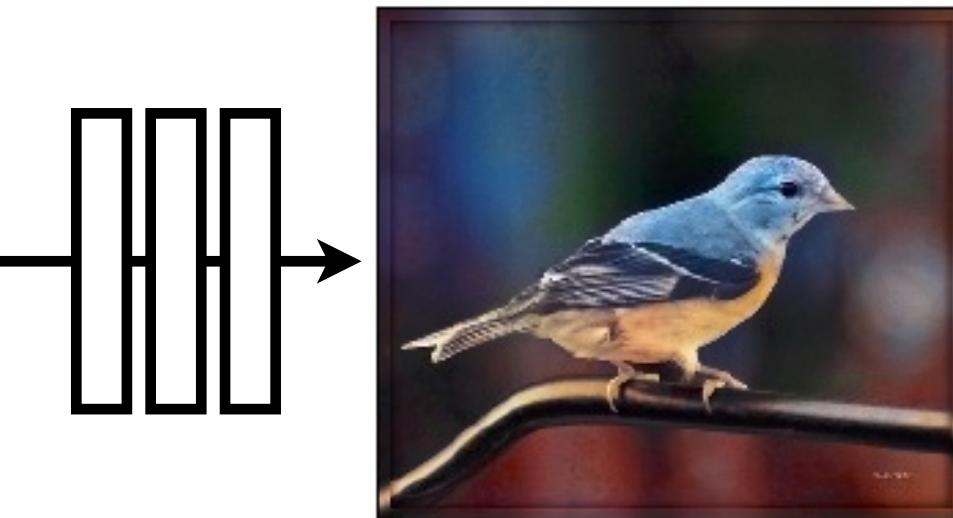
⋮

⋮



Universal loss?

Generated images

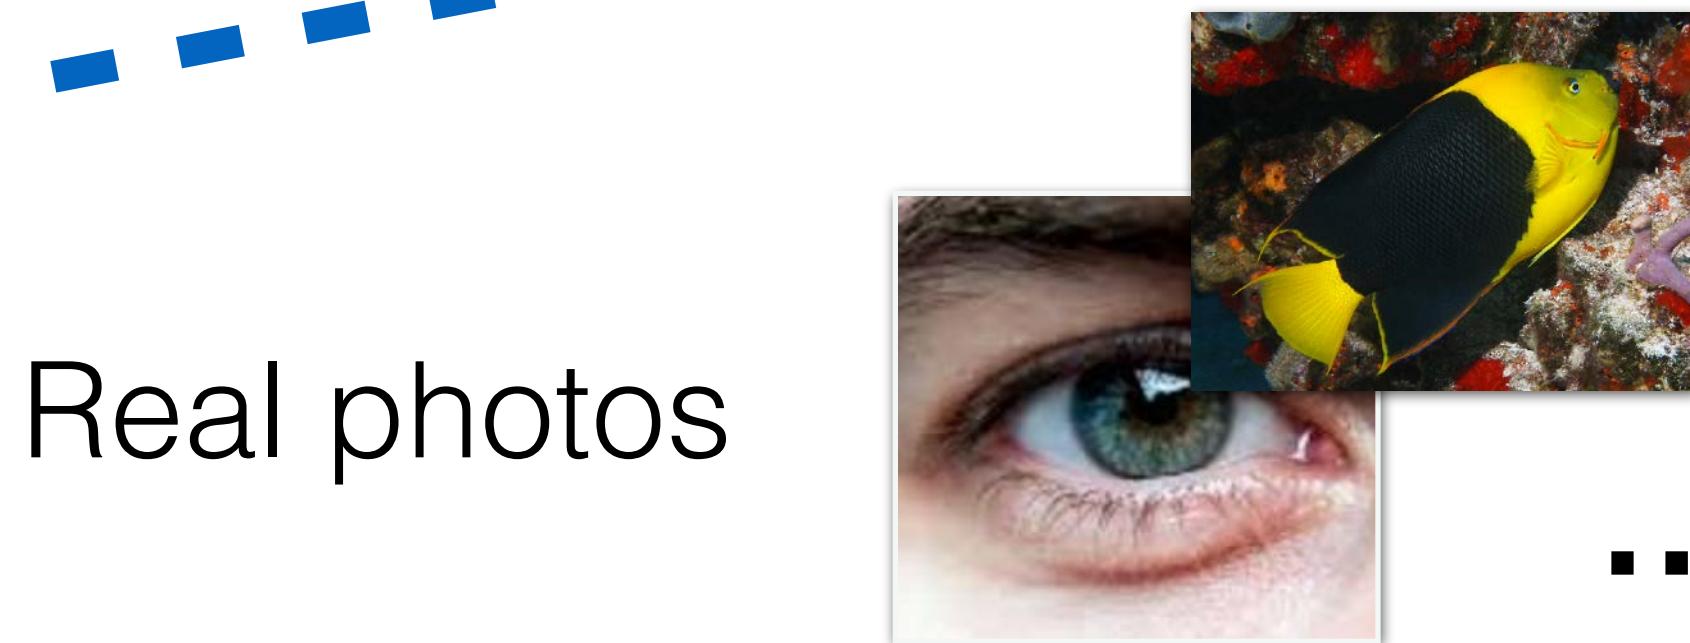


:

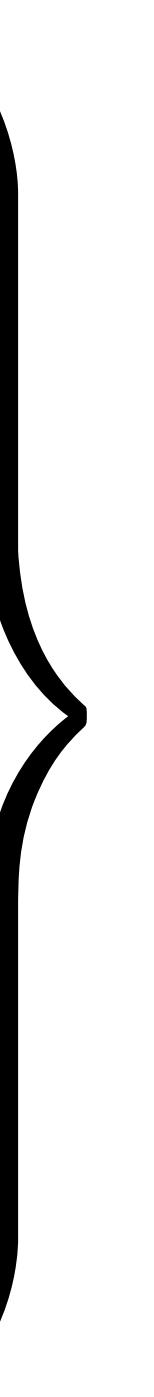
:

...

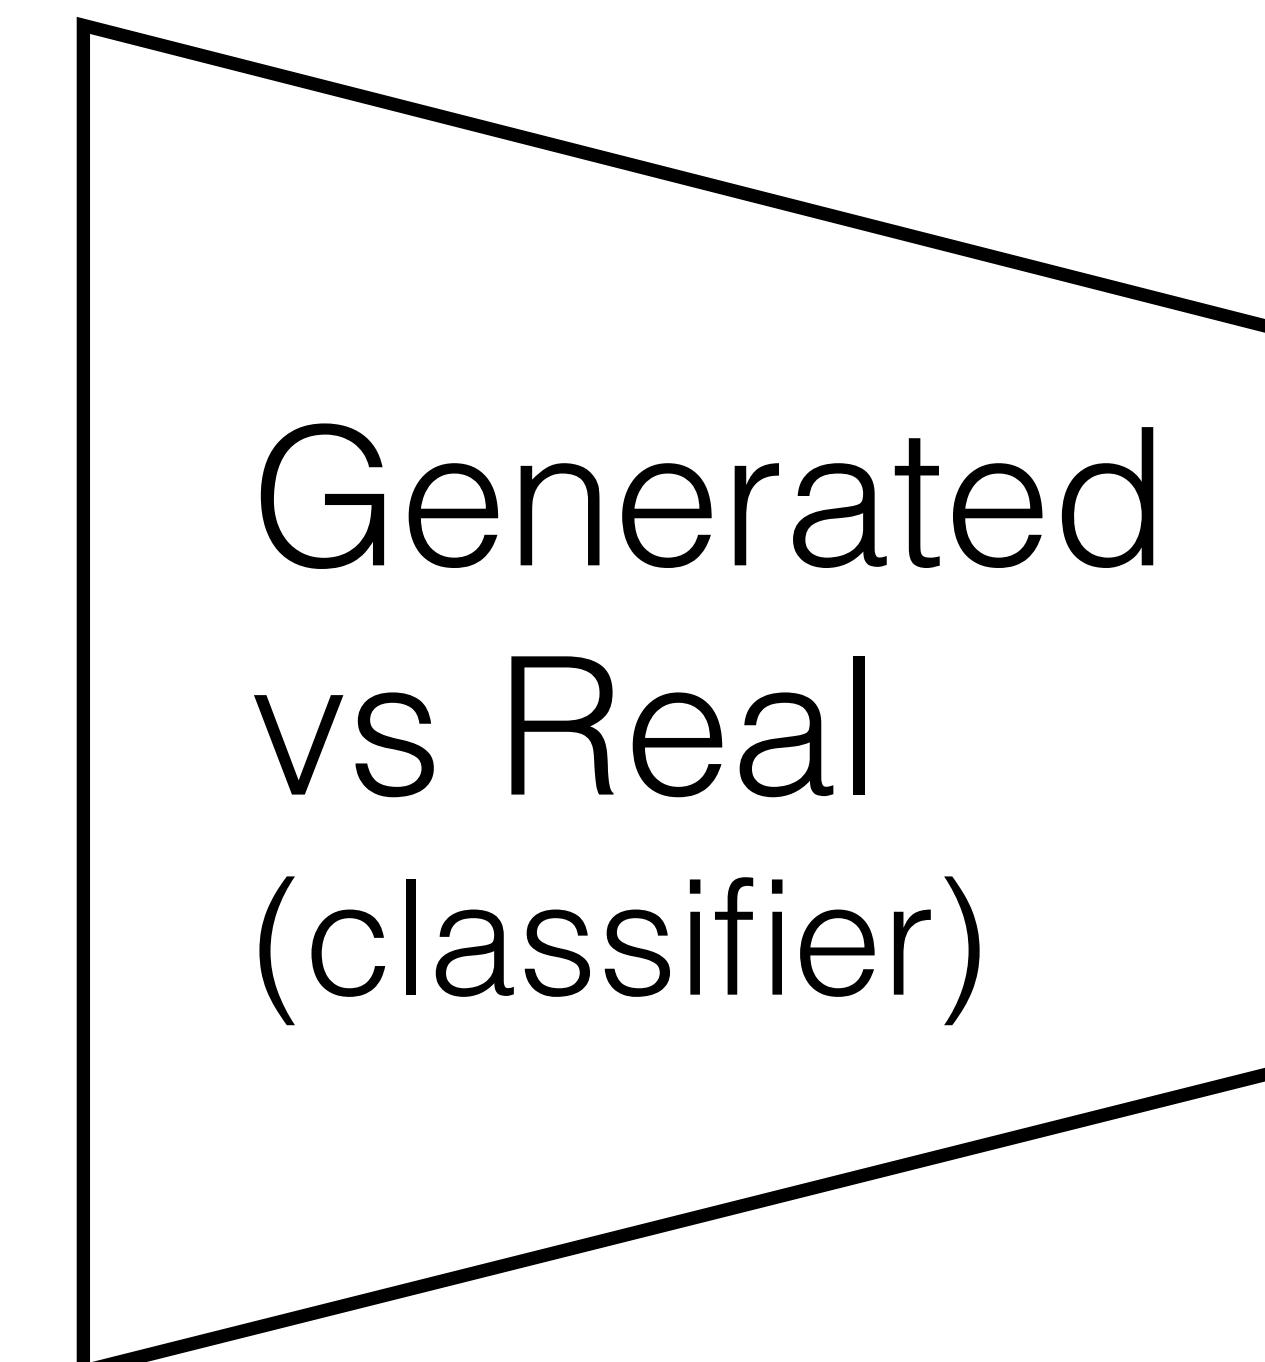
“Generative Adversarial Network” (GANs)



Real photos



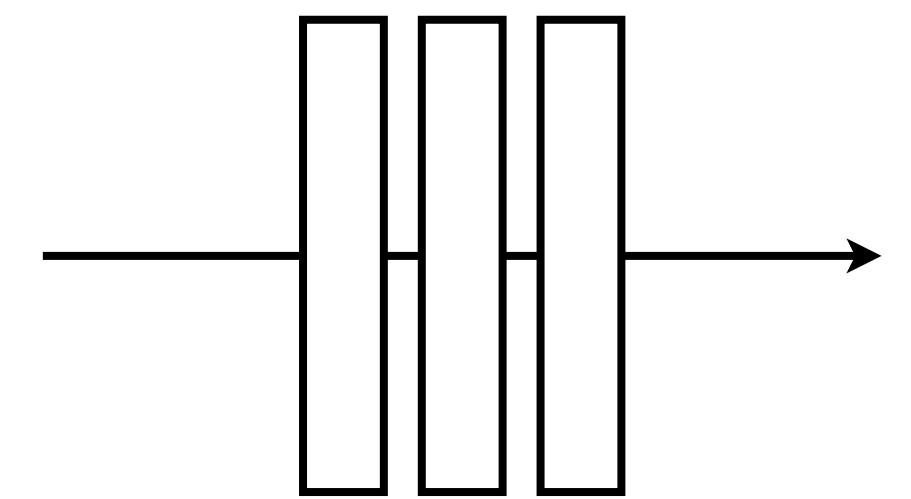
[Goodfellow, Pouget-Abadie, Mirza, Xu,
Warde-Farley, Ozair, Courville, Bengio 2014]



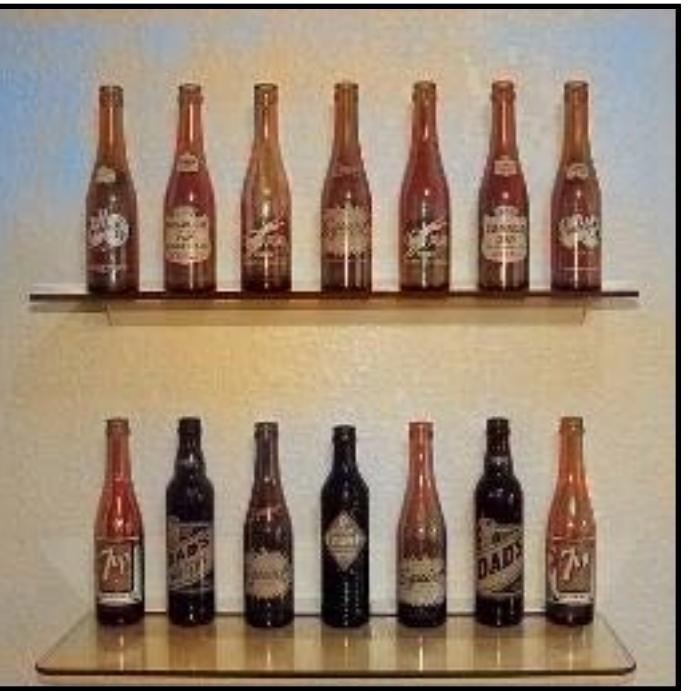
x



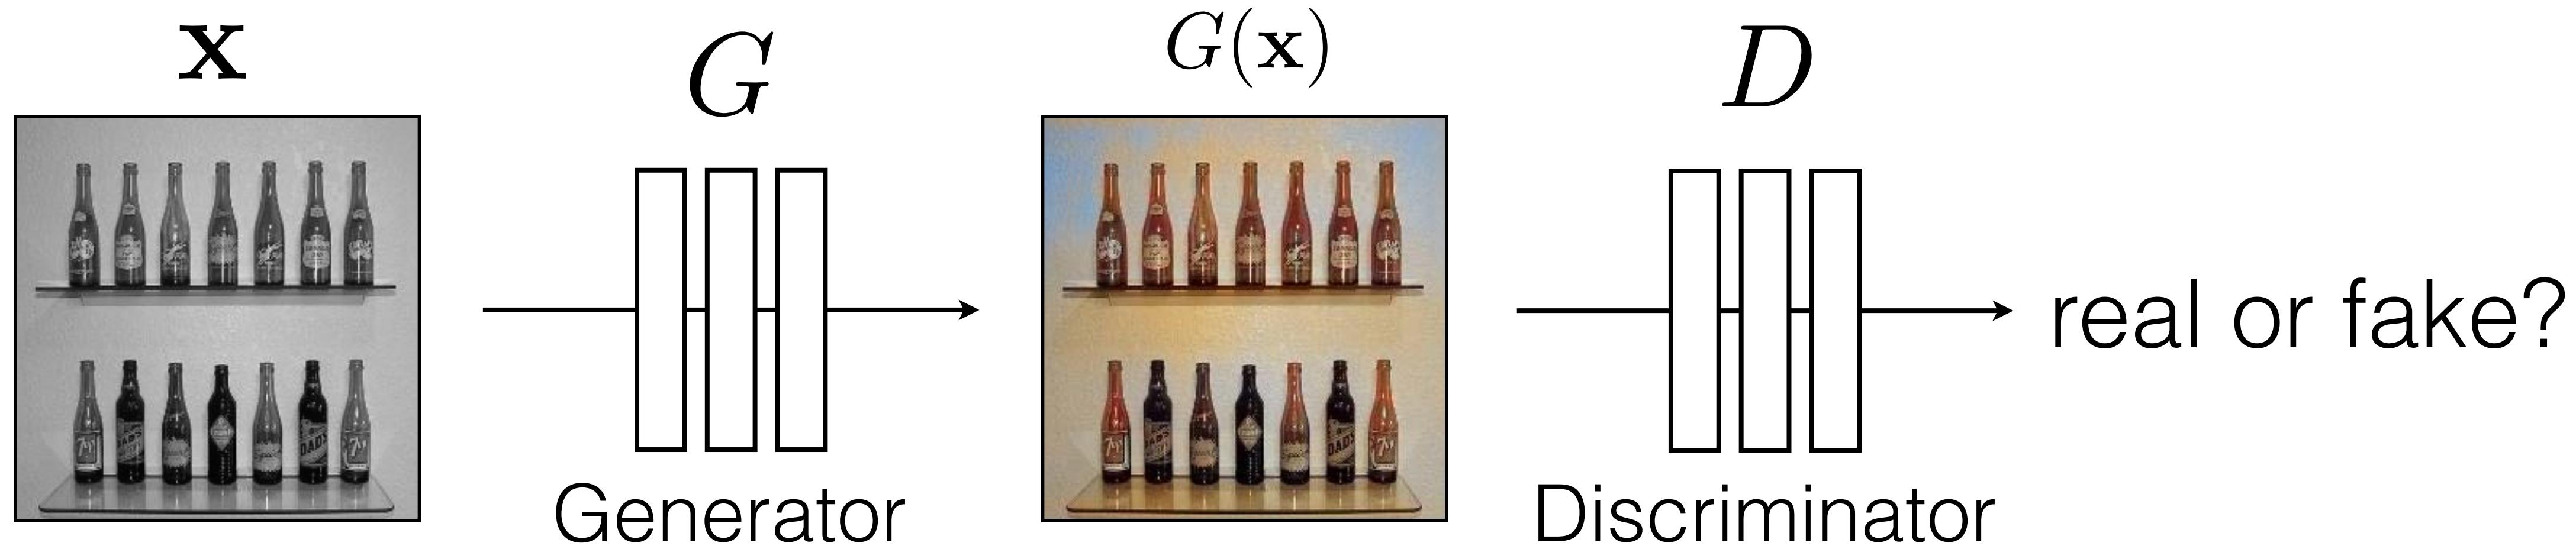
G



G(x)



Generator



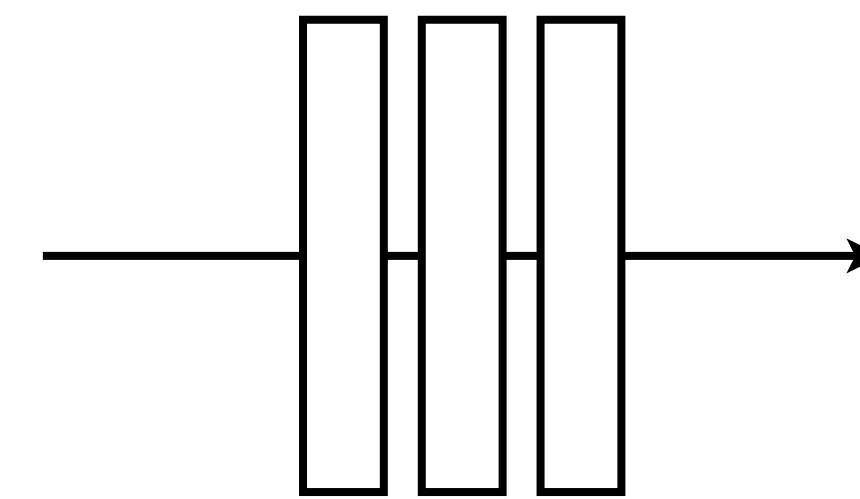
G tries to synthesize fake images that fool **D**

D tries to identify the fakes

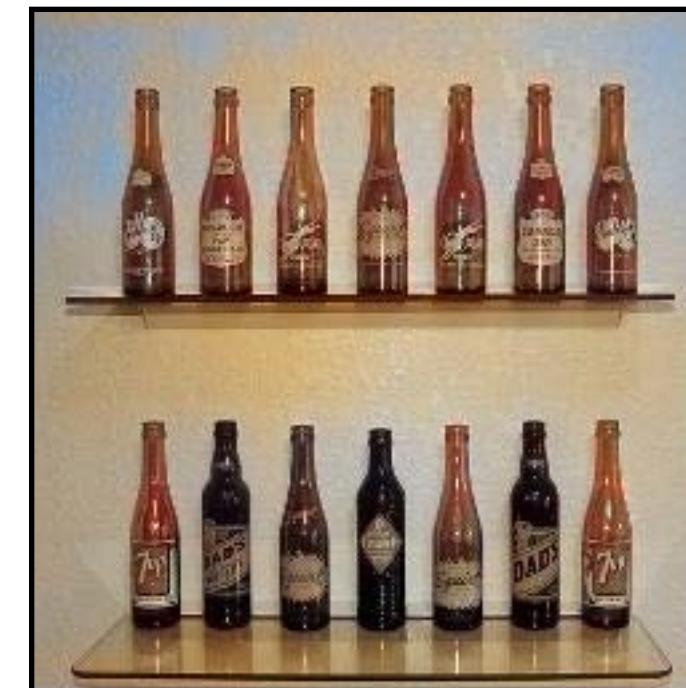
x



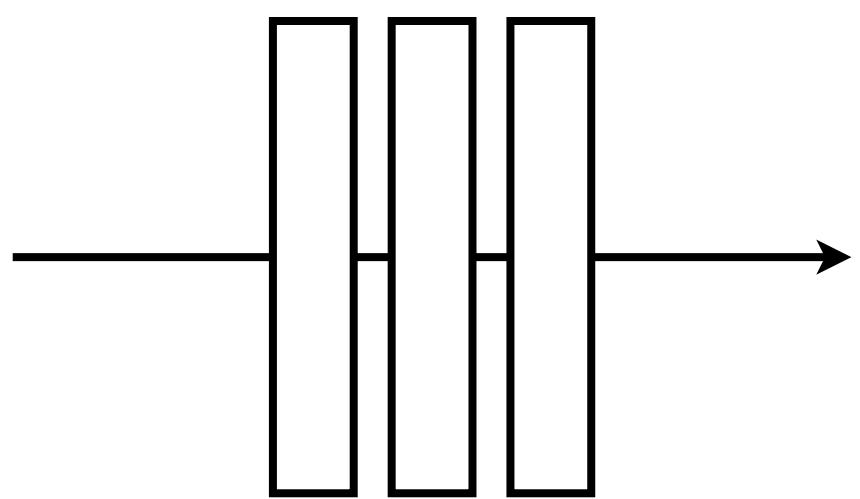
G



G(x)



D

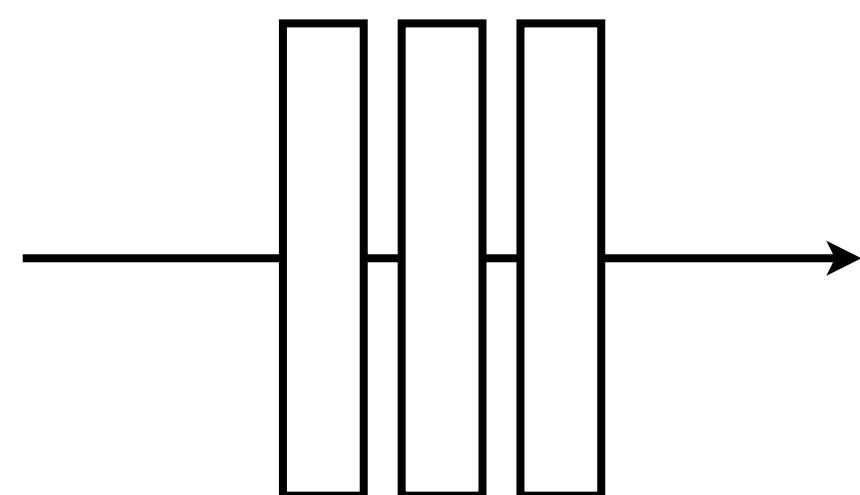


fake (0.9)

y

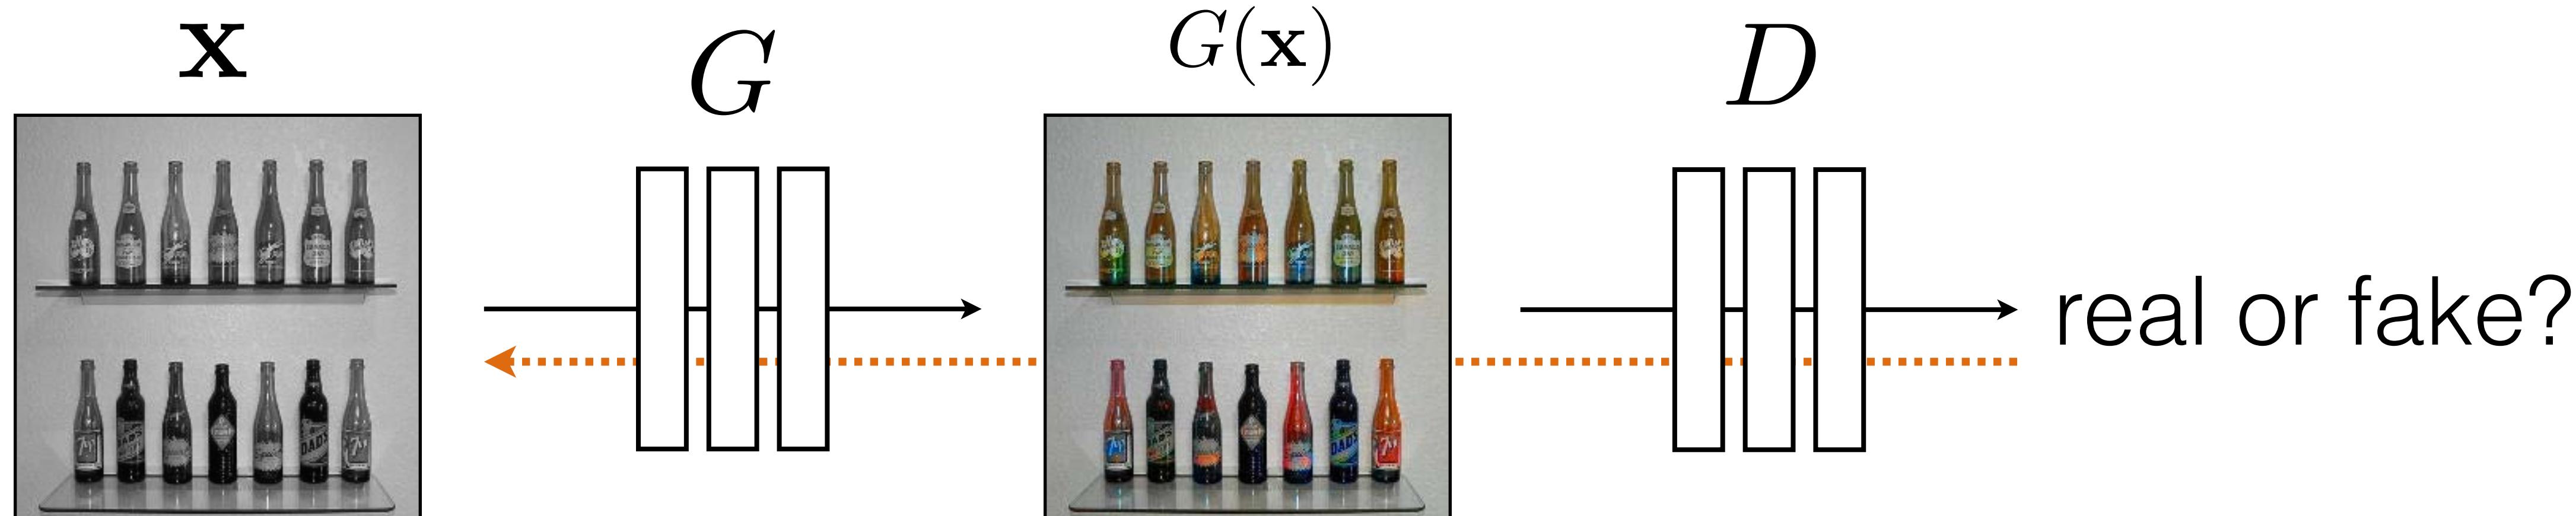


D



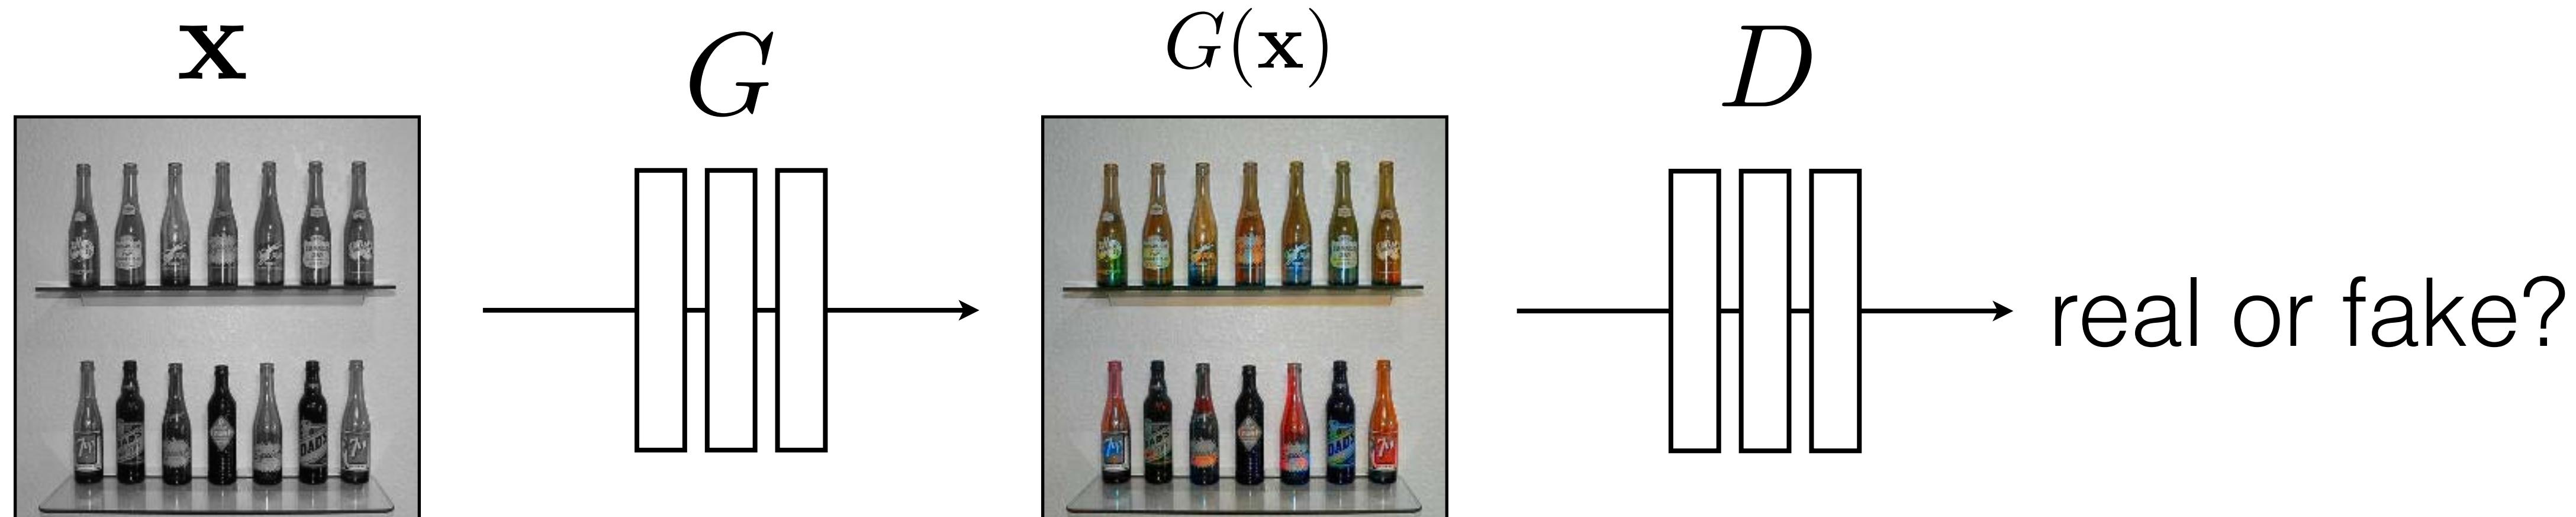
real (0.1)

$$\arg \max_D \mathbb{E}_{\mathbf{x}, \mathbf{y}} [\boxed{\log D(G(\mathbf{x}))} + \boxed{\log(1 - D(\mathbf{y}))}]$$



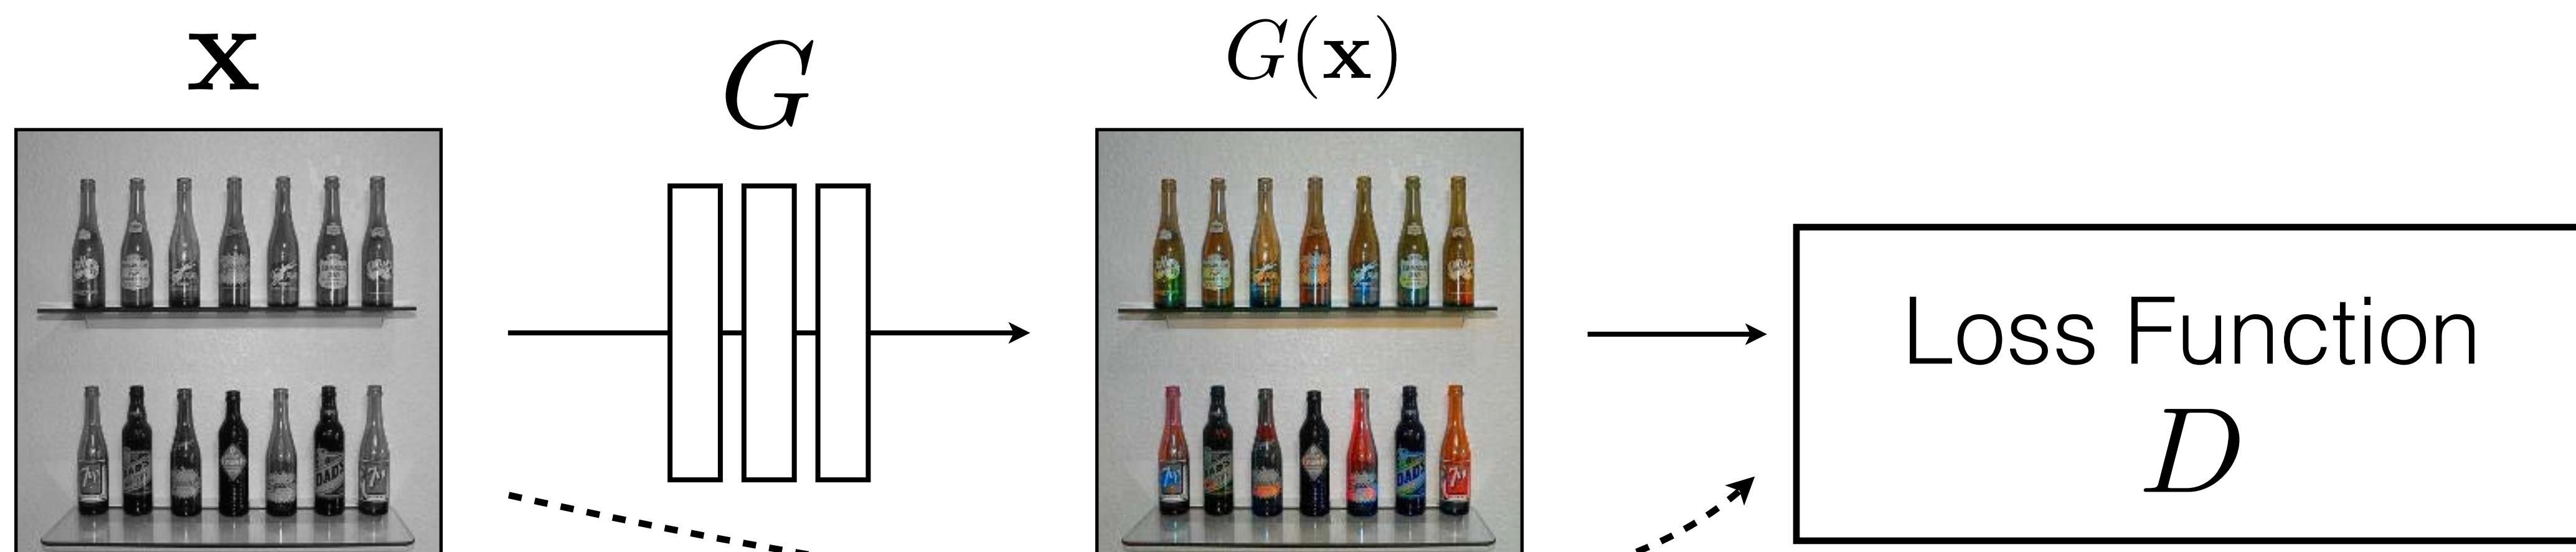
G tries to synthesize fake images that **fool** **D**:

$$\arg \min_G \mathbb{E}_{\mathbf{x}, \mathbf{y}} [\log D(G(\mathbf{x})) + \log(1 - D(\mathbf{y}))]$$



G tries to synthesize fake images that **fool** the **best** **D**:

$$\arg \min_G \max_D \mathbb{E}_{\mathbf{x}, \mathbf{y}} [\log D(G(\mathbf{x})) + \log(1 - D(\mathbf{y}))]$$



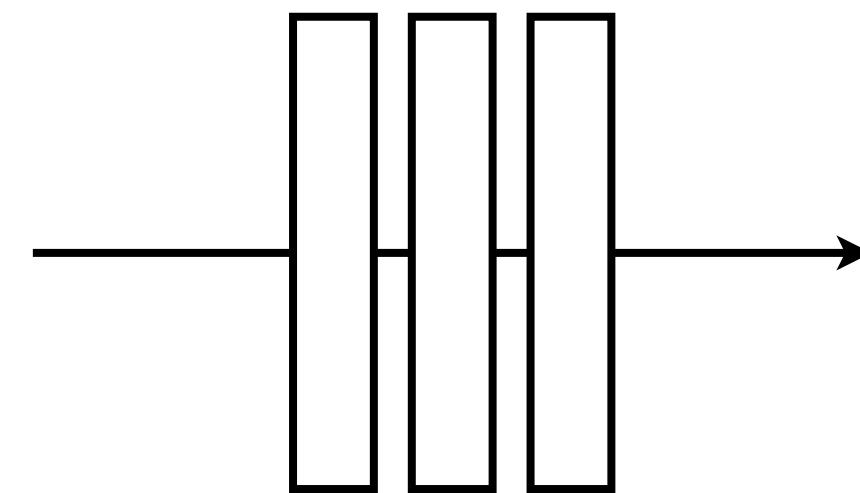
G's perspective: **D** is a loss function.

Rather than being hand-designed, it is *learned*.

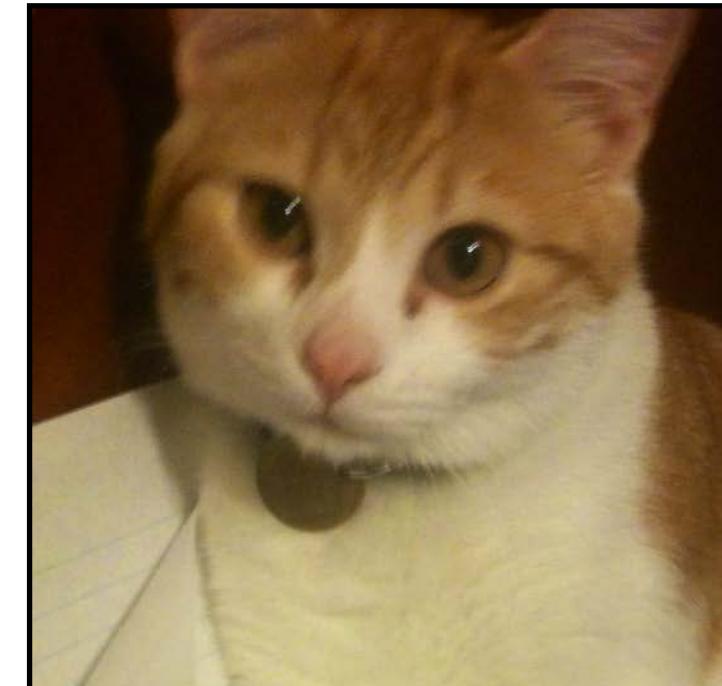
x



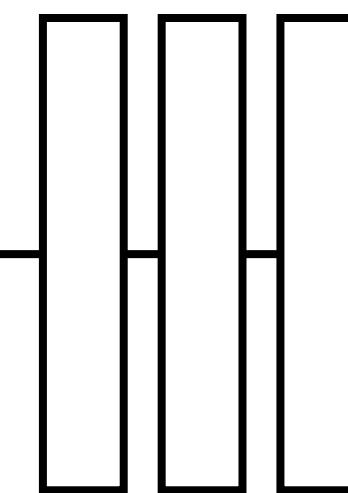
G



$G(\mathbf{x})$



D



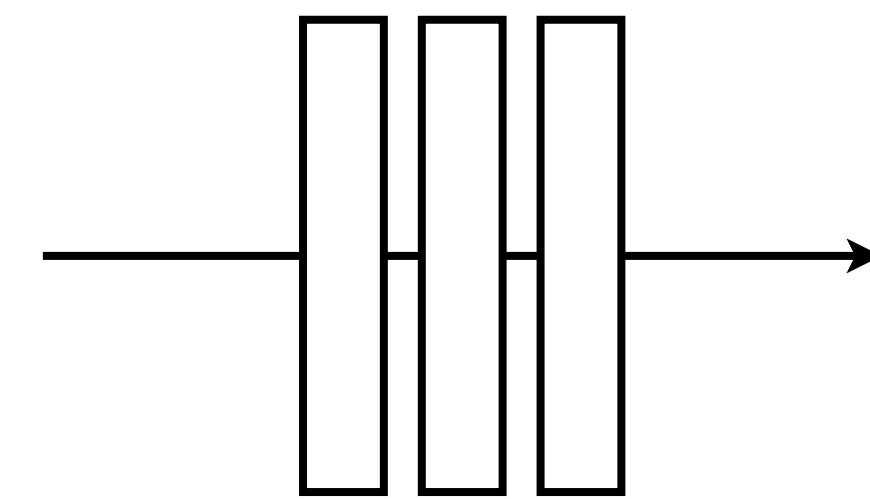
real or fake?

$$\arg \min_G \max_D \mathbb{E}_{\mathbf{x}, \mathbf{y}} [\log D(G(\mathbf{x})) + \log(1 - D(\mathbf{y}))]$$

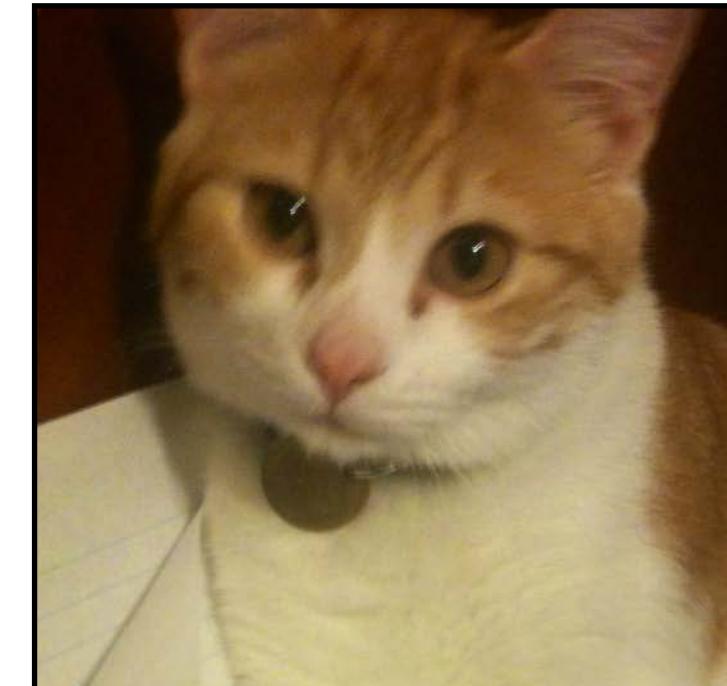
x



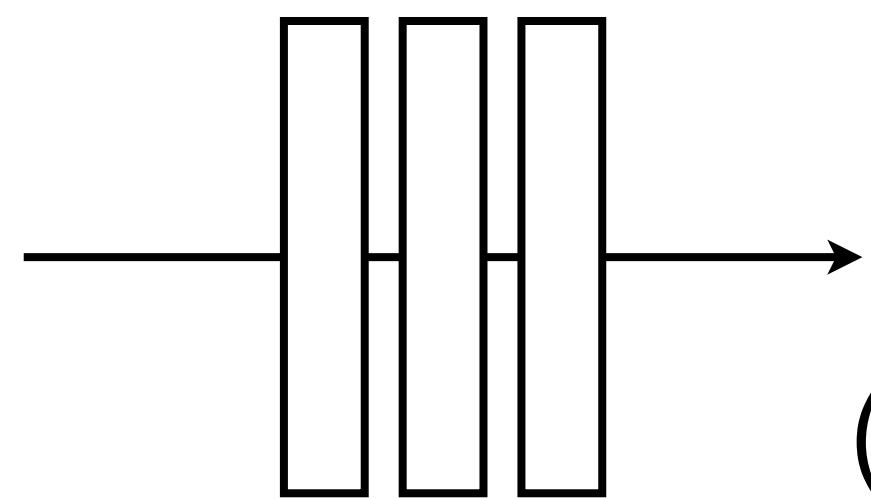
G



G(x)



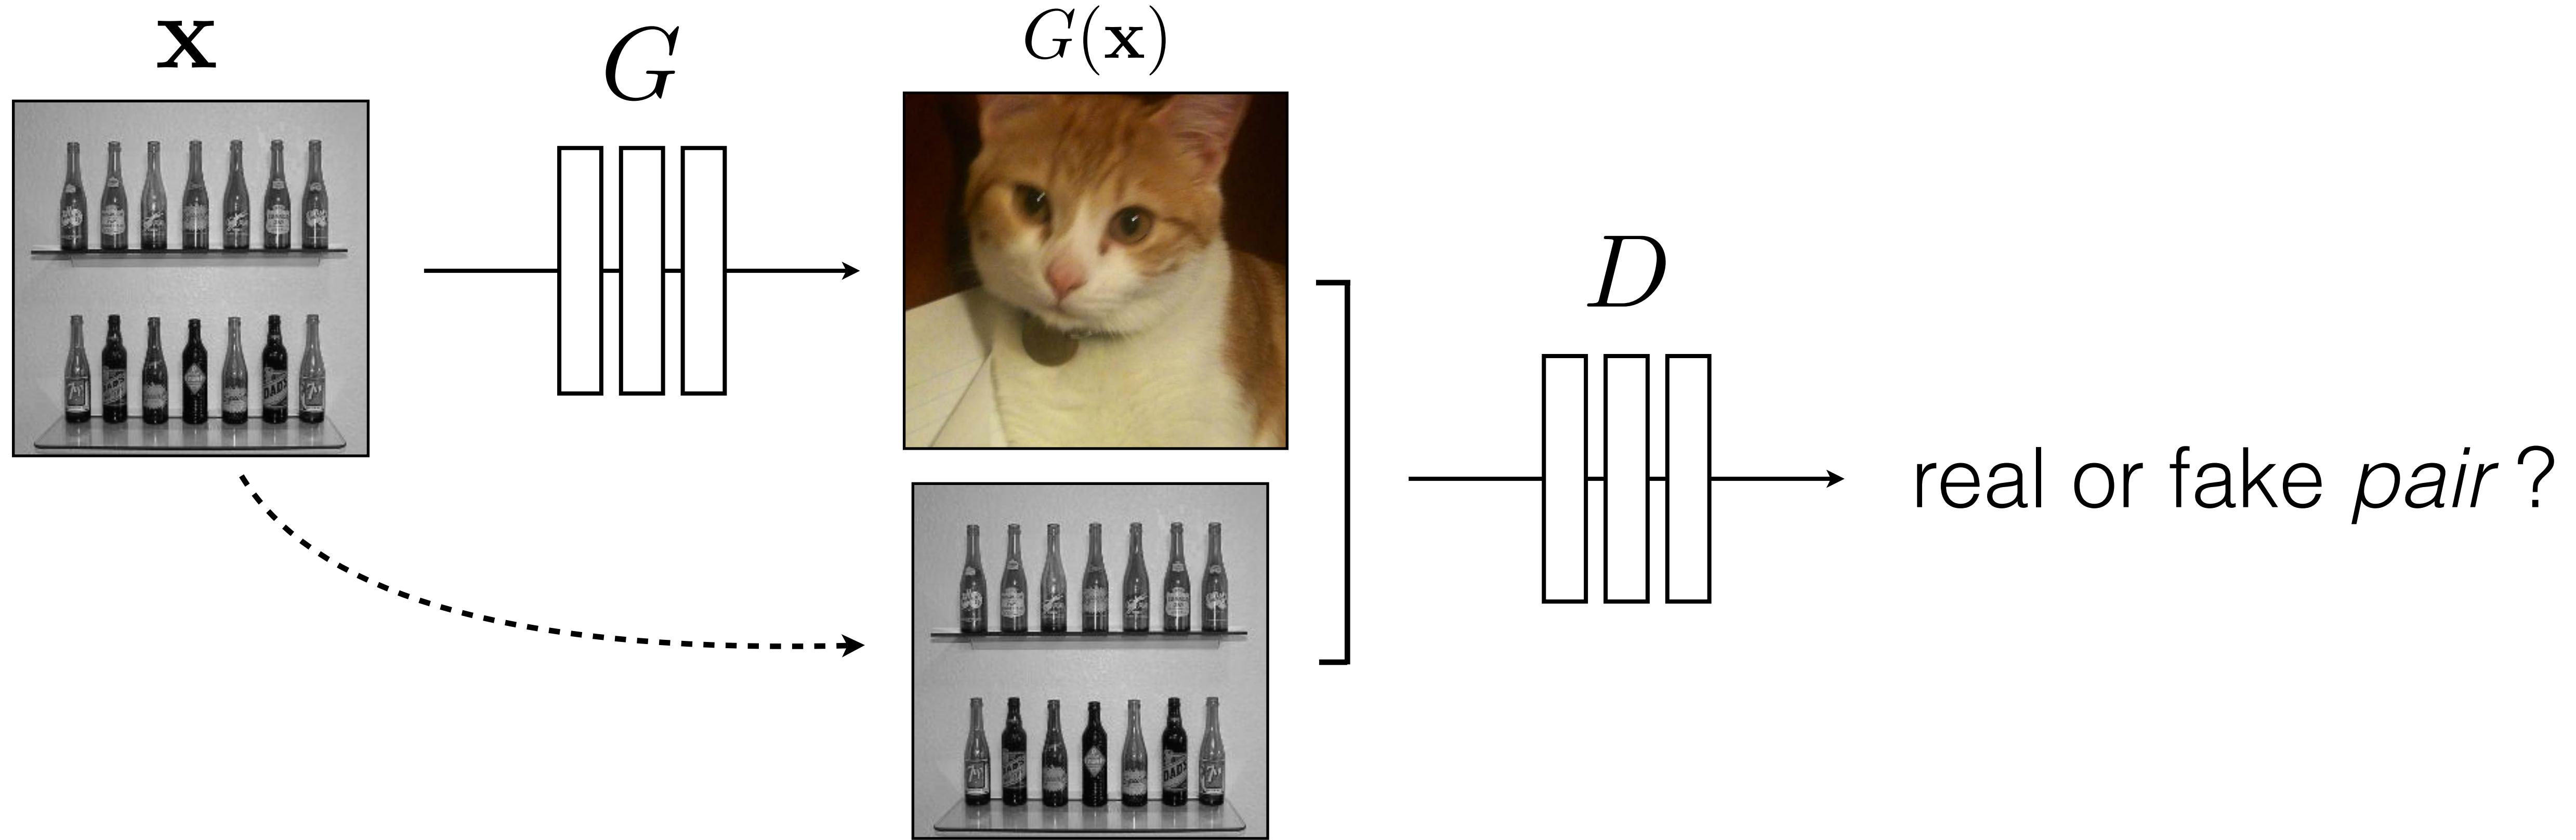
D



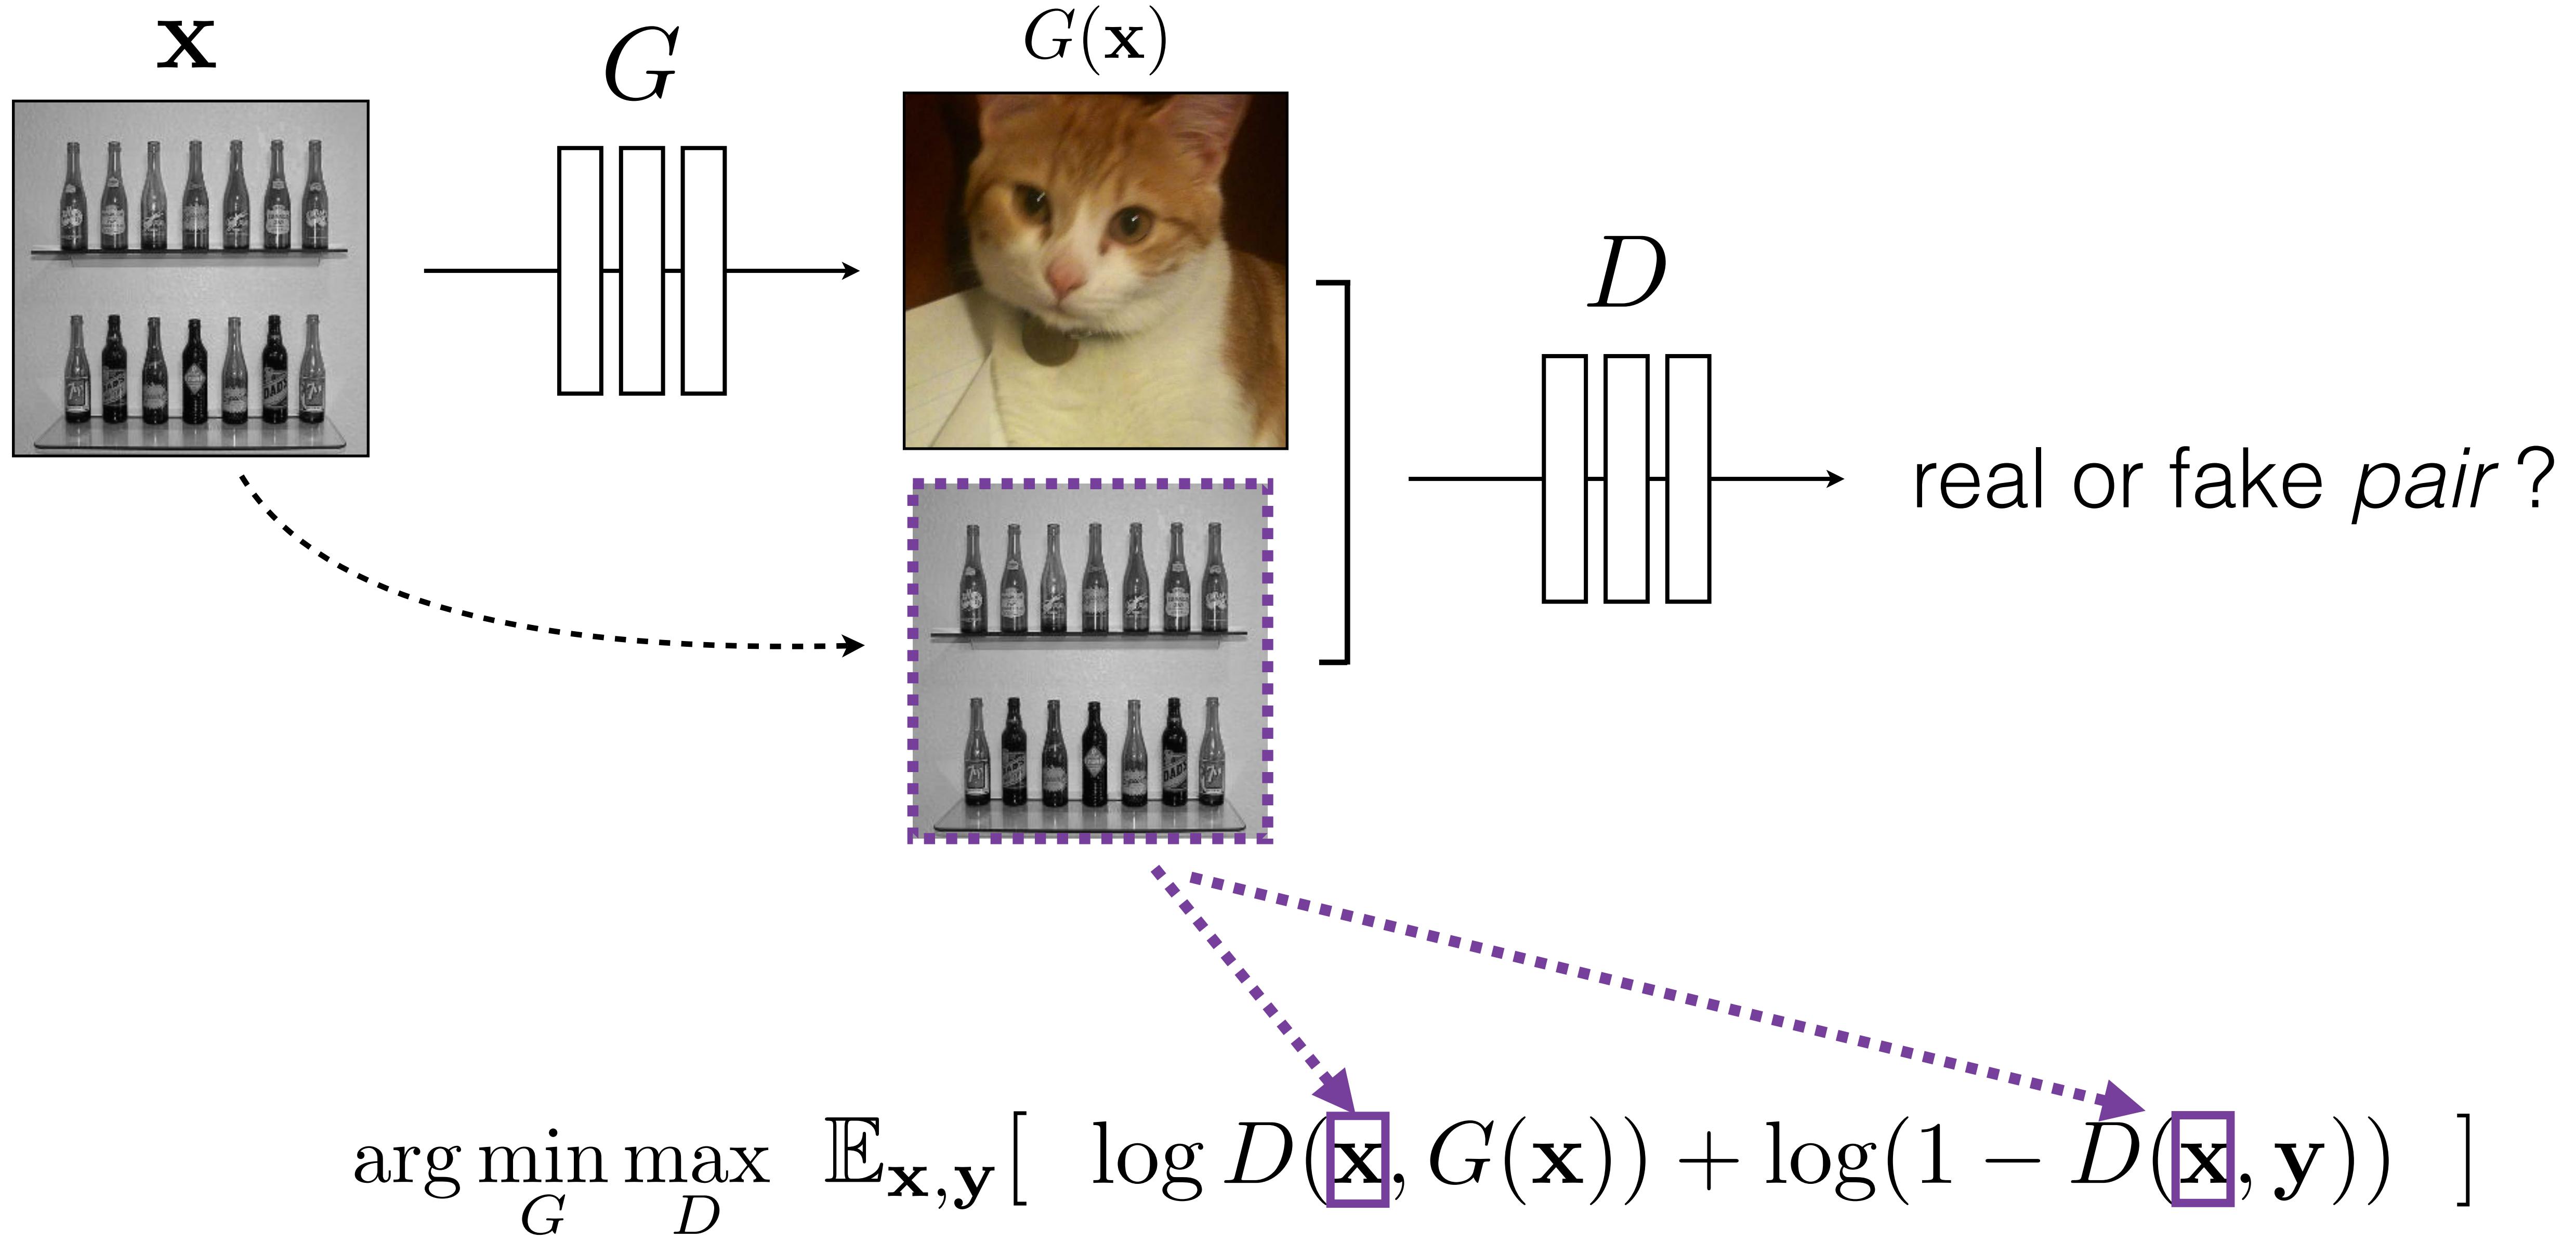
real!

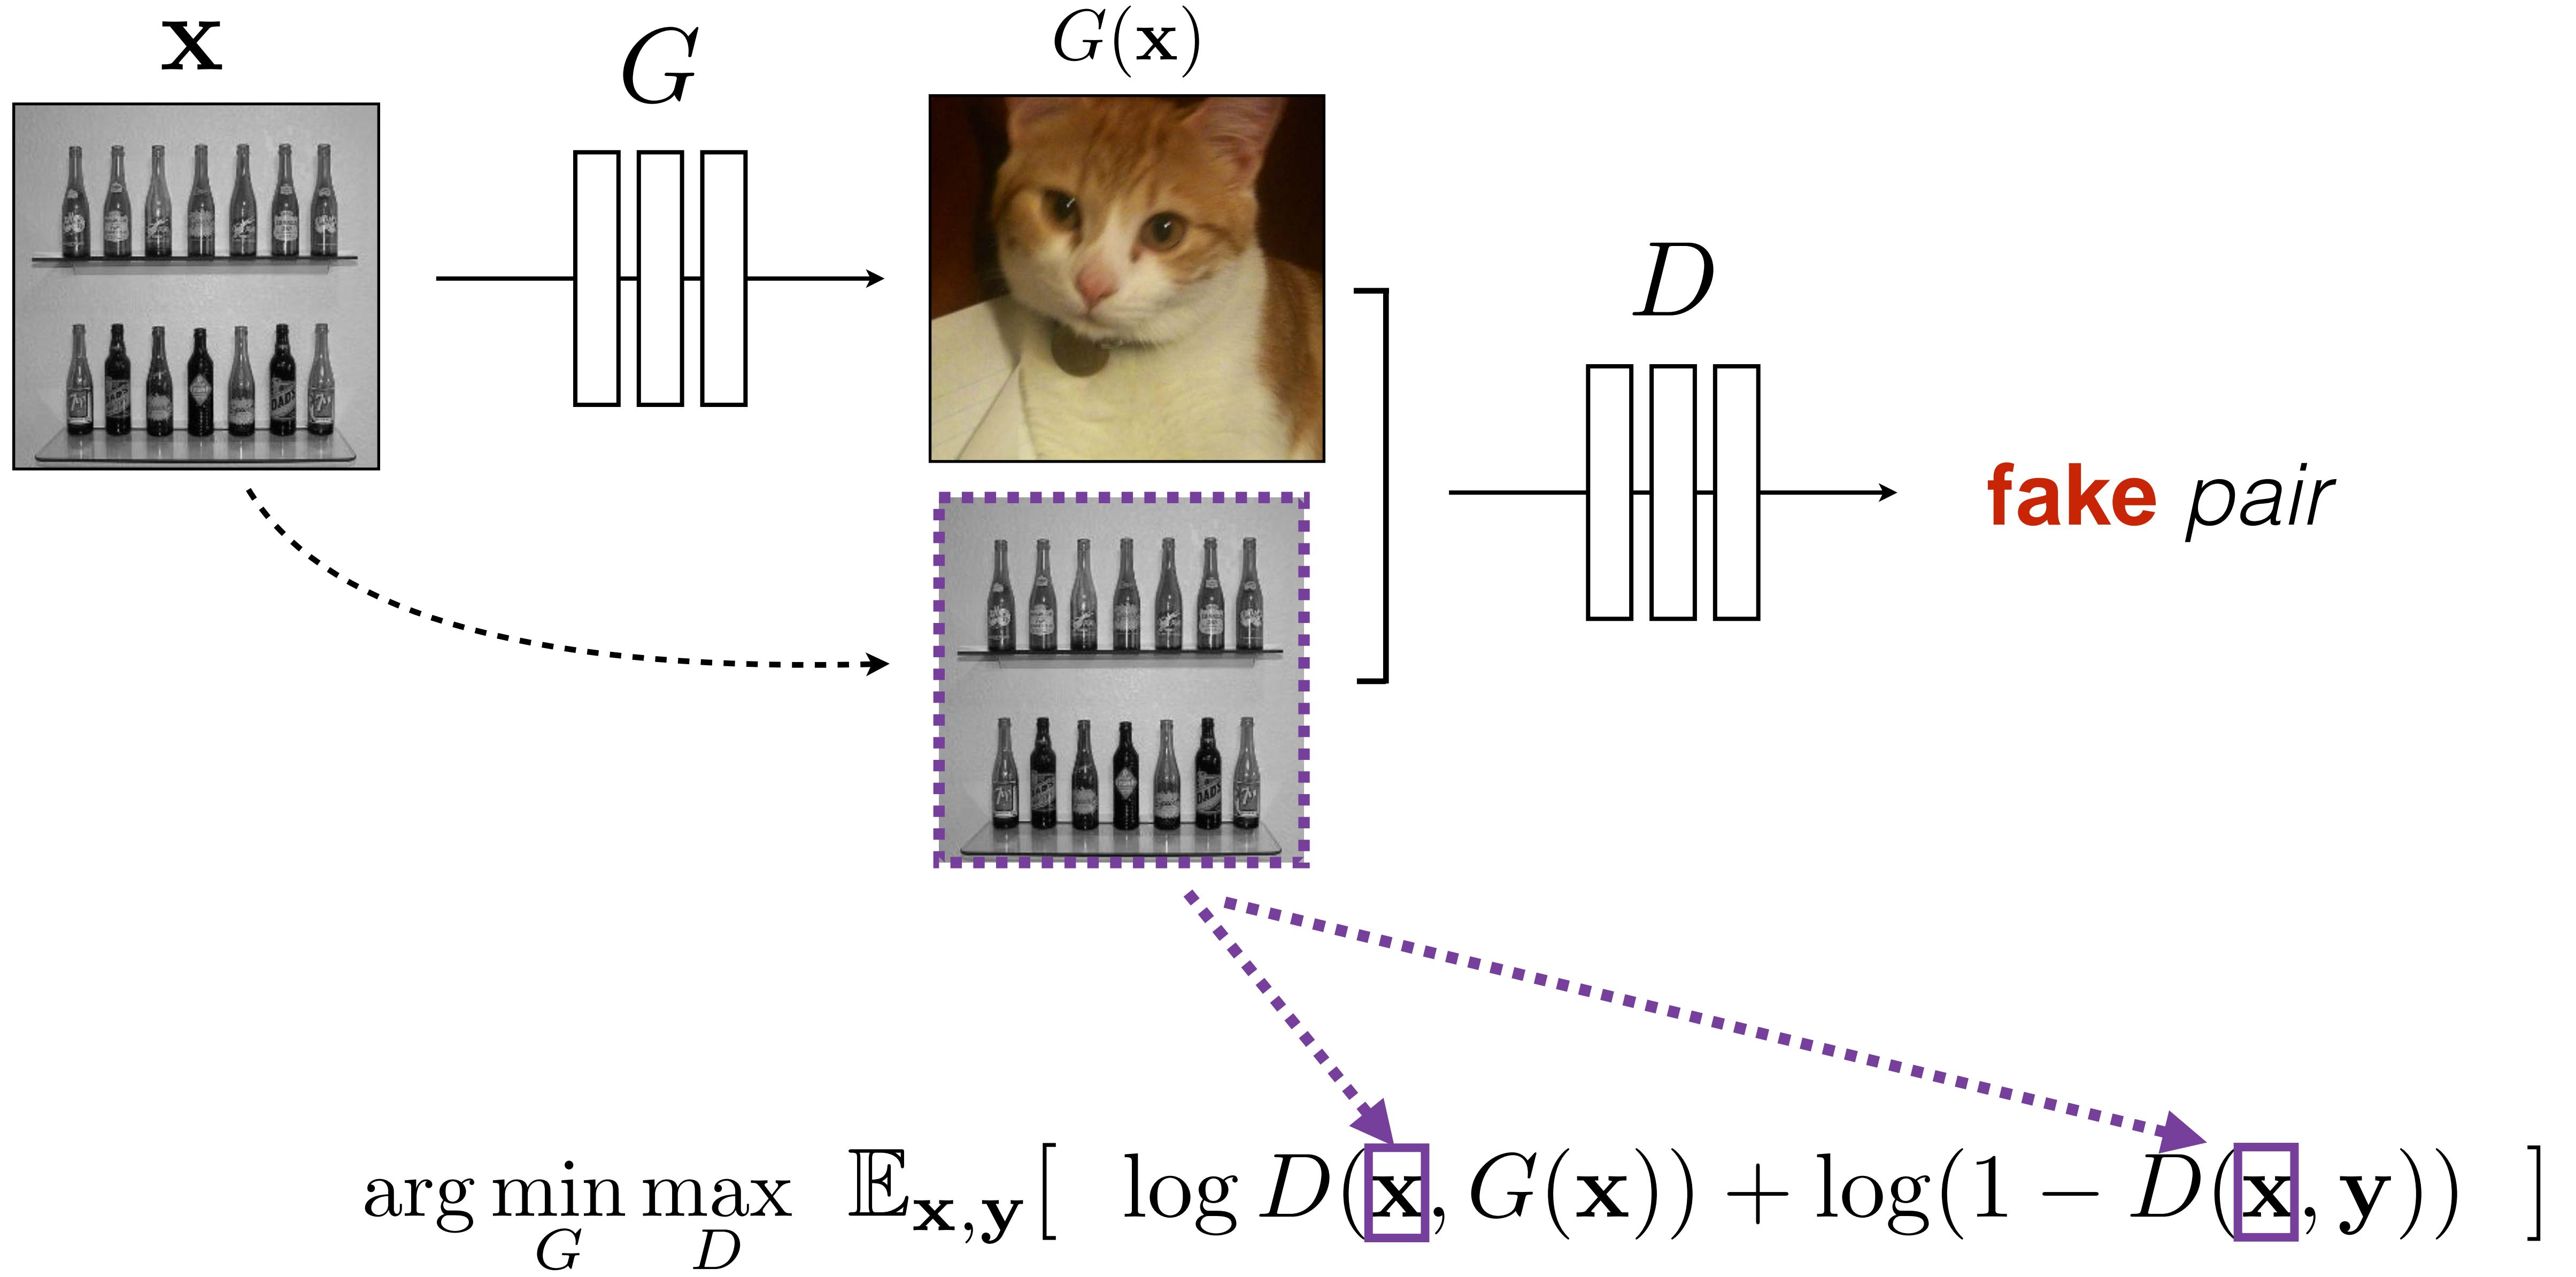
("Aquarius")

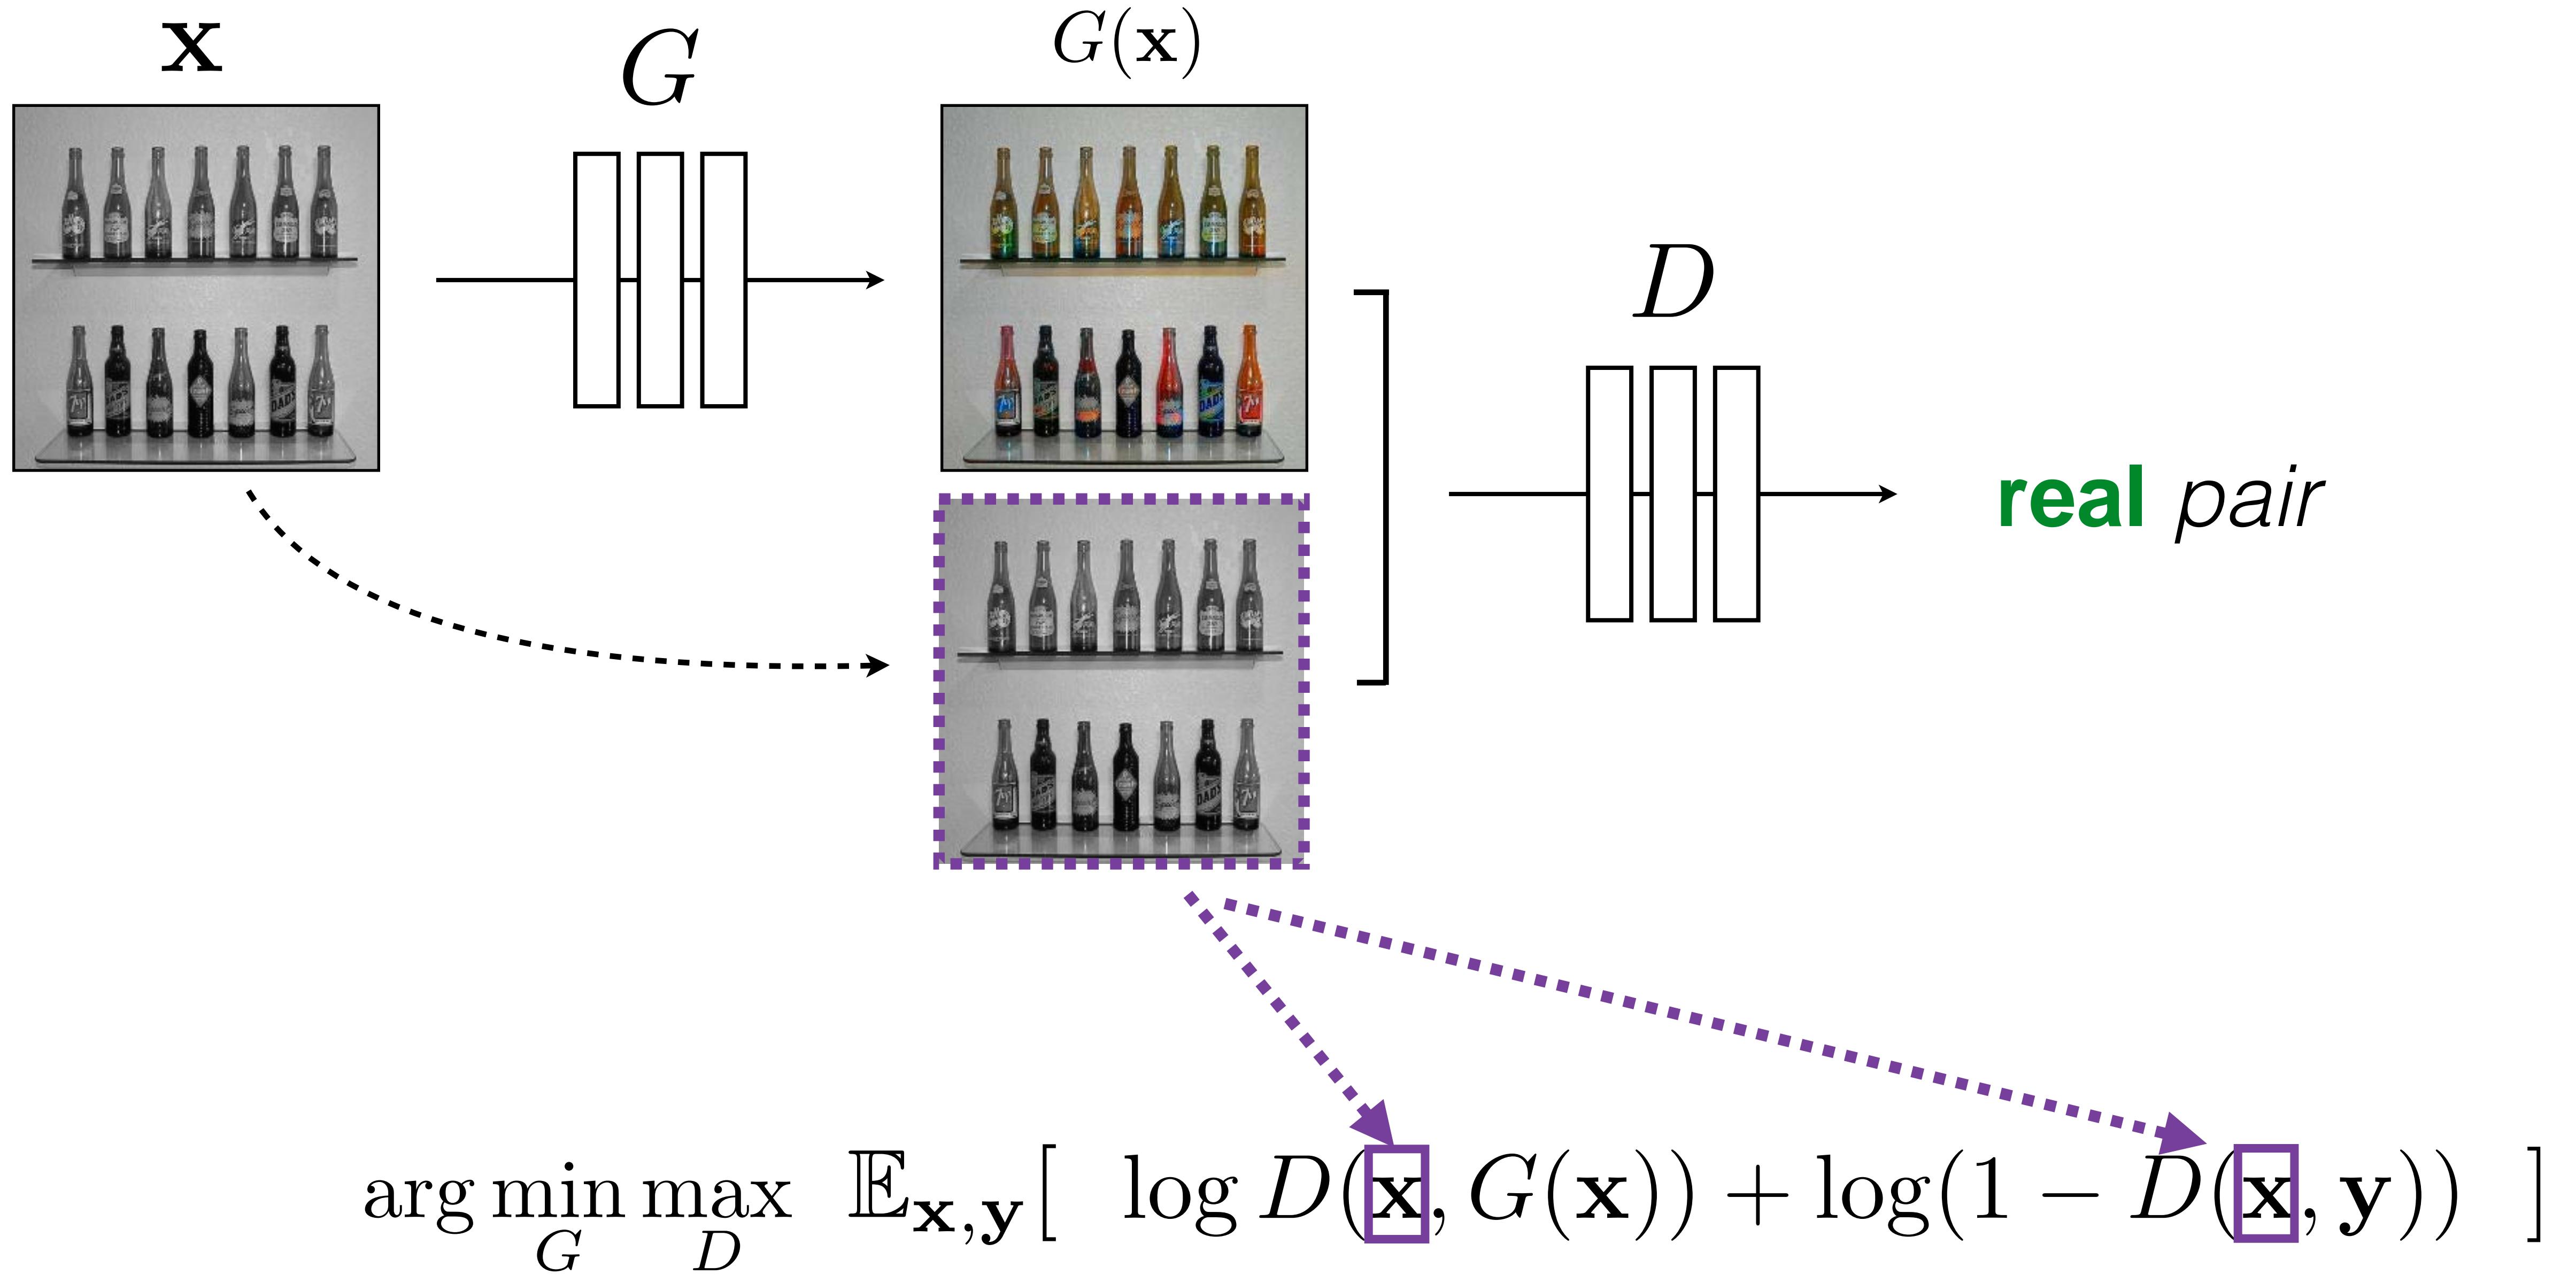
$$\arg \min_G \max_D \mathbb{E}_{\mathbf{x}, \mathbf{y}} [\log D(G(\mathbf{x})) + \log(1 - D(\mathbf{y}))]$$

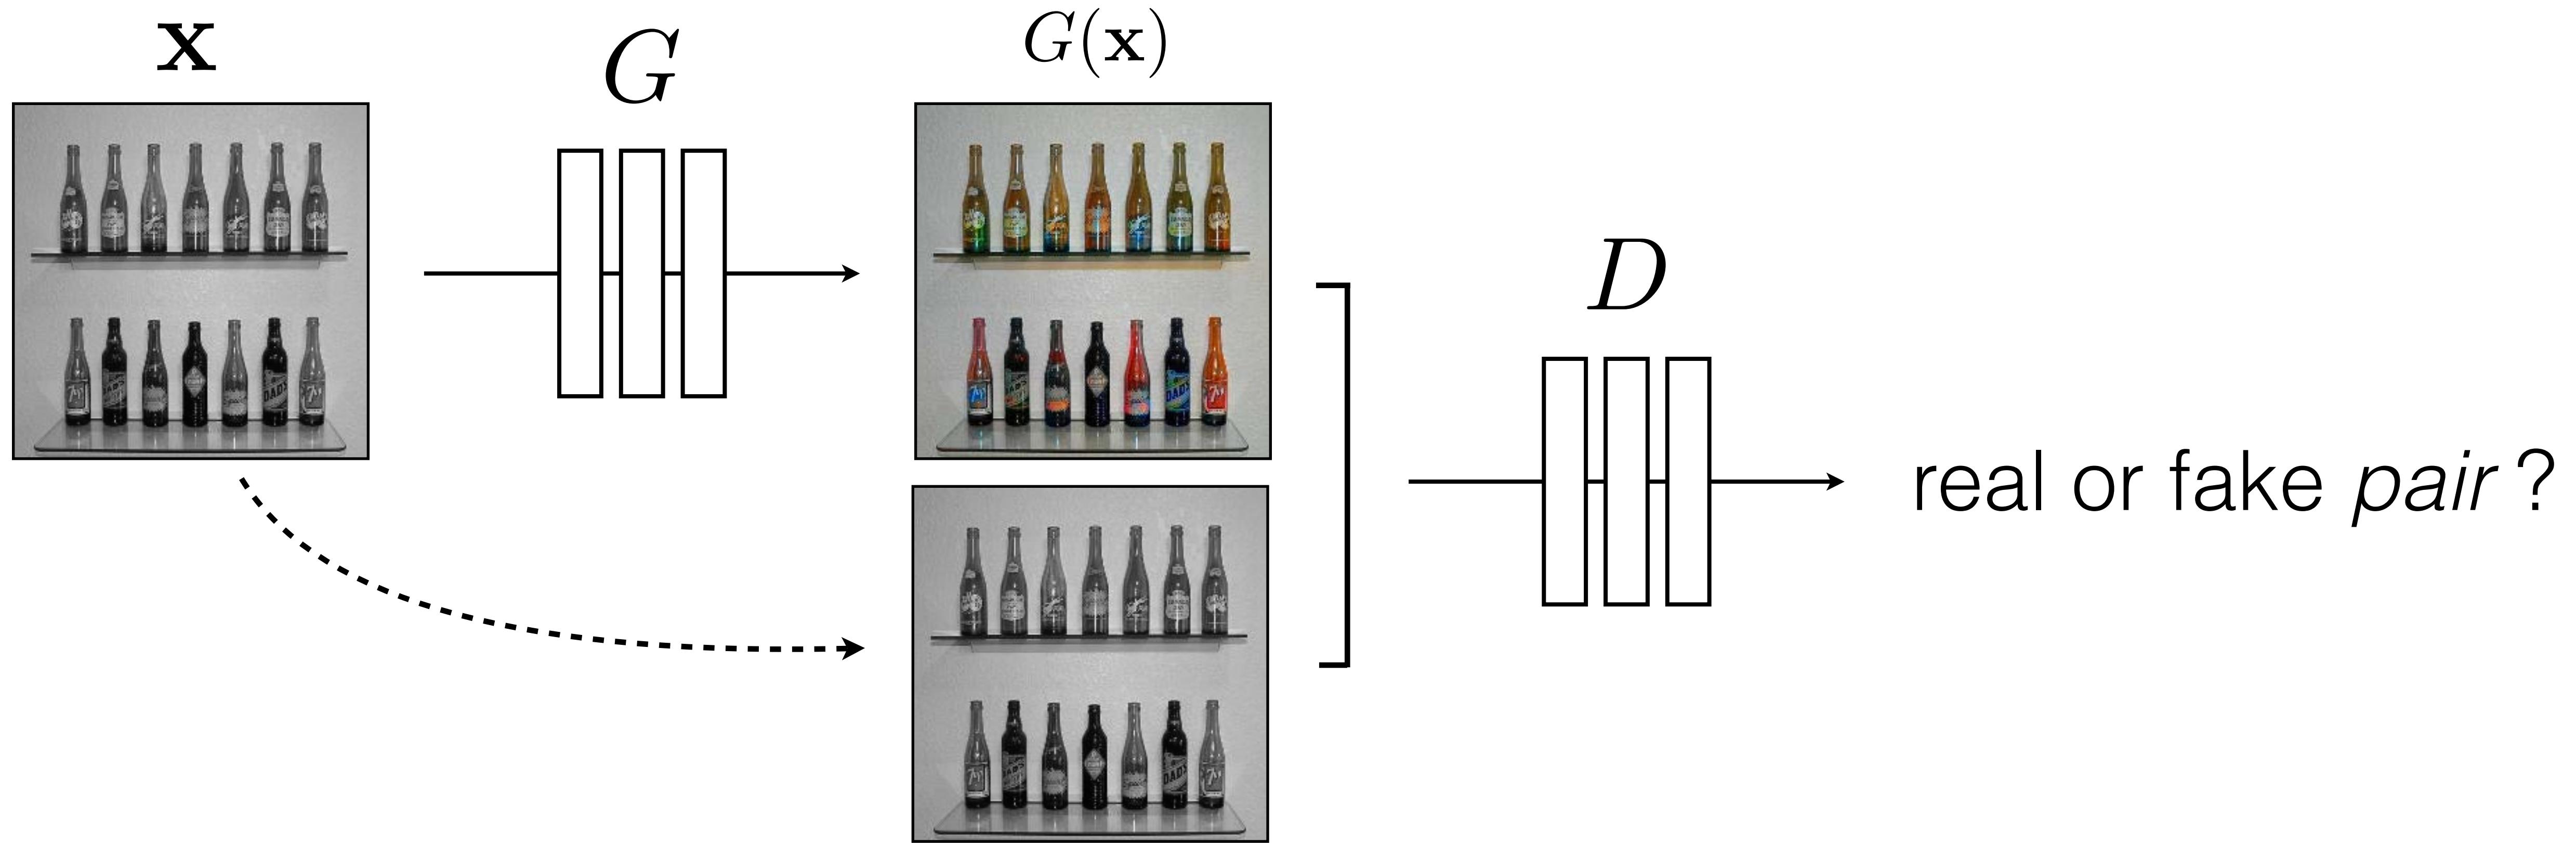


$$\arg \min_G \max_D \mathbb{E}_{\mathbf{x}, \mathbf{y}} [\log D(G(\mathbf{x})) + \log(1 - D(\mathbf{y}))]$$









$$\arg \min_G \max_D \mathbb{E}_{\mathbf{x}, \mathbf{y}} [\log D(\mathbf{x}, G(\mathbf{x})) + \log(1 - D(\mathbf{x}, \mathbf{y}))]$$

Training Details: Loss function

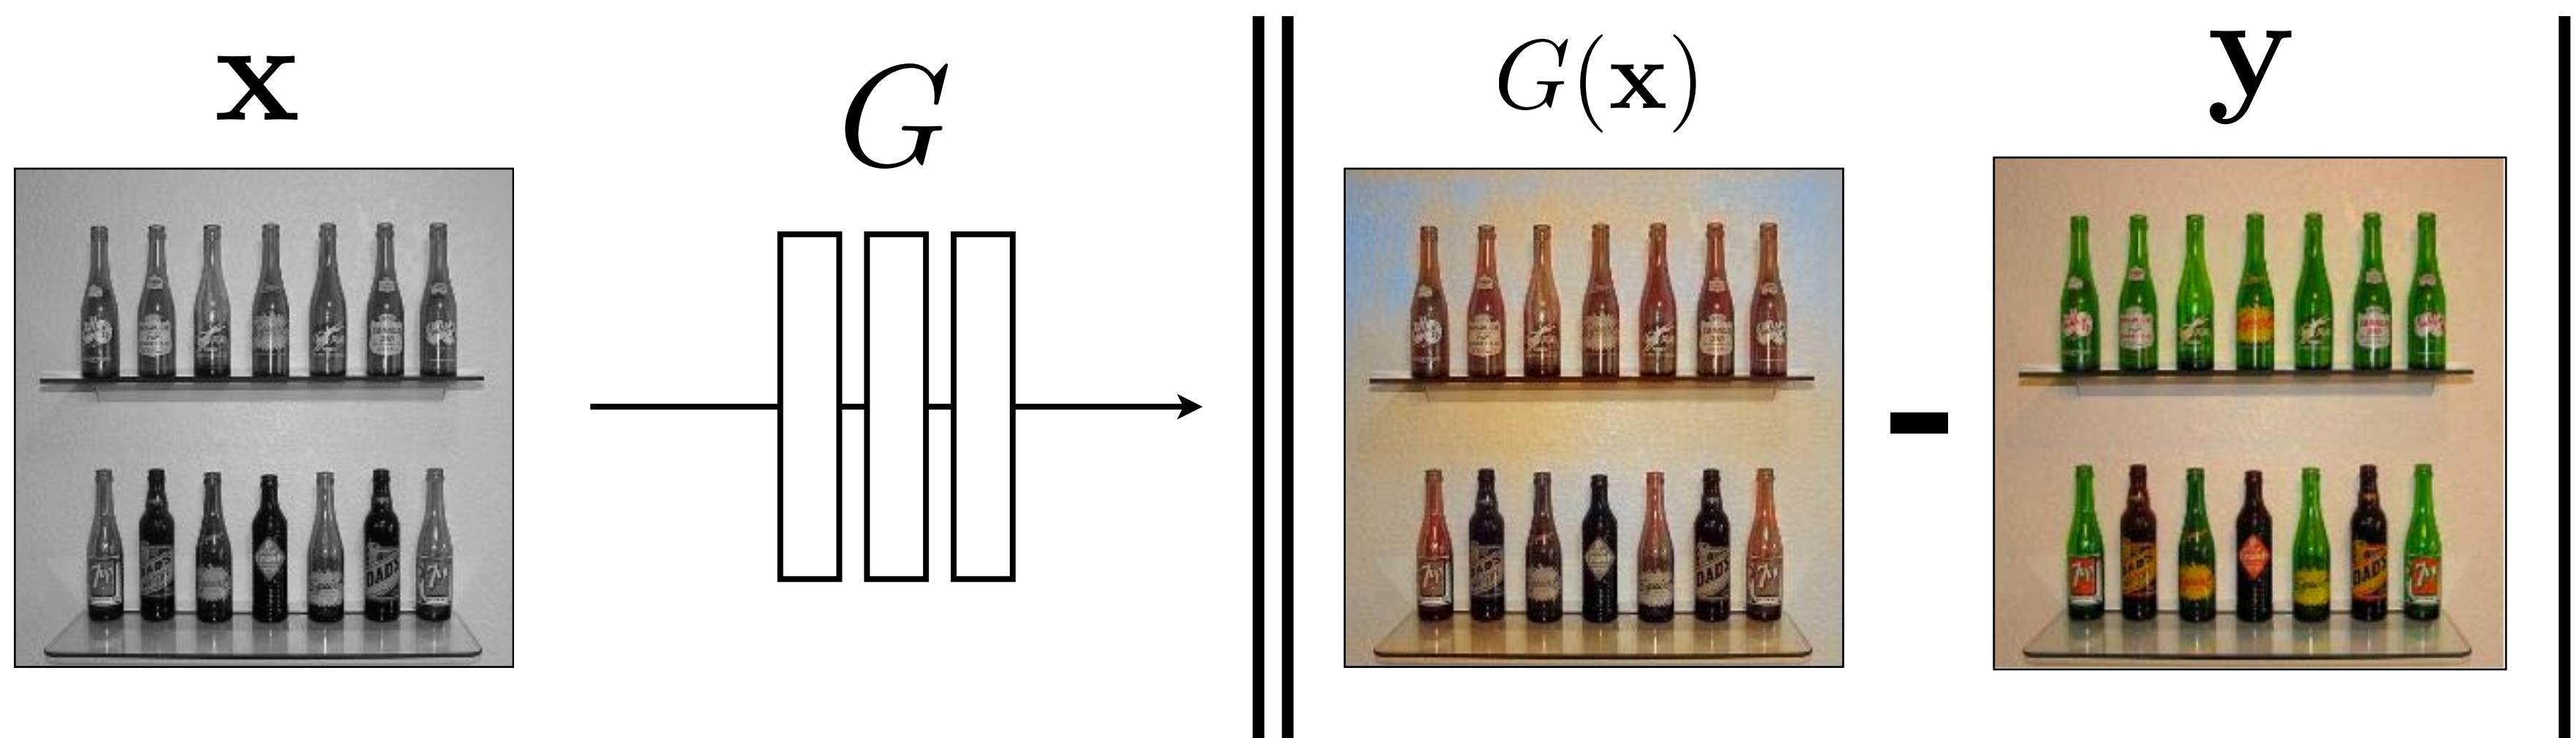
Conditional GAN

$$G^* = \arg \min_G \max_D \mathcal{L}_{cGAN}(G, D) + \lambda \mathcal{L}_{L1}(G).$$

Training Details: Loss function

Conditional GAN

$$G^* = \arg \min_G \max_D \mathcal{L}_{cGAN}(G, D) + \lambda \mathcal{L}_{L1}(G).$$



Stable training + fast convergence

[c.f. Pathak et al. CVPR 2016]

BW → Color

Input



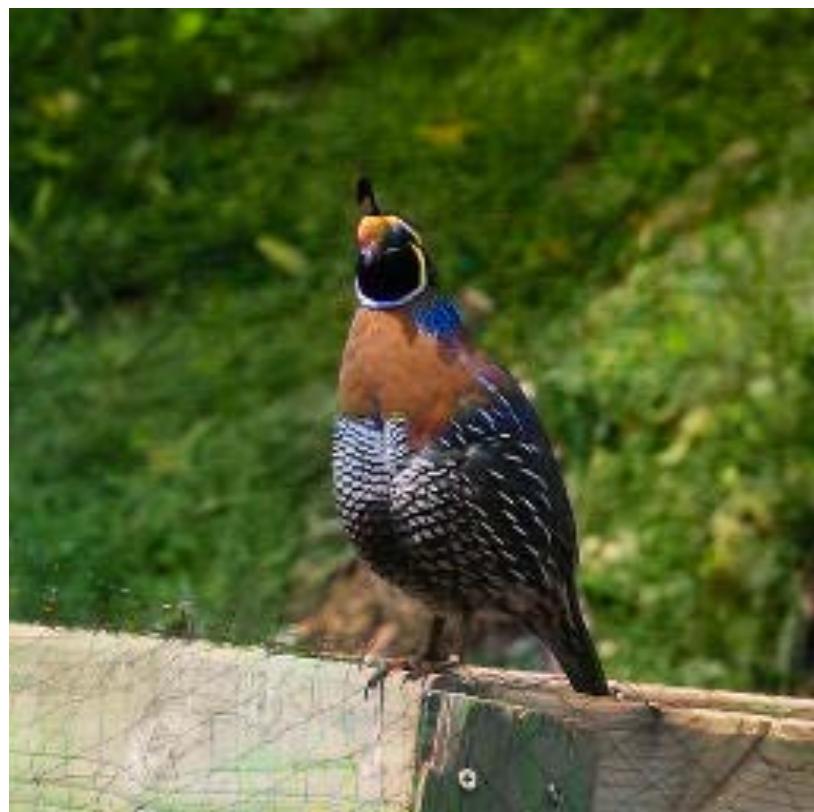
Output



Input



Output



Input

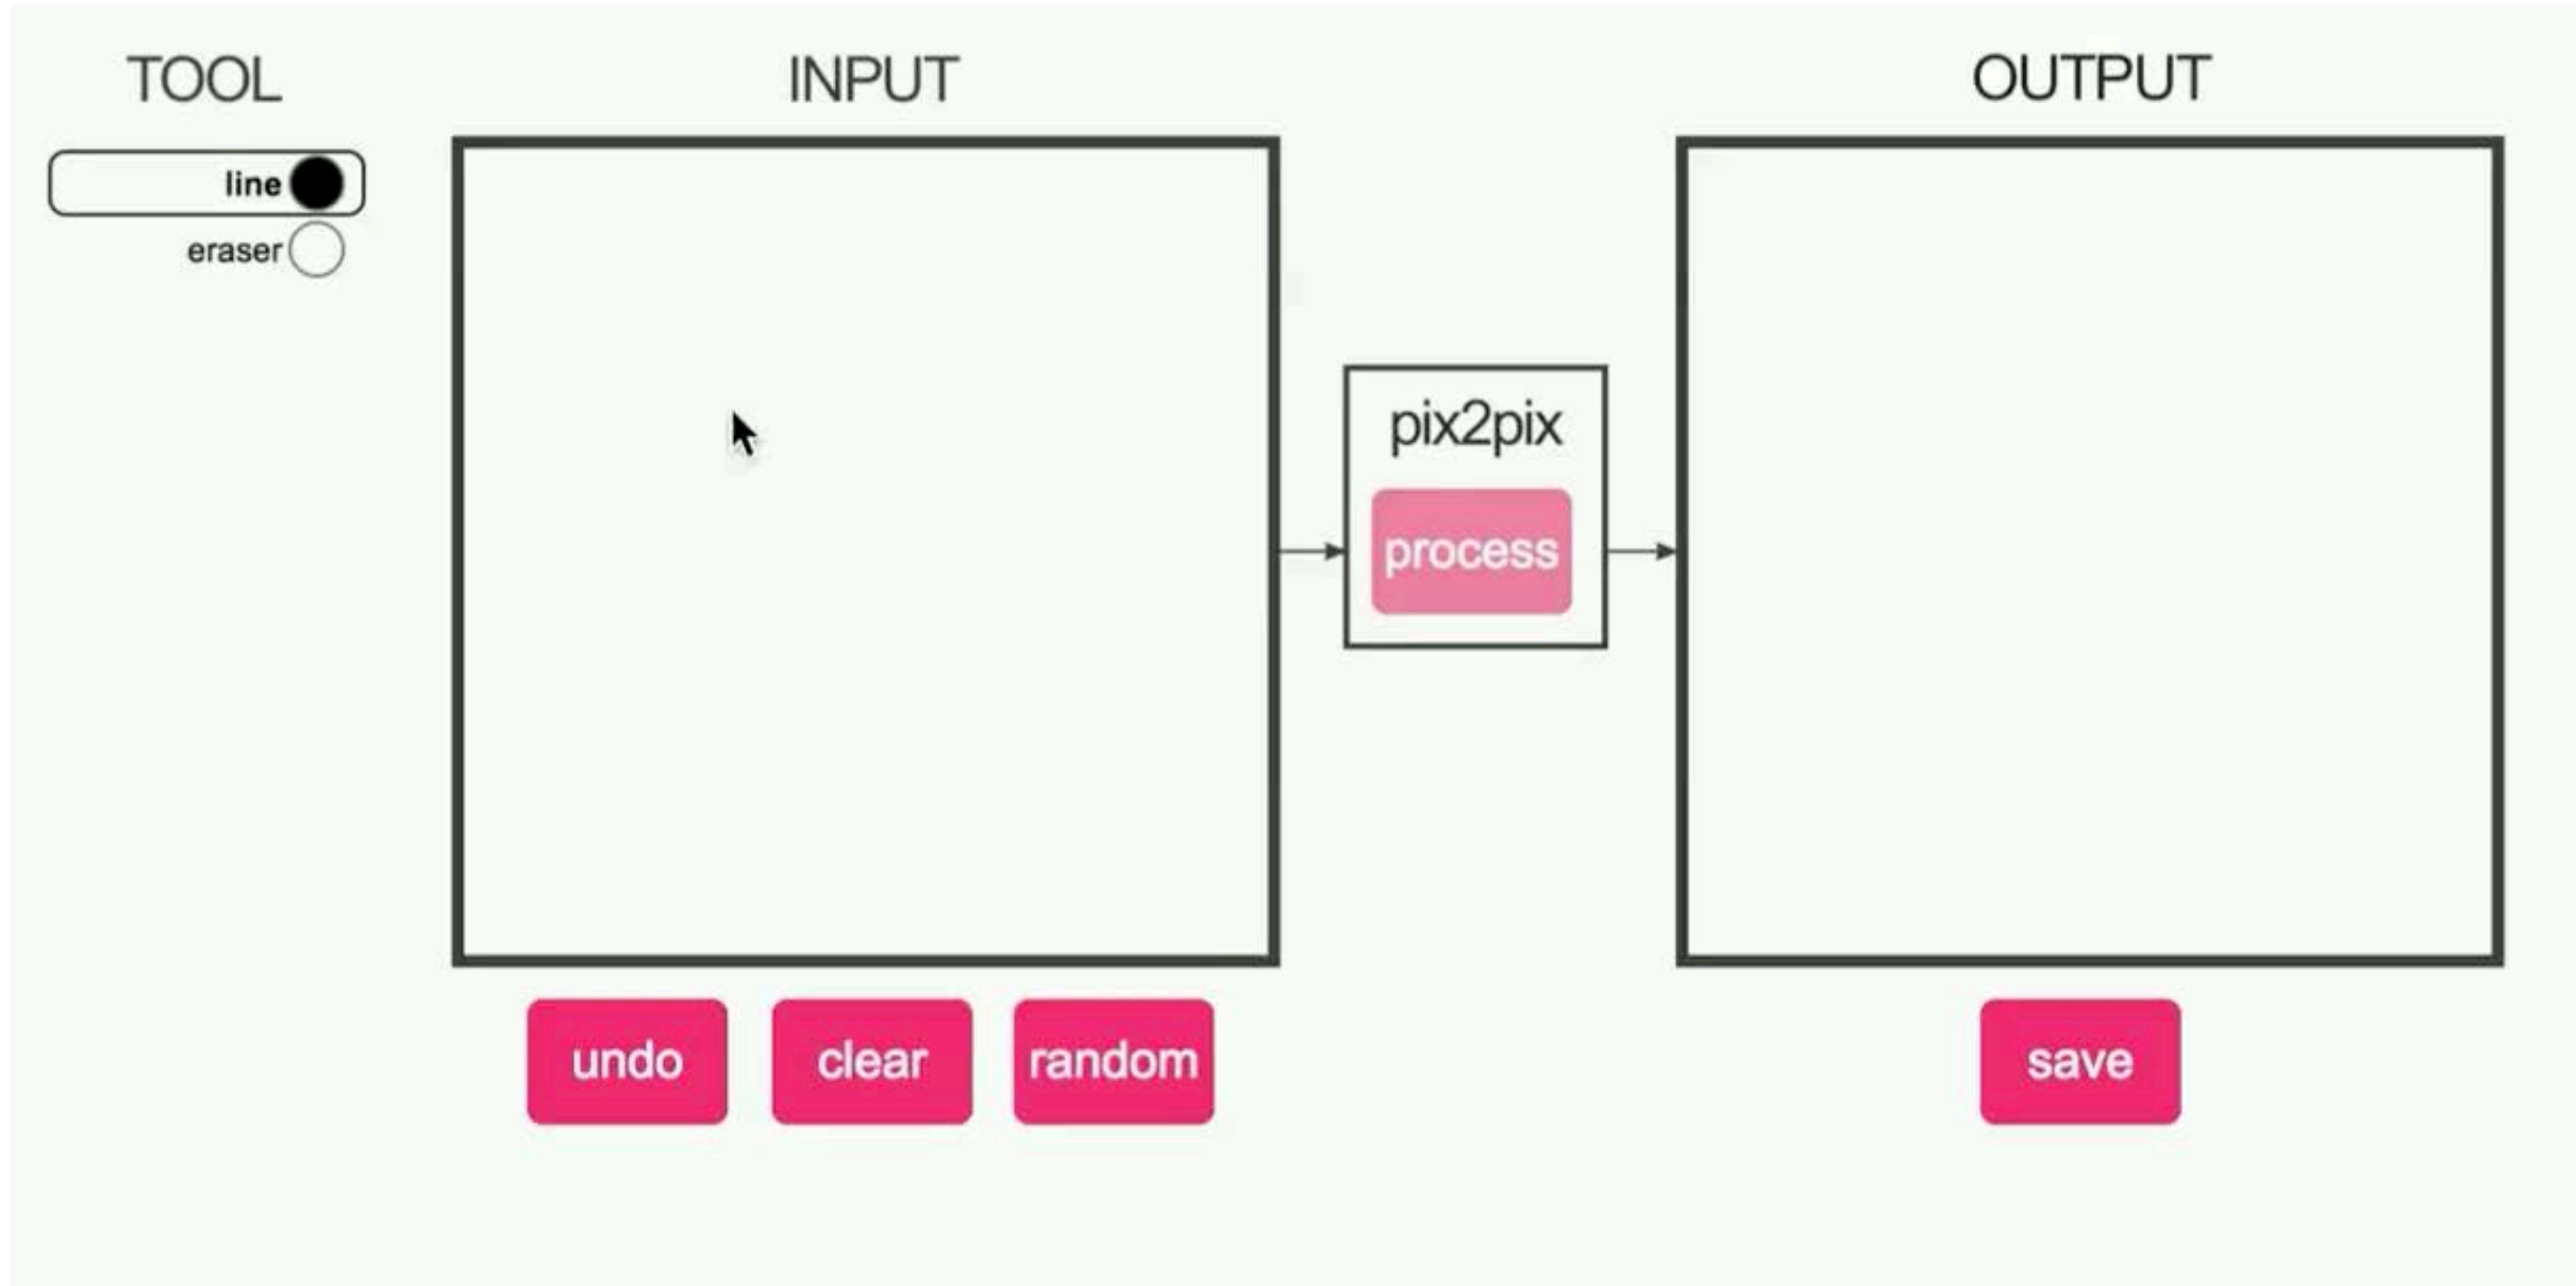


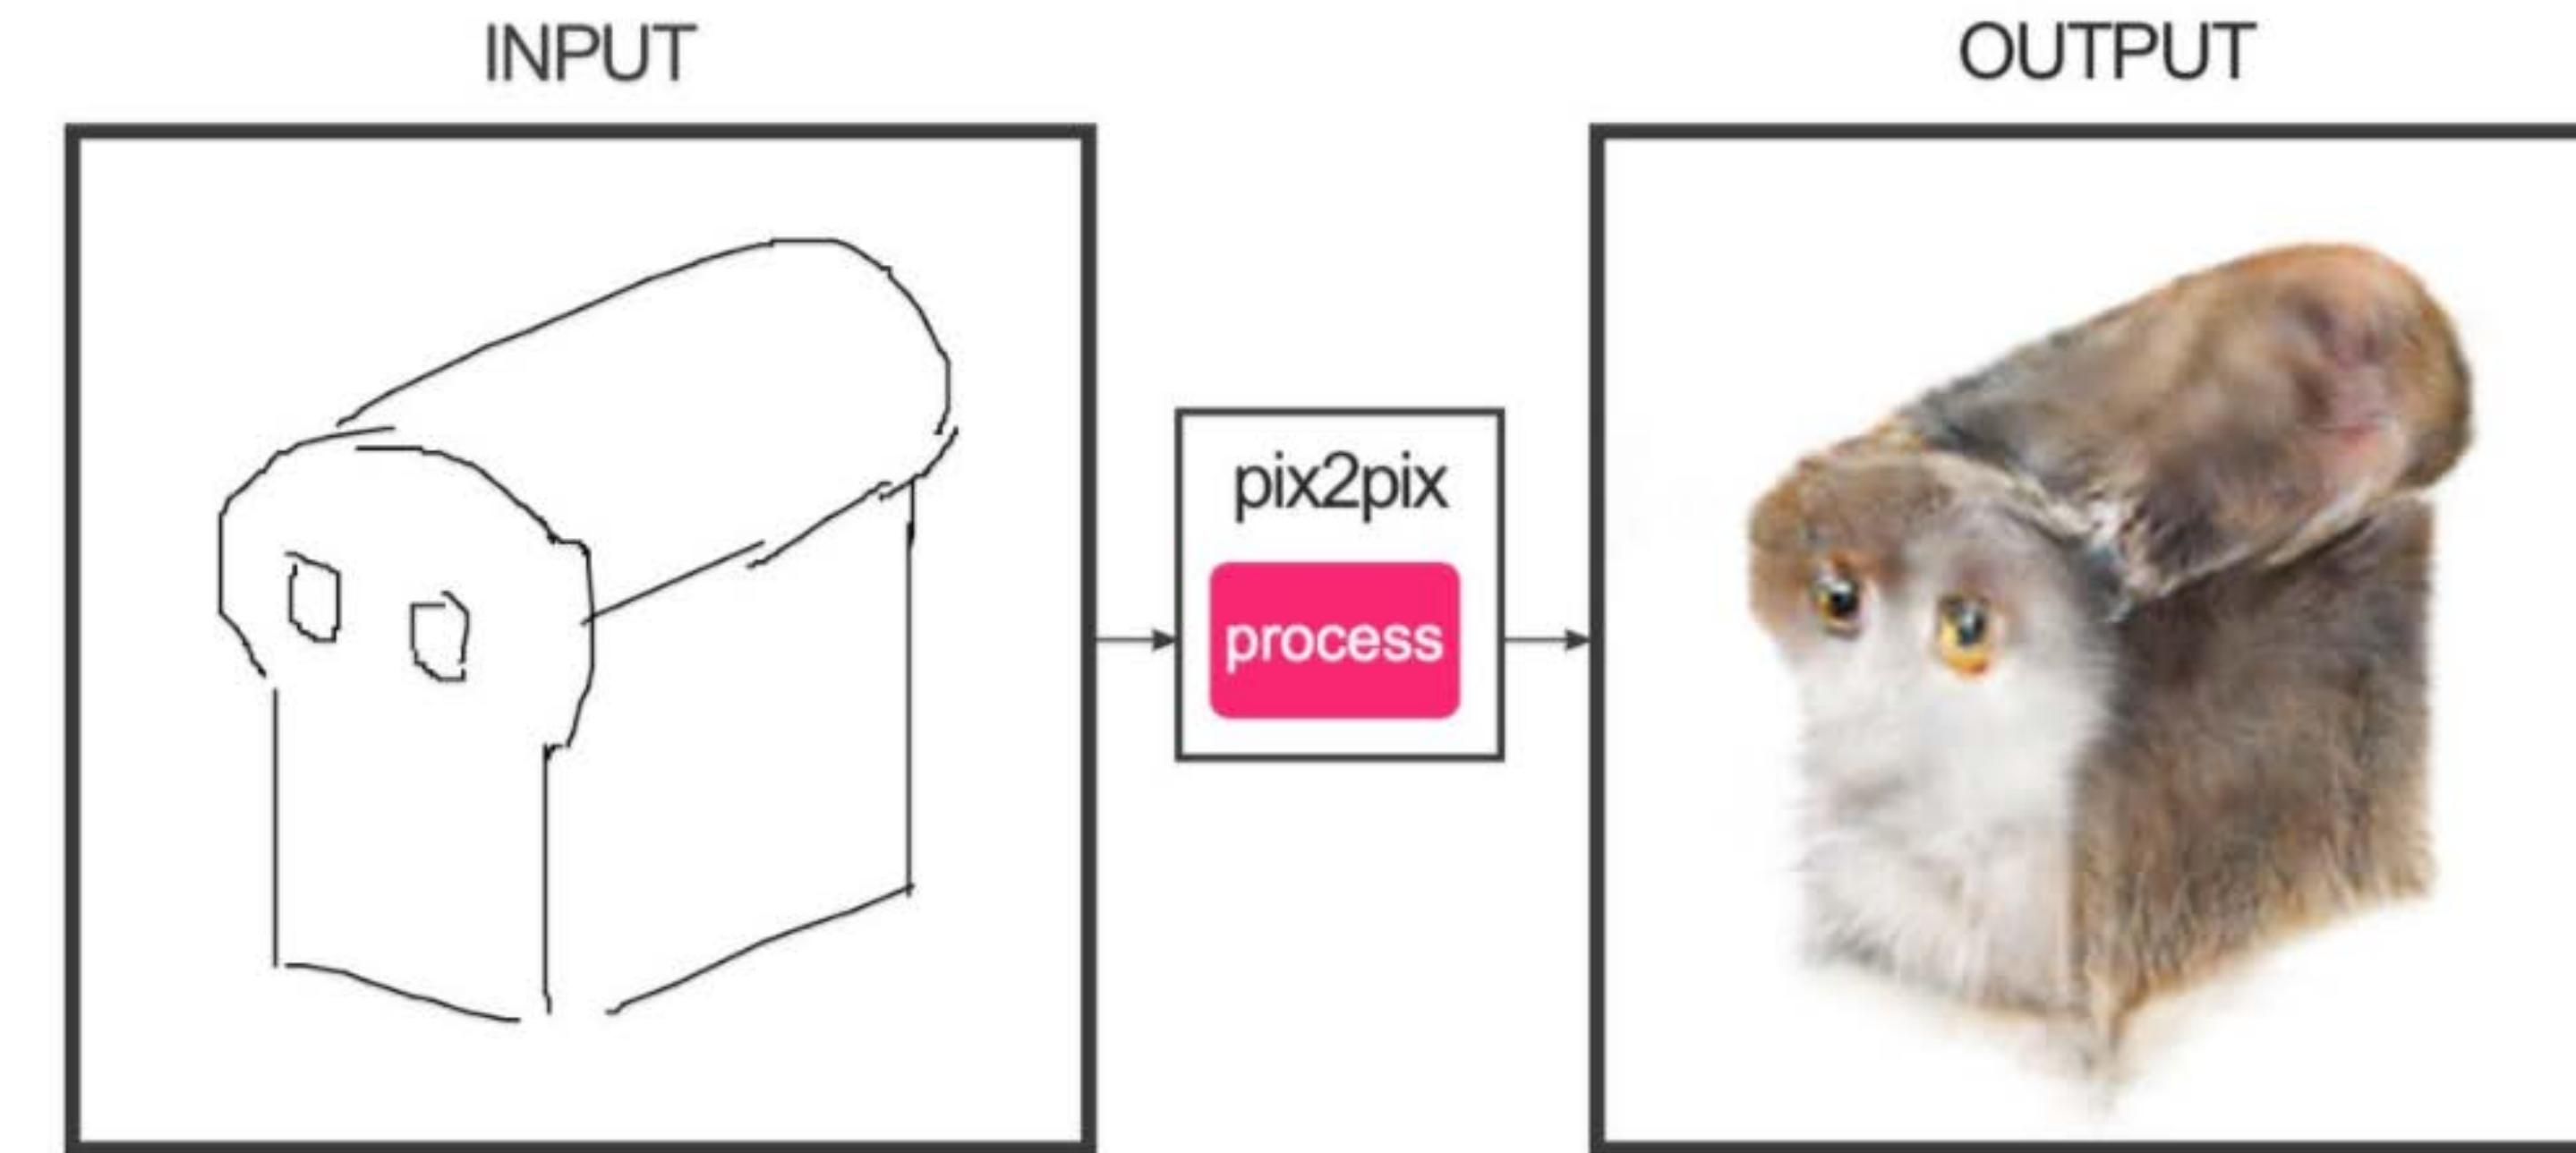
Output



Data from [Russakovsky et al. 2015]

#edges2cats [Chris Hesse]





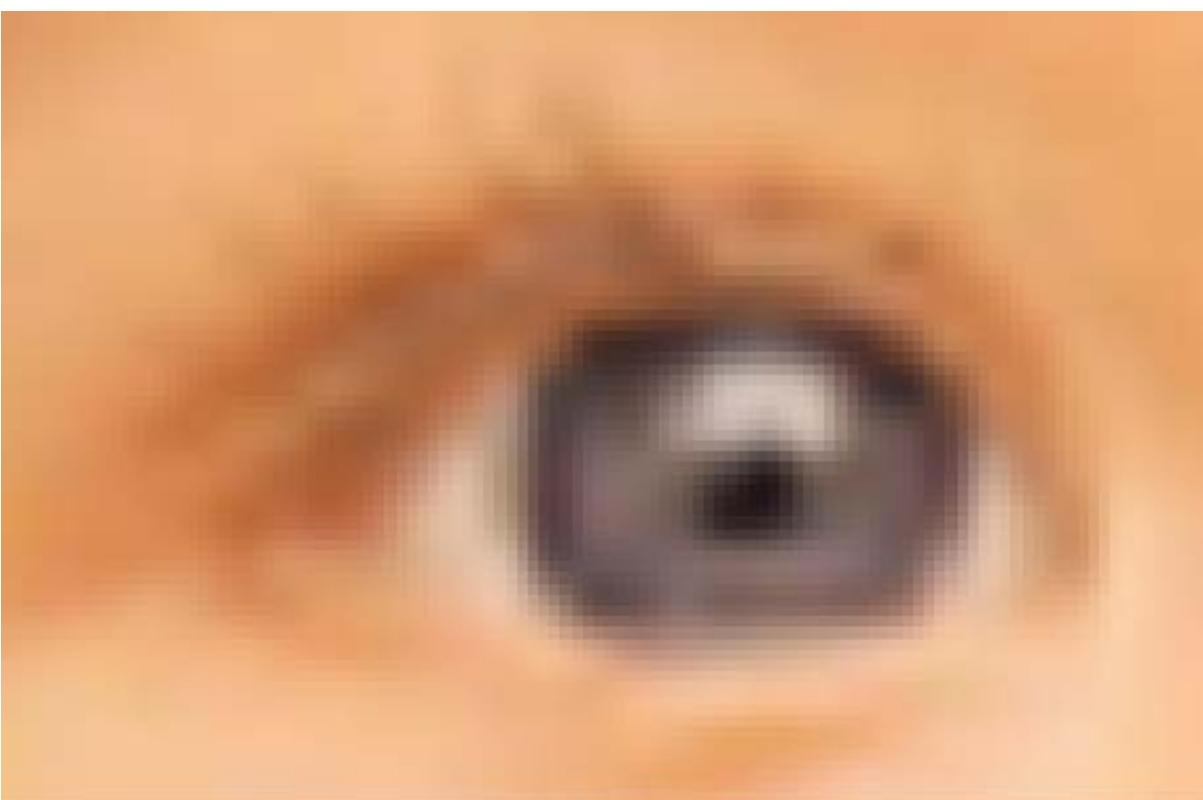
Ivy Tasi @ivymyt



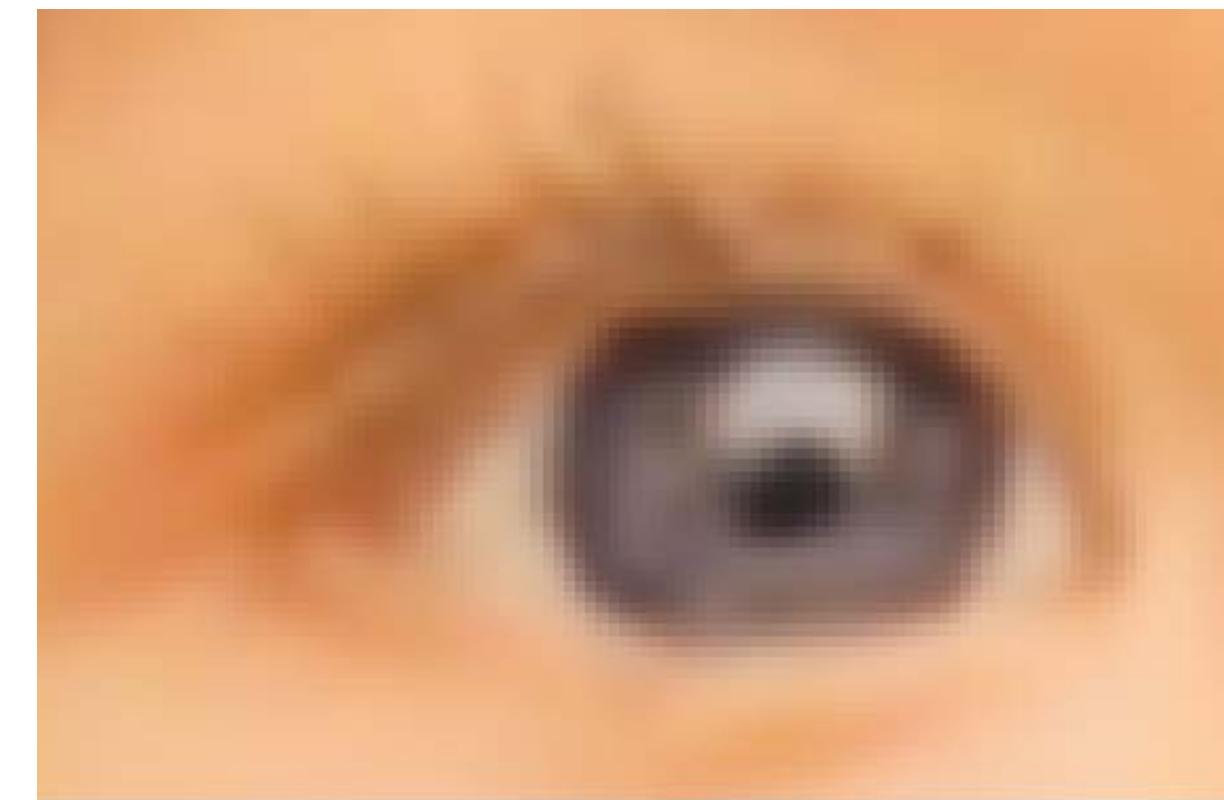
Vitaly Vidmirov @vvid

Structured Prediction

Input
 \mathbf{x}



Output
 $\hat{\mathbf{y}}$

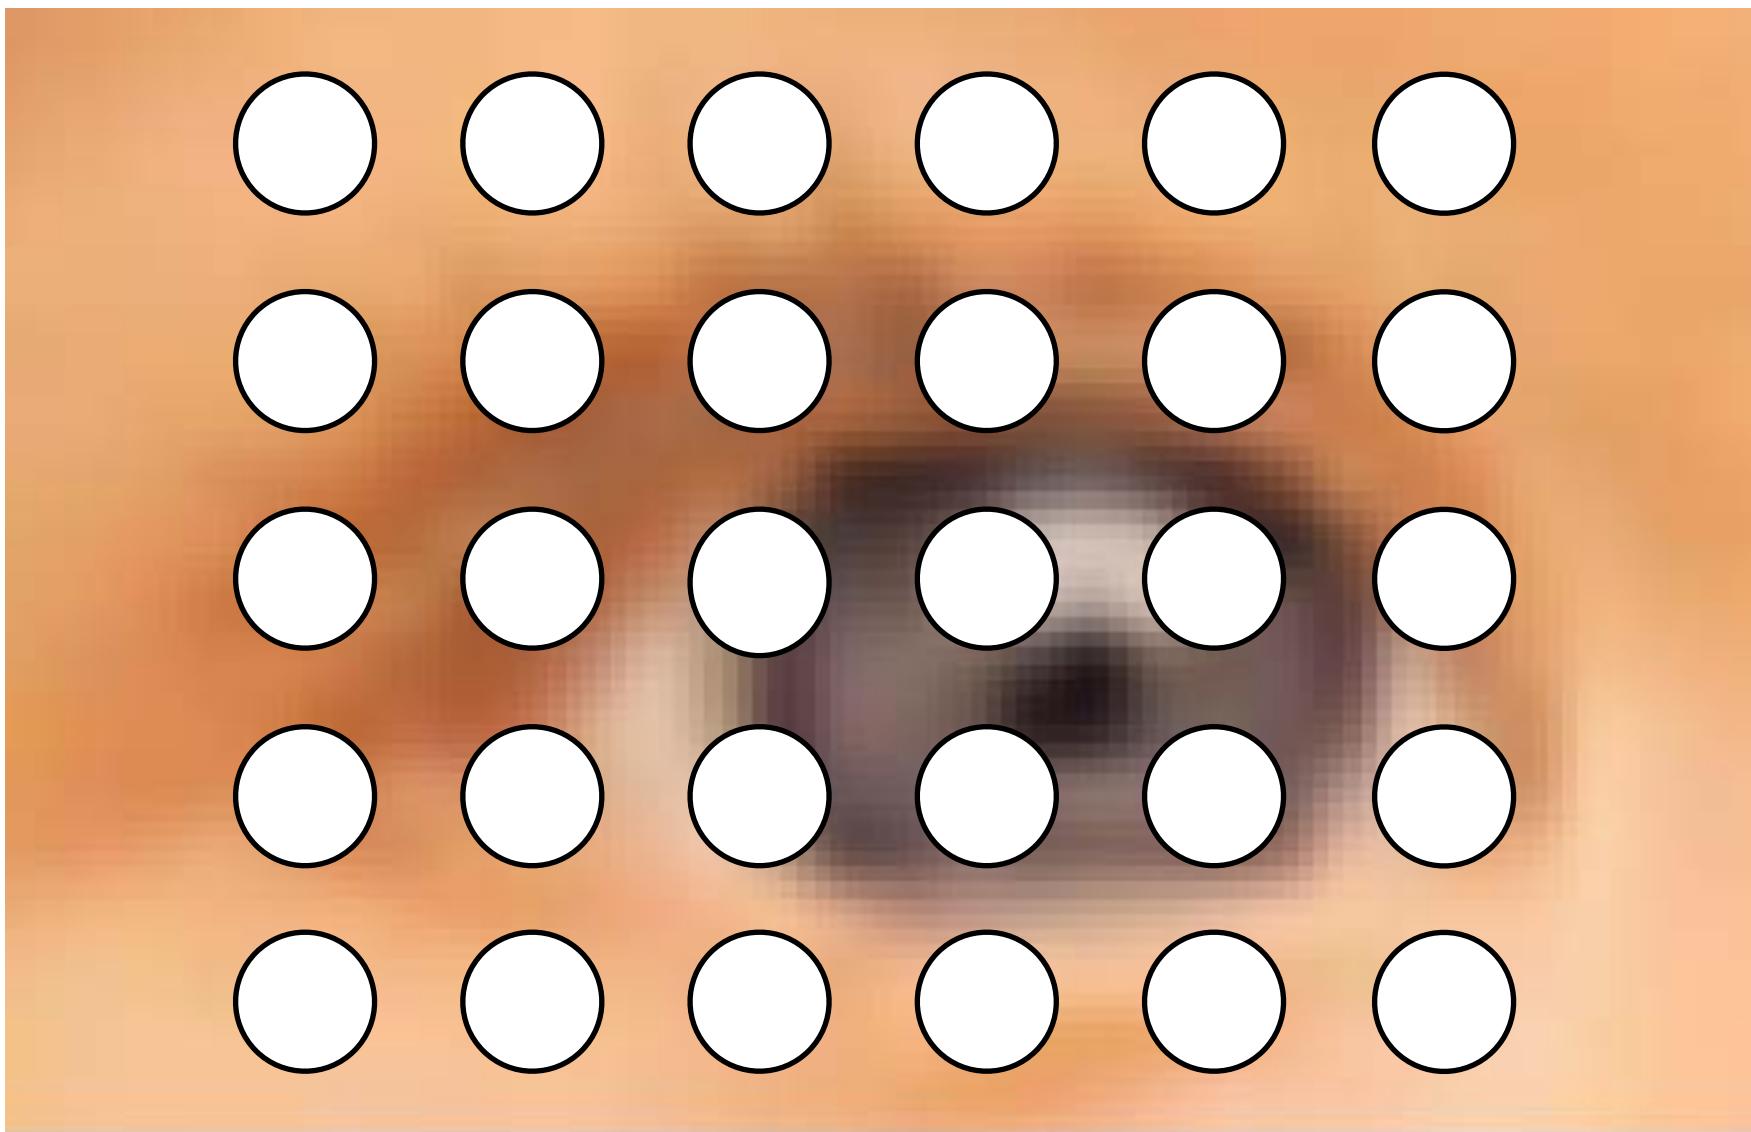


Target
 \mathbf{y}



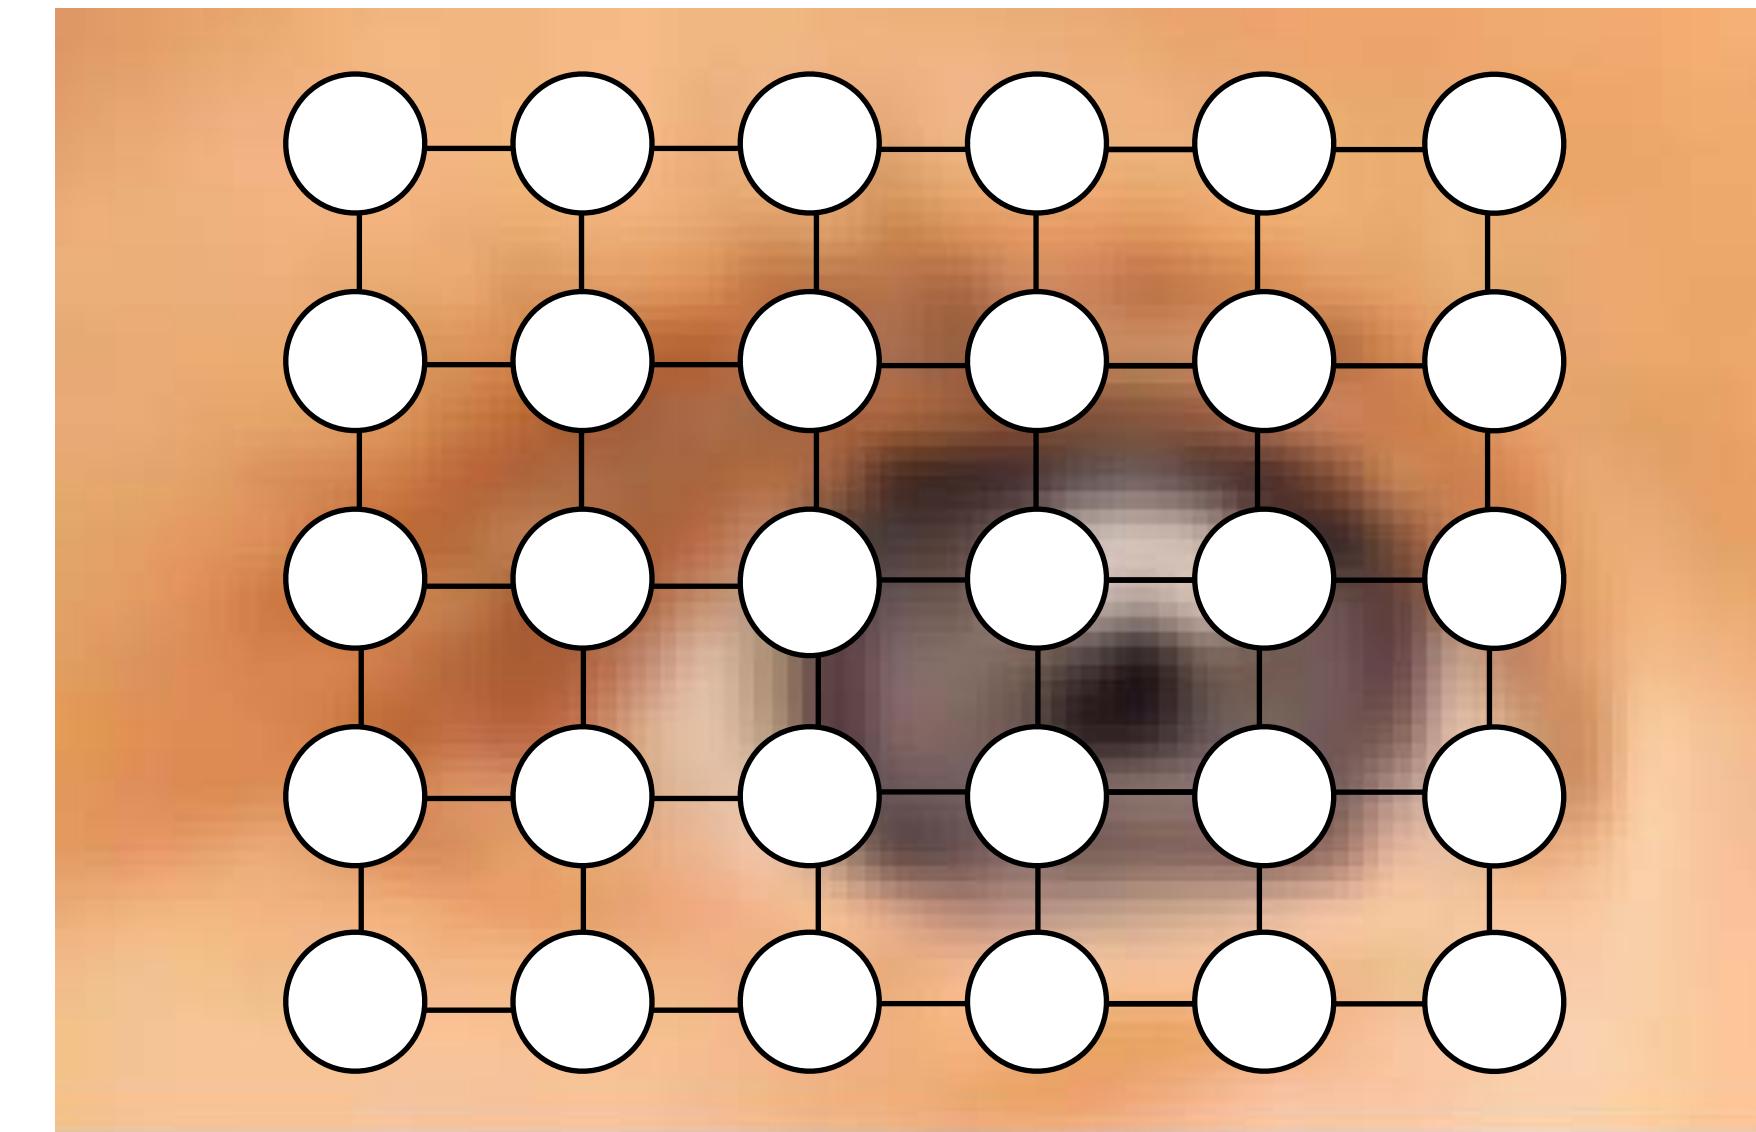
$$L(\hat{\mathbf{y}}, \mathbf{y}) = \|\hat{\mathbf{y}} - \mathbf{y}\|_2$$

Structured Prediction



Each pixel treated as
independent

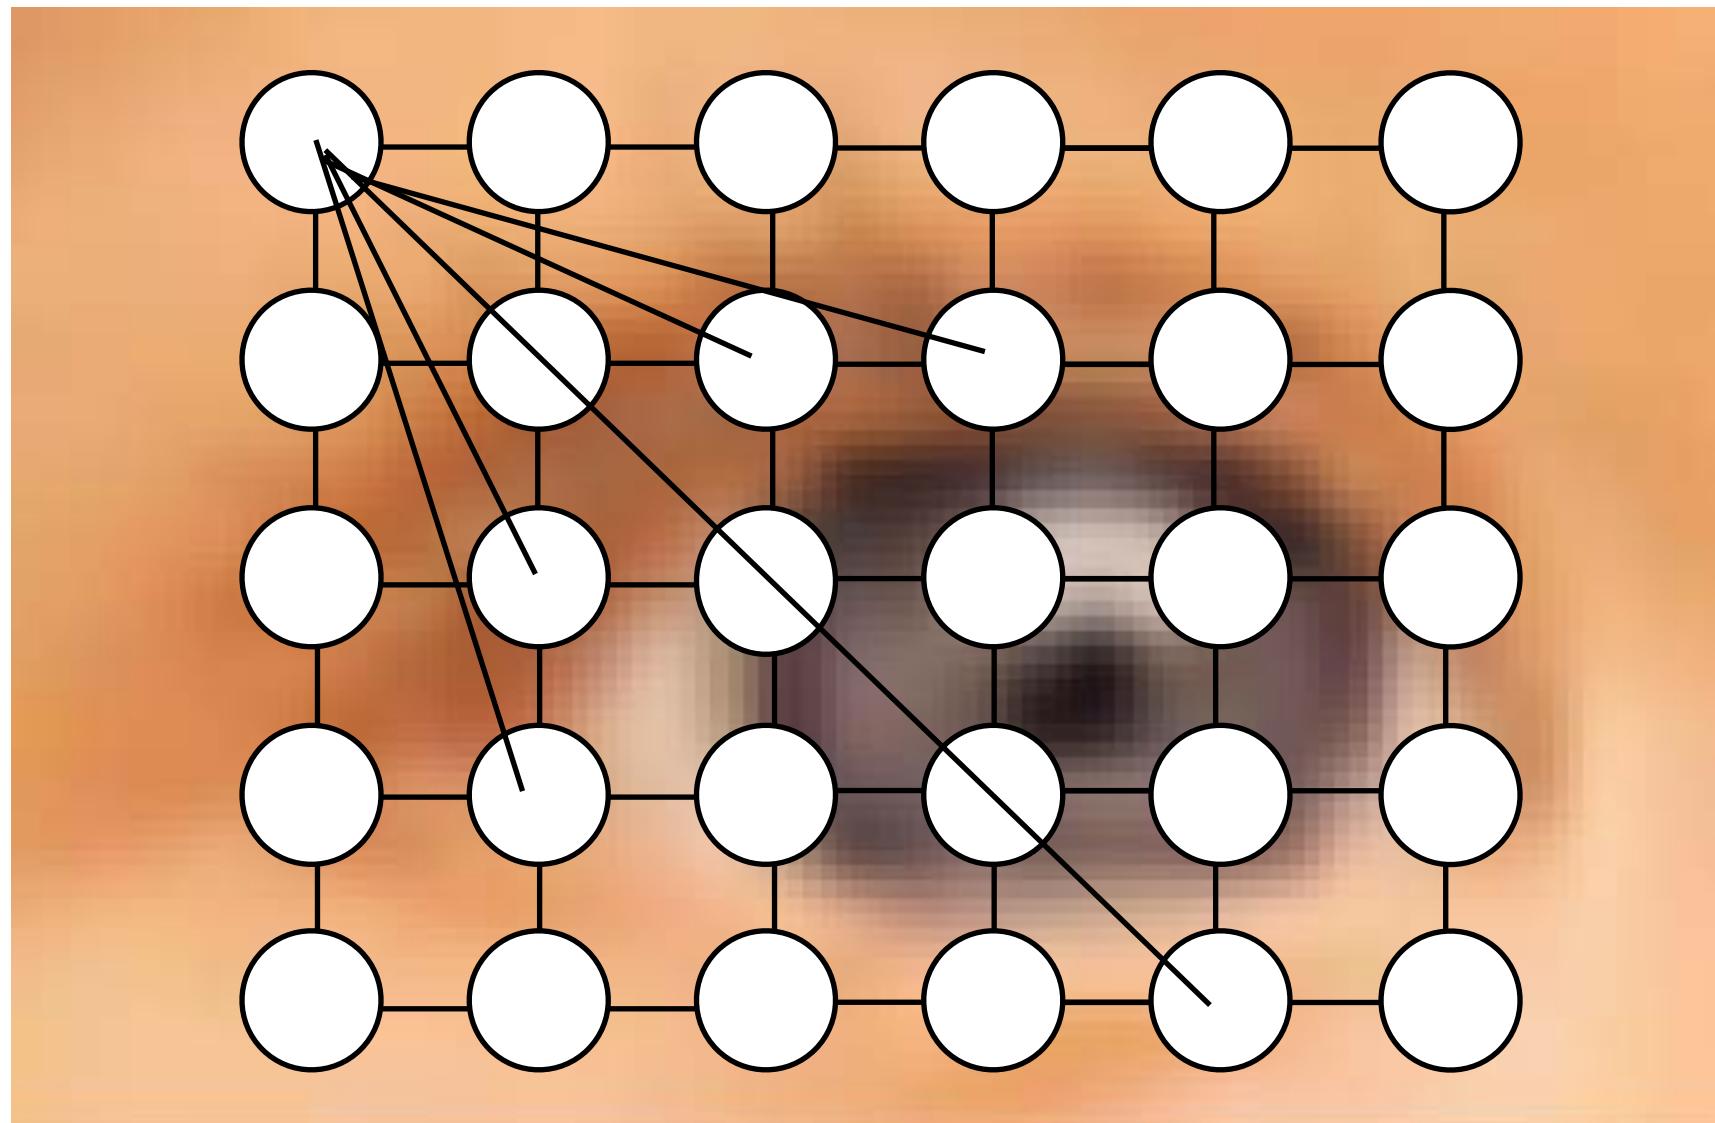
$$\prod_i p(y_i | \mathbf{x})$$



Models at pairwise configuration
of pixels

$$\frac{1}{Z} \prod_{i,j} p(y_i, y_j | \mathbf{x})$$

Structured Prediction



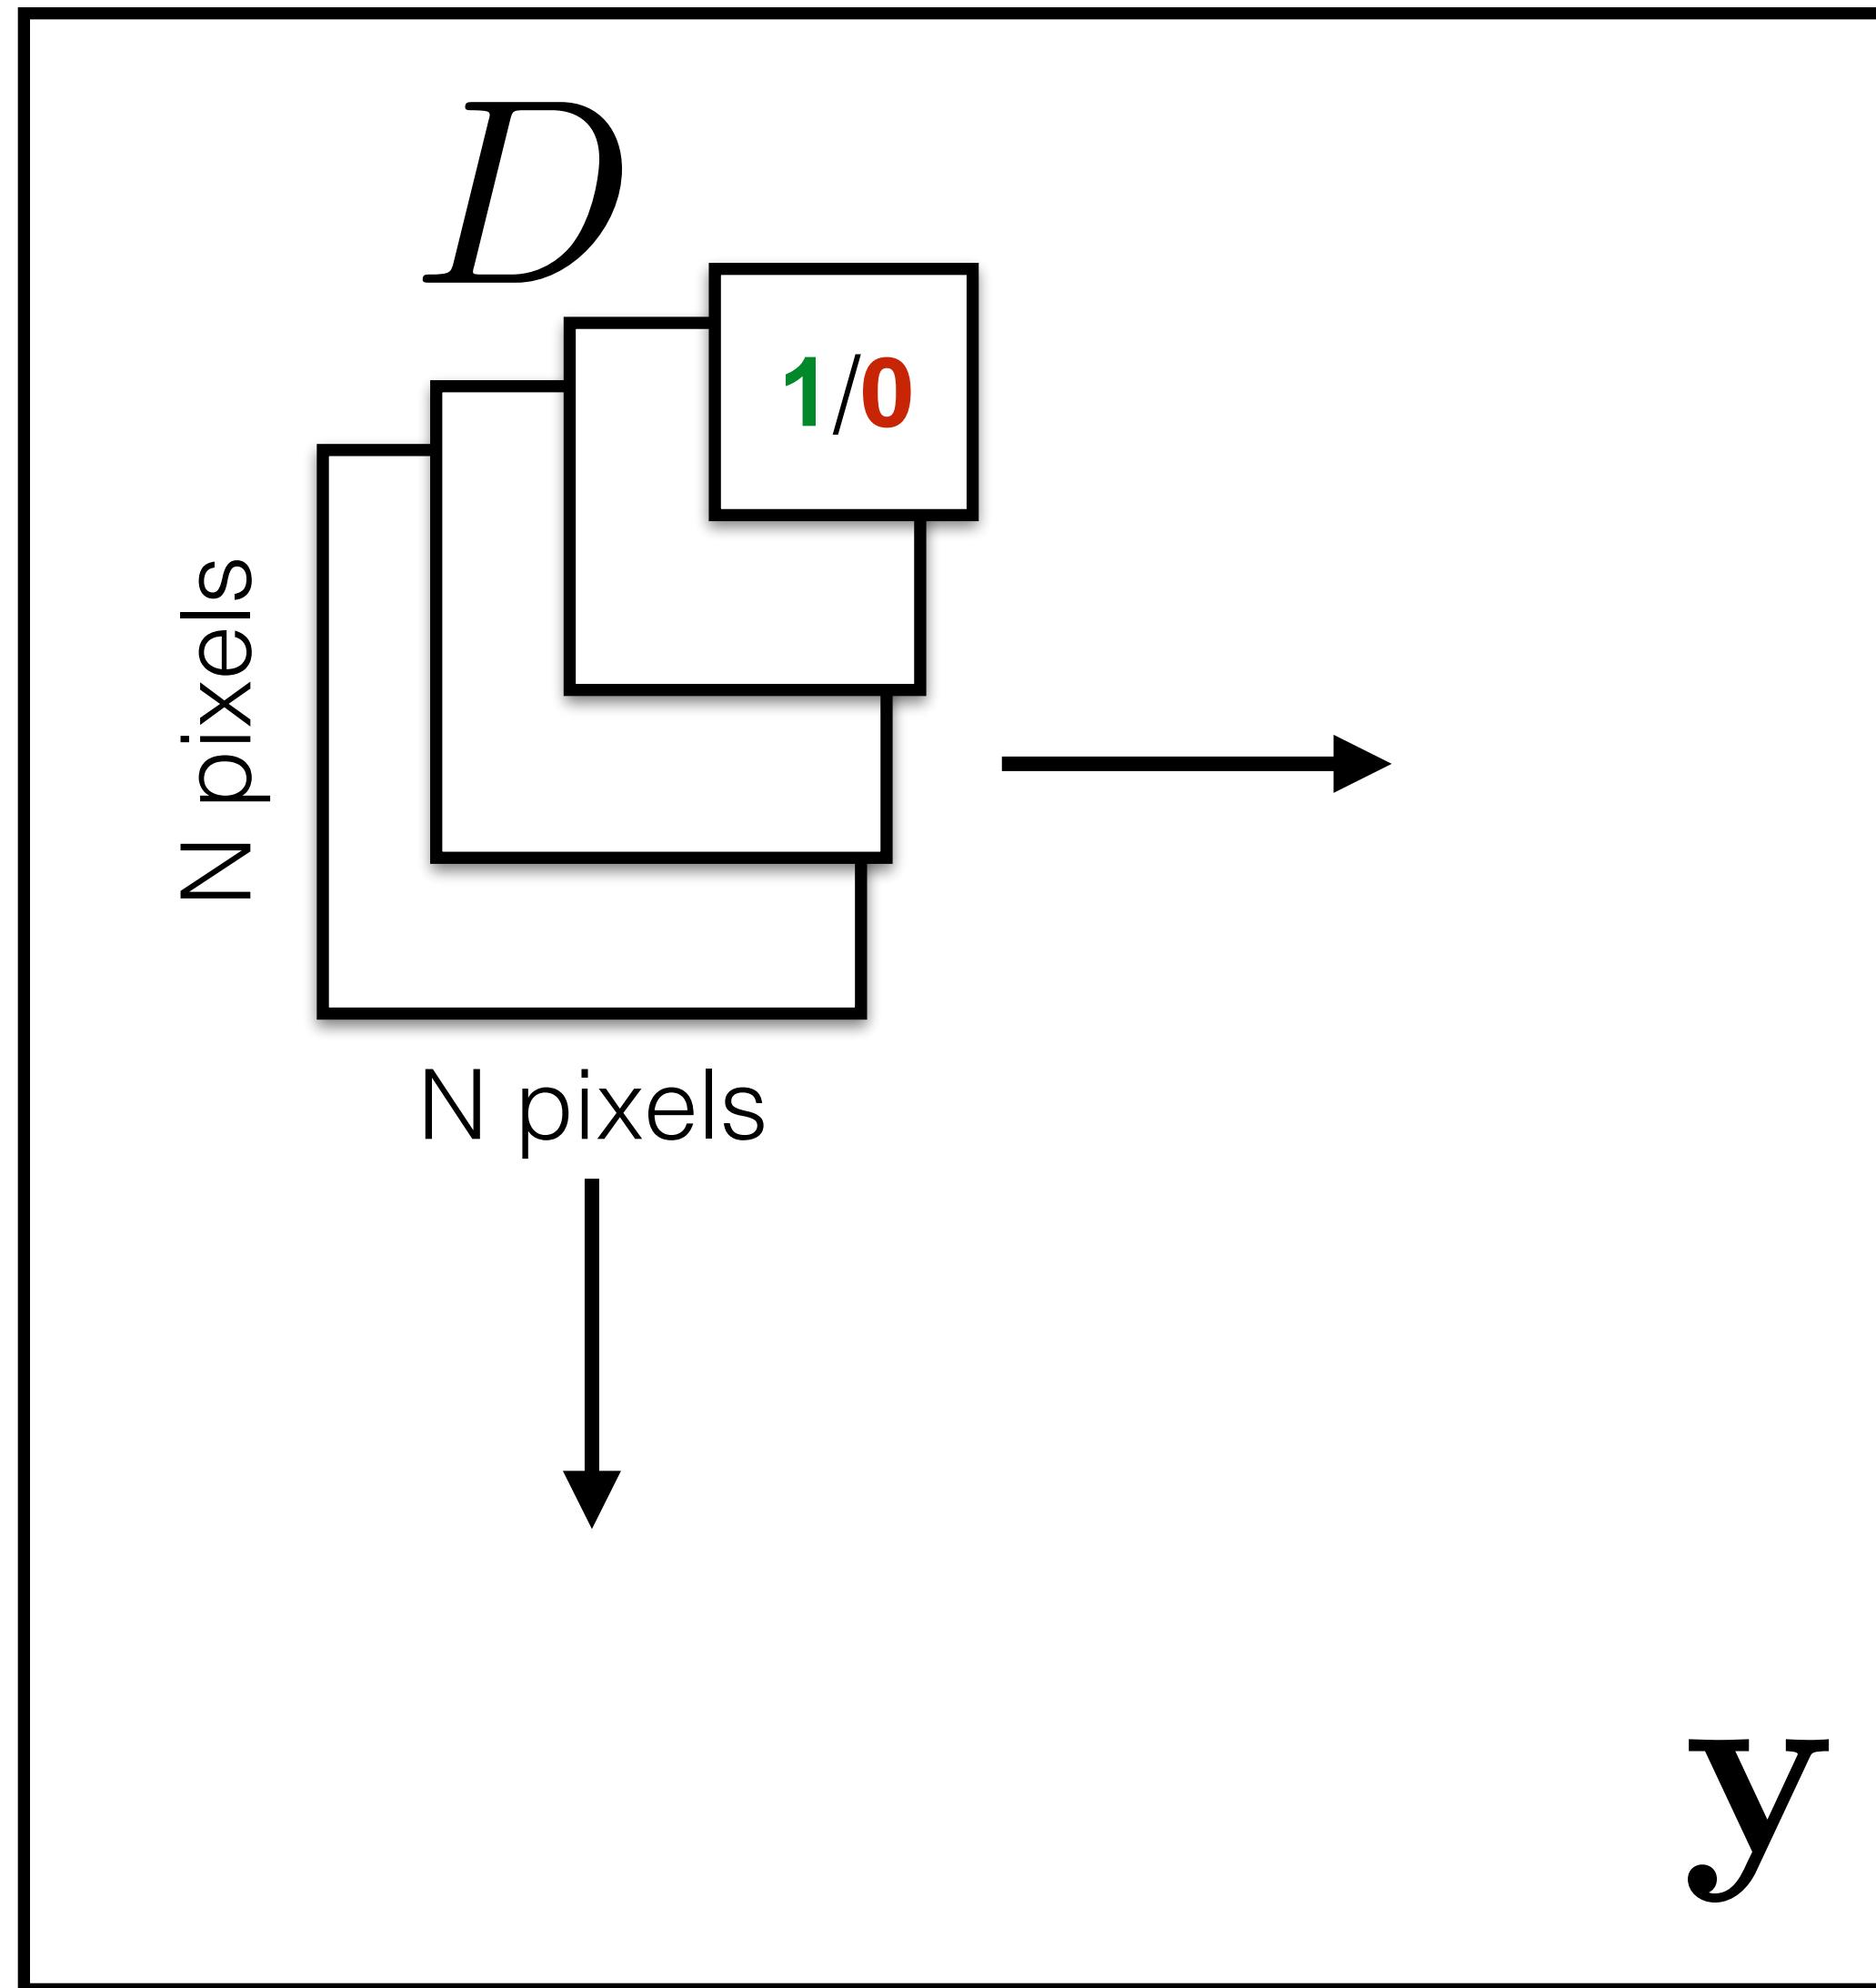
Model *joint* configuration
of all pixels

$$p(\mathbf{y}|\mathbf{x})$$

A GAN, with sufficient capacity,
samples from the full joint distribution
(at equilibrium)

Most generative models have this
property! Give them **sufficient**
capacity and **infinite data**, and they
are the complete solution to
prediction problems.

Shrinking the capacity: Patch Discriminator



Rather than penalizing if output *image* looks fake, penalize if each overlapping *patch* in output looks fake

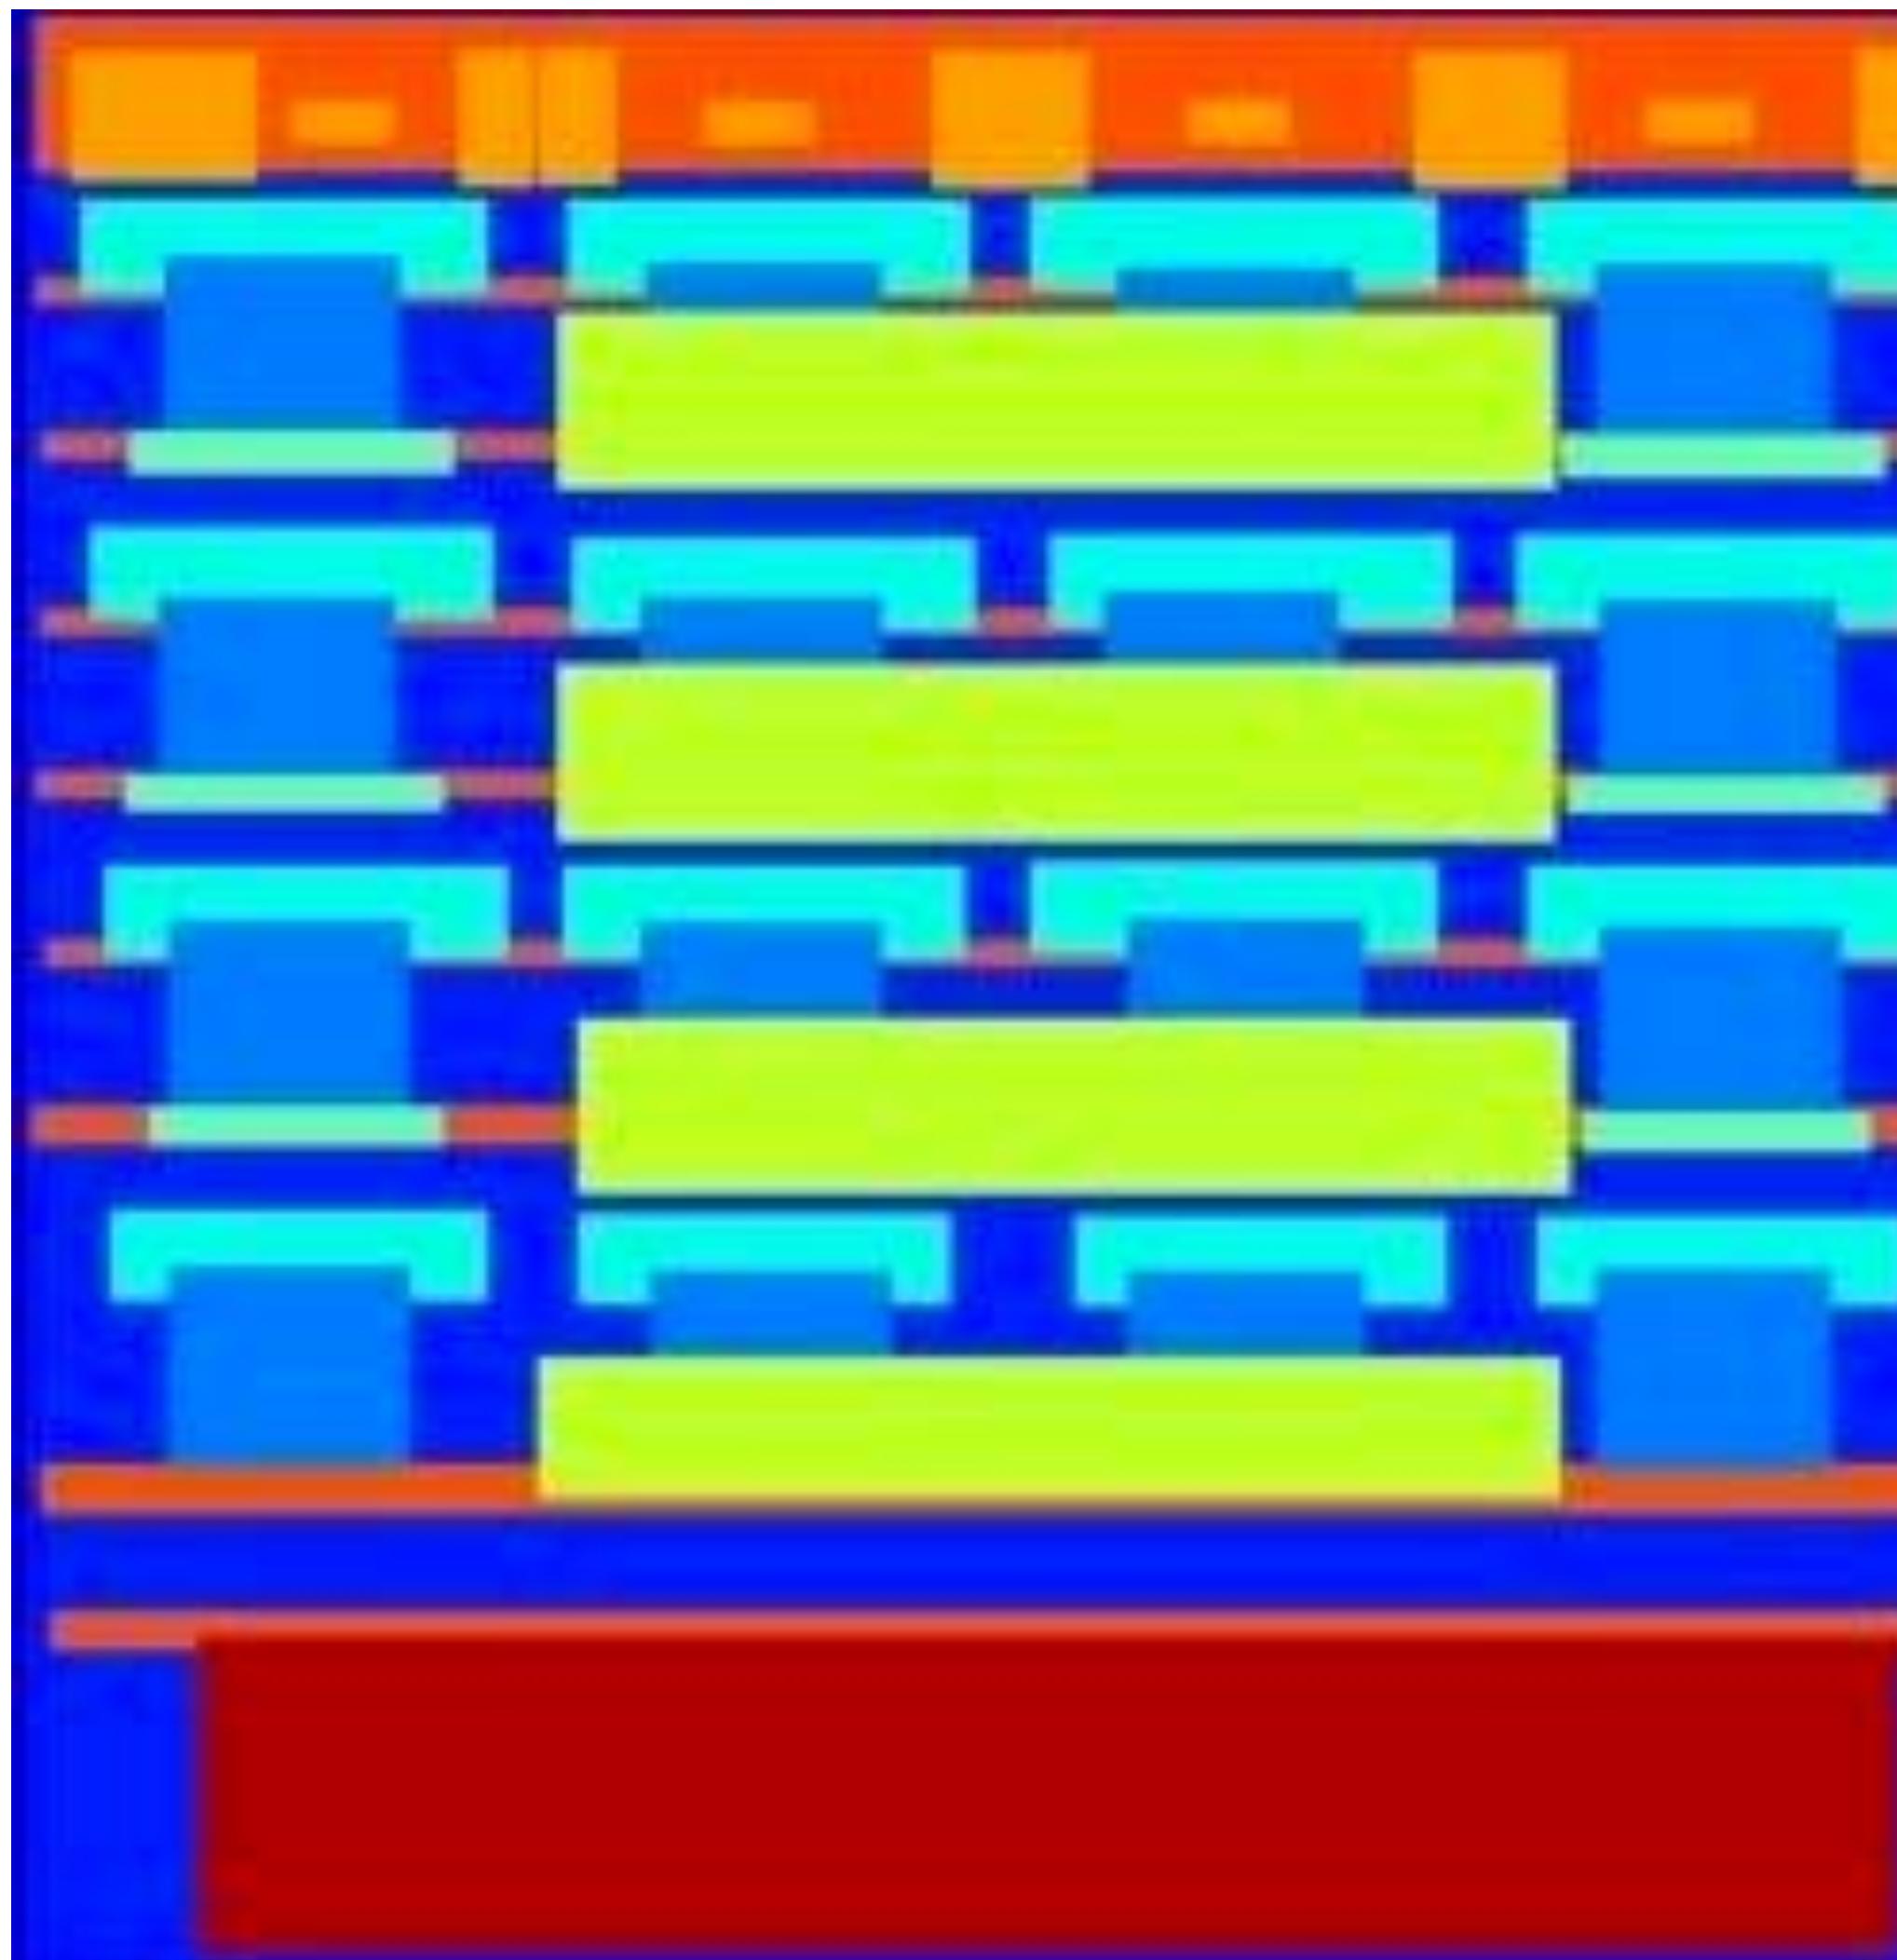
[Li & Wand 2016]

[Shrivastava et al. 2017]

[Isola et al. 2017]

Labels → Facades

Input



1x1 Discriminator



Data from [Tylecek, 2013]

Labels → Facades

Input



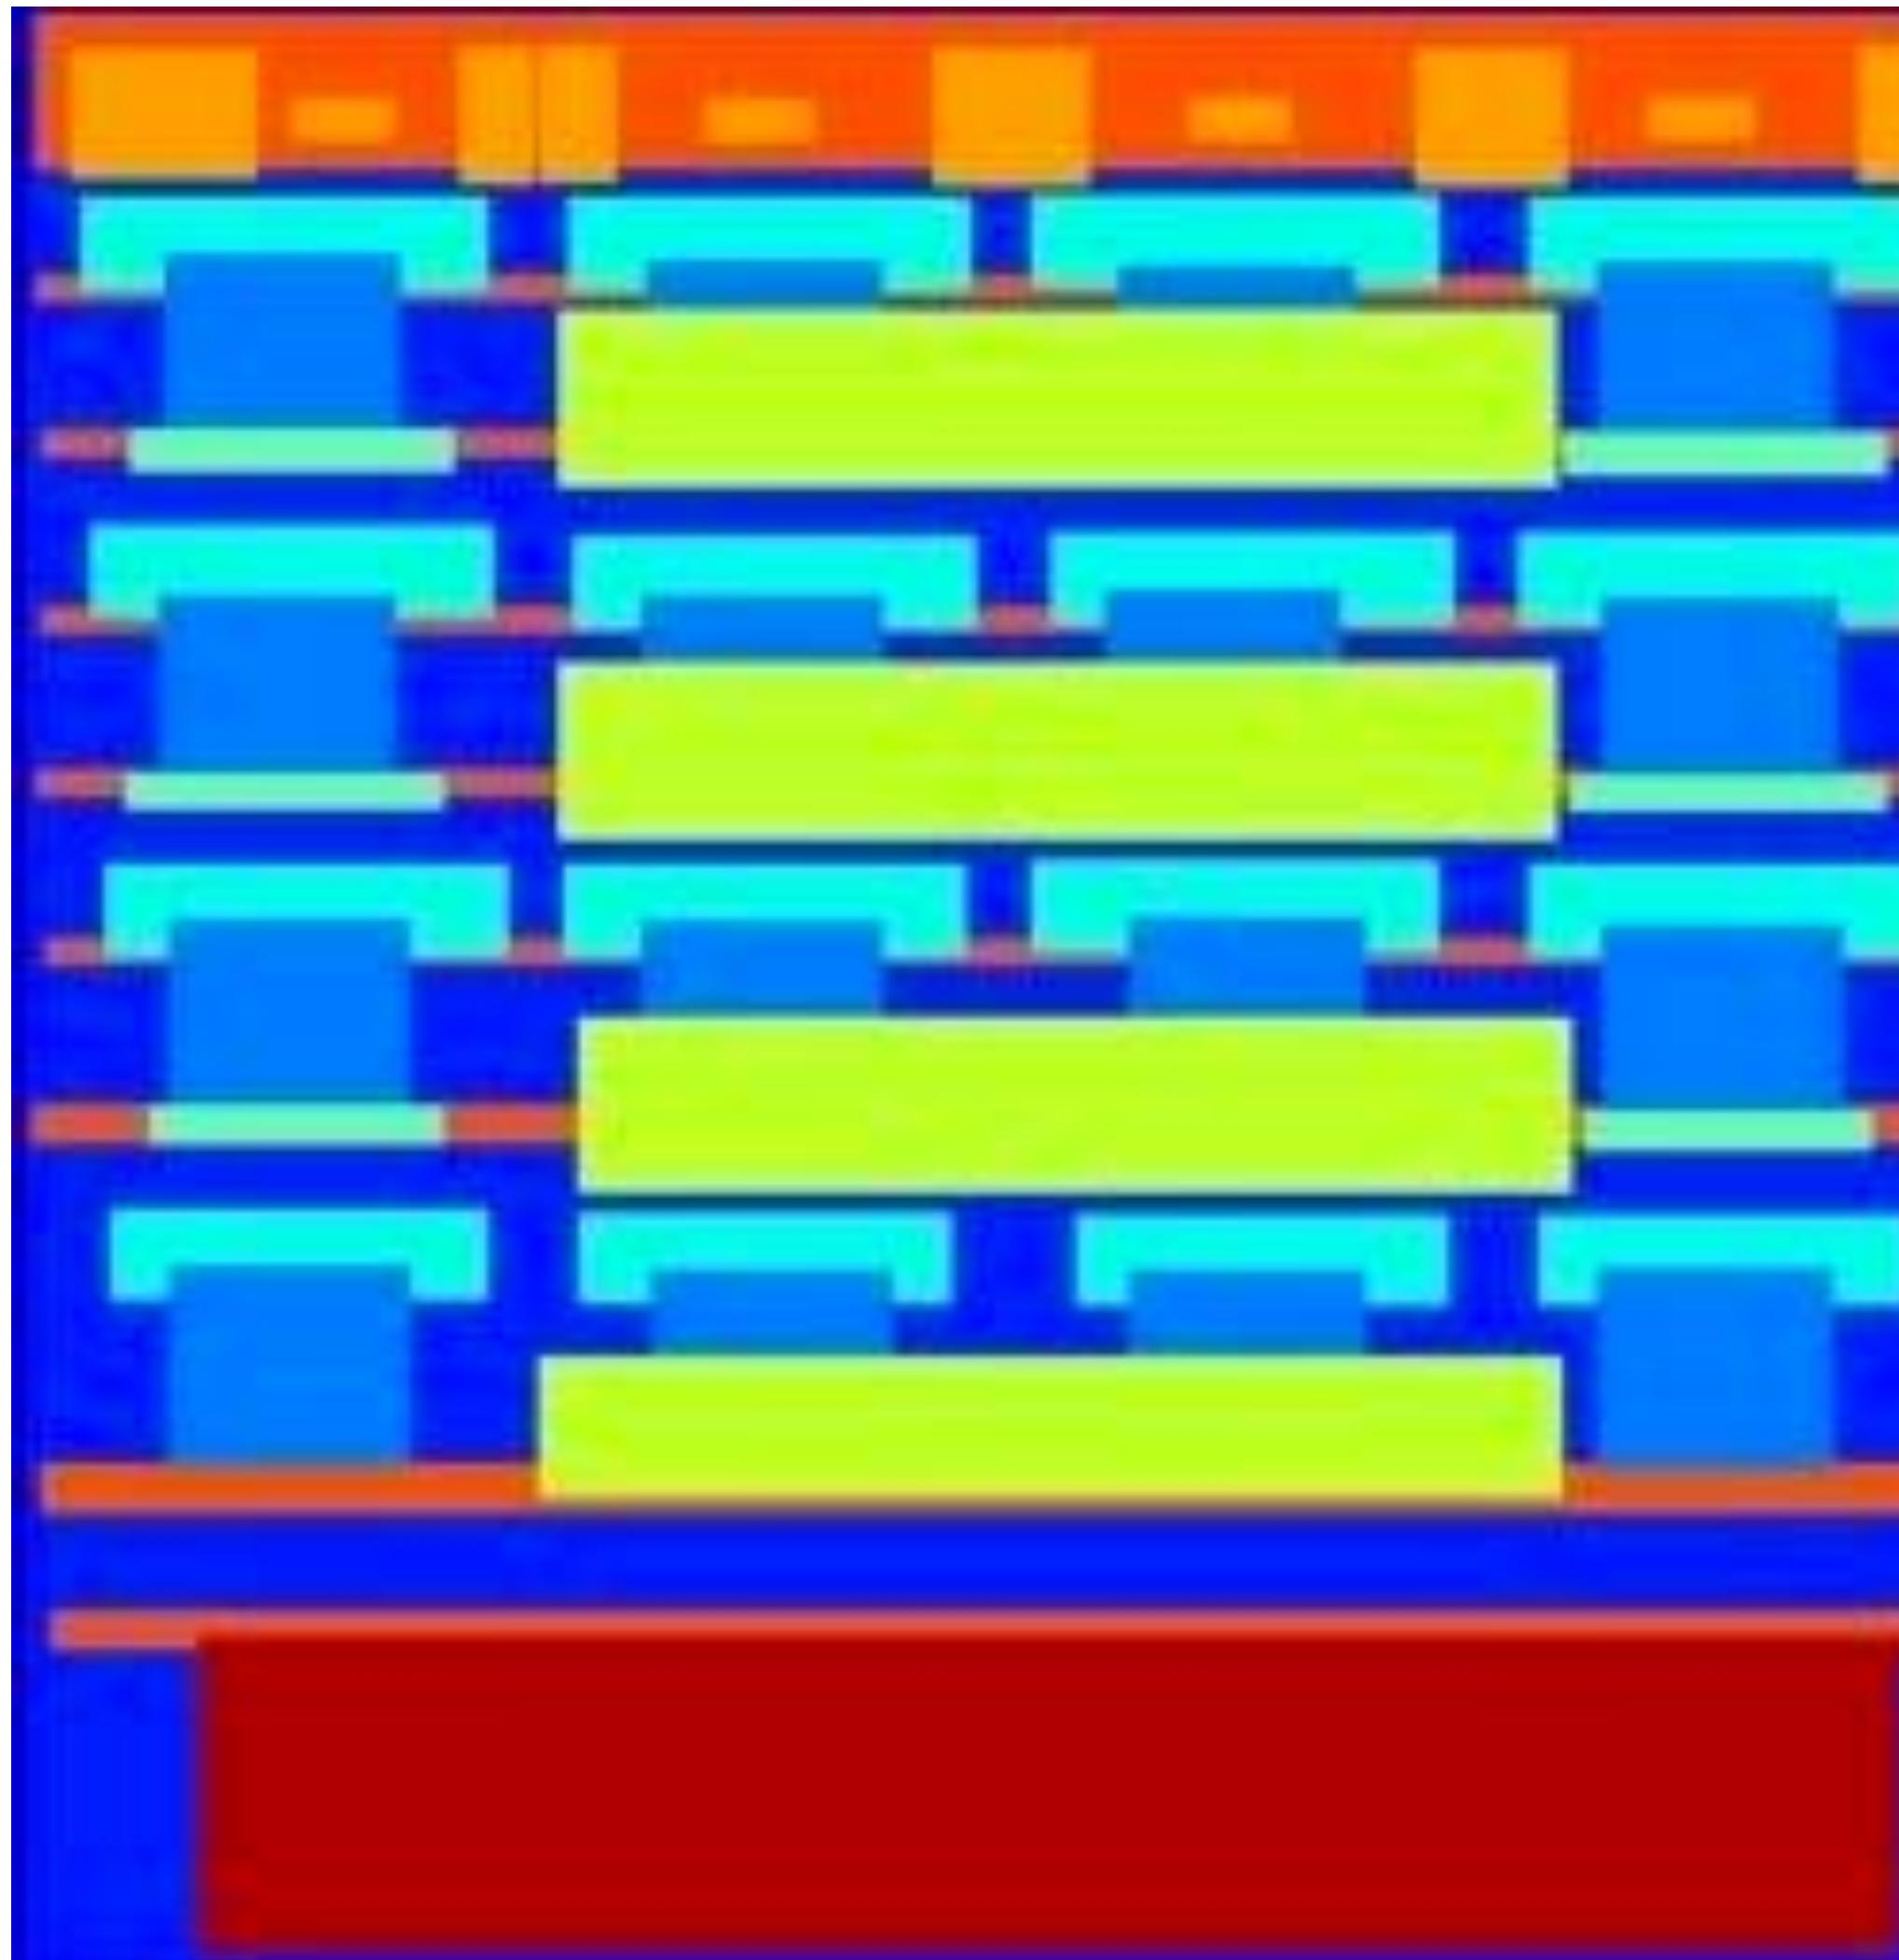
16x16 Discriminator



Data from [Tylecek, 2013]

Labels → Facades

Input



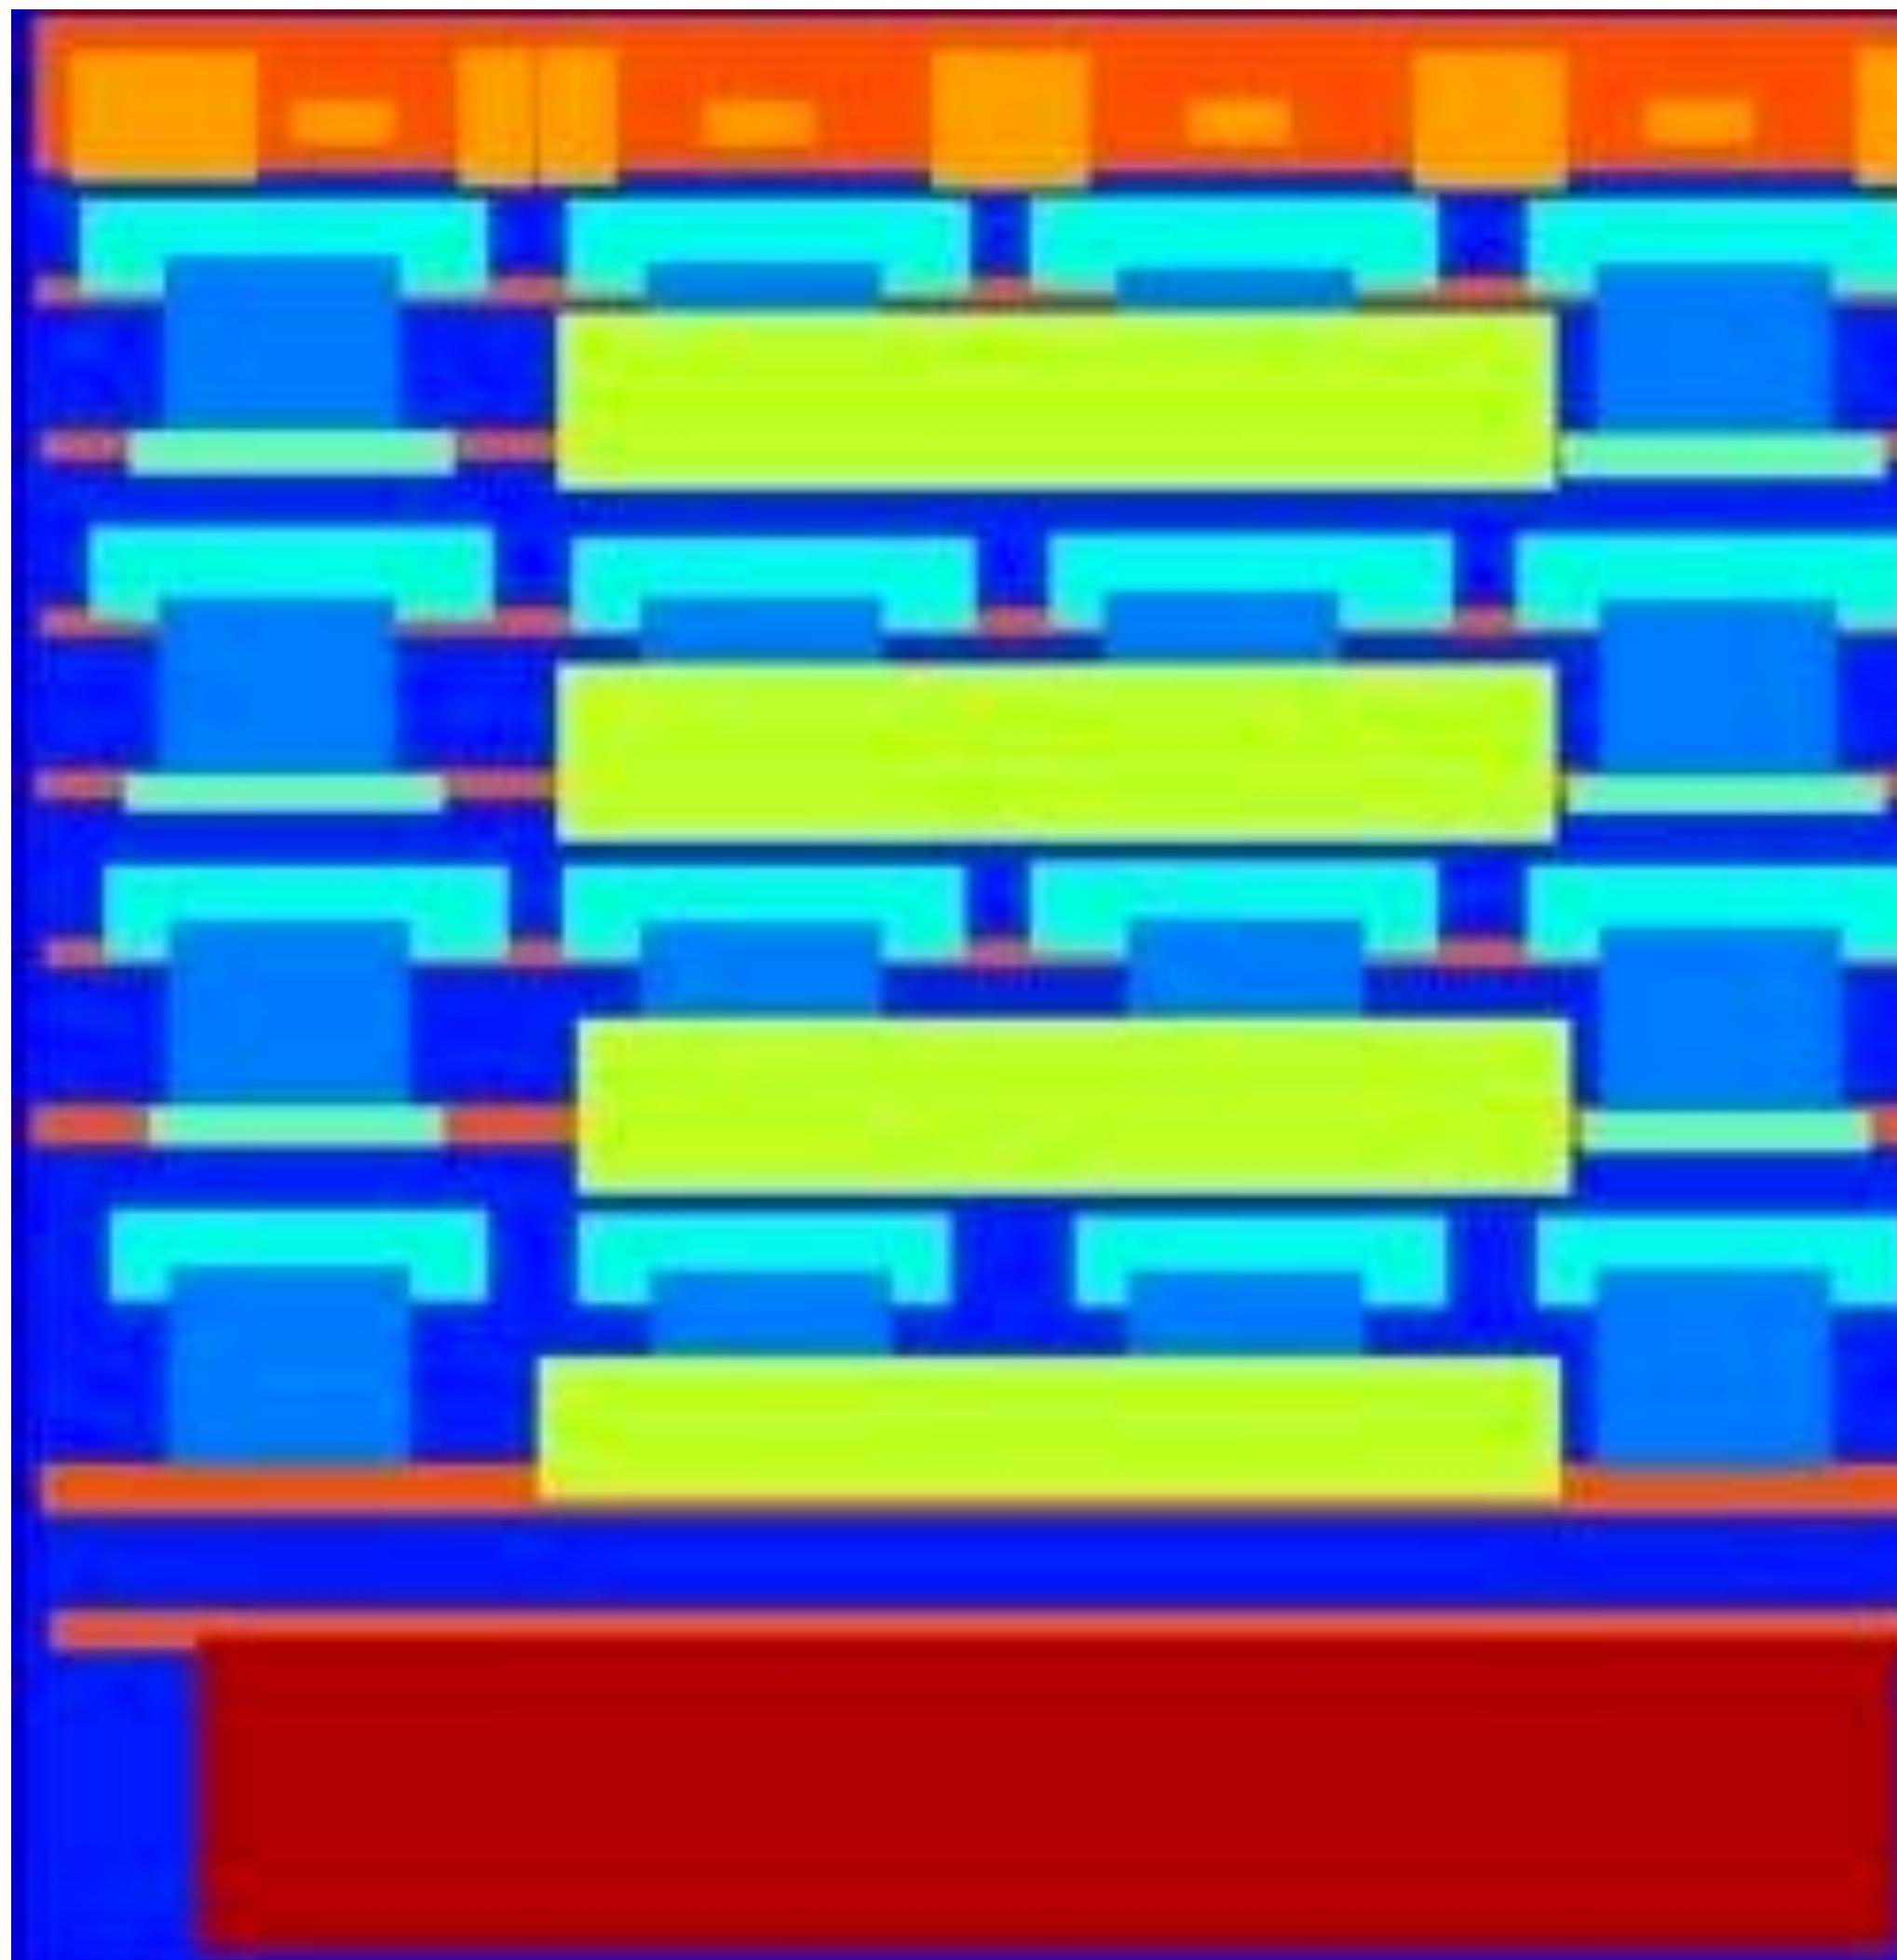
70x70 Discriminator



Data from [Tylecek, 2013]

Labels → Facades

Input

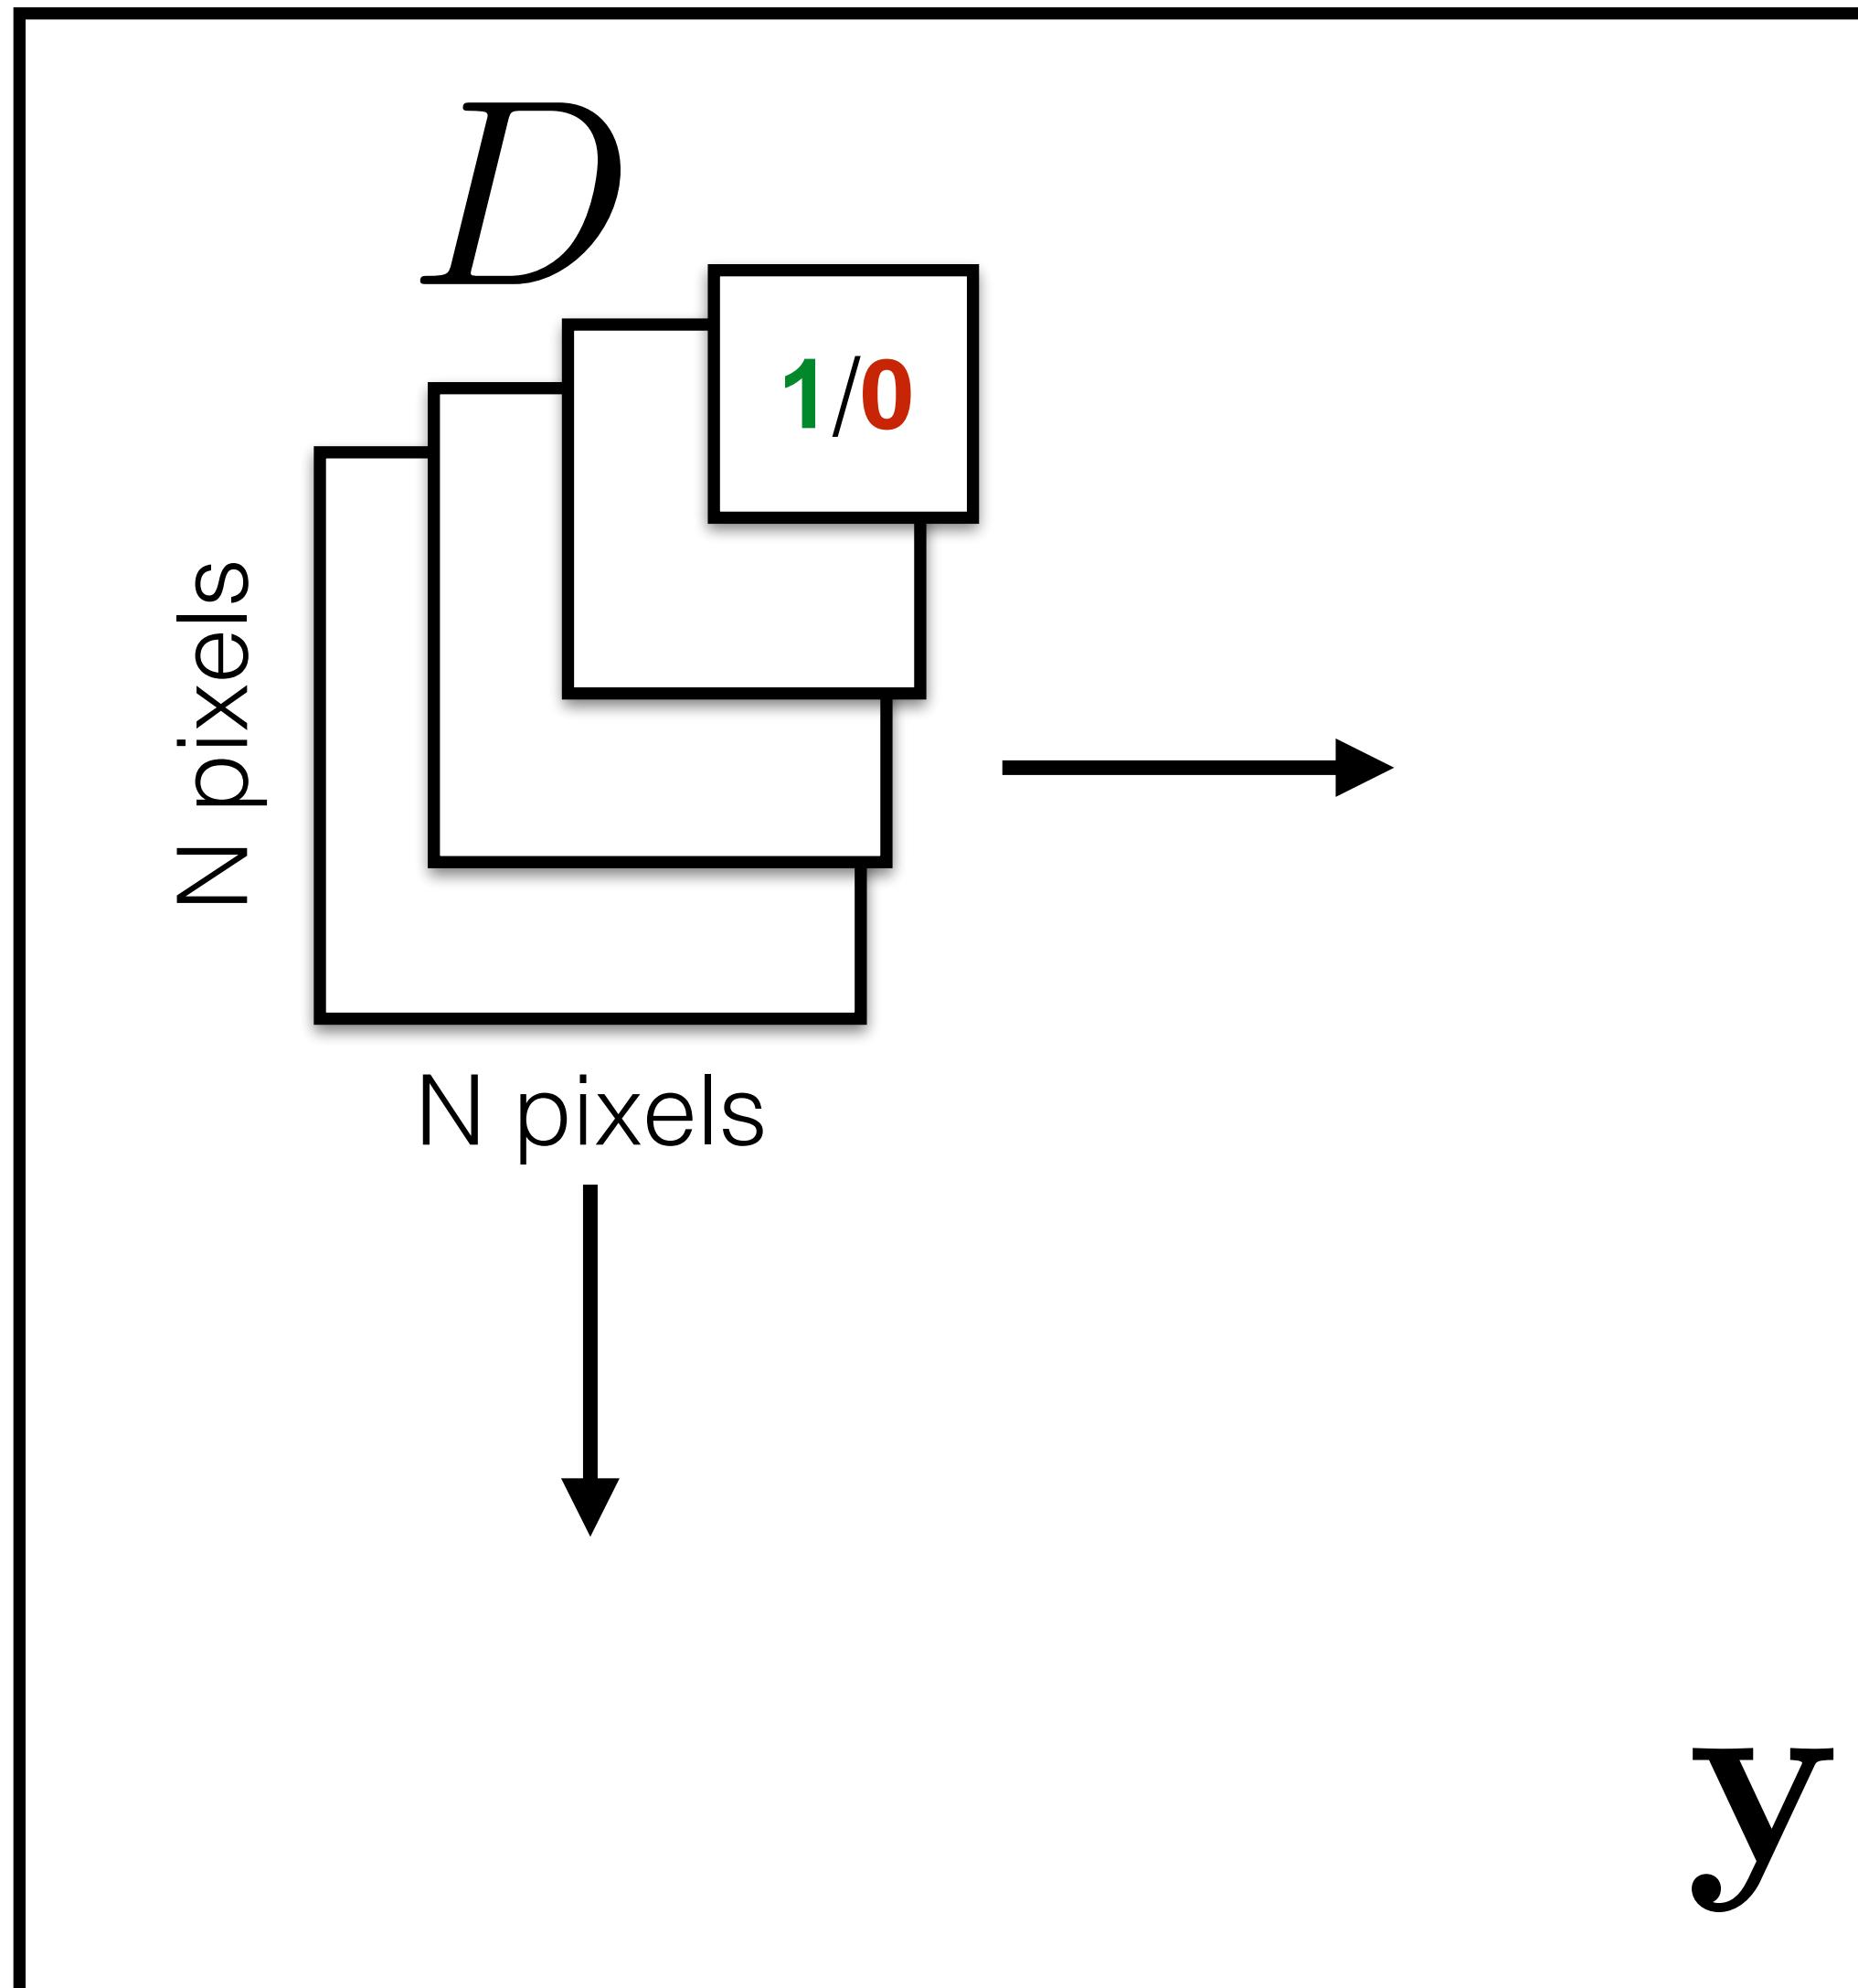


Full image Discriminator



Data from [Tylecek, 2013]

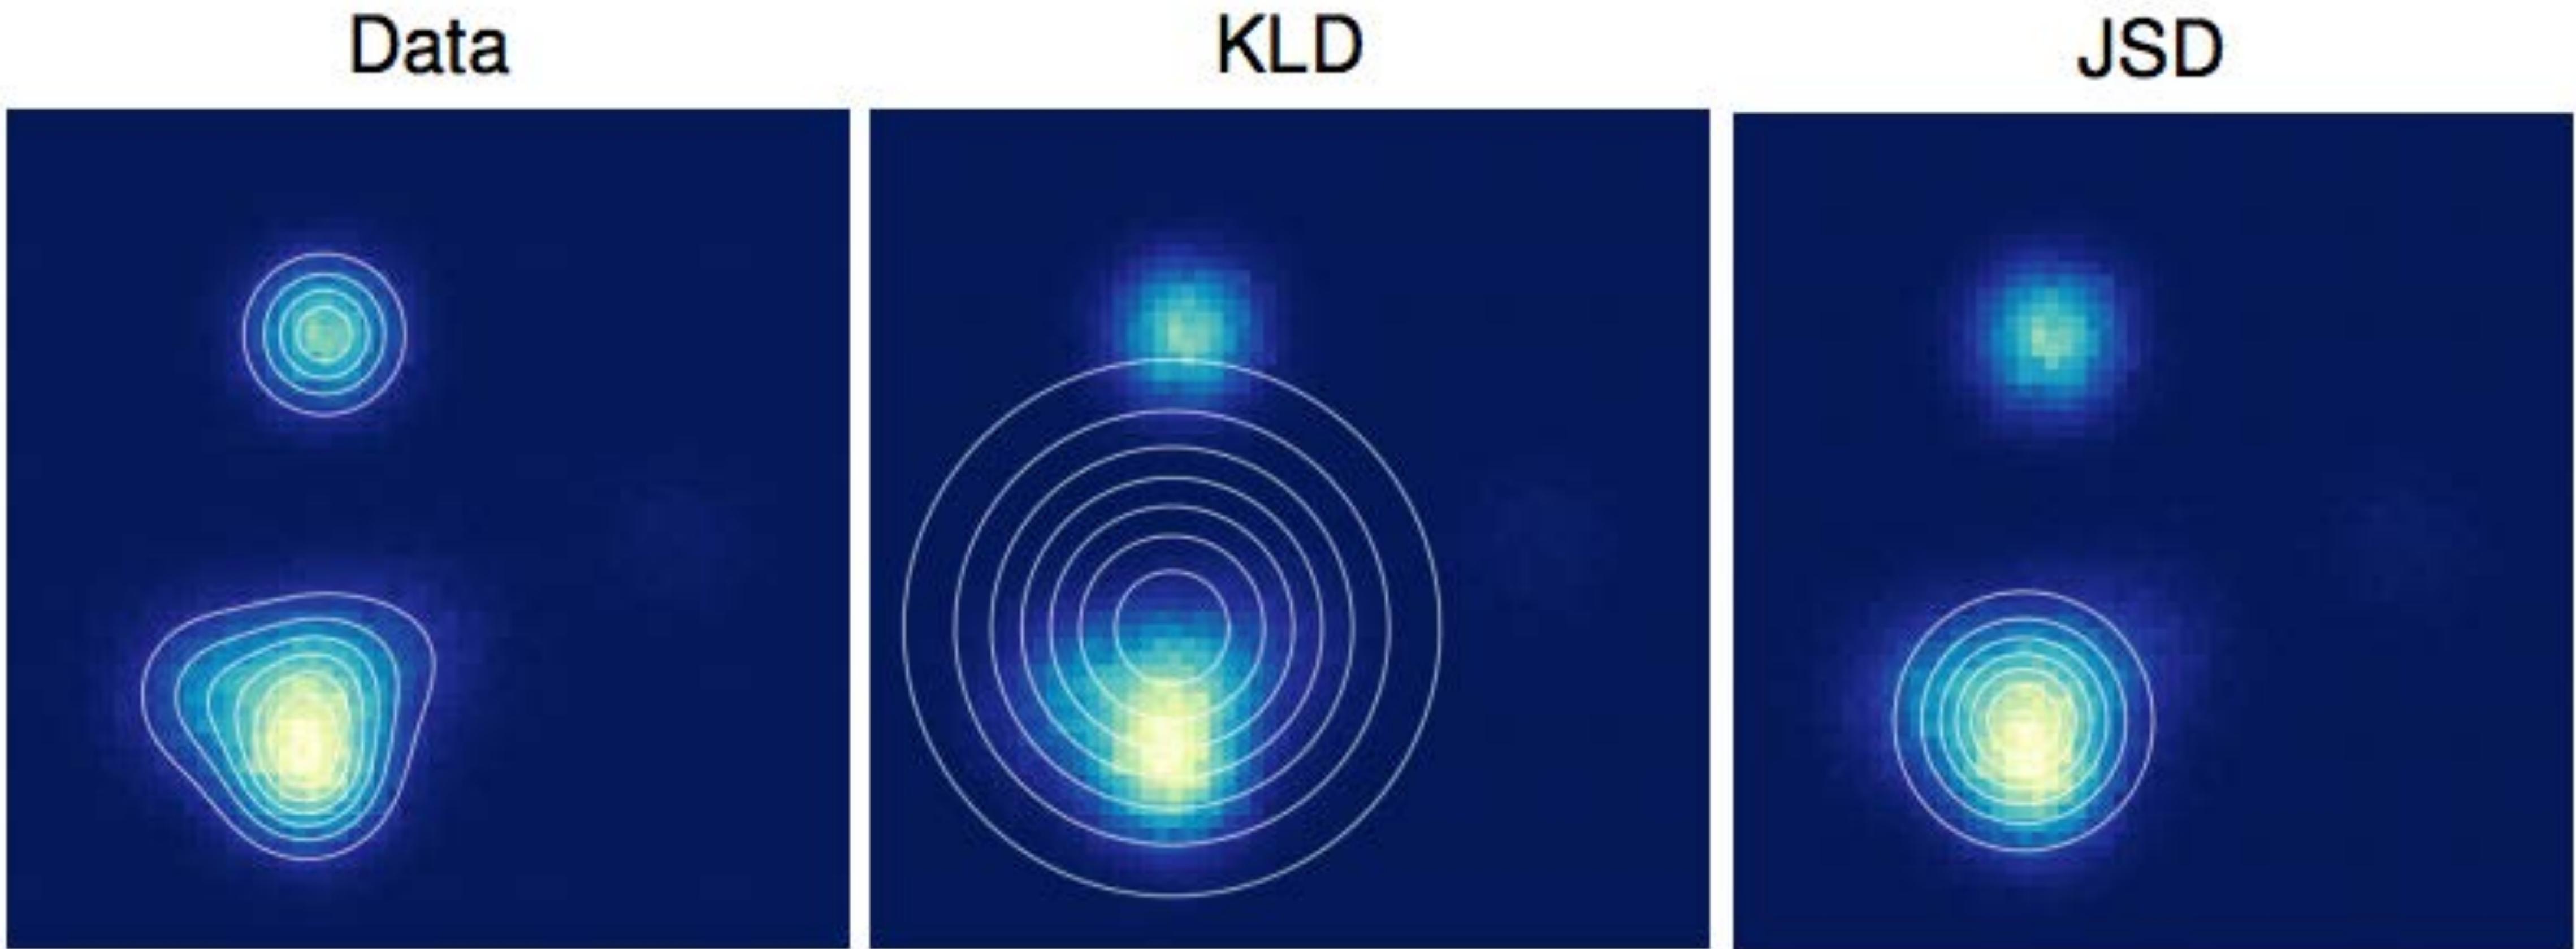
Patch Discriminator



Rather than penalizing if output *image* looks fake, penalize if each overlapping *patch* in output looks fake

- Faster, fewer parameters
- More supervised observations
- Applies to arbitrarily large images

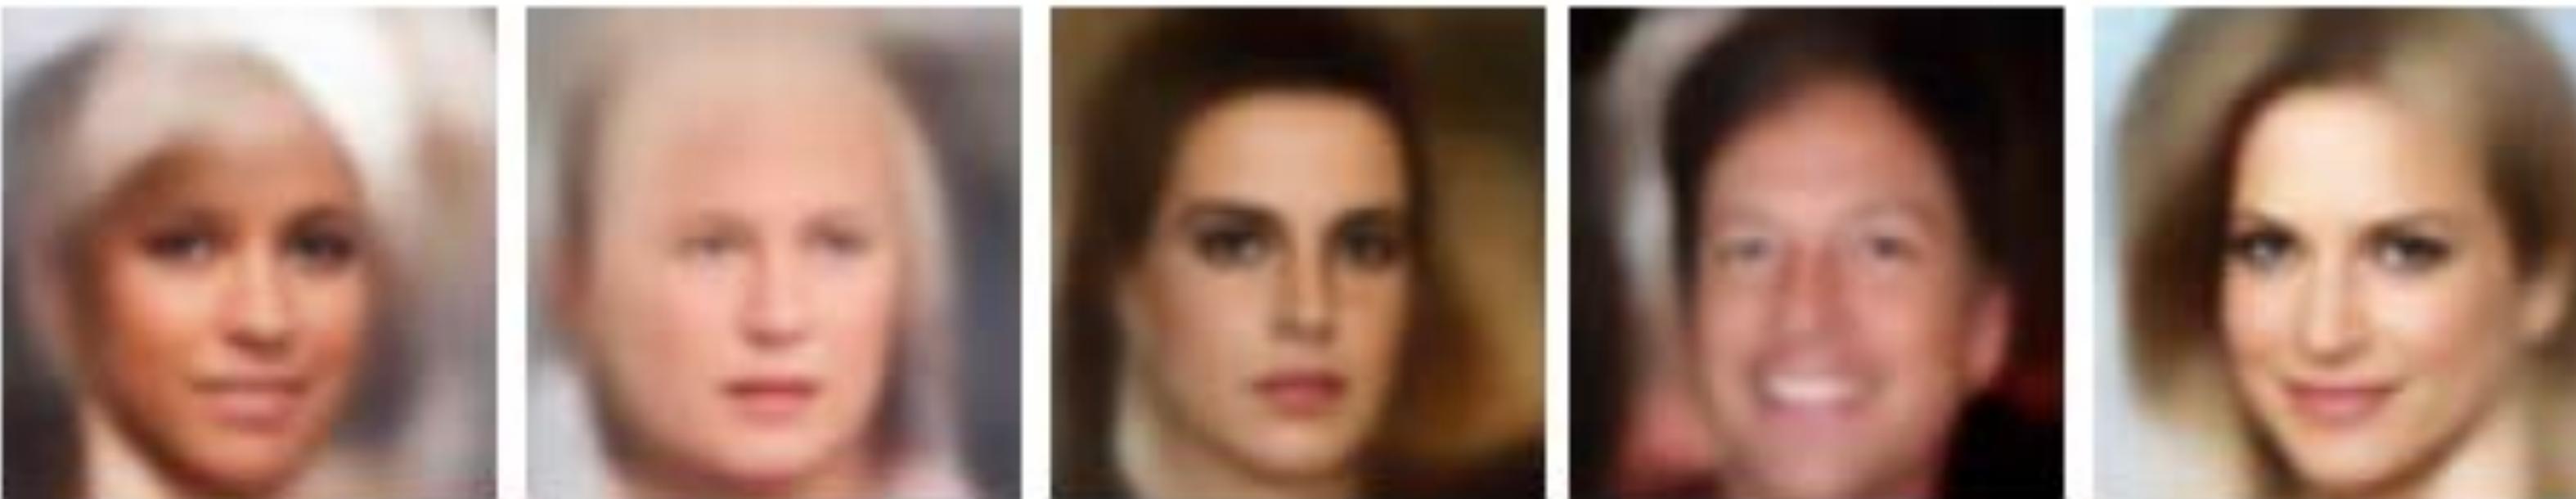
Shrinking the capacity: model misspecification



[Theis et al. 2016]

Mode covering versus mode seeking

VAE

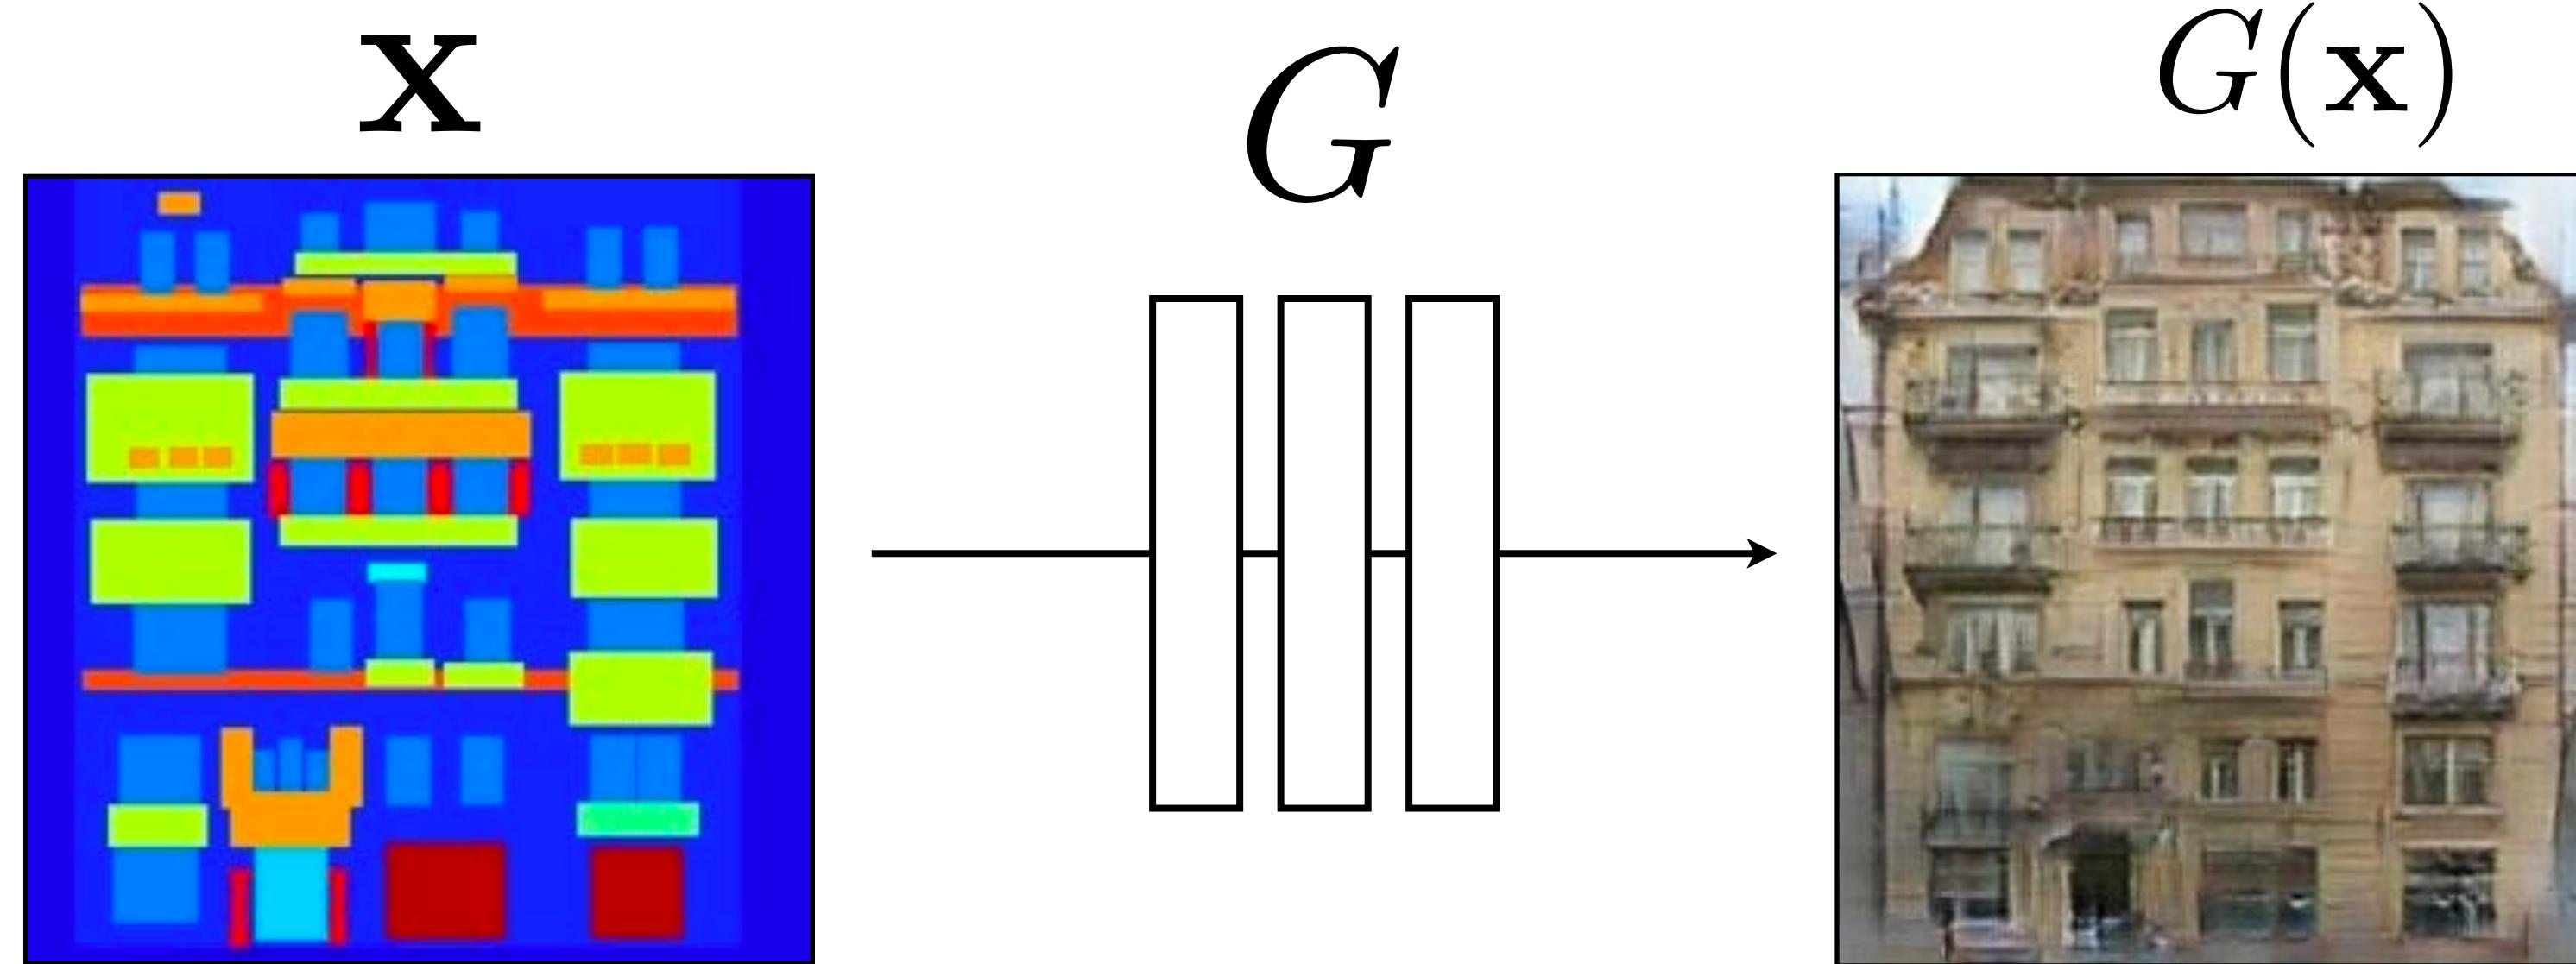


GAN

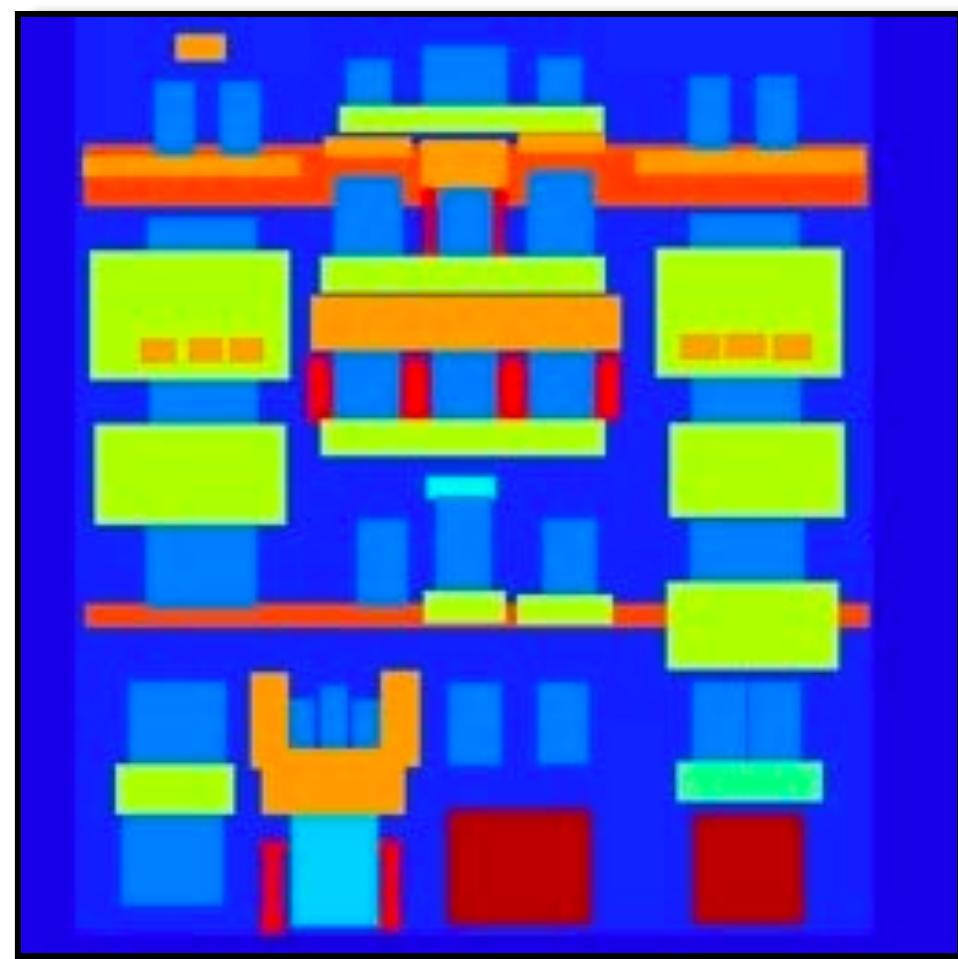


[Larsen et al. 2016]

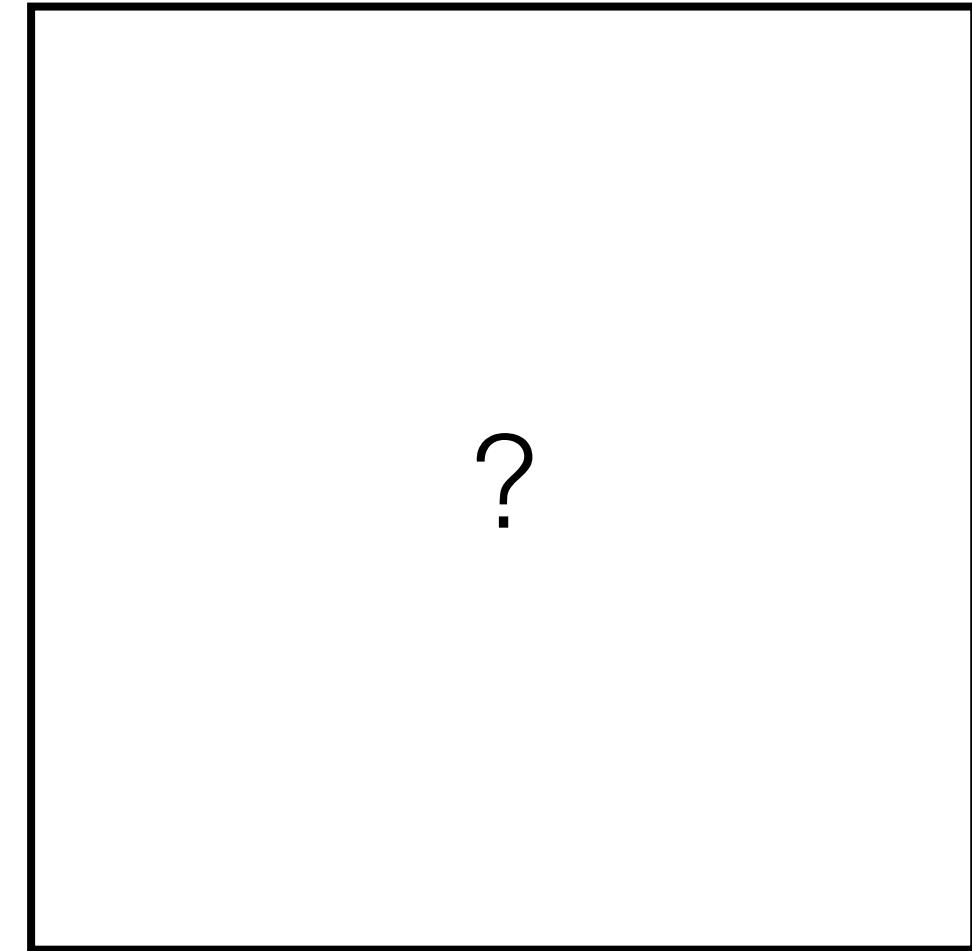
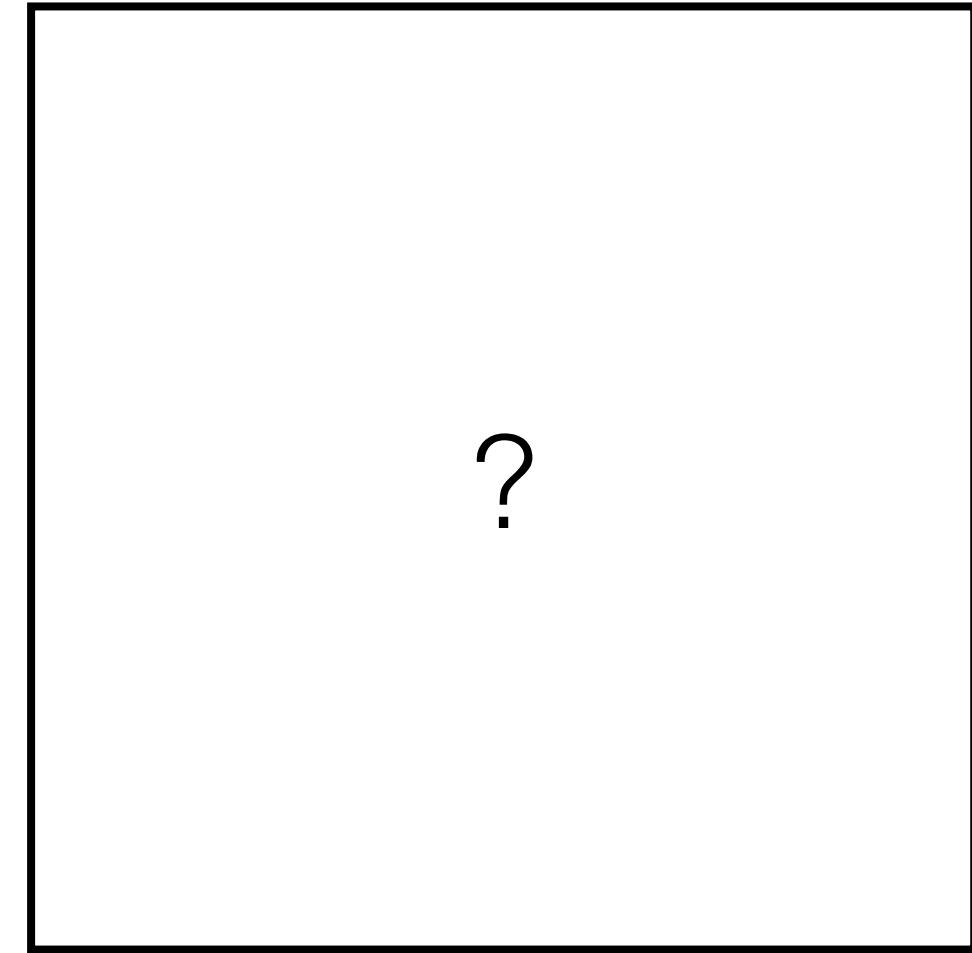
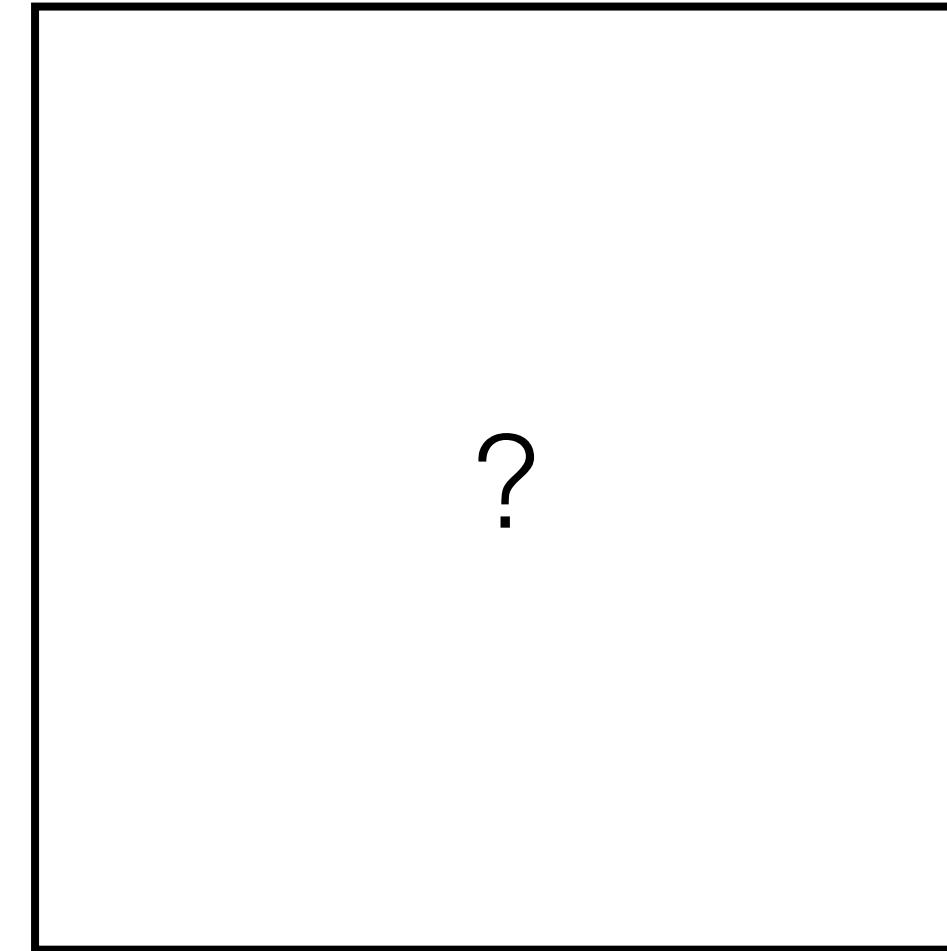
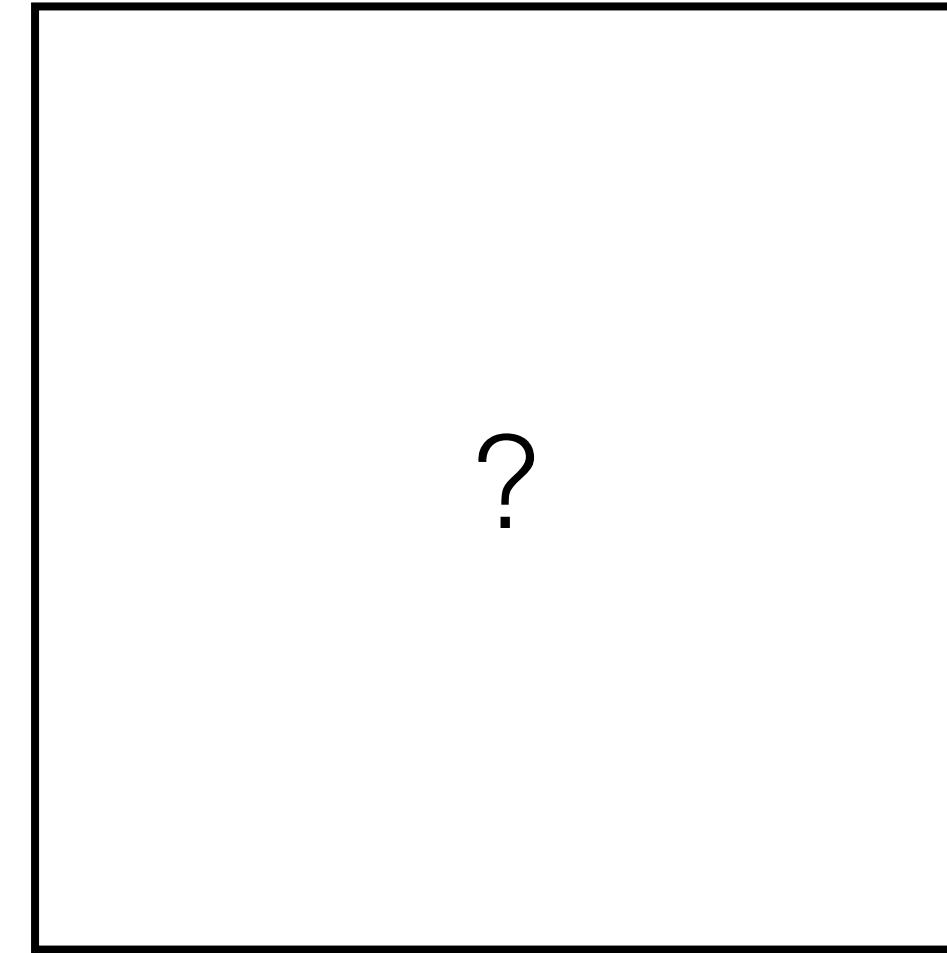
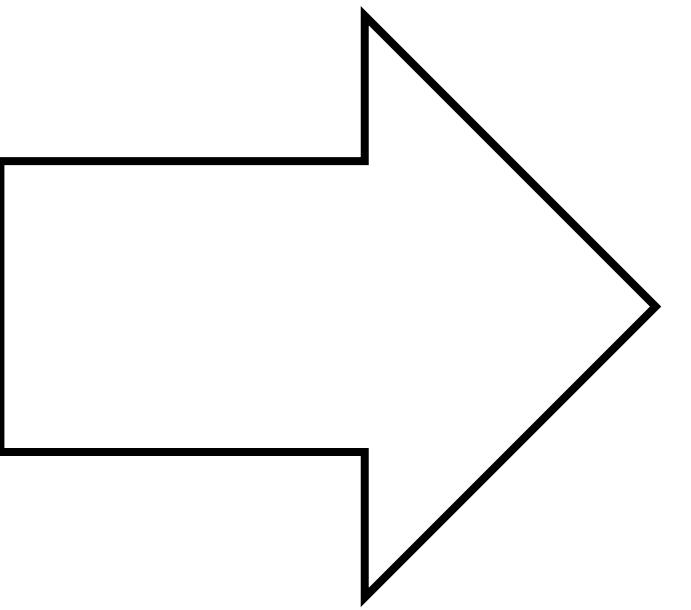
Modeling multiple possible outputs



Modeling multiple possible outputs

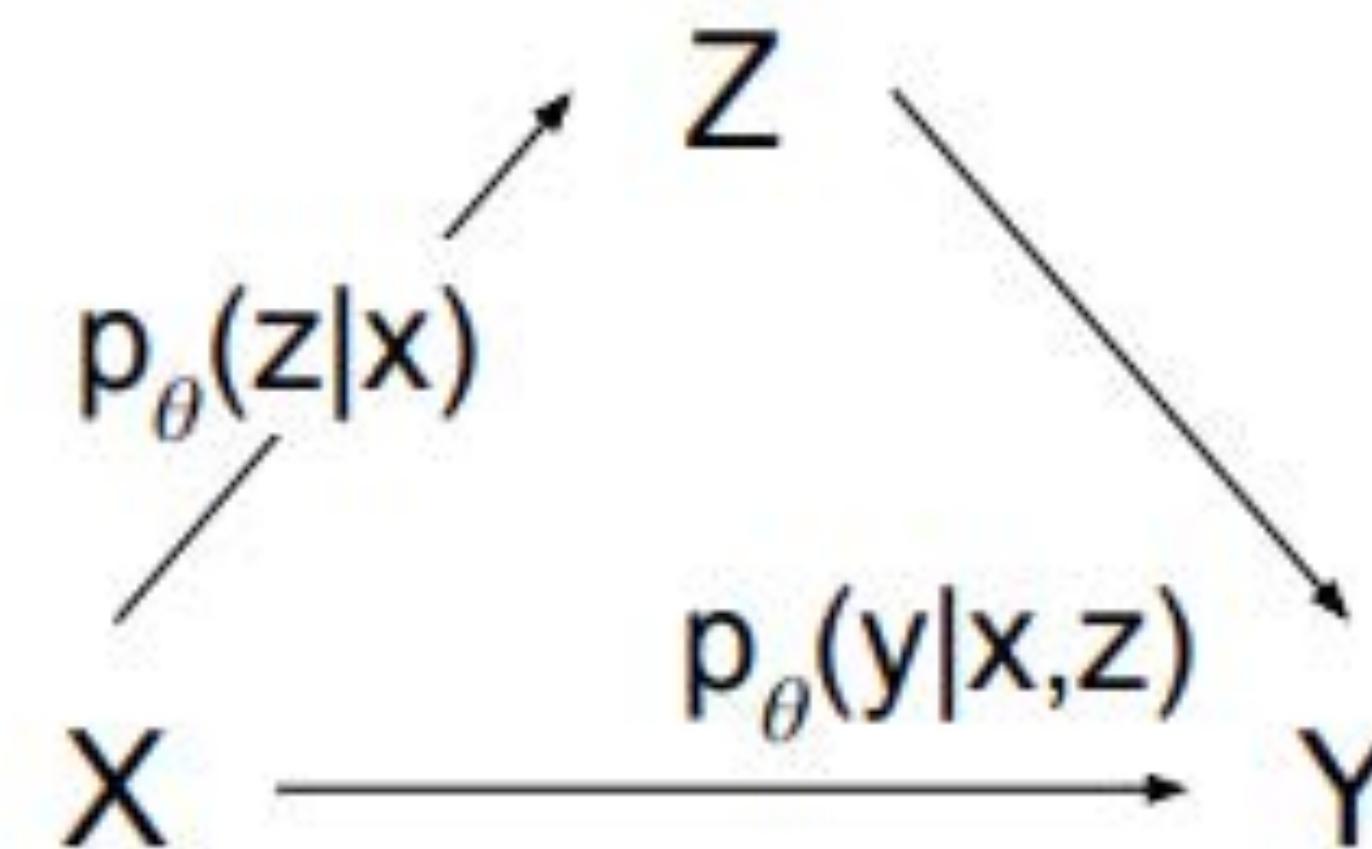


Input



Possible outputs

Conditional Variational Autoencoders



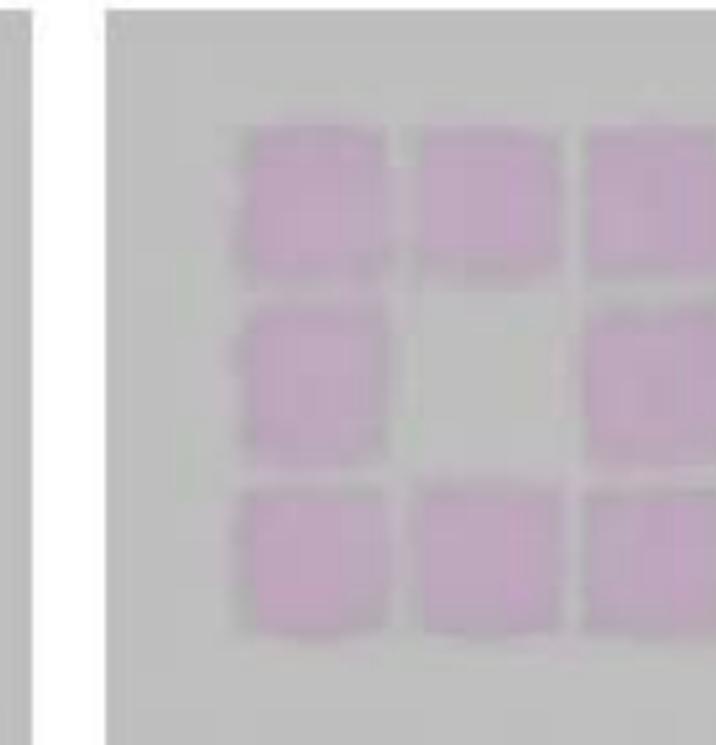
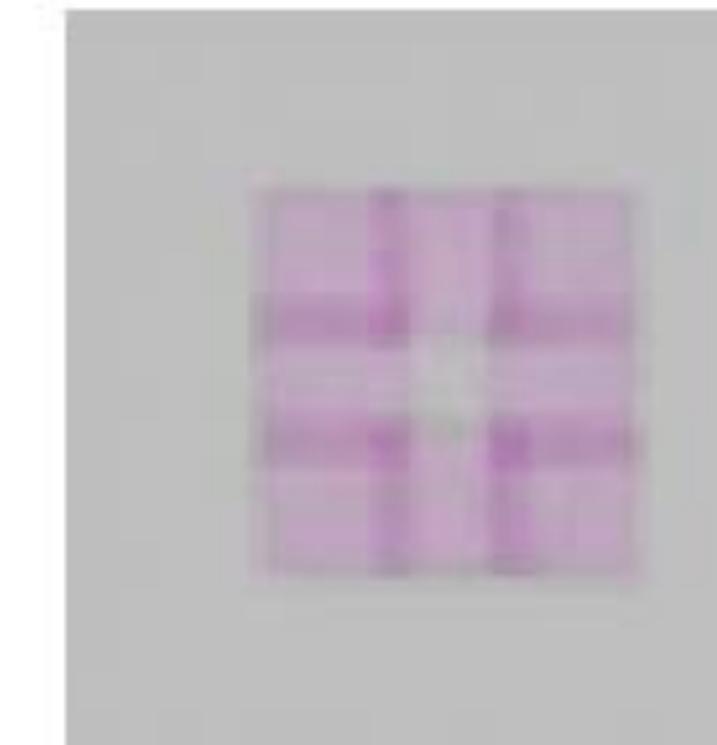
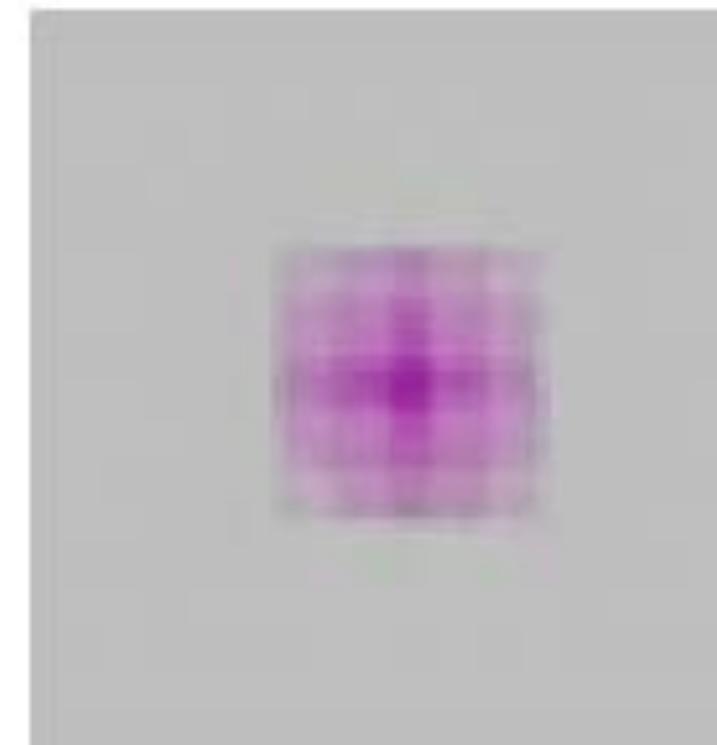
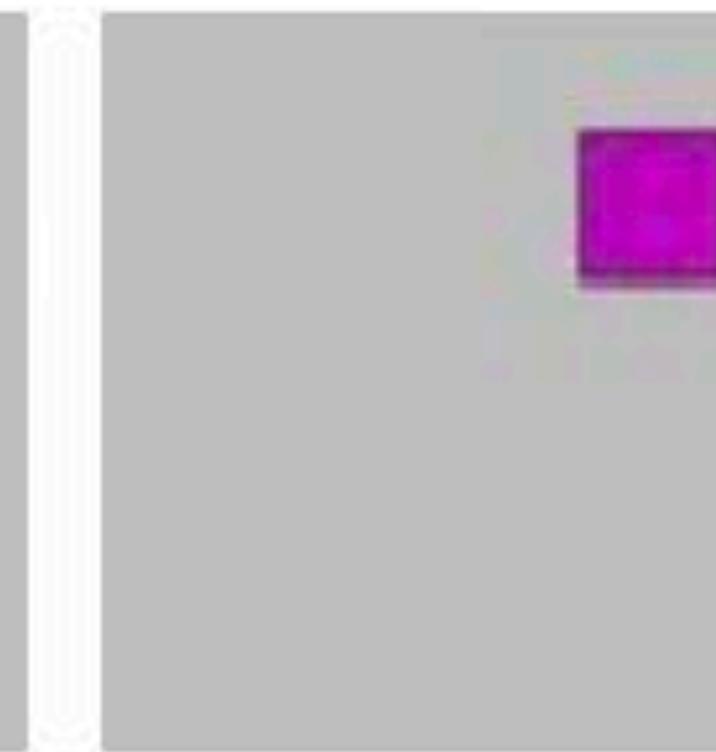
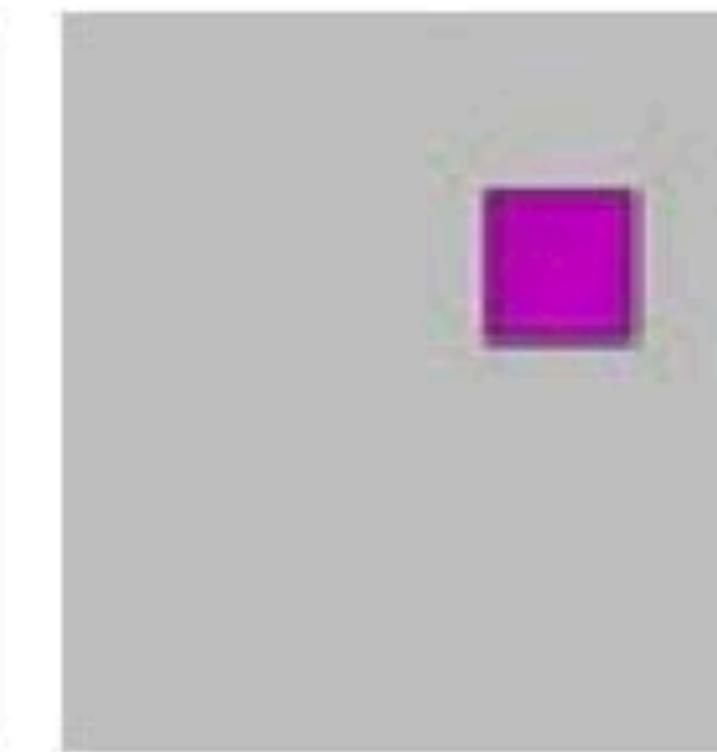
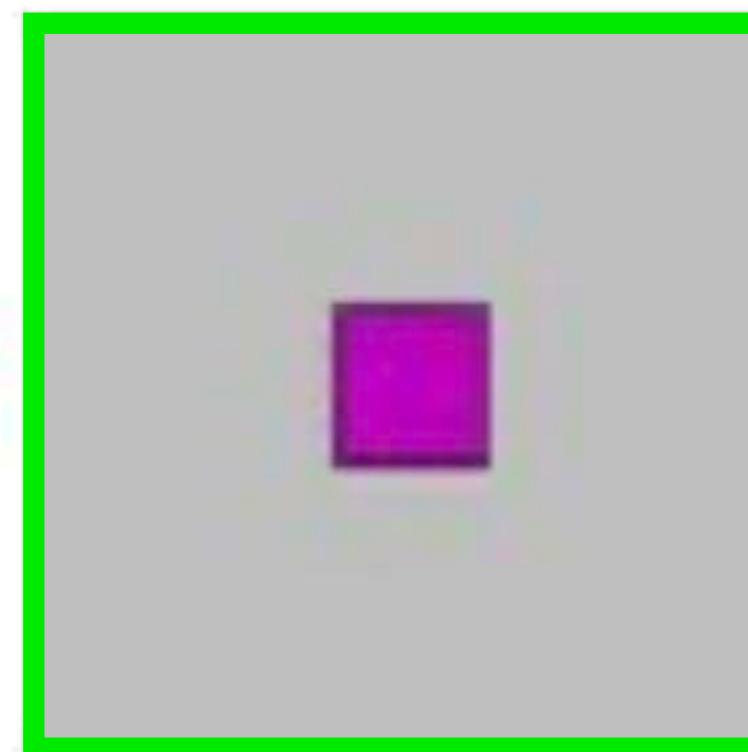
$$\log p_\theta(\mathbf{y}|\mathbf{x}) \geq -KL(q_\phi(\mathbf{z}|\mathbf{x}, \mathbf{y})\|p_\theta(\mathbf{z}|\mathbf{x})) + \mathbb{E}_{q_\phi(\mathbf{z}|\mathbf{x}, \mathbf{y})} [\log p_\theta(\mathbf{y}|\mathbf{x}, \mathbf{z})]$$

[Sohn, Yan, Lee, NIPS 2015]

Context
Frame

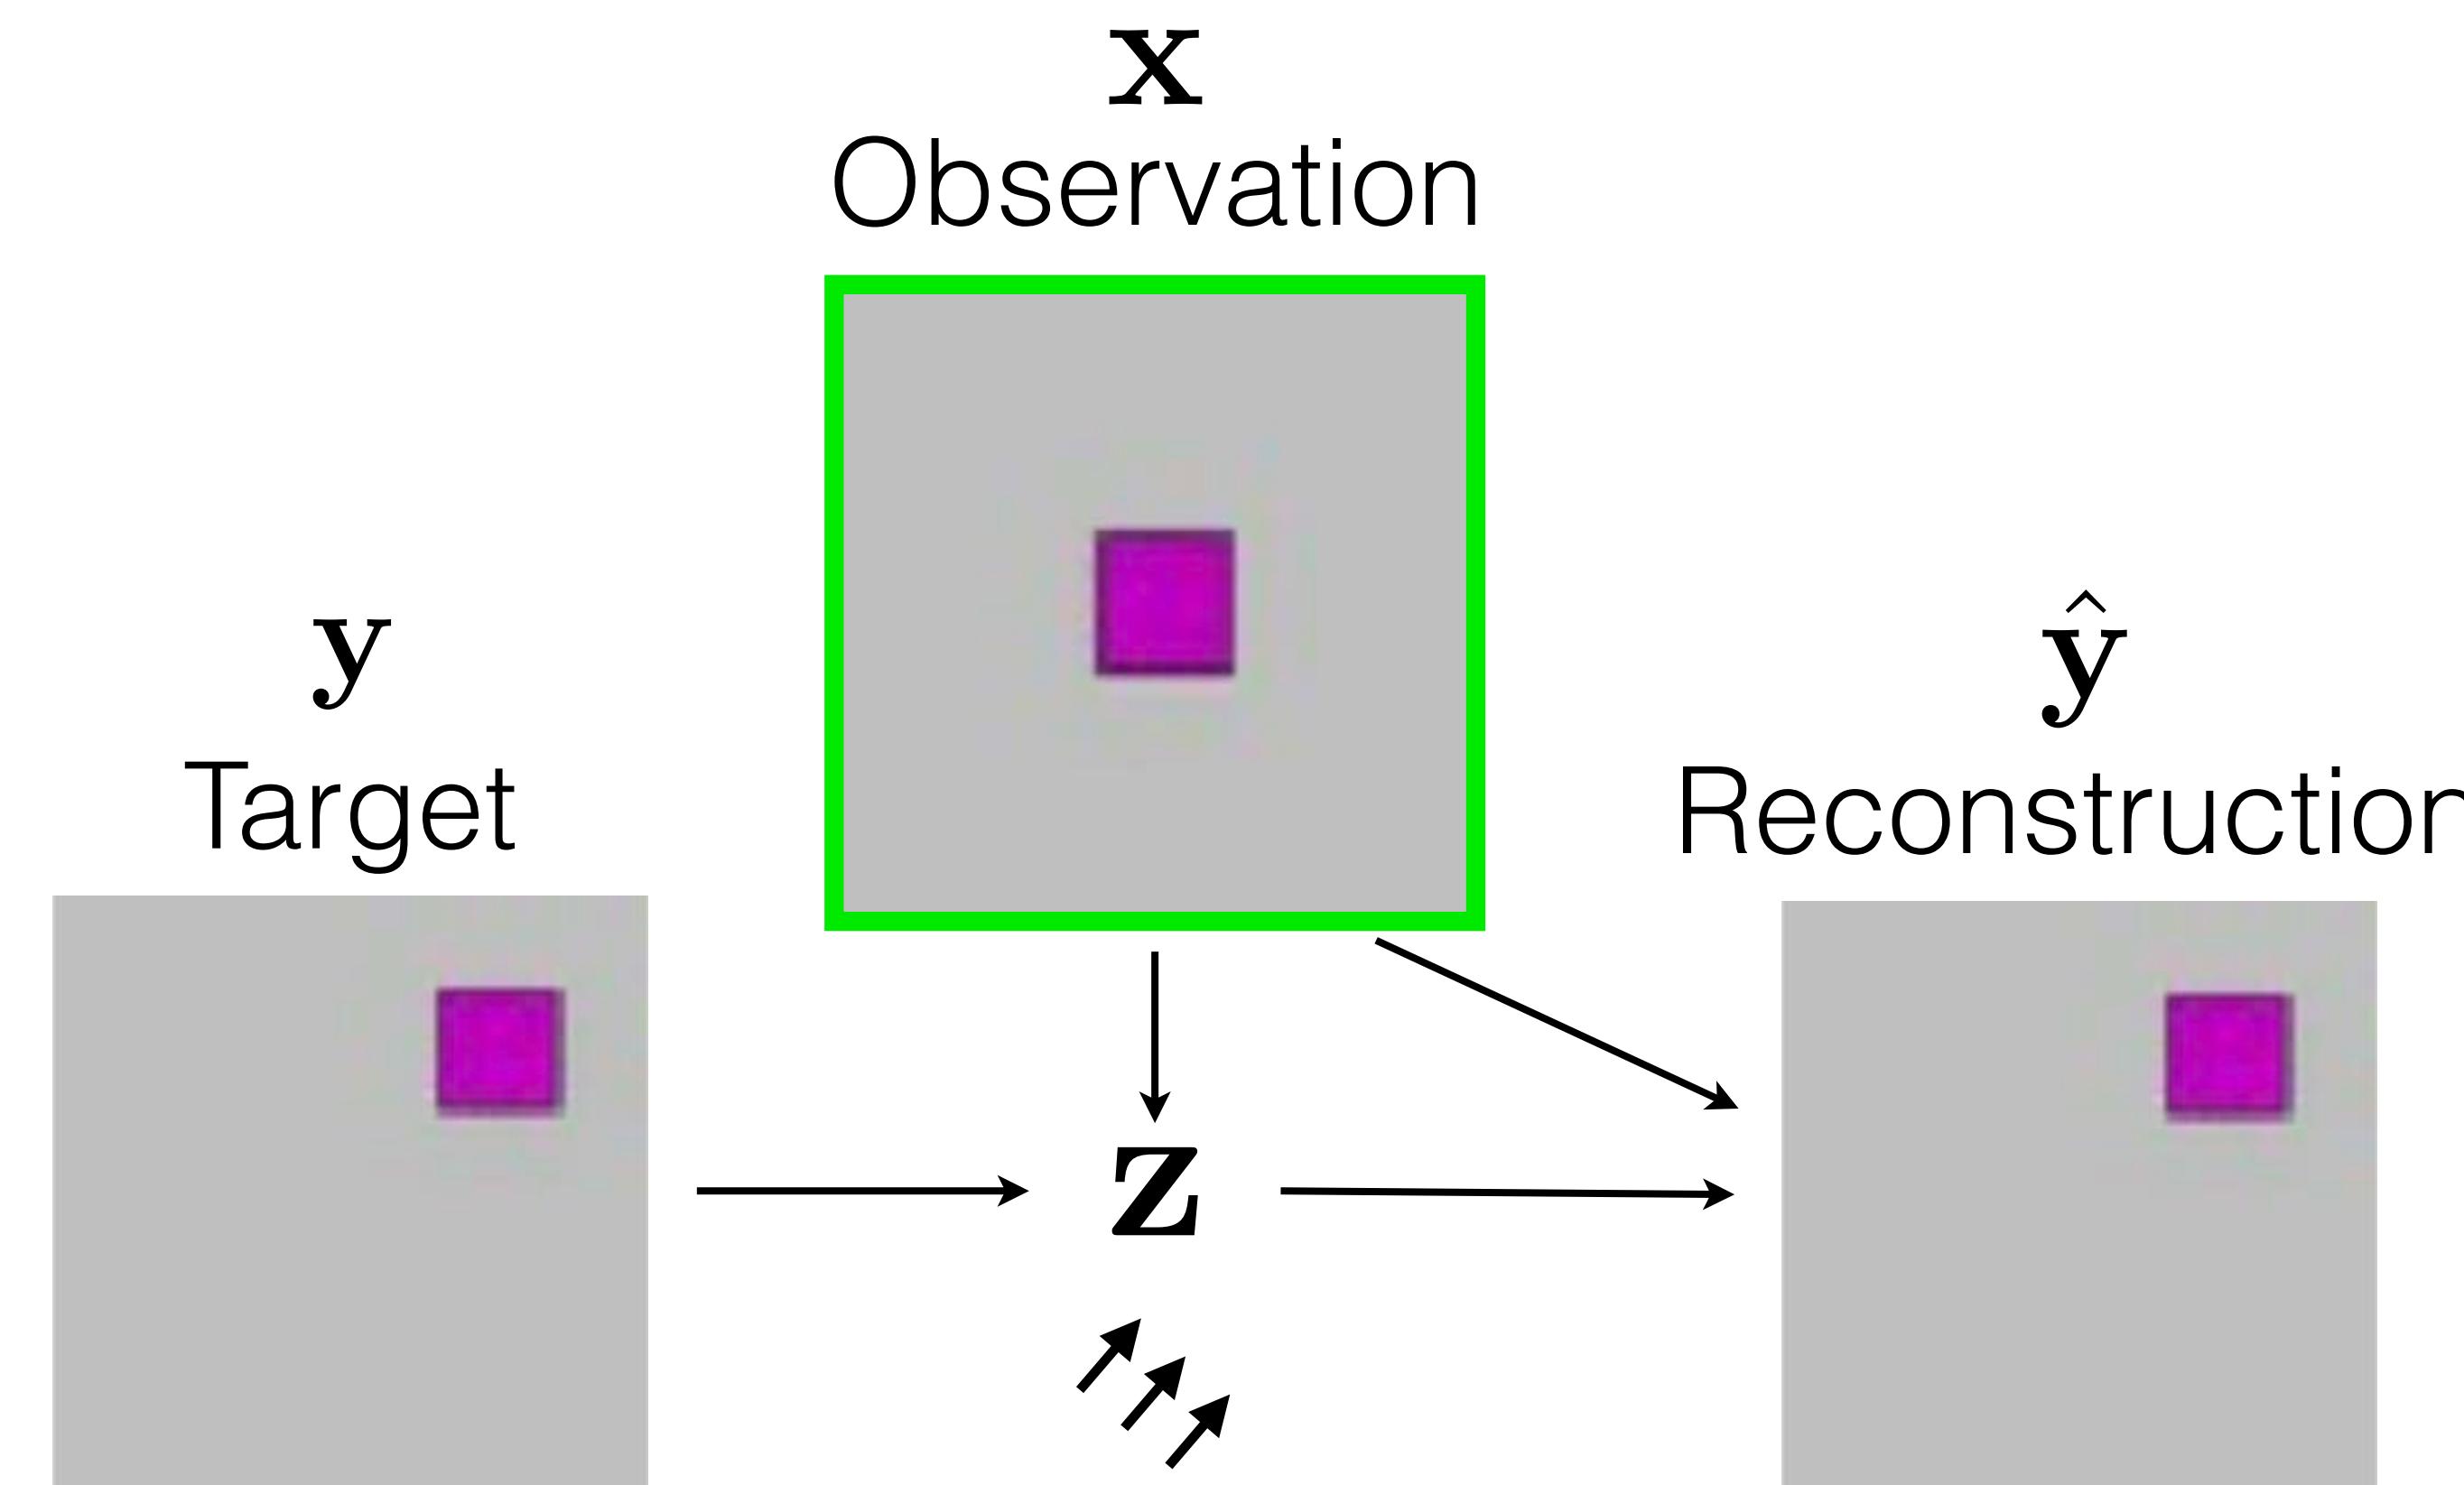
Time

GT

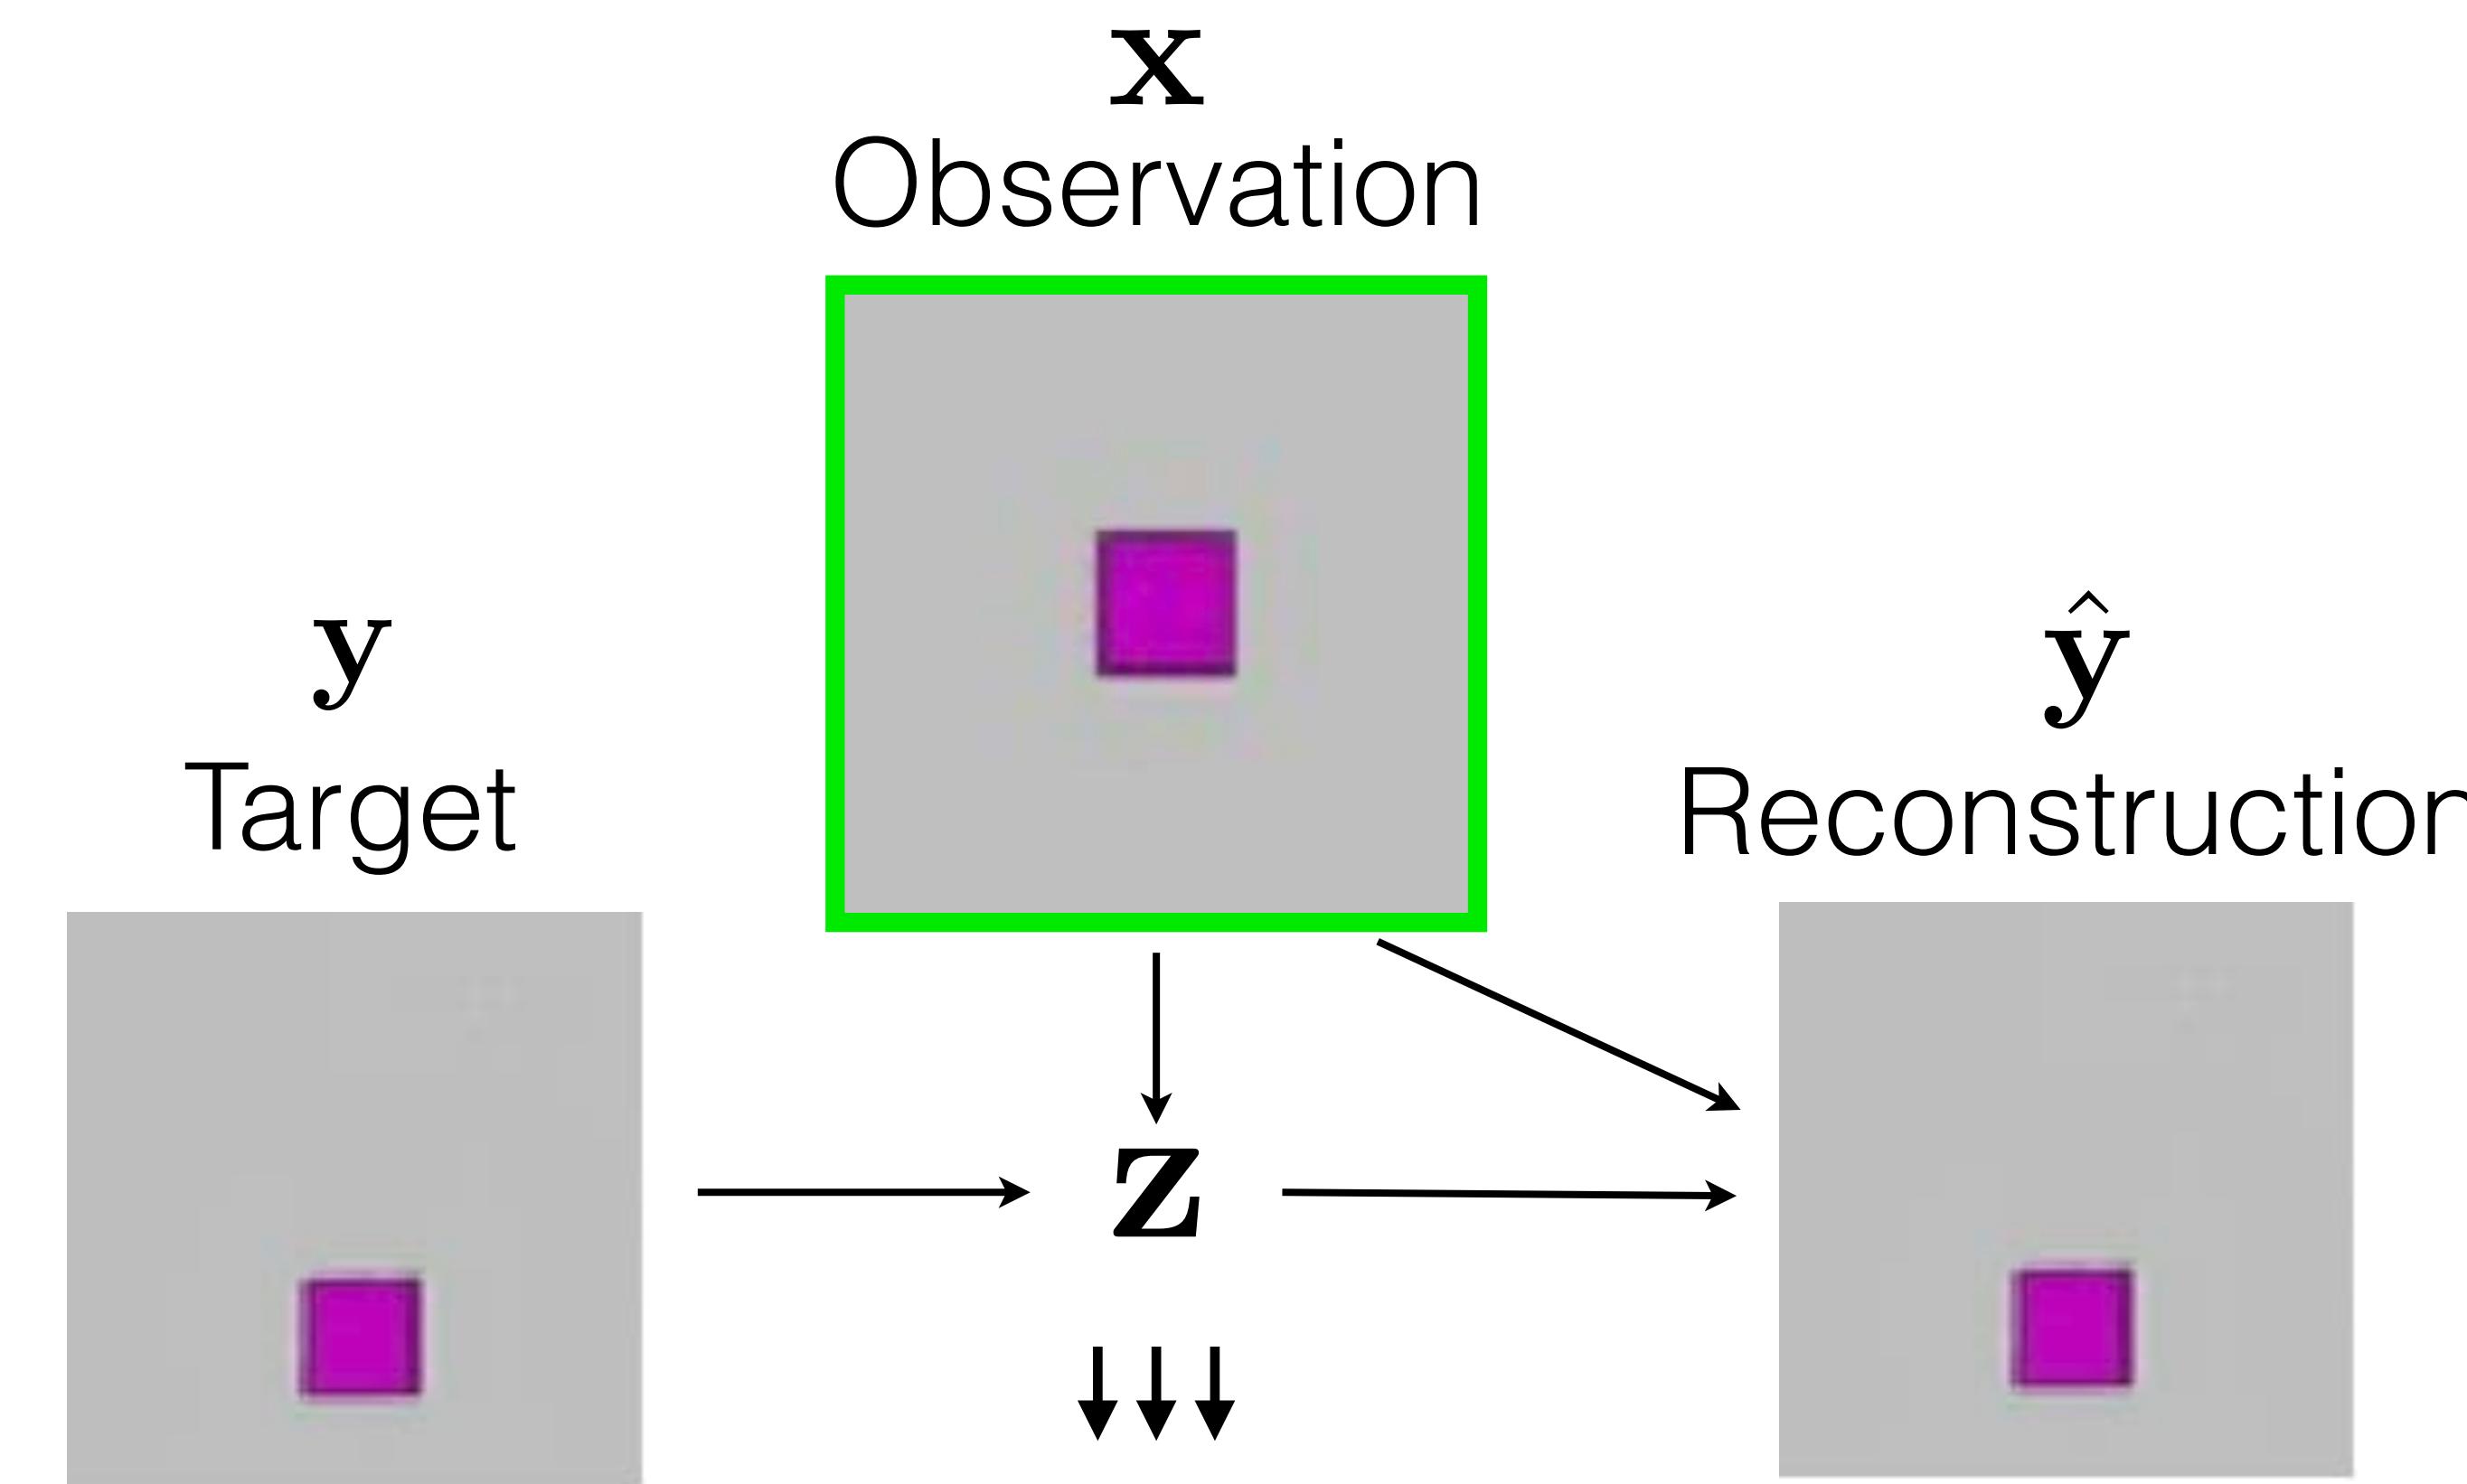


Finn et al.
(2016)

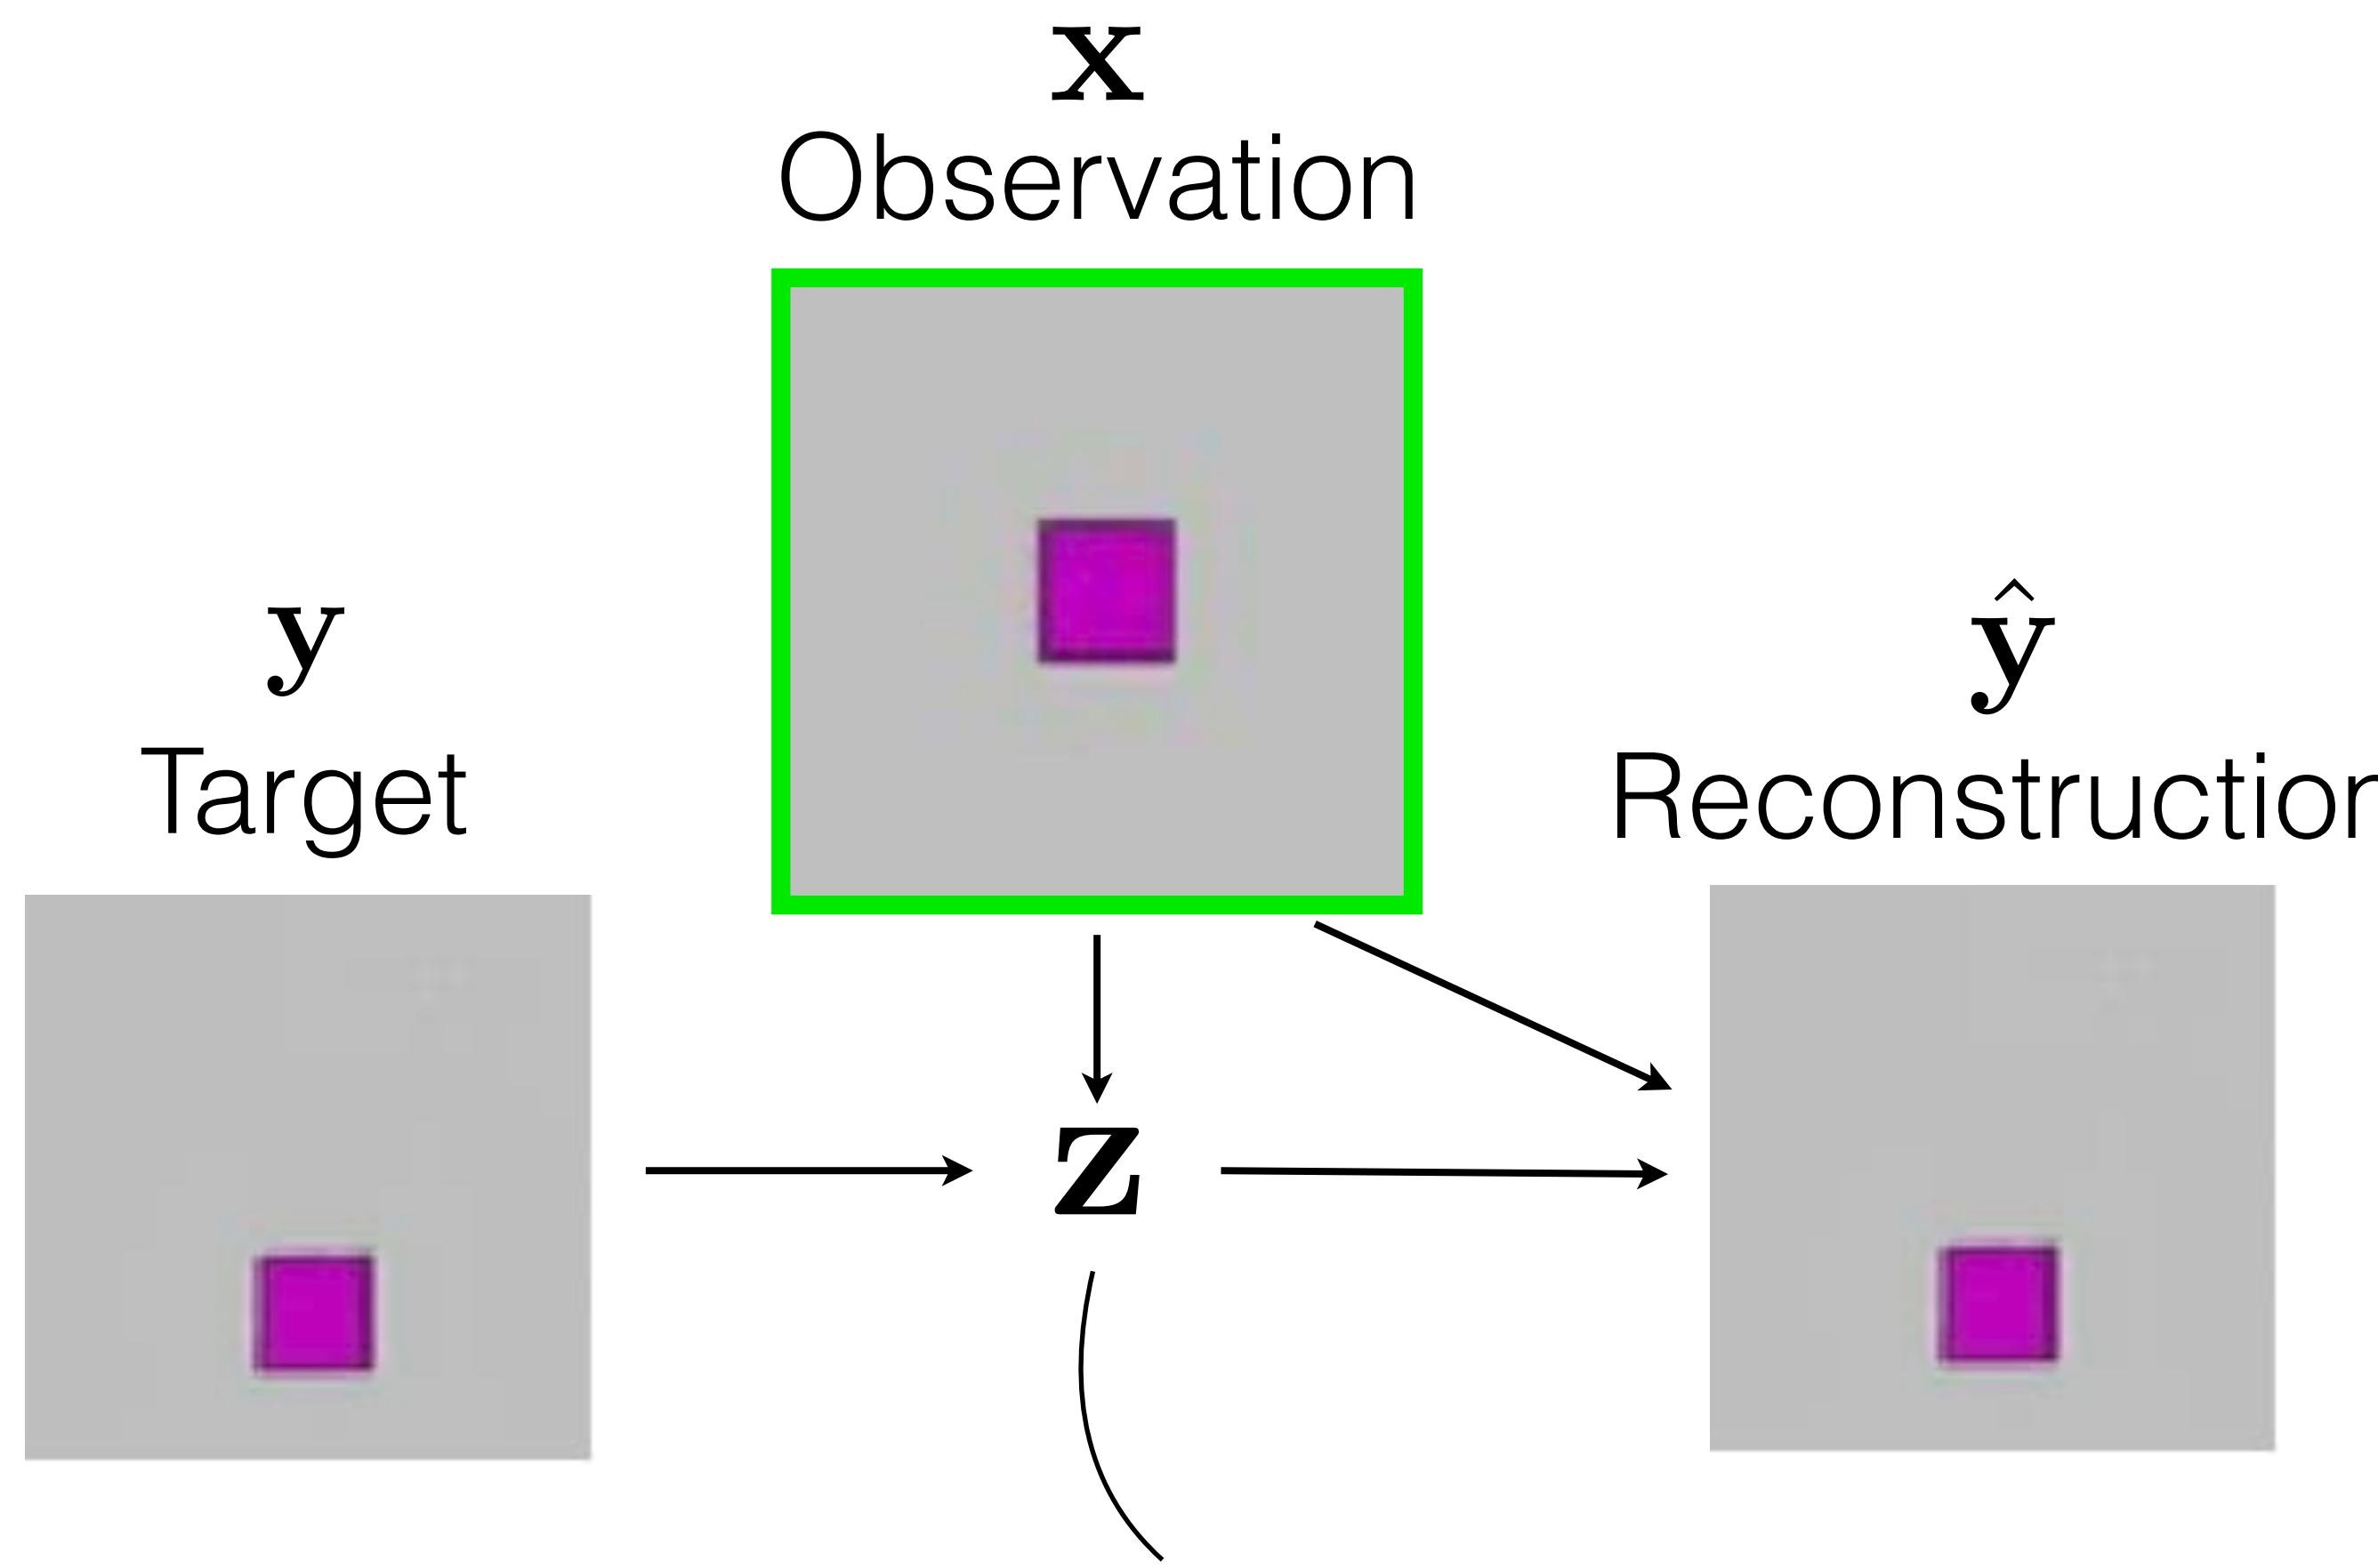
Example from [Babaeizadeh et al., ICLR 2018]
see also [Walker et al., ECCV 2016], [Xue*, Wu*, et al., NIPS 2016]



Example from [Babaeizadeh et al., ICLR 2018]
see also [Walker et al., ECCV 2016], [Xue*, Wu*, et al., NIPS 2016]

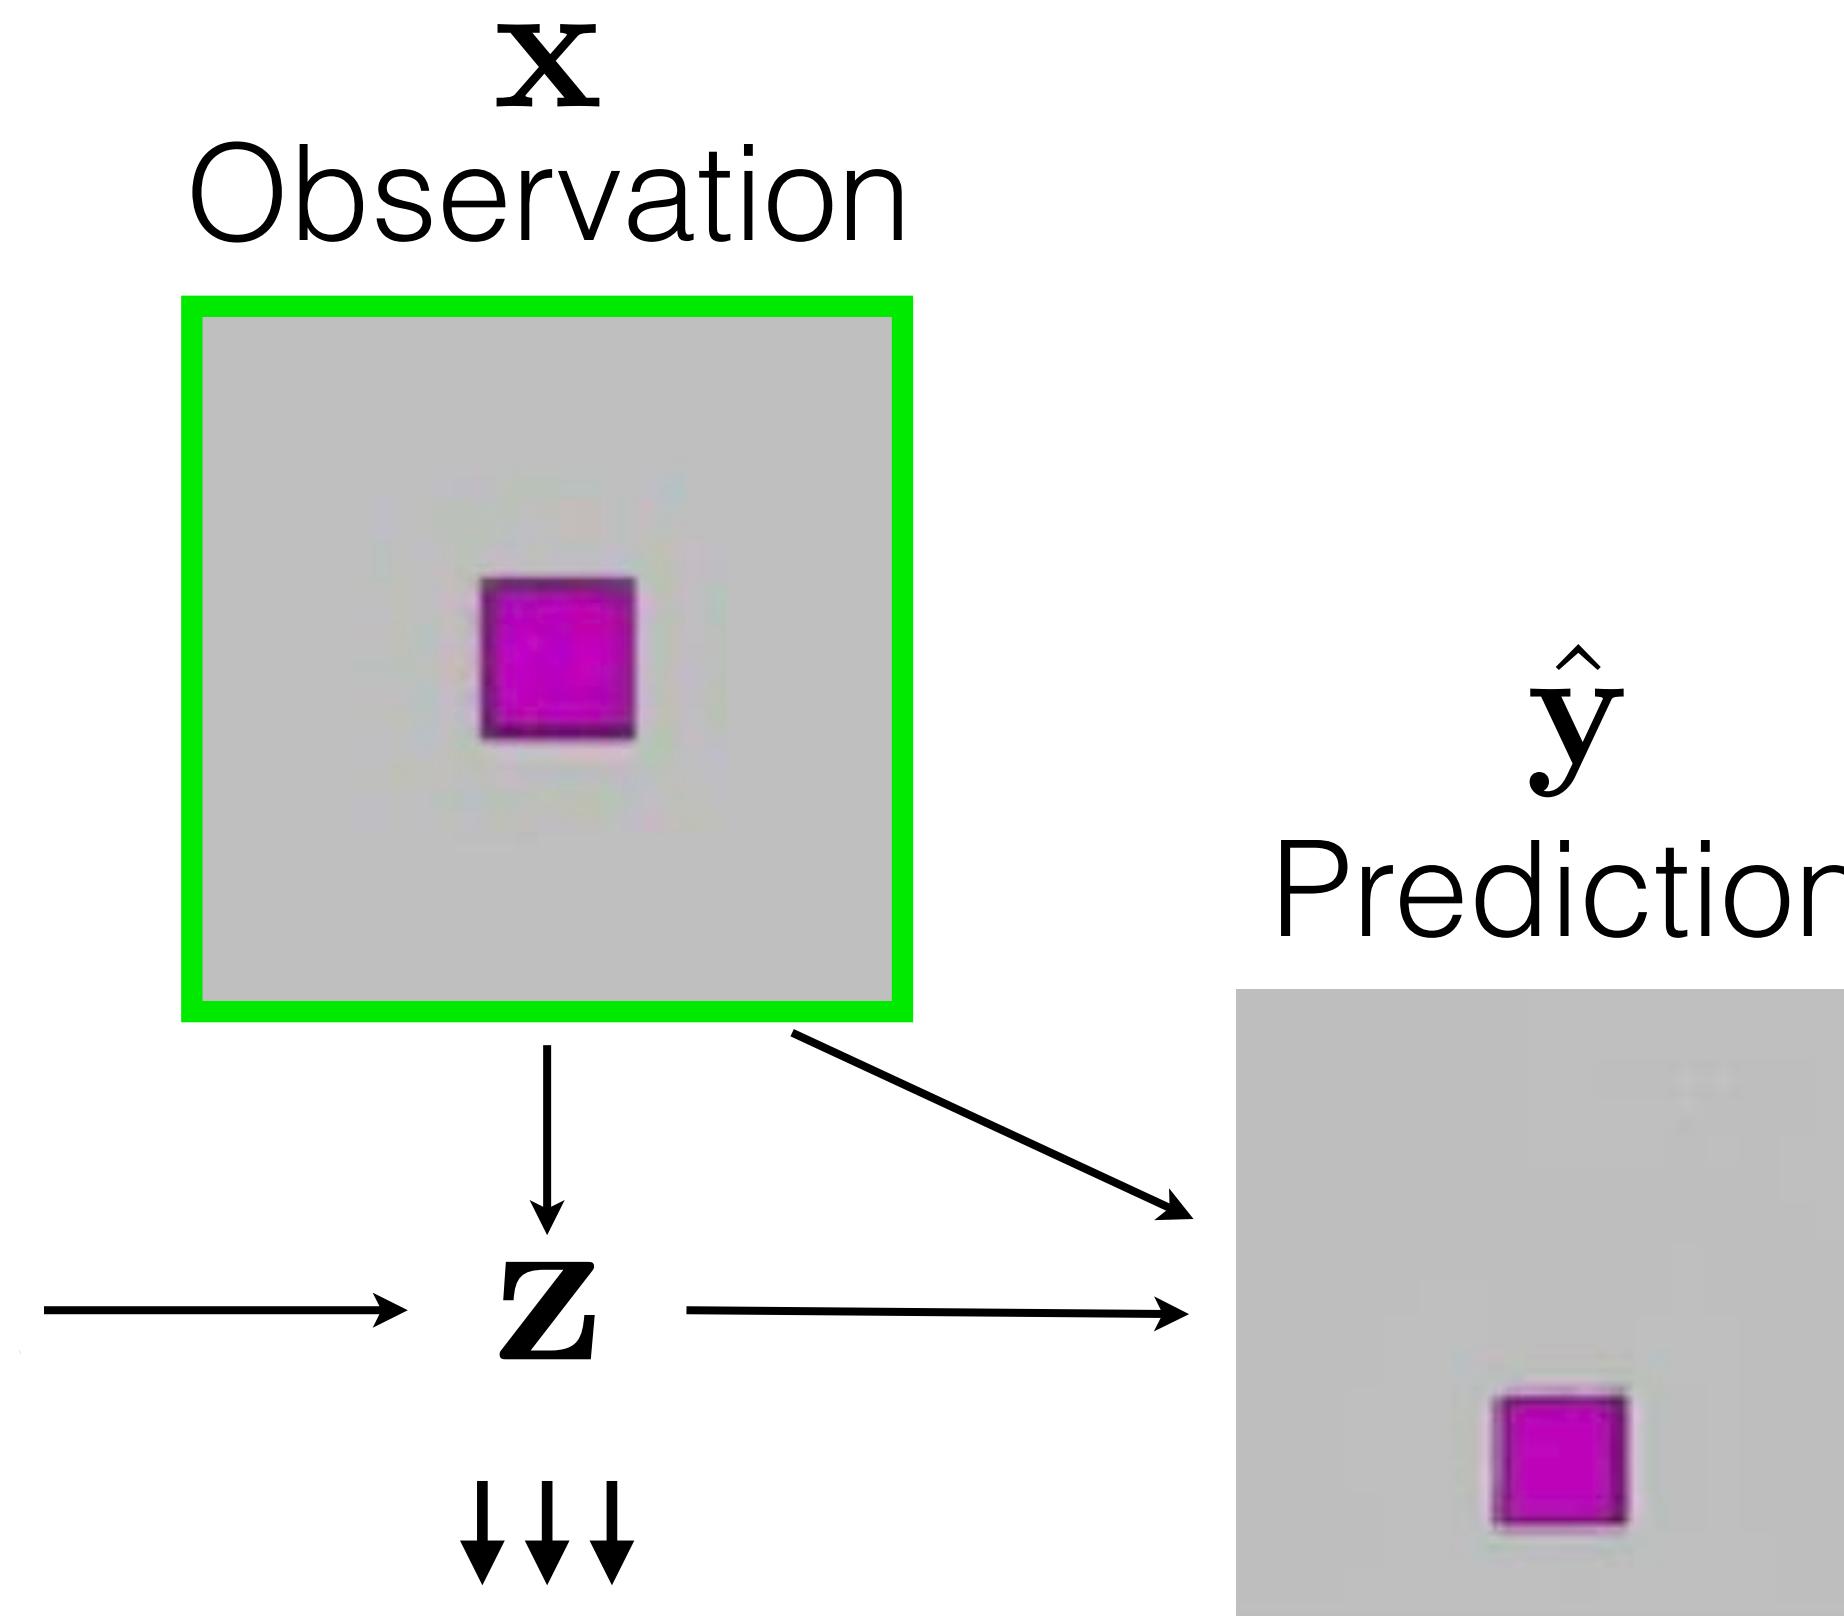
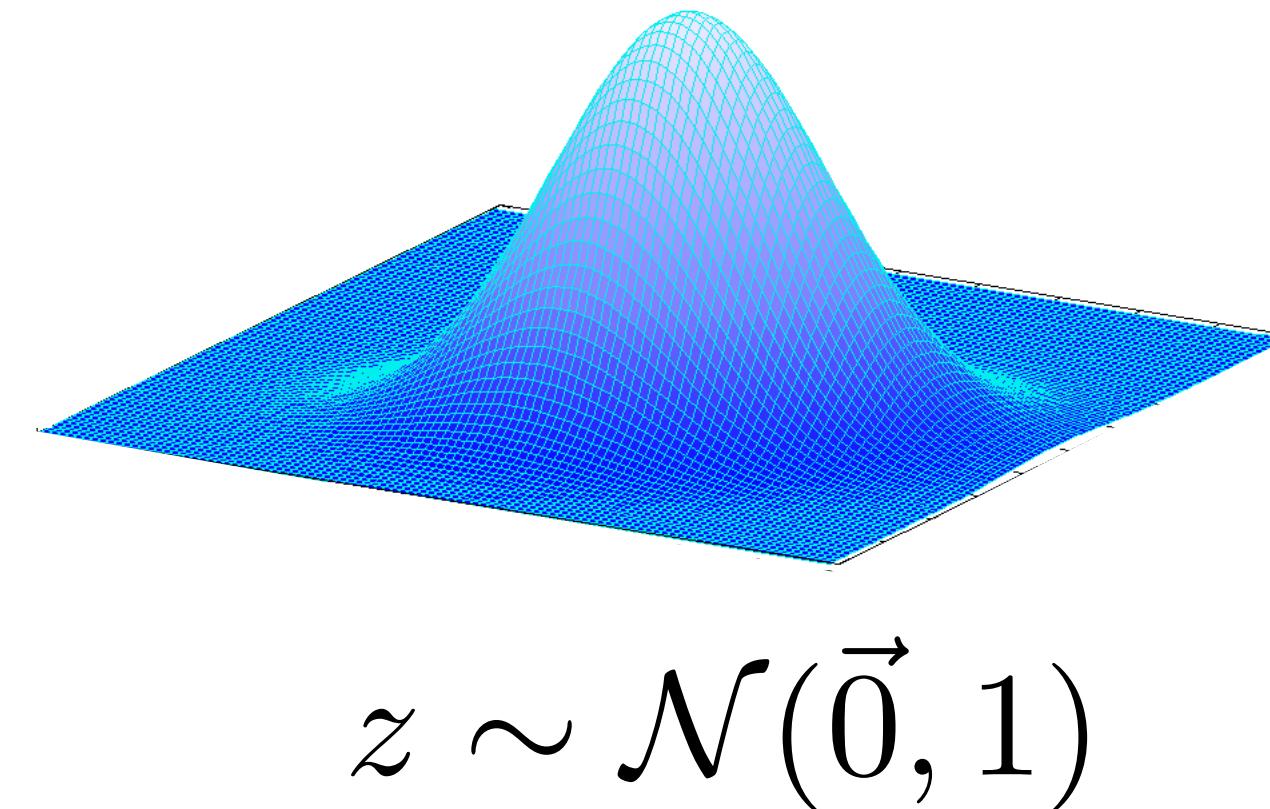


Example from [Babaeizadeh et al., ICLR 2018]
see also [Walker et al., ECCV 2016], [Xue*, Wu*, et al., NIPS 2016]

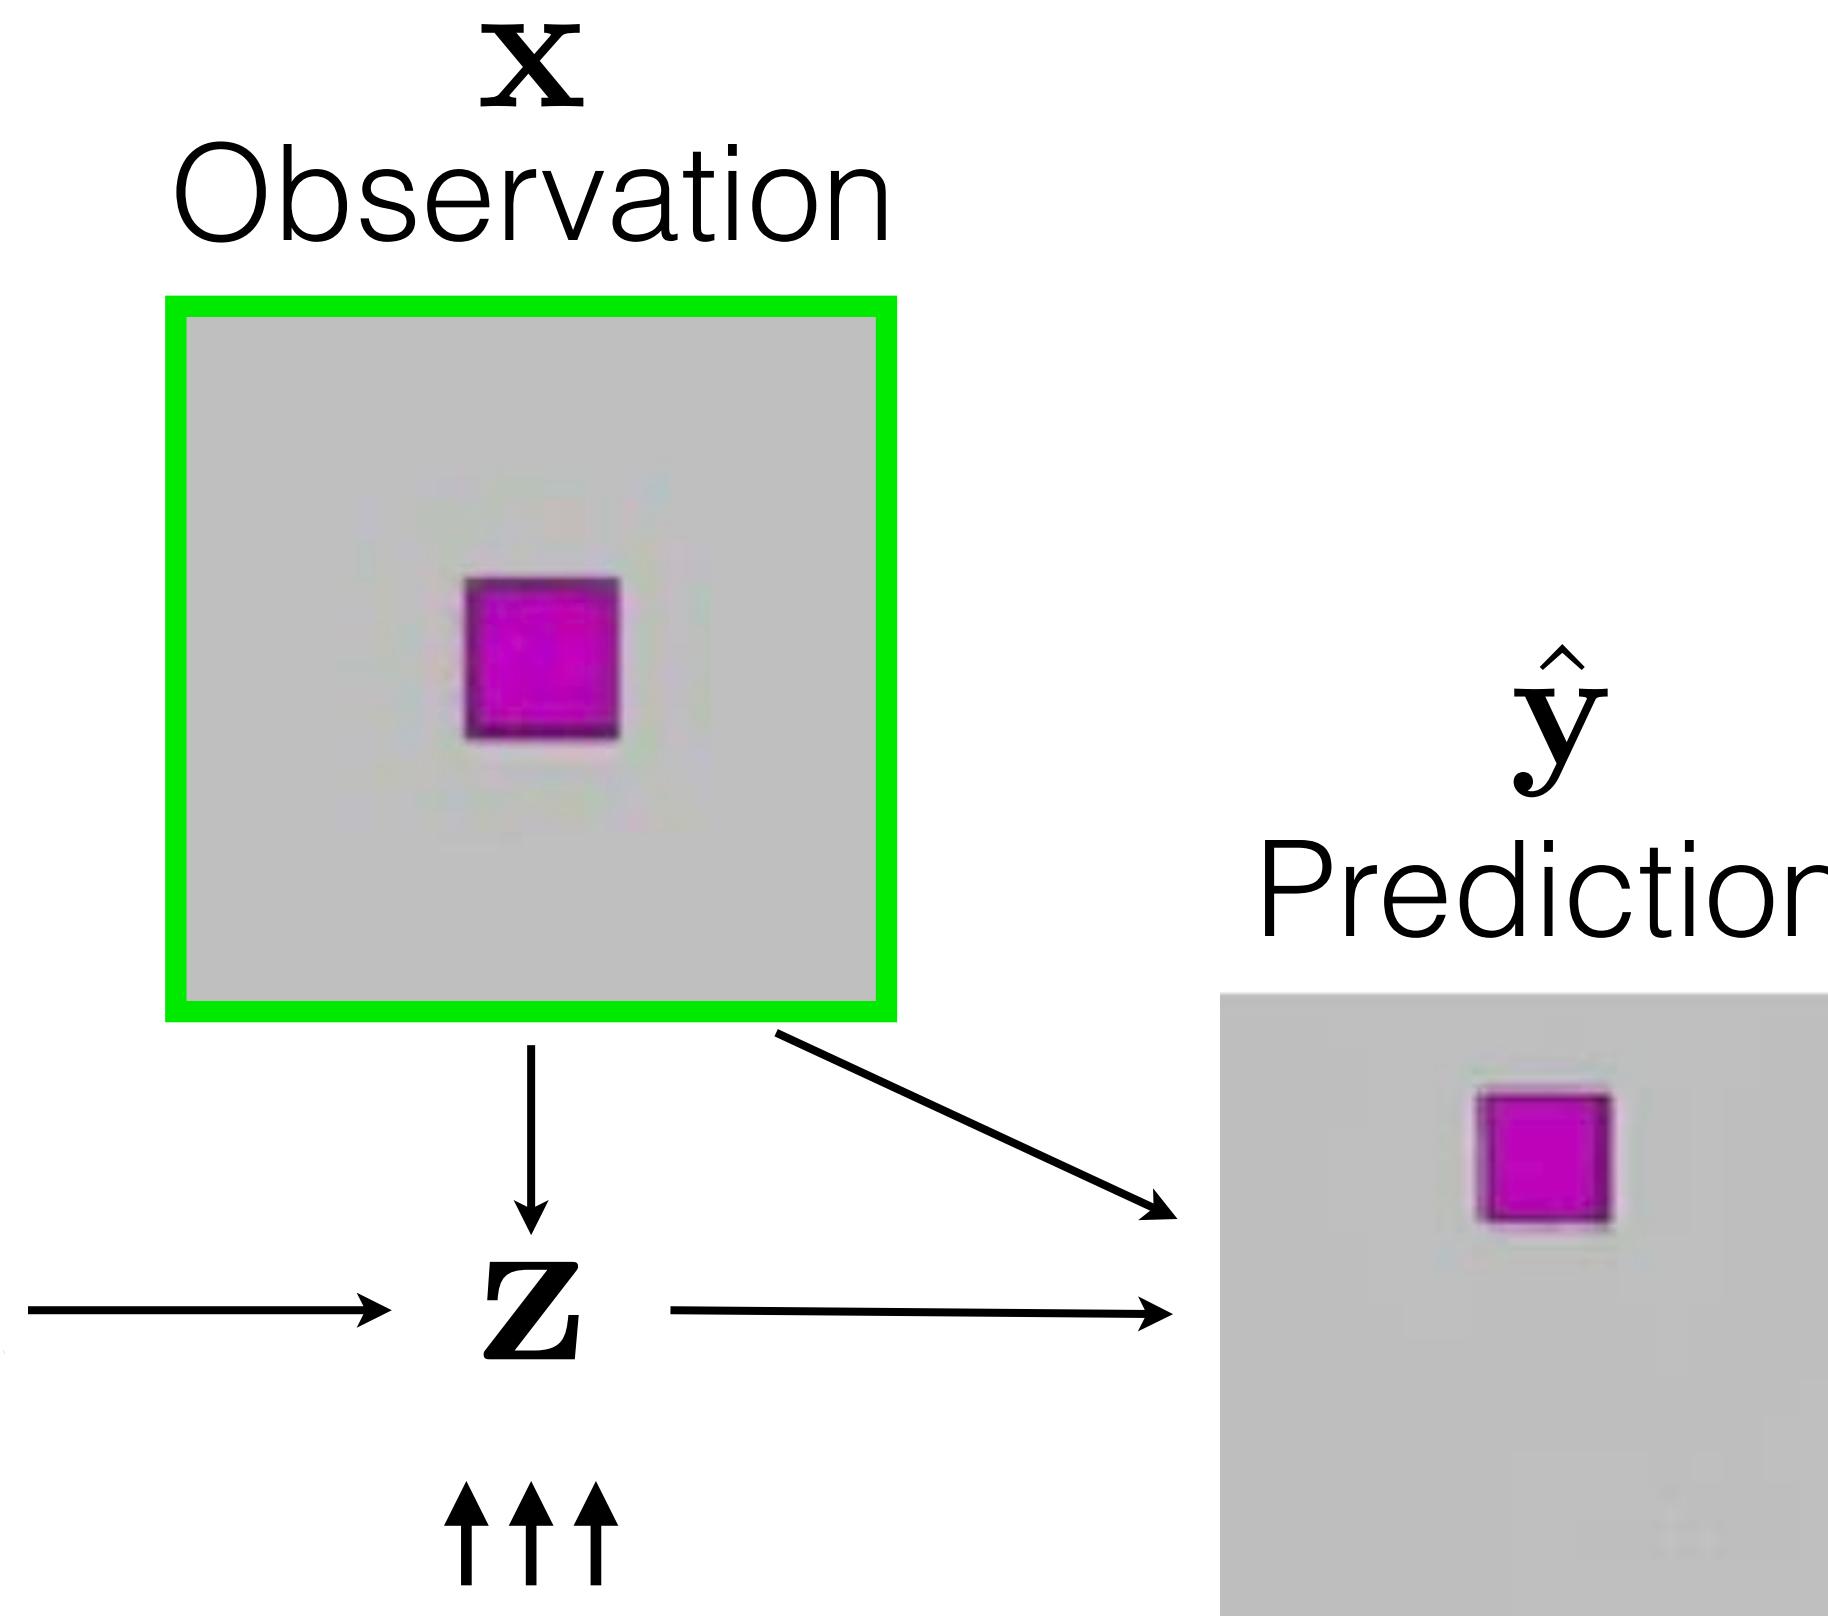
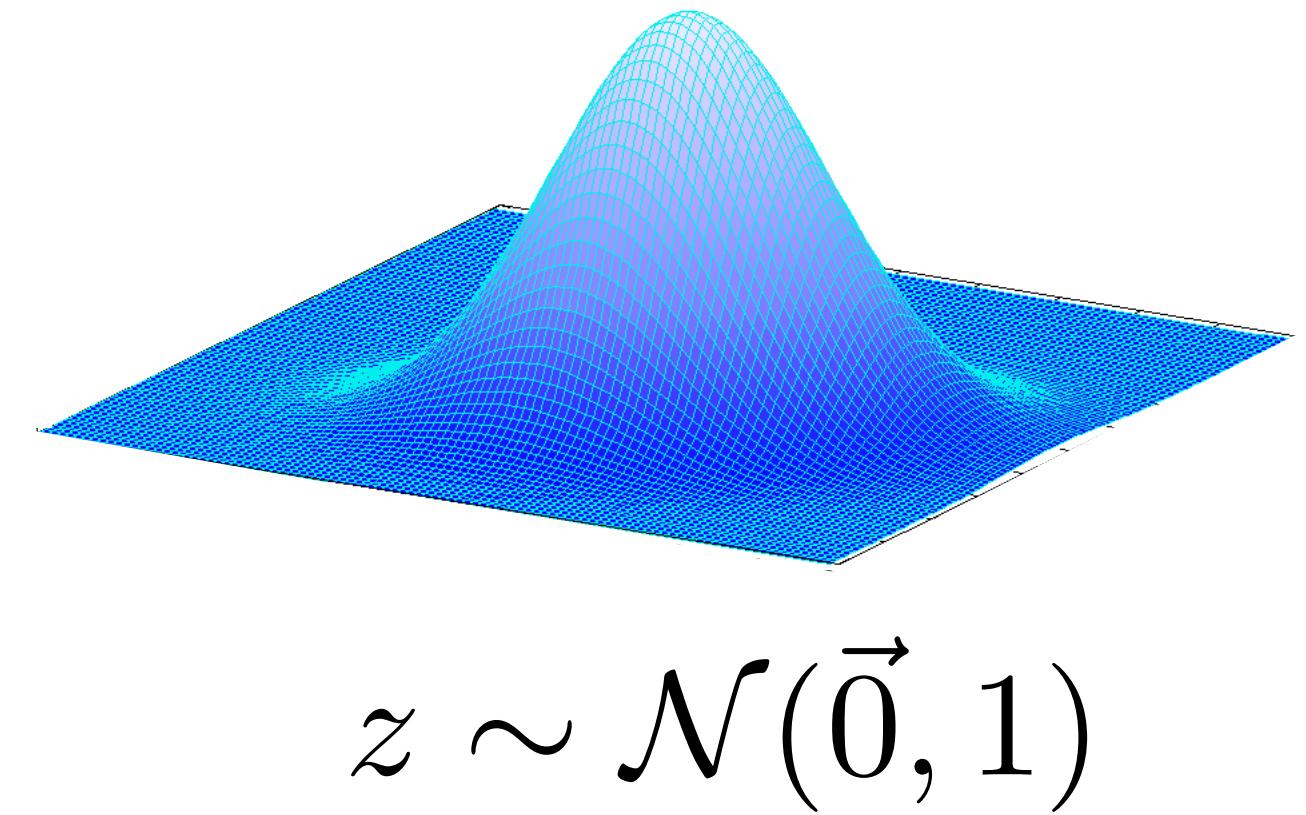


z learns to encode the missing information necessary
to predict y from x , i.e. the direction in which the
purple box moves

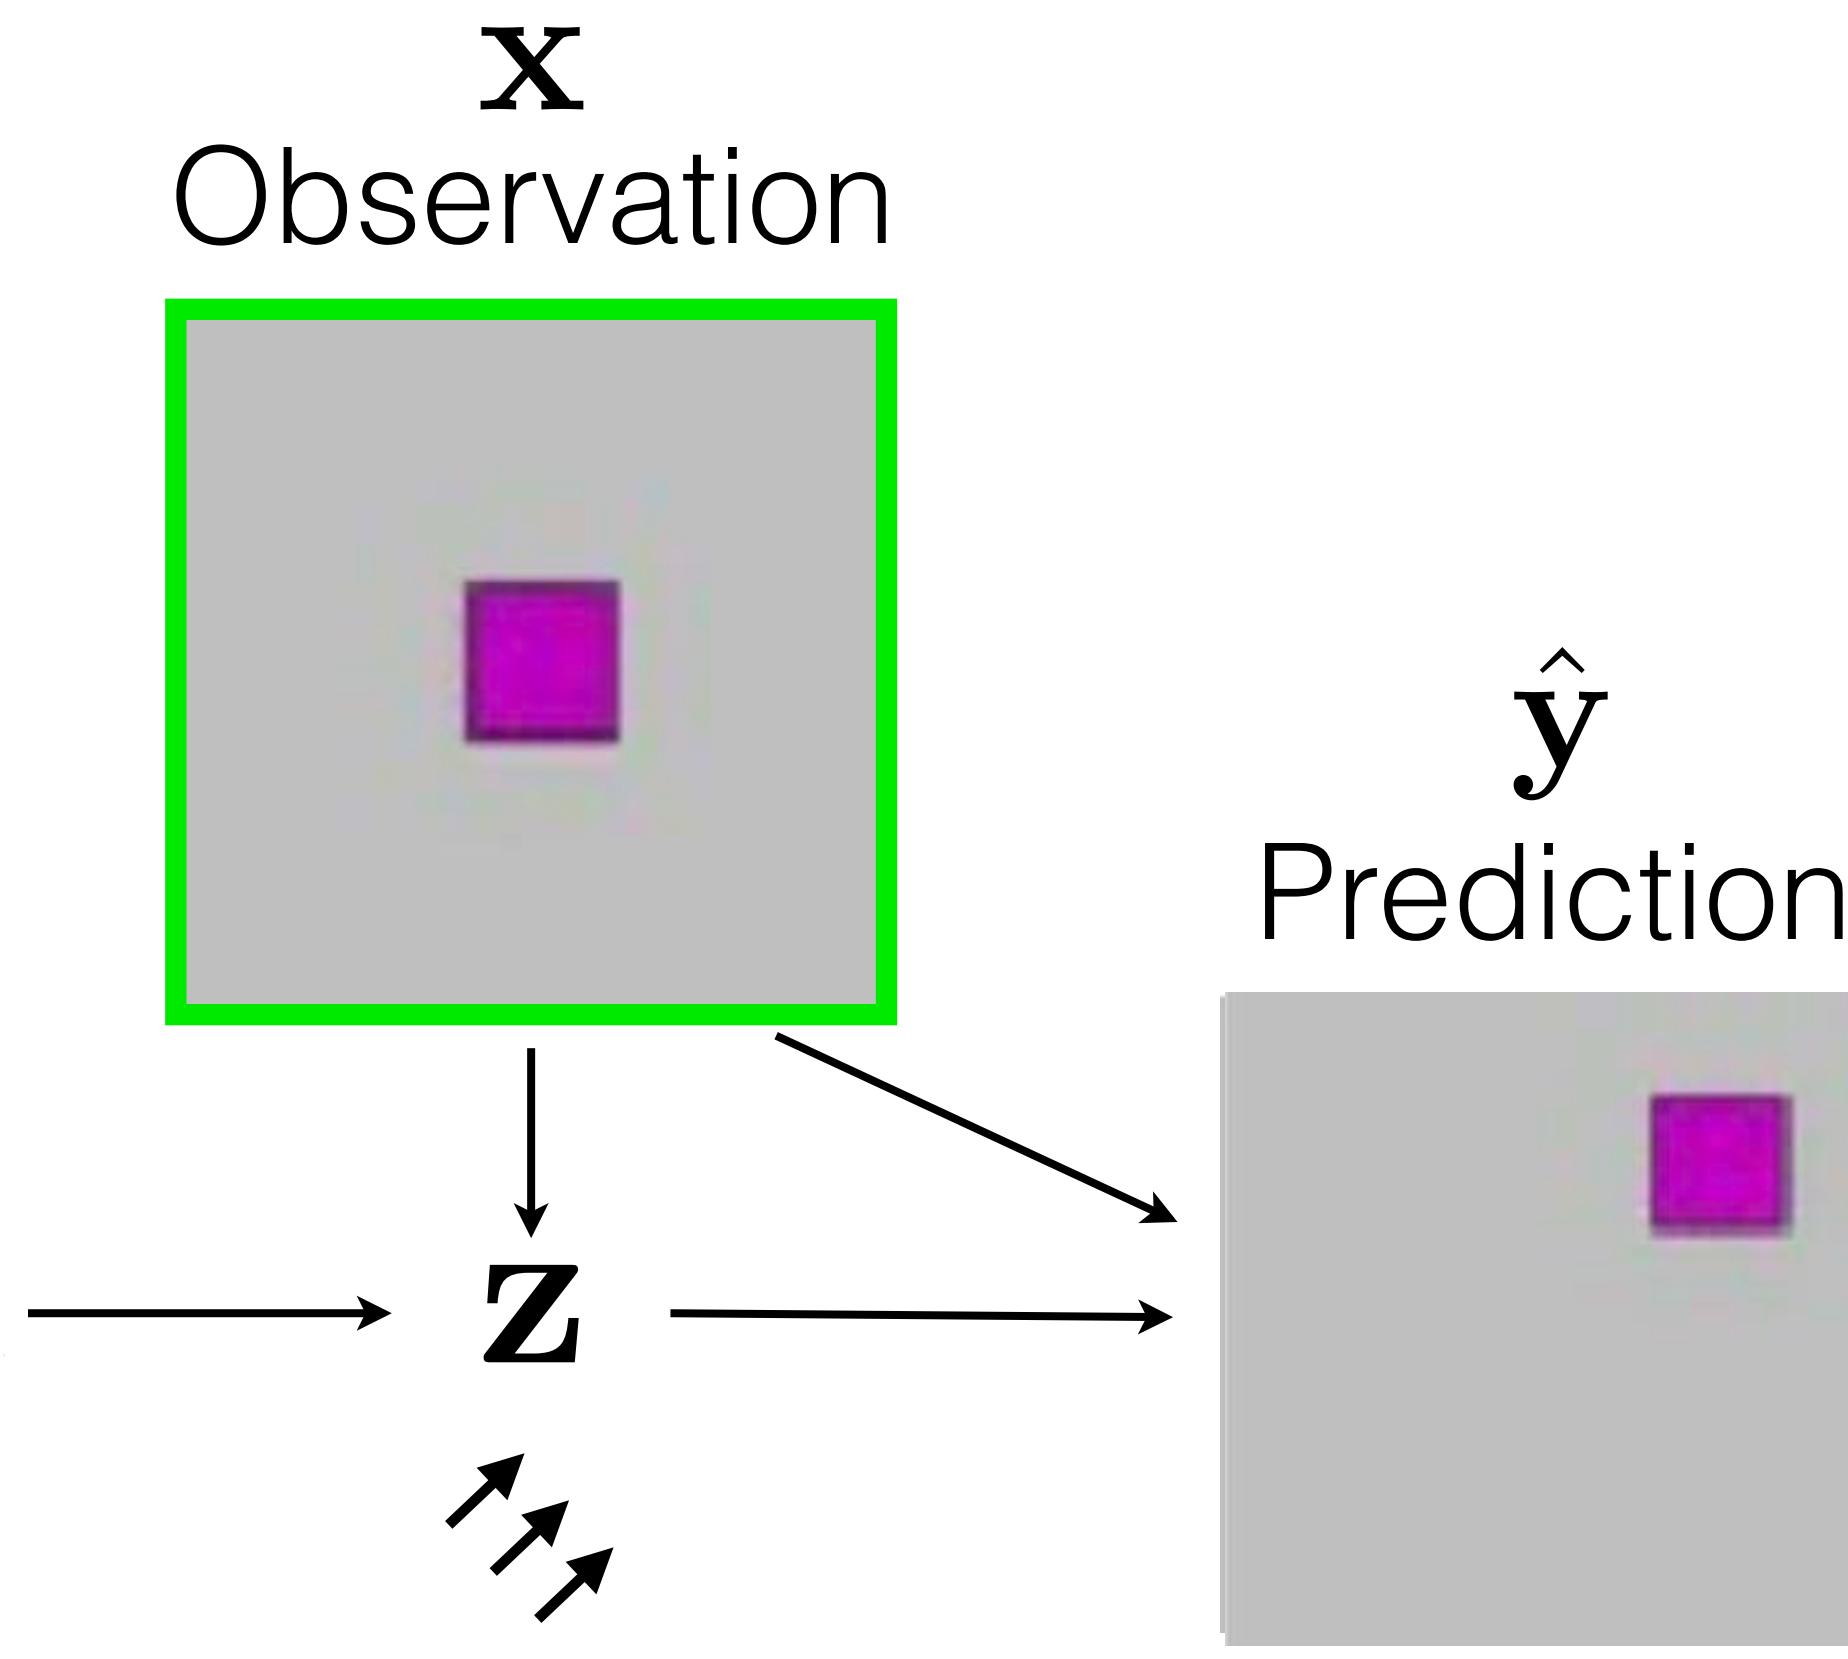
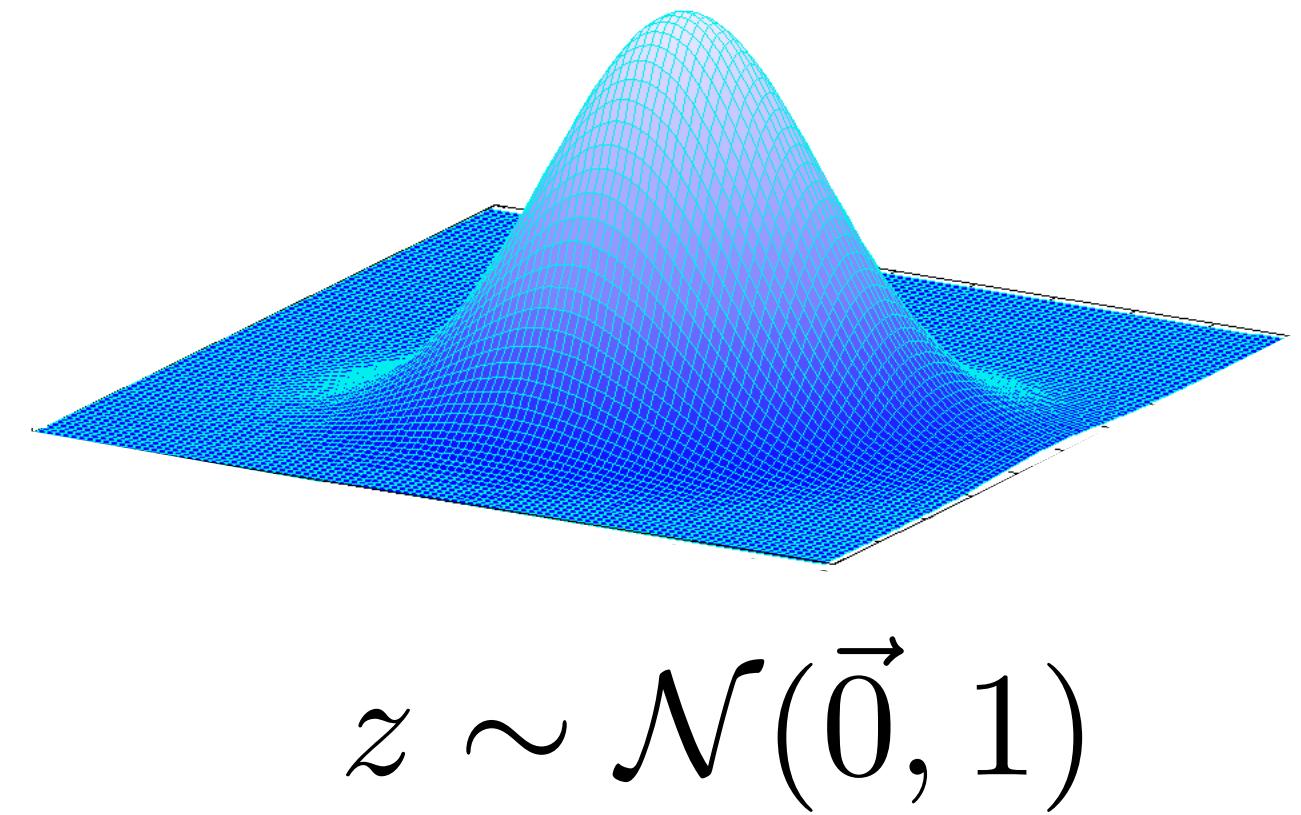
Example from [Babaeizadeh et al., ICLR 2018]
see also [Walker et al., ECCV 2016], [Xue*, Wu*, et al., NIPS 2016]



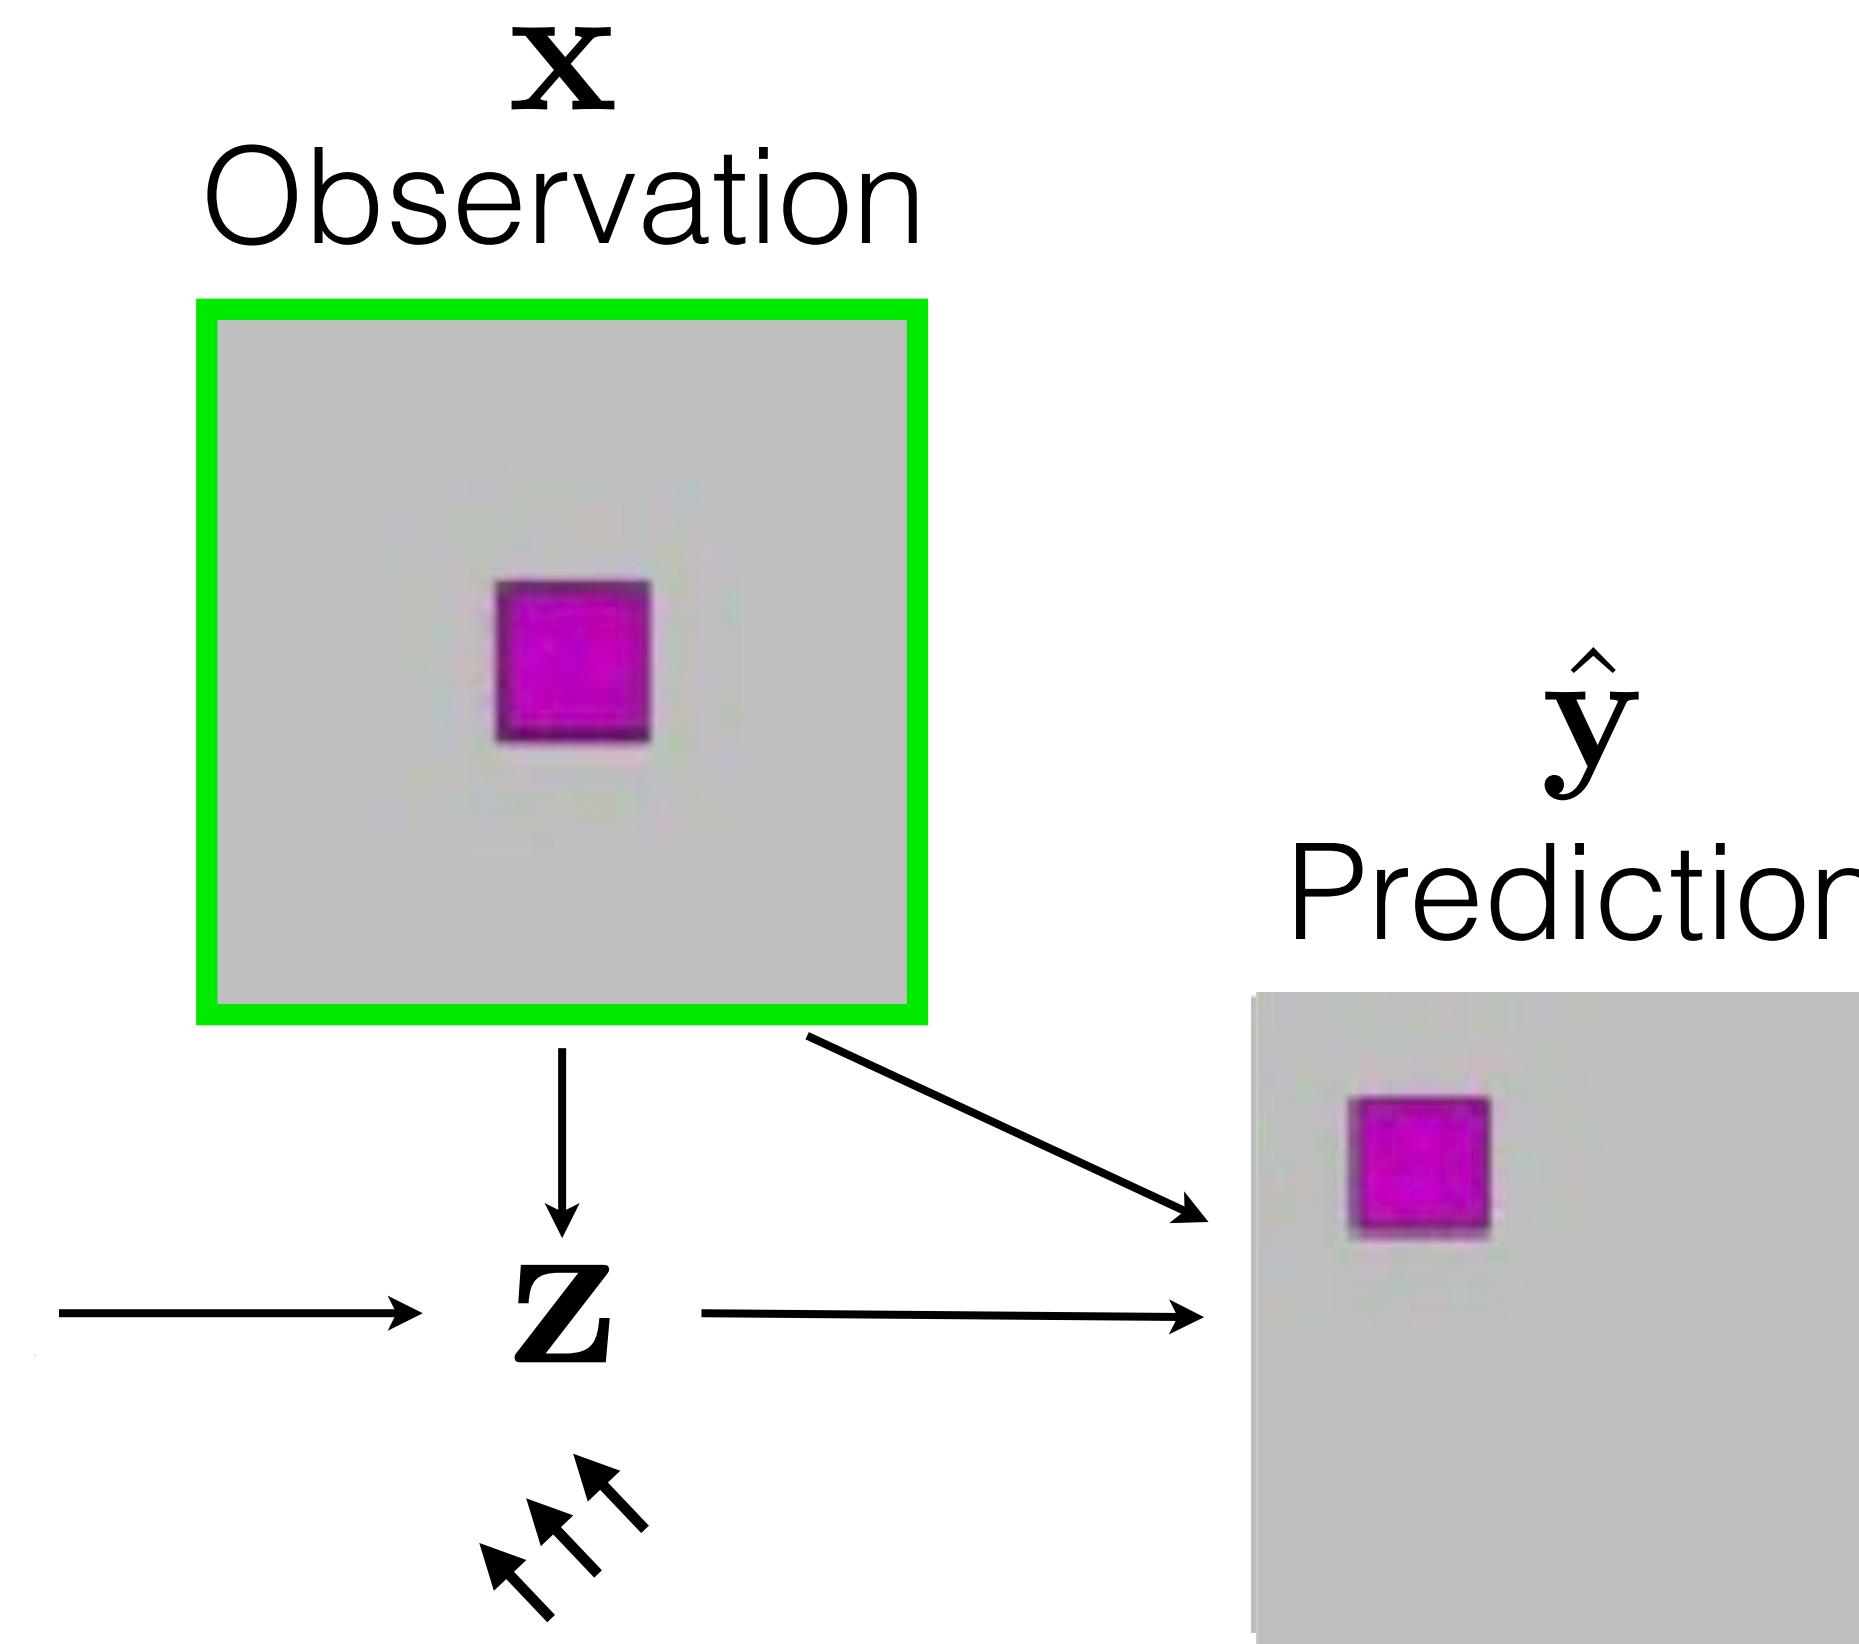
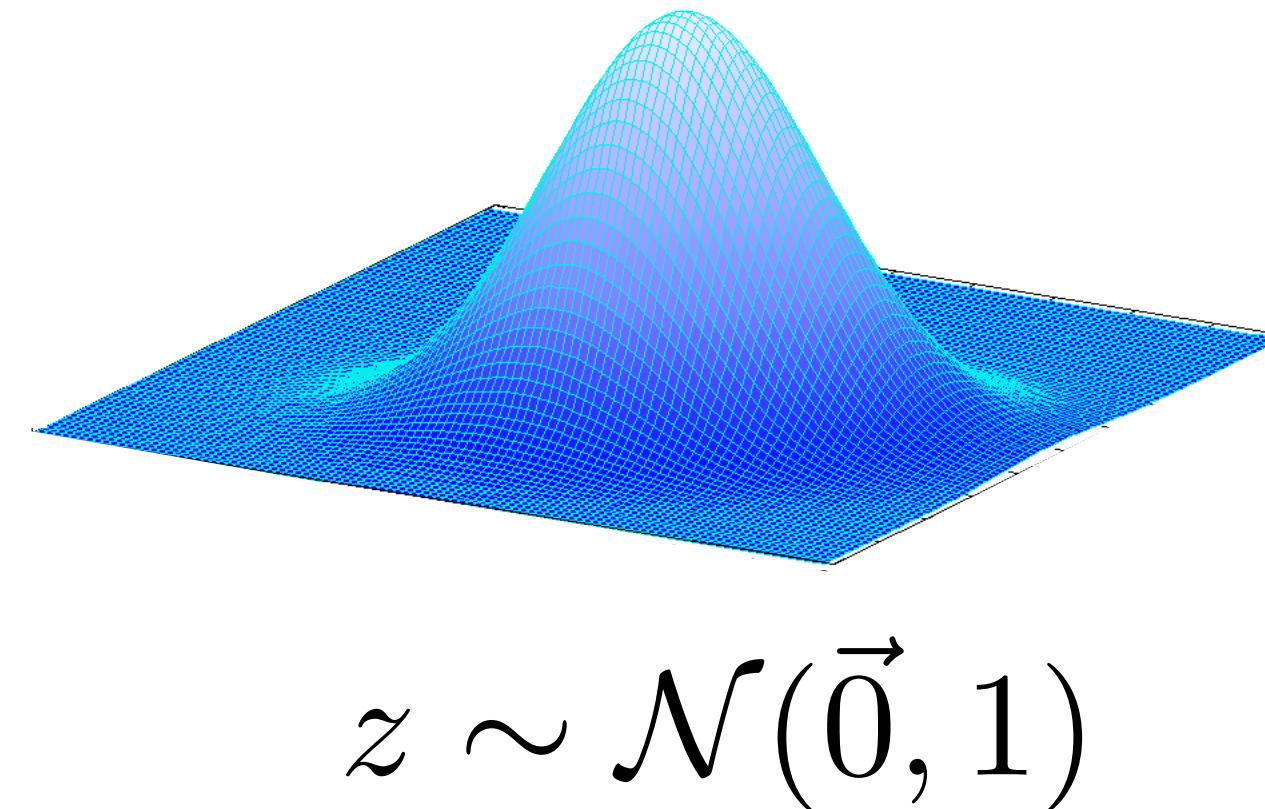
see also [Walker et al., ECCV 2016], [Xue*, Wu*, et al., NIPS 2016]



see also [Walker et al., ECCV 2016], [Xue*, Wu*, et al., NIPS 2016]



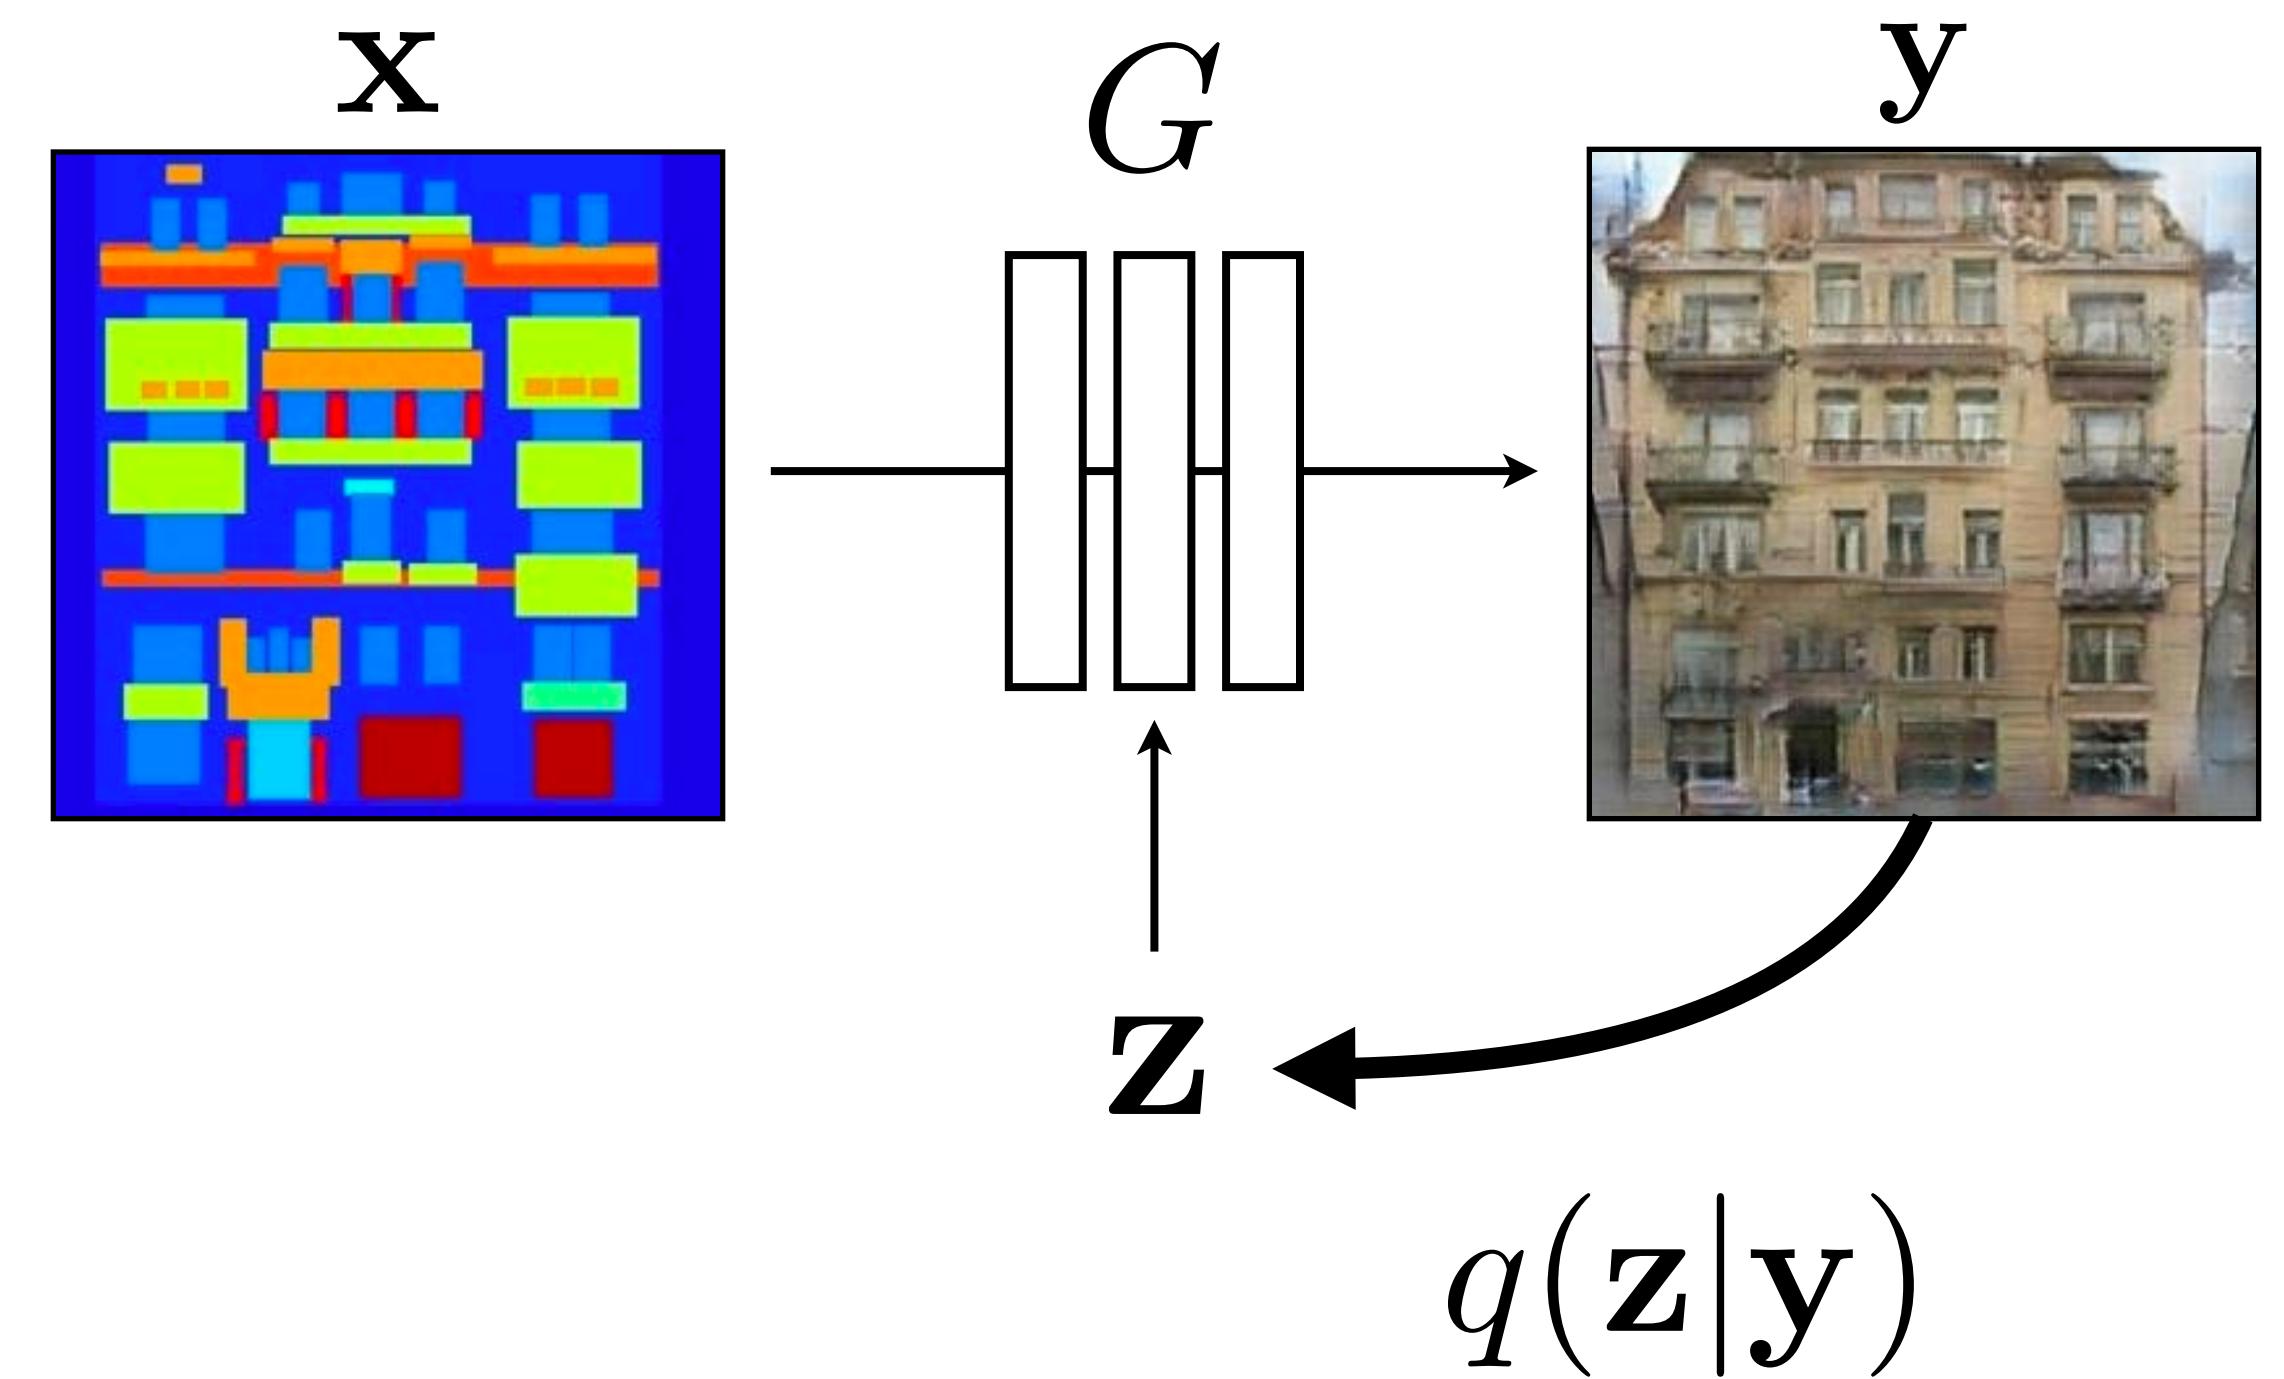
see also [Walker et al., ECCV 2016], [Xue*, Wu*, et al., NIPS 2016]



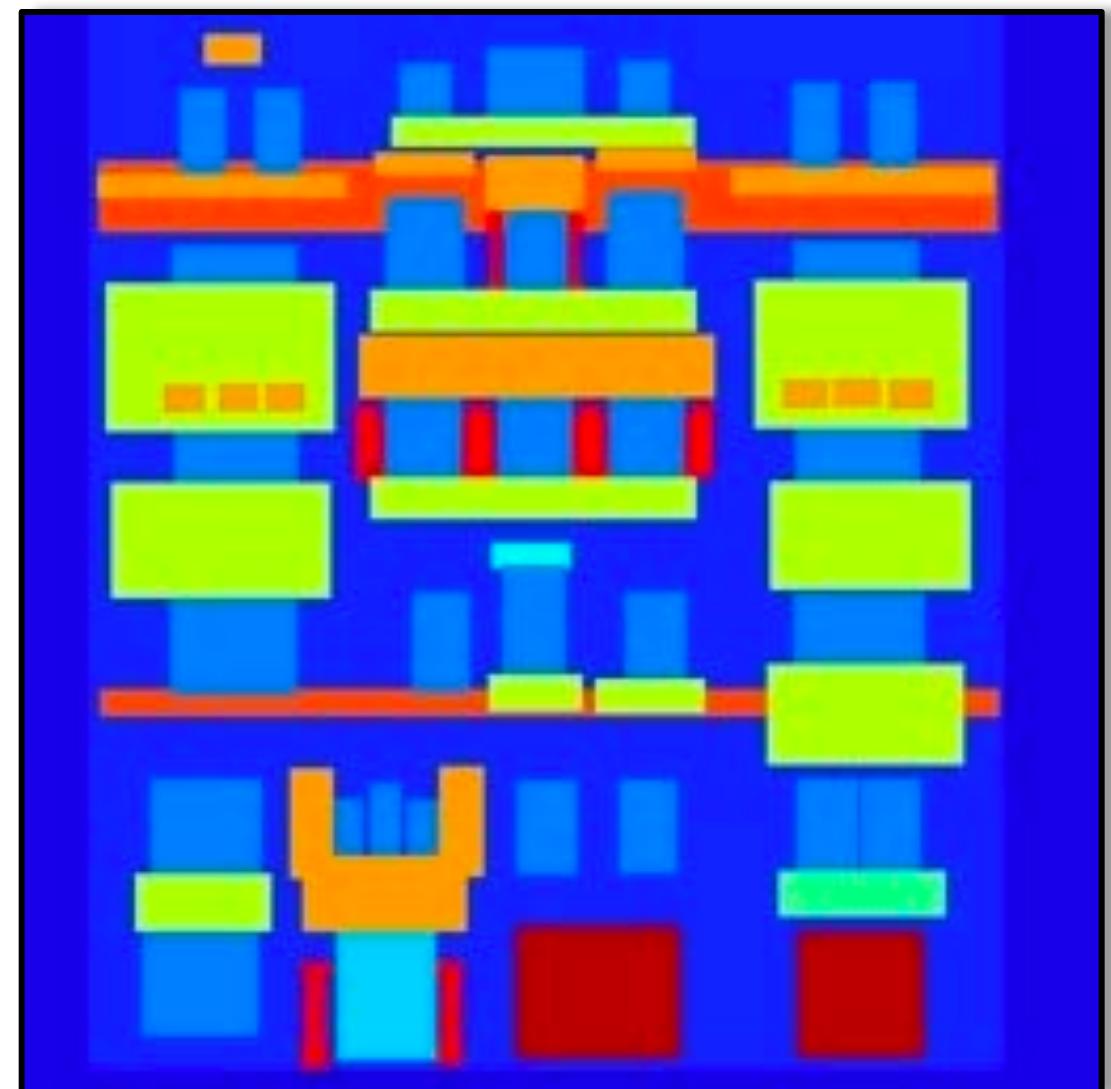
see also [Walker et al., ECCV 2016], [Xue*, Wu*, et al., NIPS 2016]

InfoGAN [Chen et al. 2016]

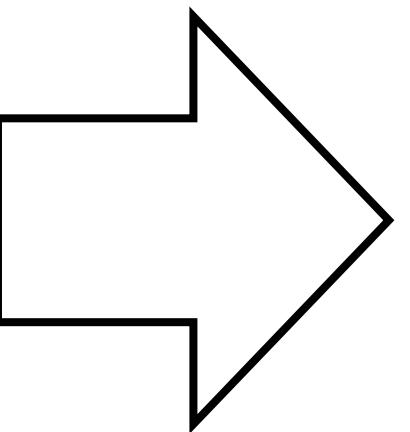
BiCycleGAN [Zhu et al., NIPS 2017]



Encourages z to relay information about the target.



Labels



Randomly generated facades

[BiCycleGAN, Zhu et al., NIPS 2017]

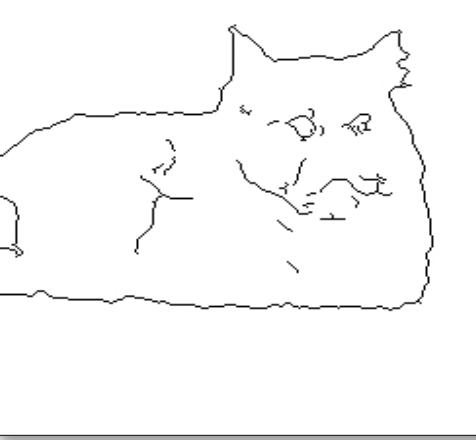
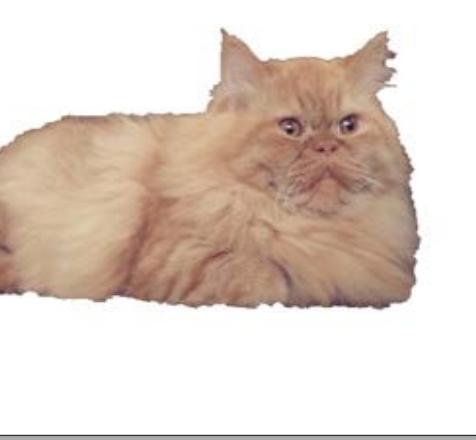
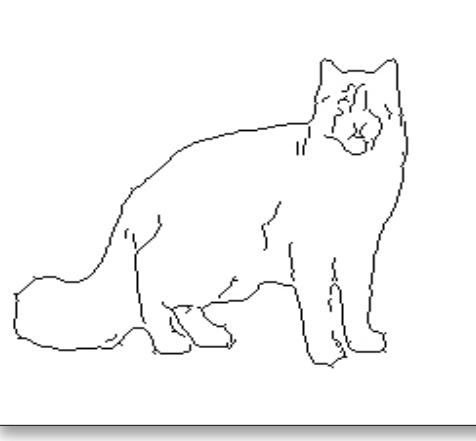
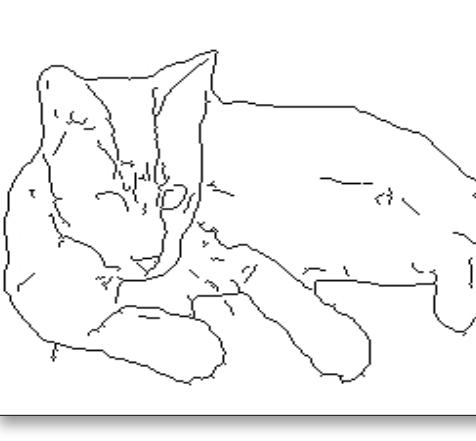
What can you do with generative models?

- 1. Data prediction**
2. Domain mapping
3. Representation learning
4. Model-based intelligence

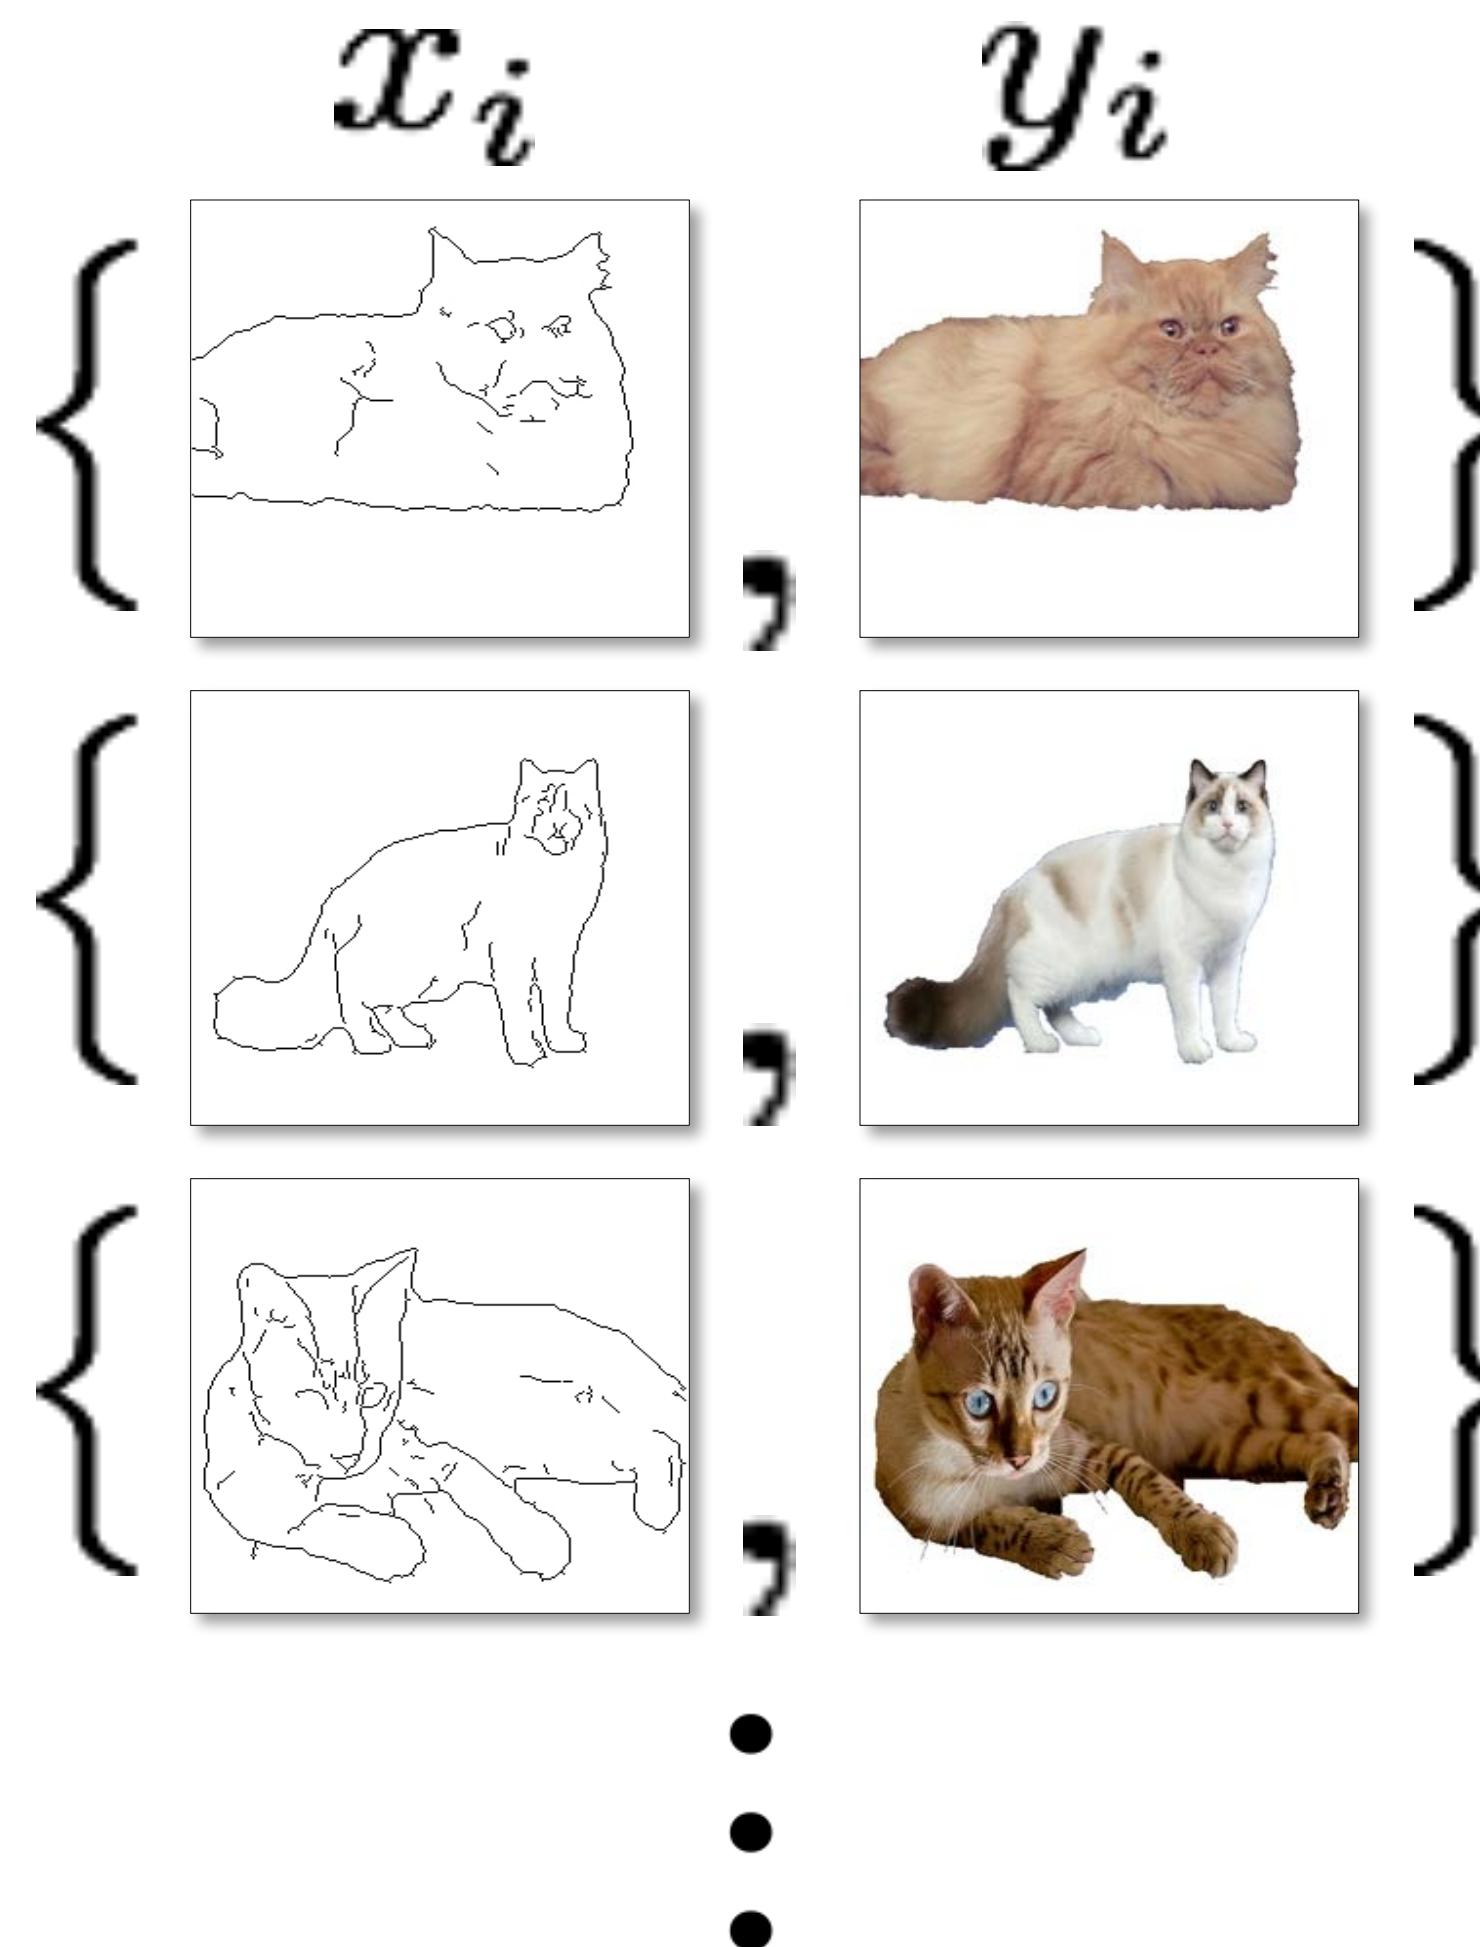
What can you do with generative models?

1. Data prediction
- 2. Domain mapping**
3. Representation learning
4. Model-based intelligence

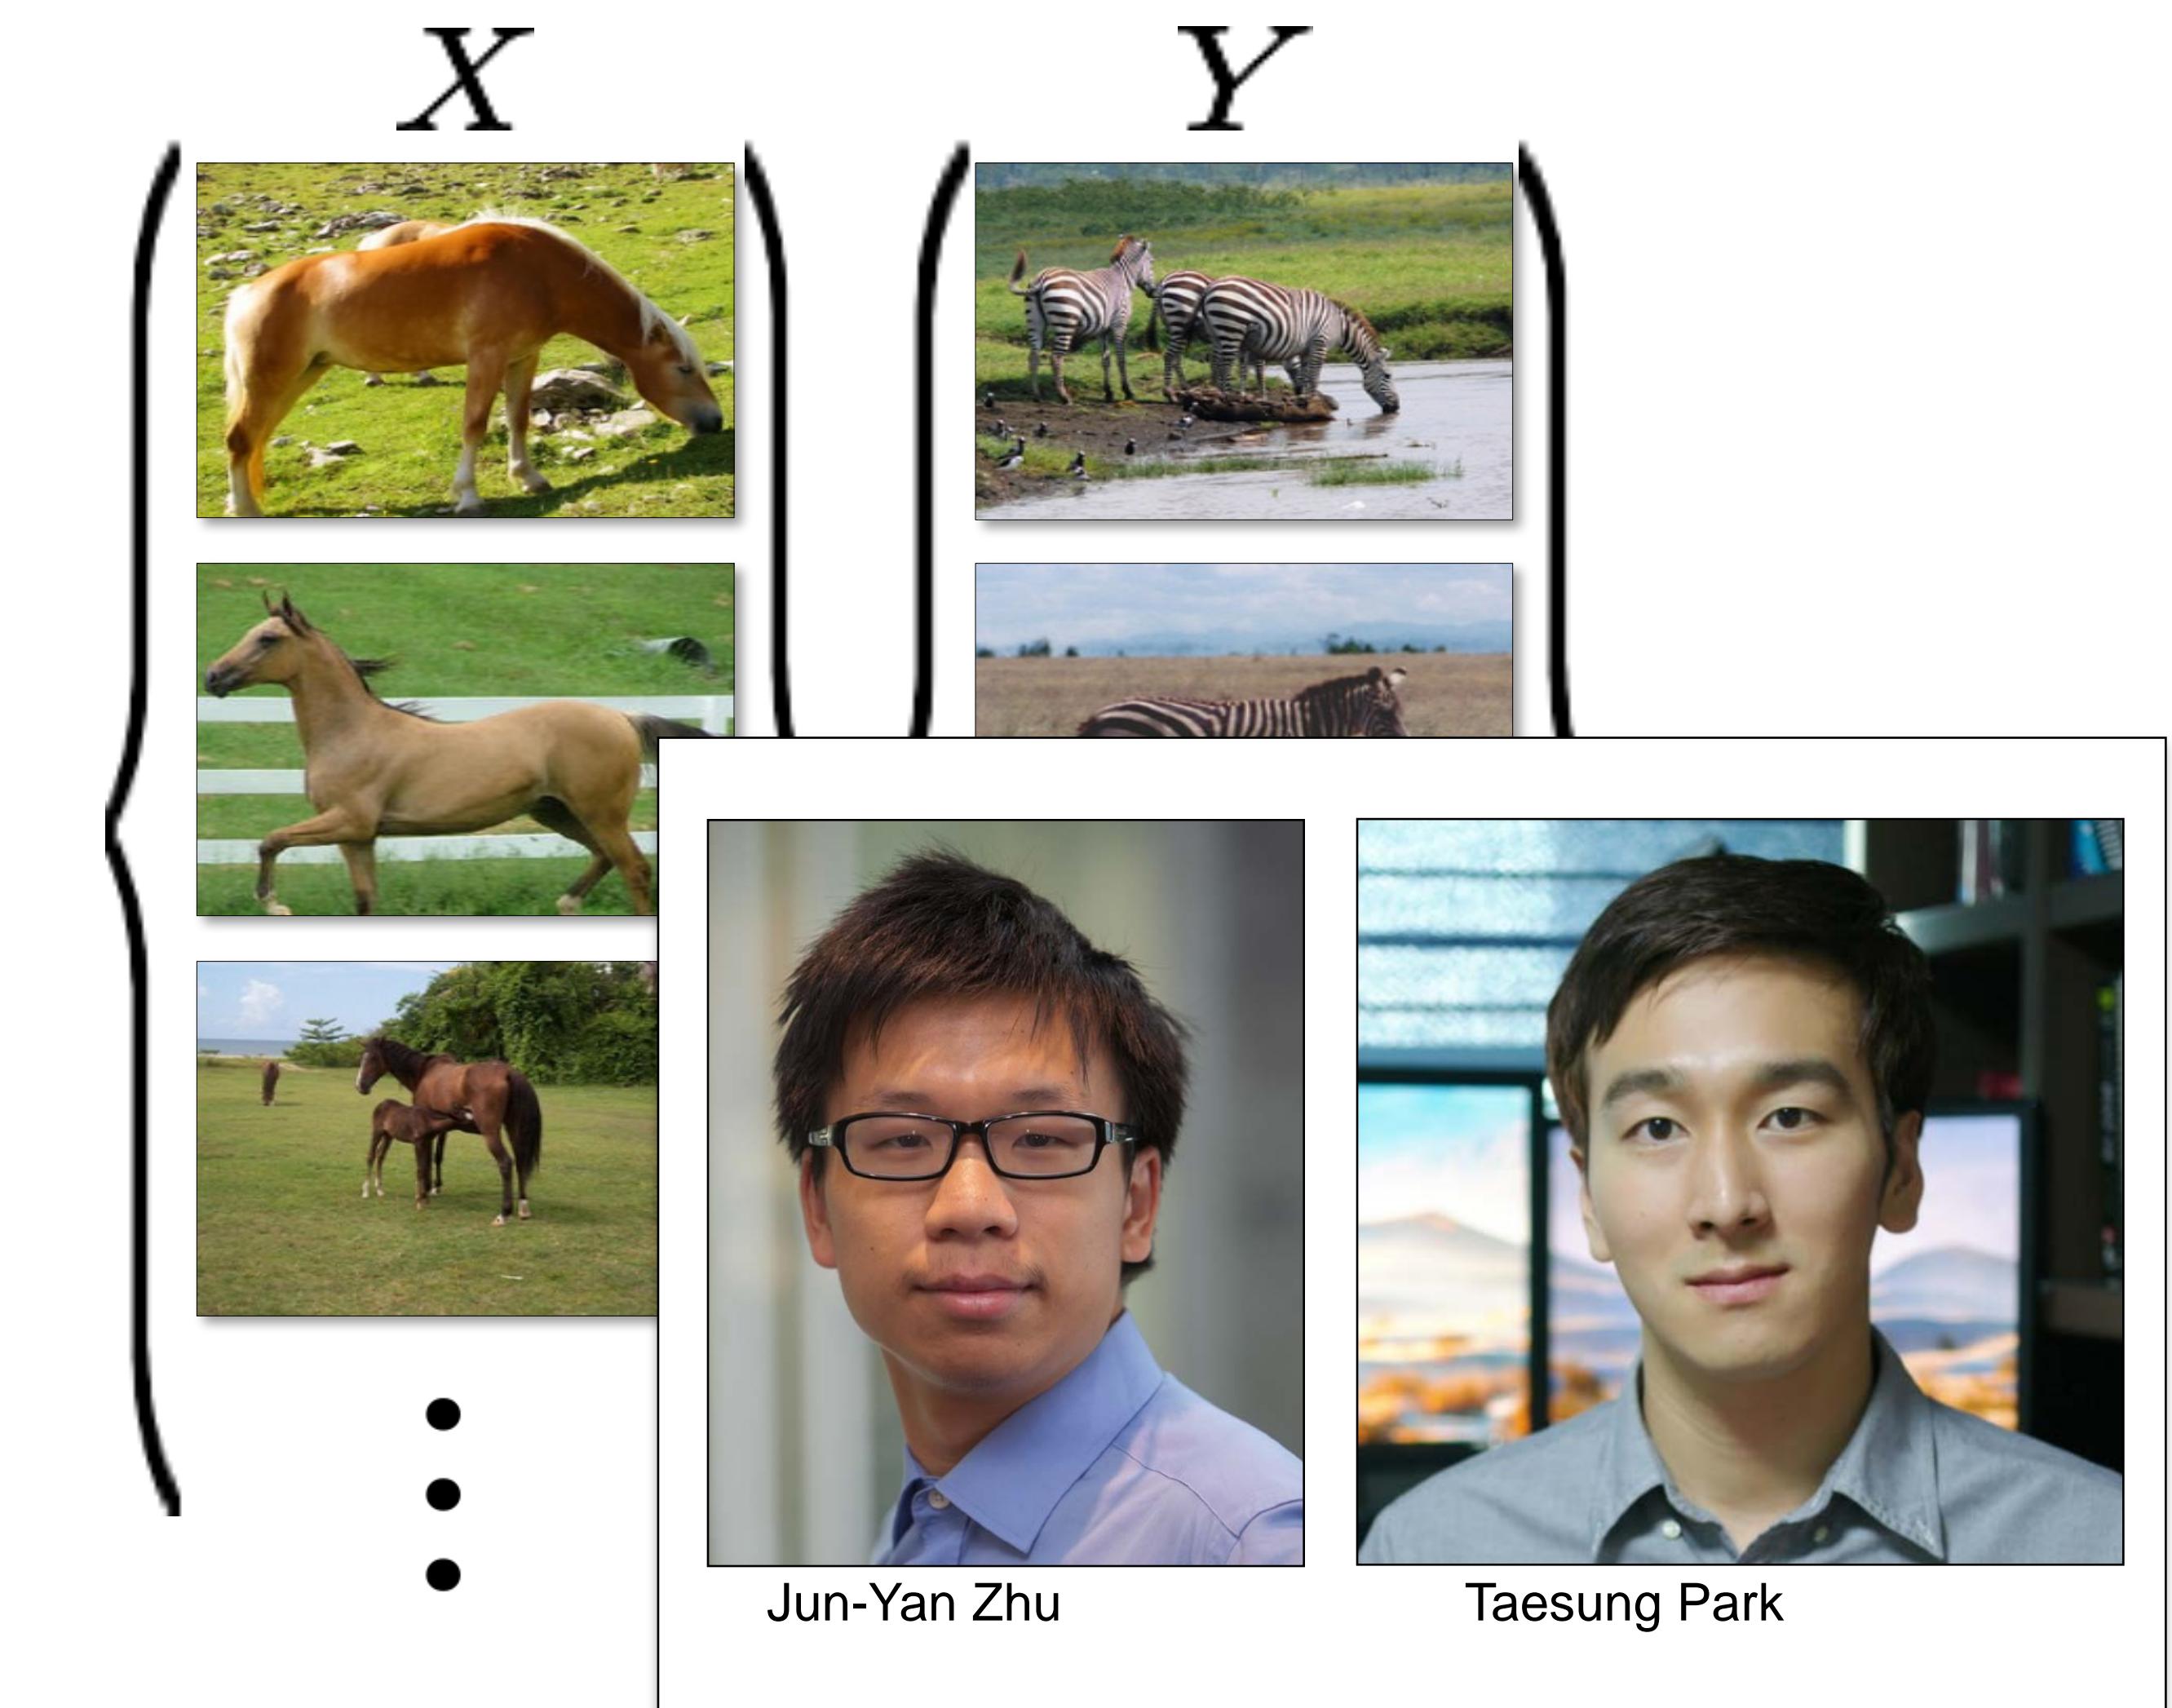
Paired data

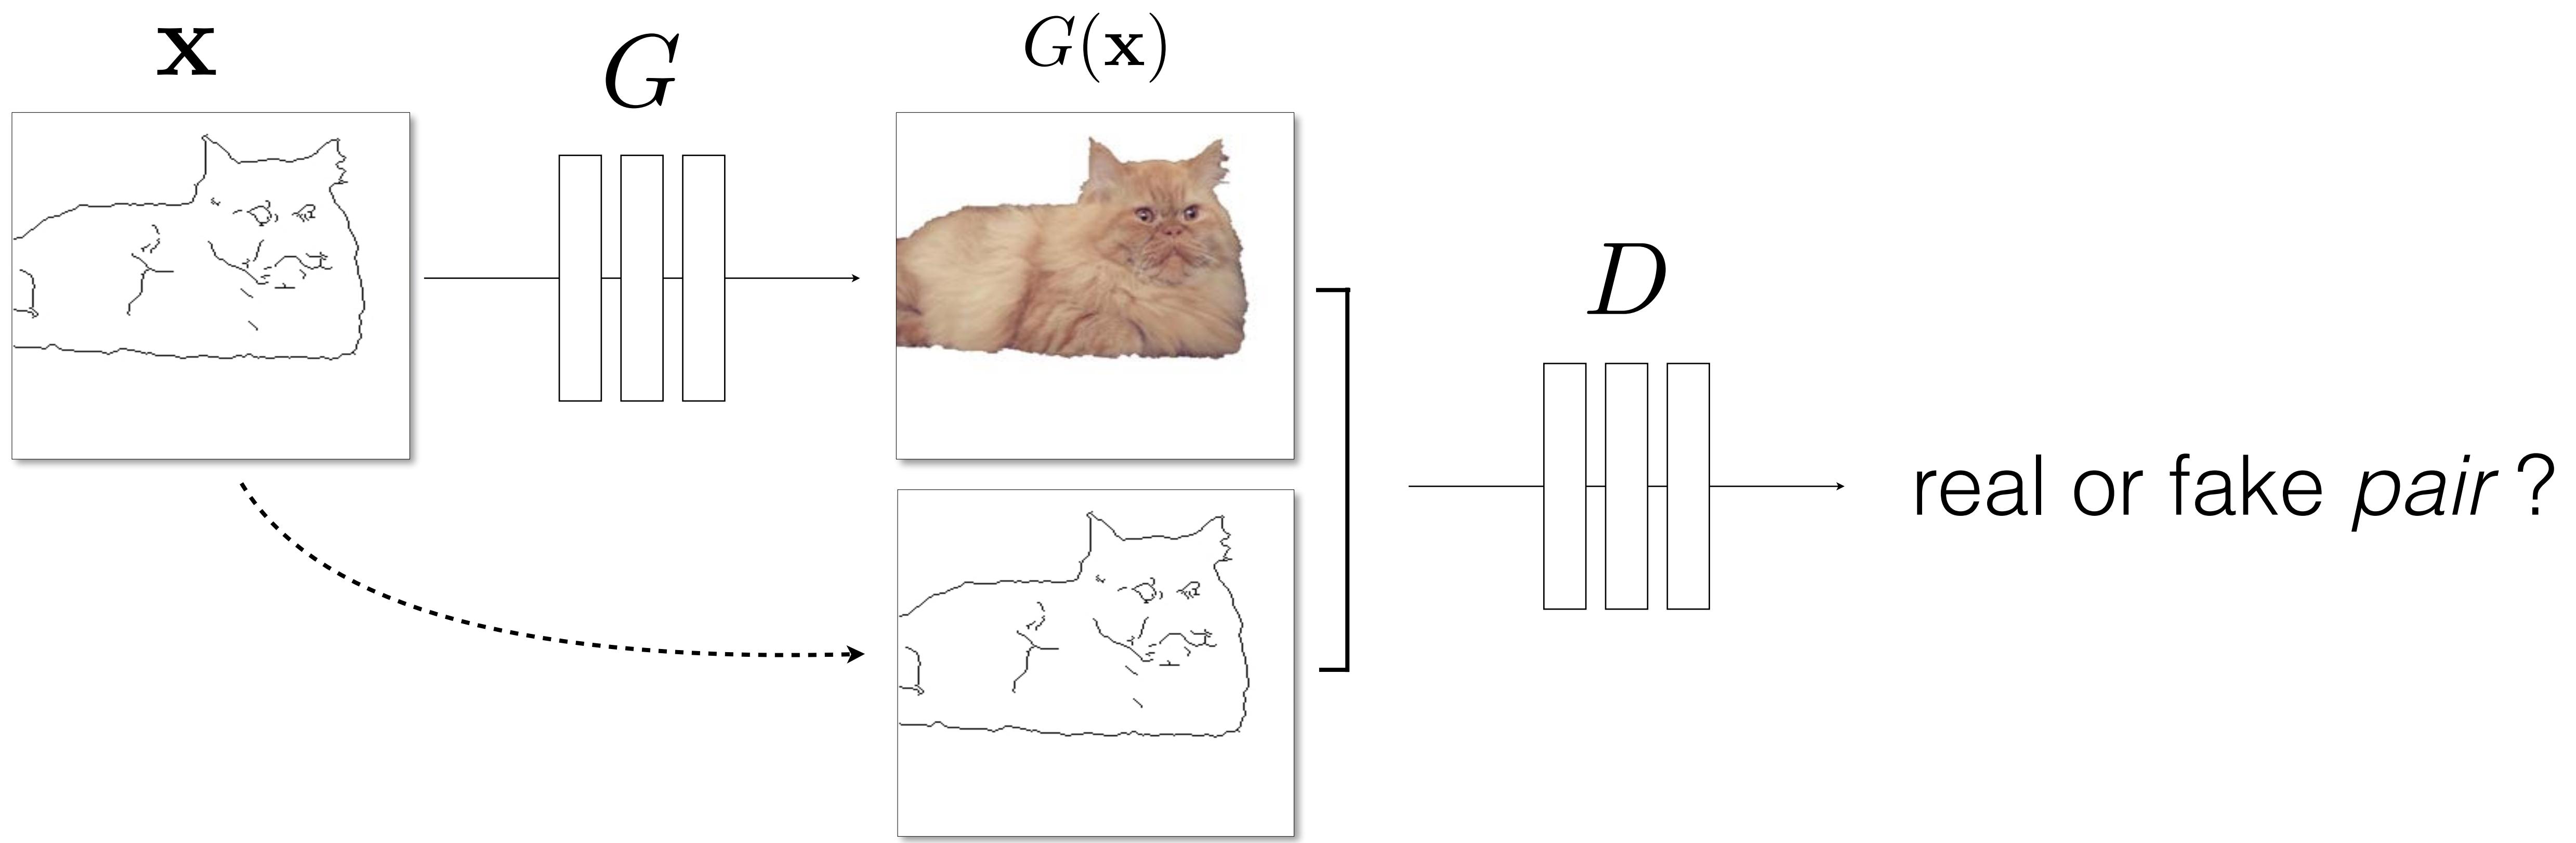
x_i	y_i	
{ 		}
{ 		}
{ 		}
⋮		

Paired data

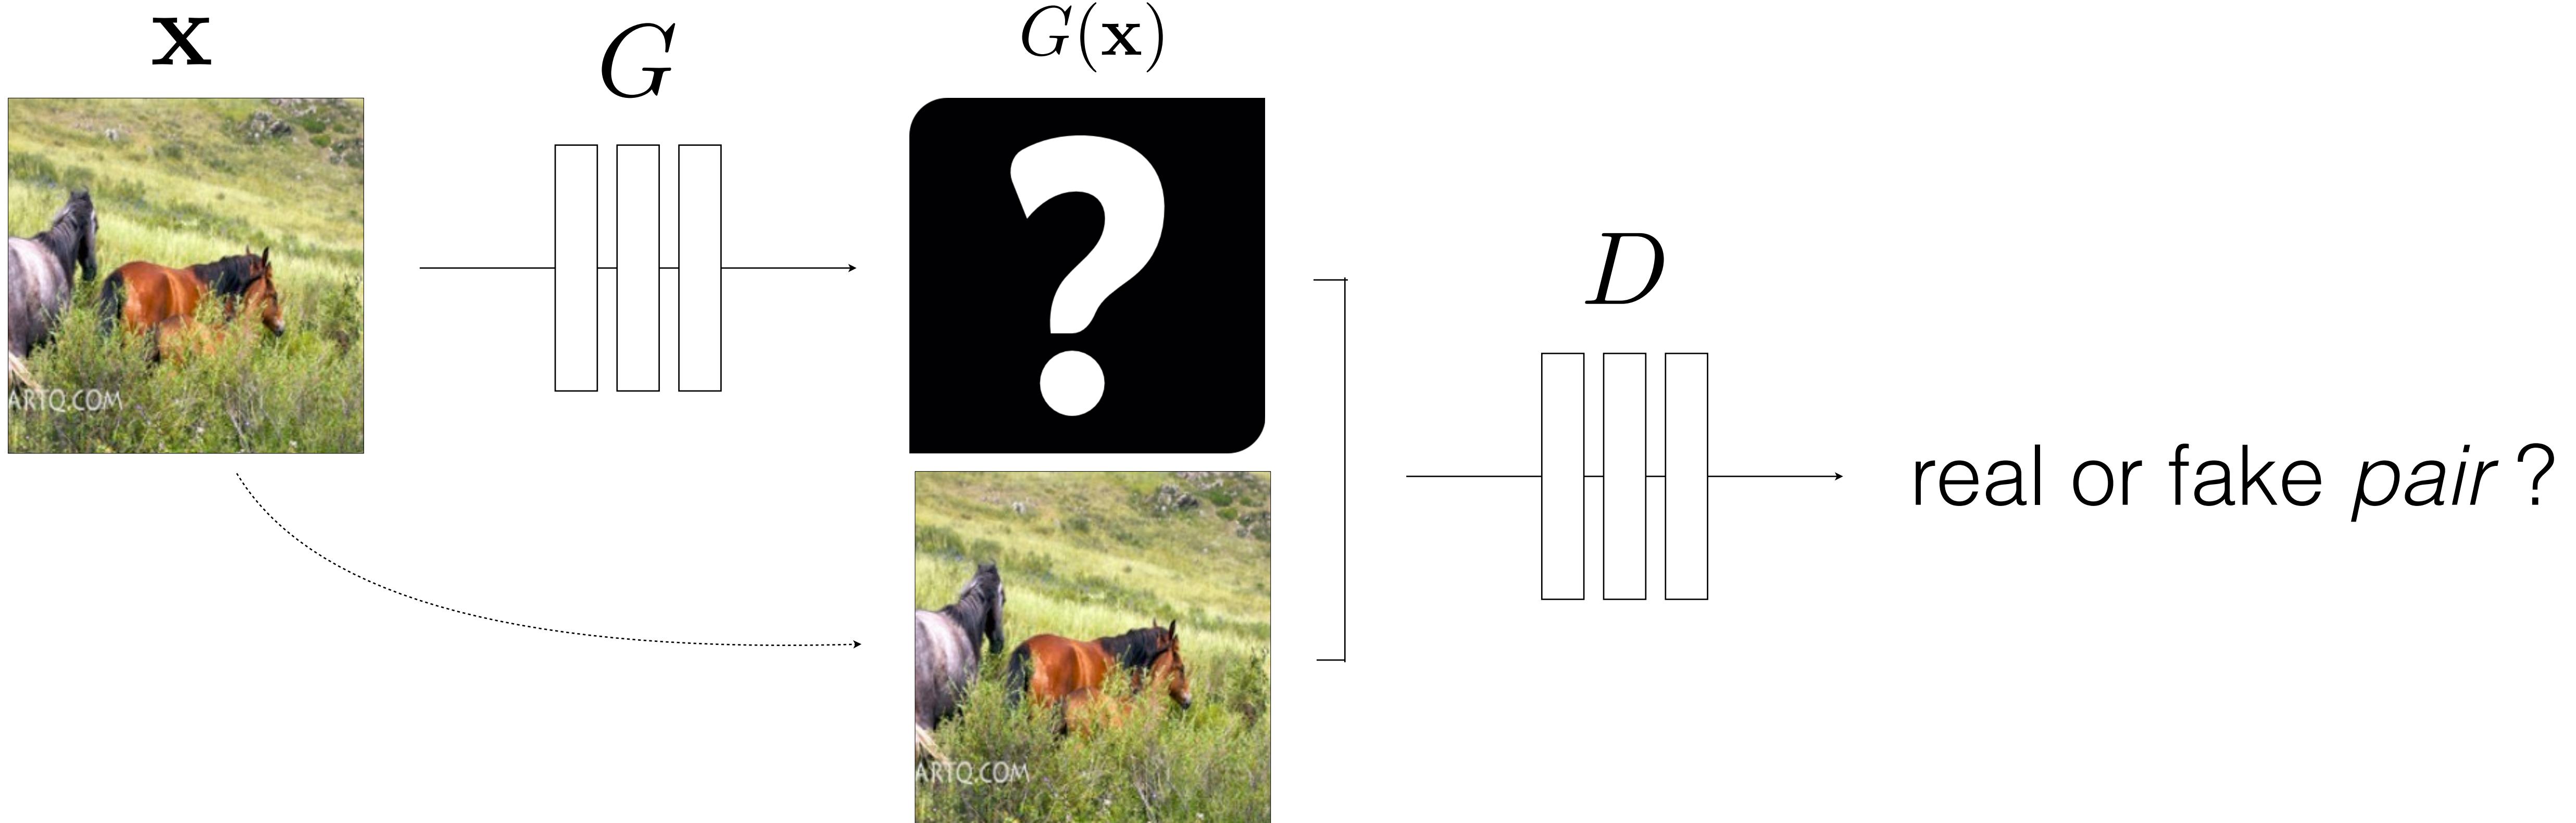


Unpaired data





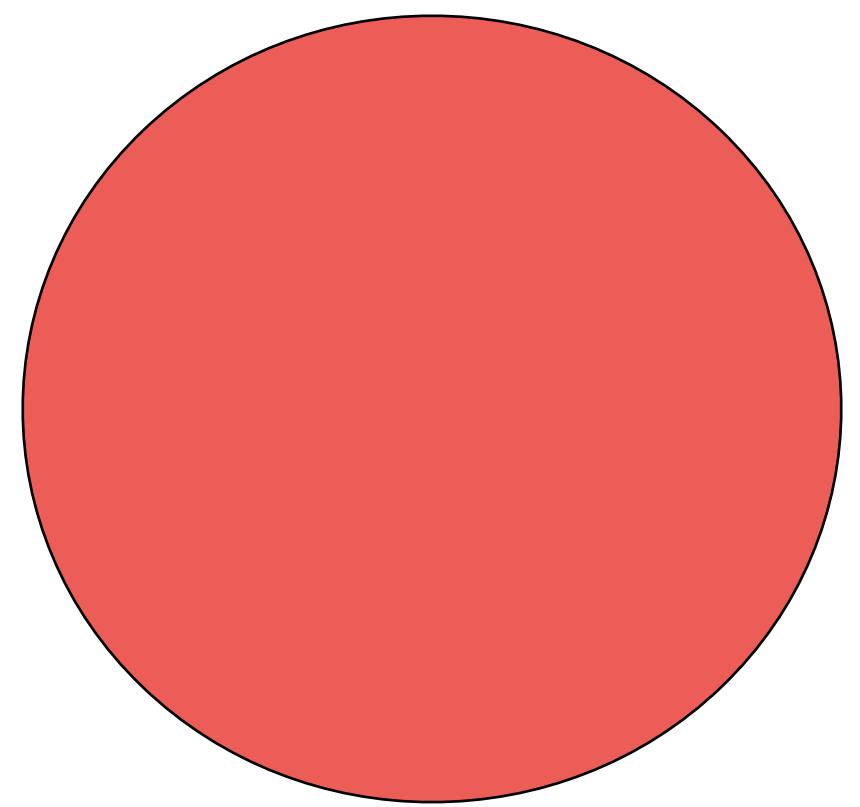
$$\arg \min_G \max_D \mathbb{E}_{\mathbf{x}, \mathbf{y}} [\log D(\mathbf{x}, G(\mathbf{x})) + \log(1 - D(\mathbf{x}, \mathbf{y}))]$$



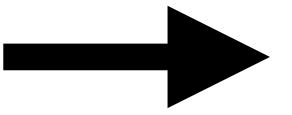
$$\arg \min_G \max_D \mathbb{E}_{\mathbf{x}, \mathbf{y}} [\log D(\mathbf{x}, G(\mathbf{x})) + \log(1 - D(\mathbf{x}, \mathbf{y}))]$$

No input-output pairs!

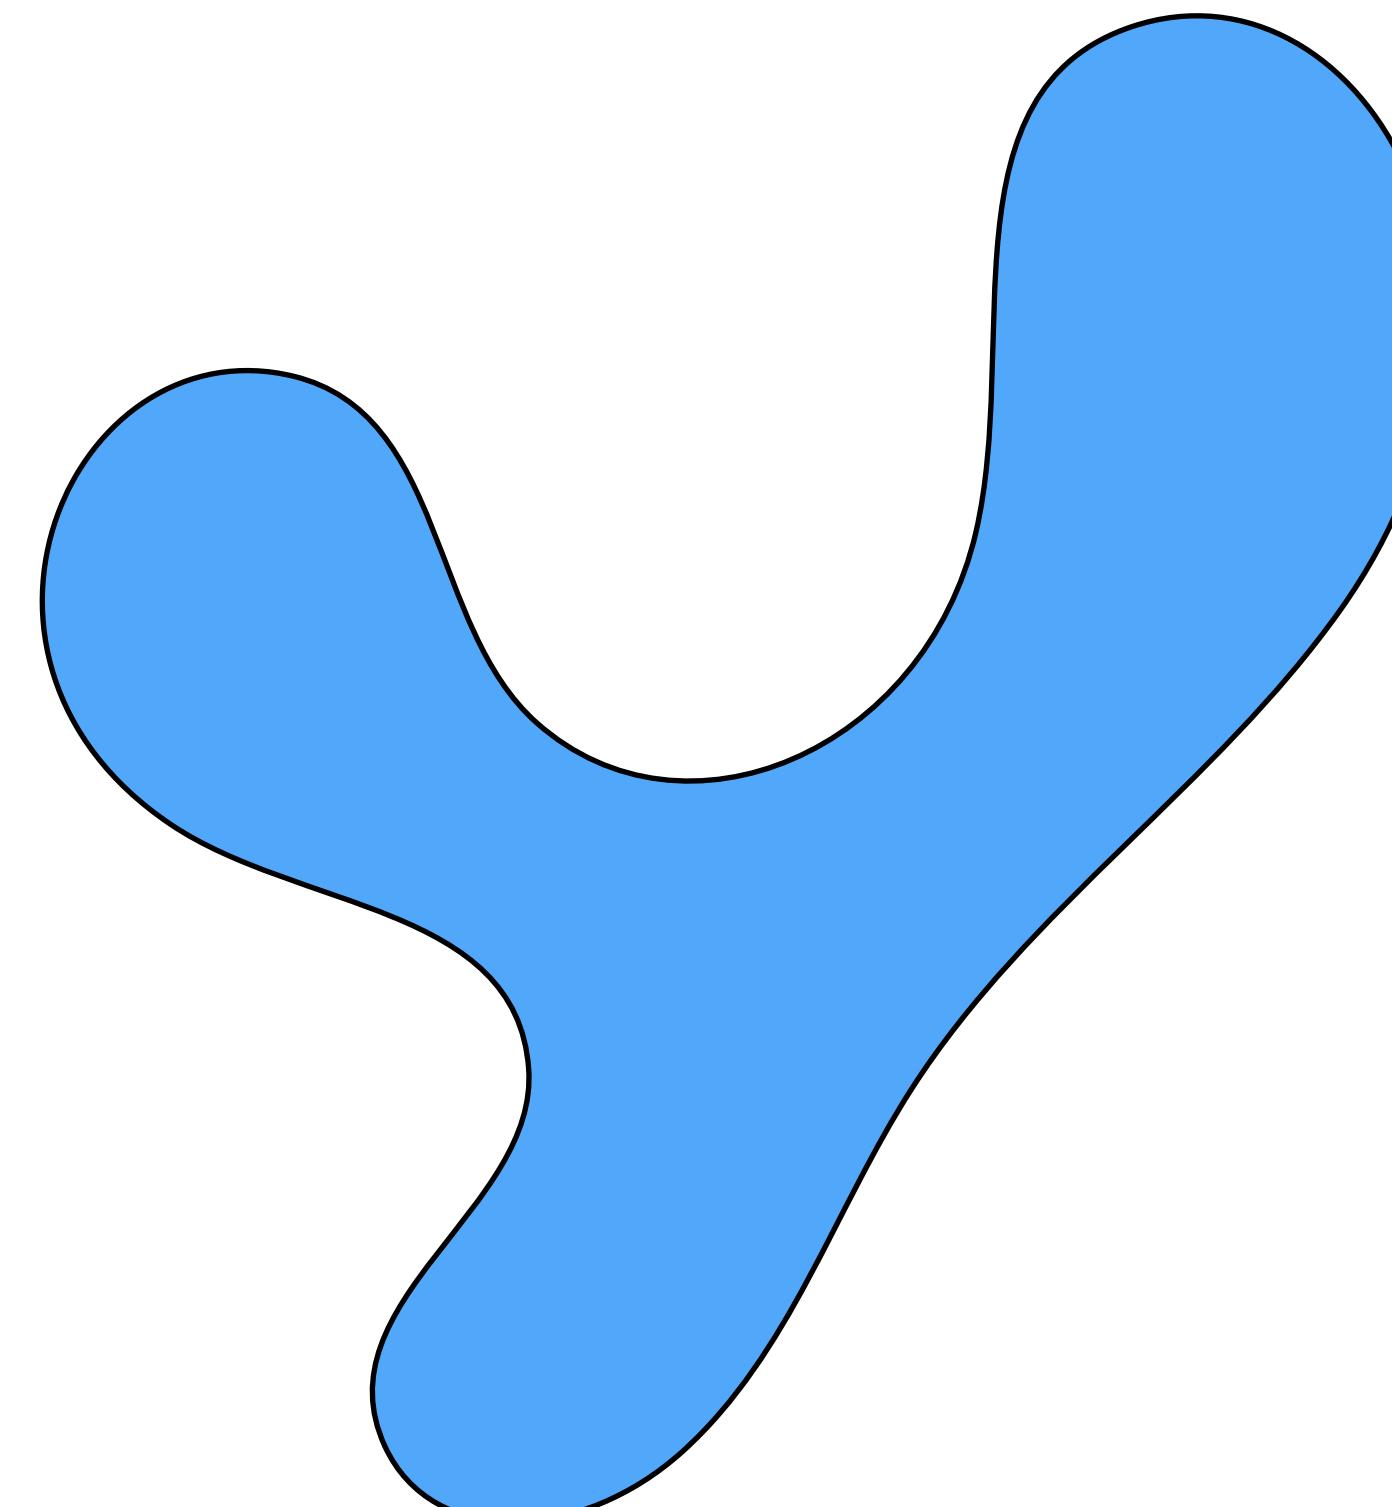
Gaussian



z

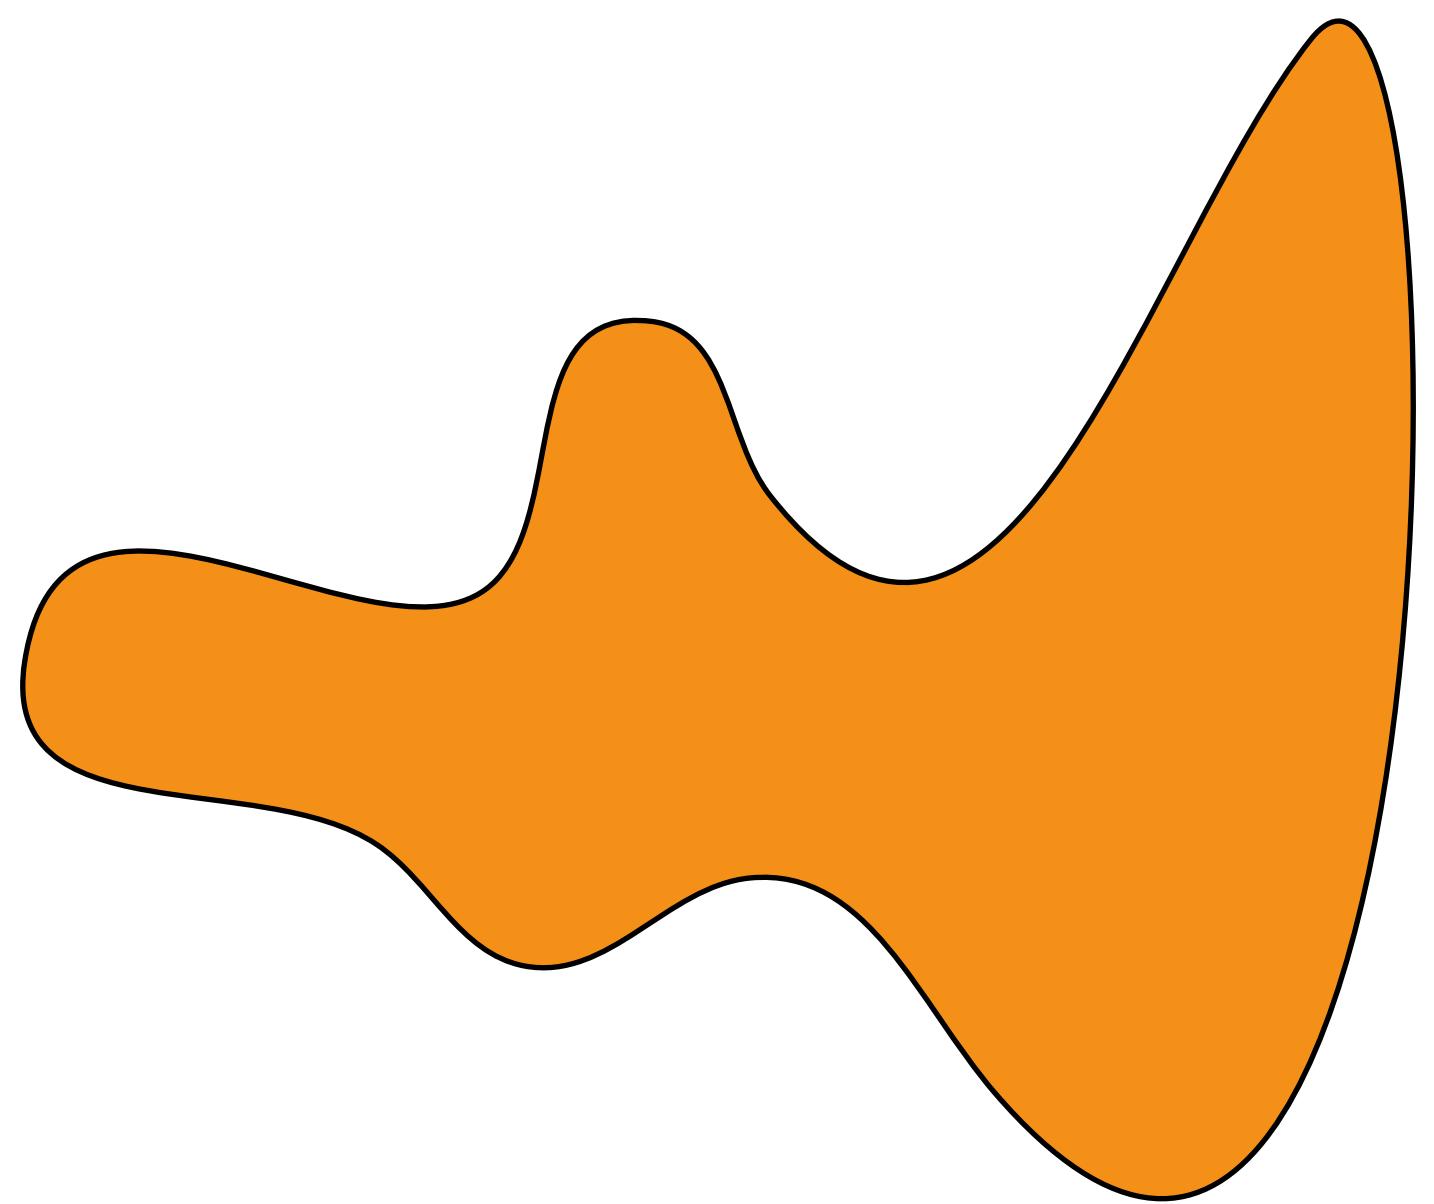


Target distribution



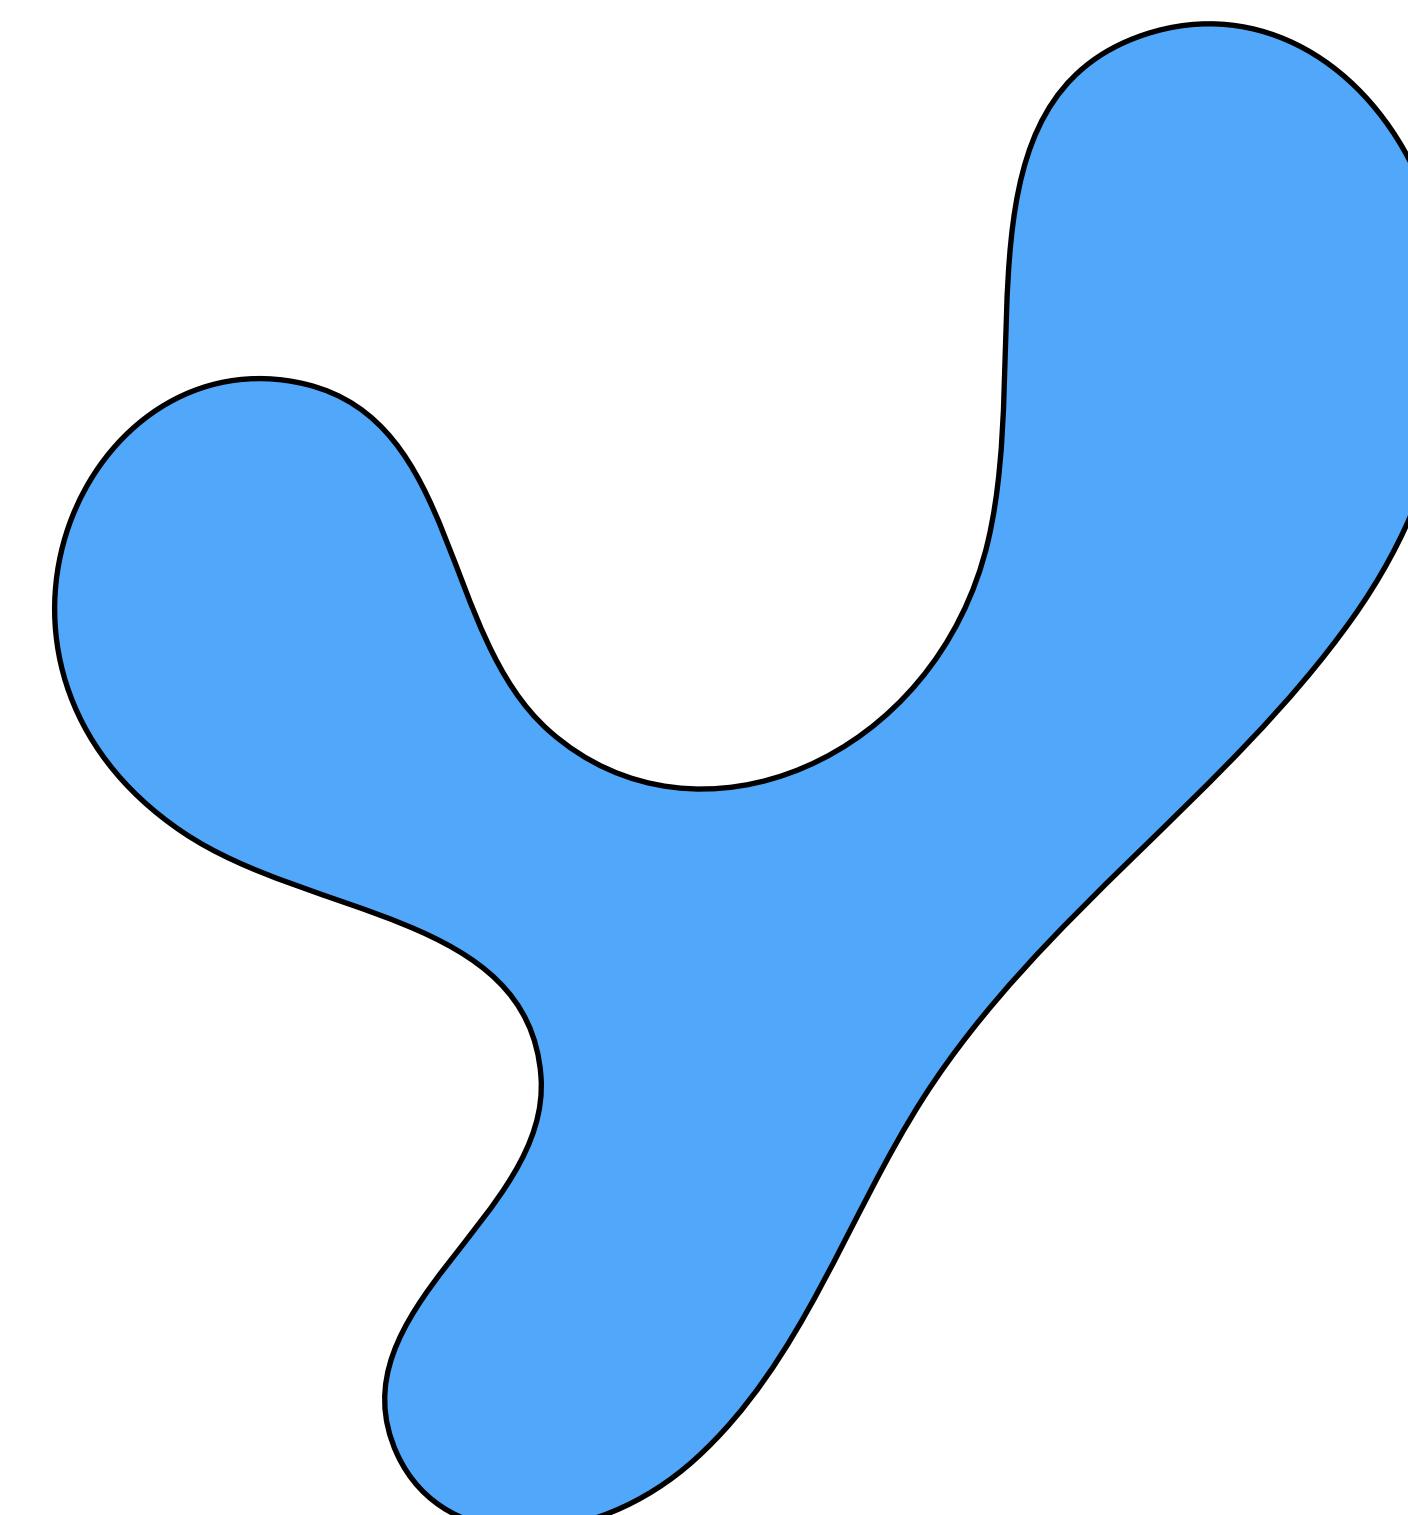
Y

Horses

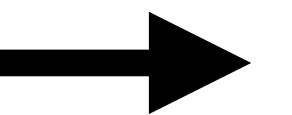


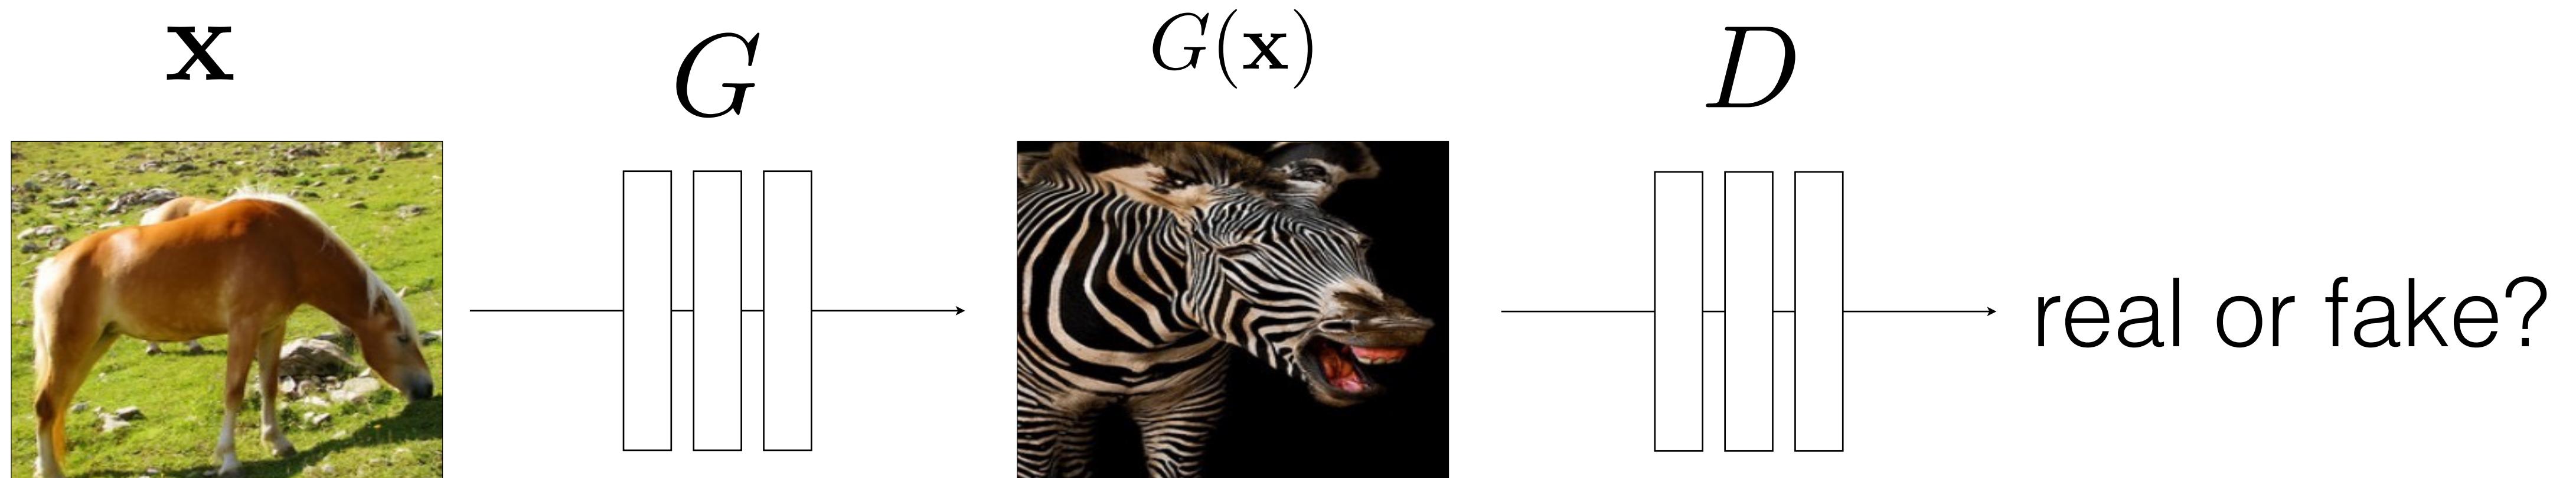
X

Zebras



Y





$$\arg \min_G \max_D \mathbb{E}_{\mathbf{x}, \mathbf{y}} [\log D(G(\mathbf{x})) + \log(1 - D(\mathbf{y}))]$$

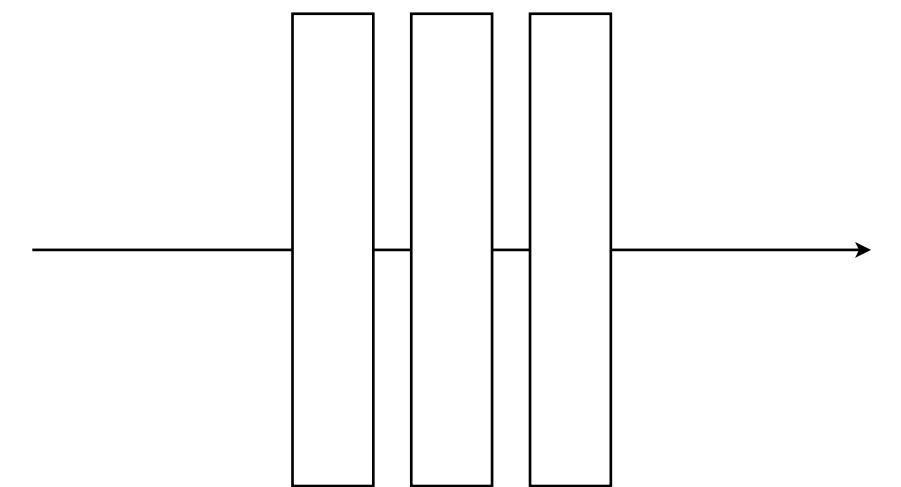
Usually loss functions check if output matches a target *instance*

GAN loss checks if output is part of an admissible set

x



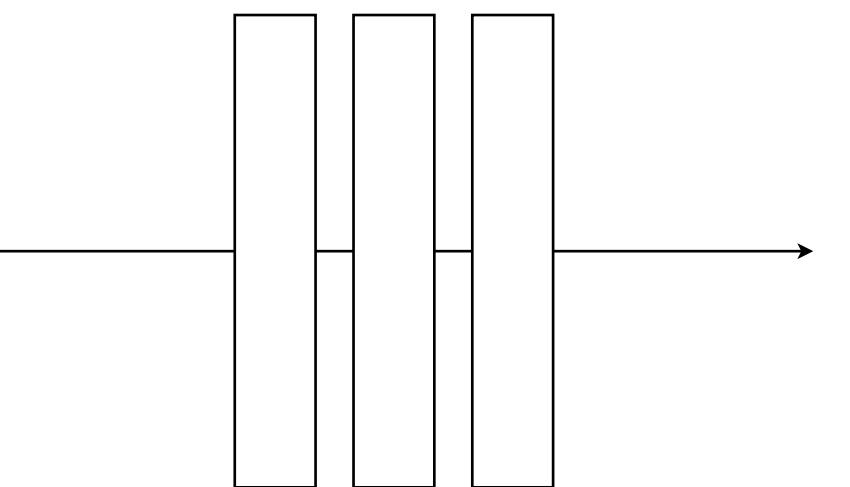
G



G(x)



D

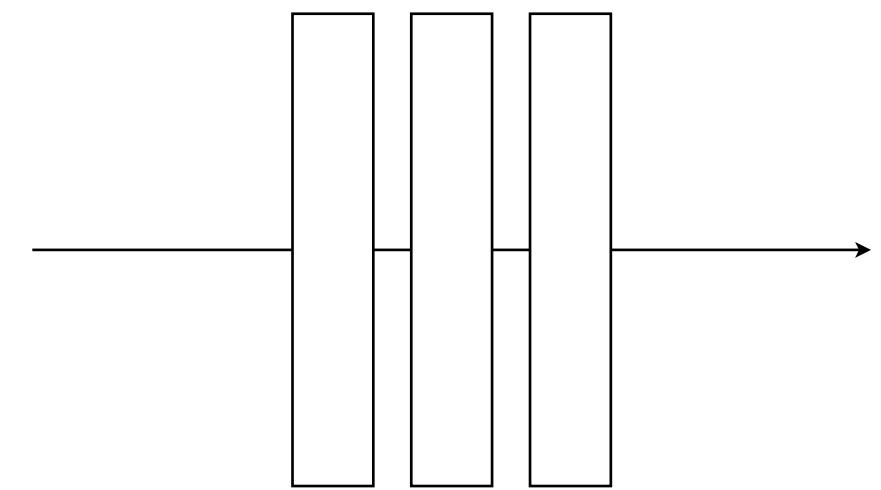


Real!

x



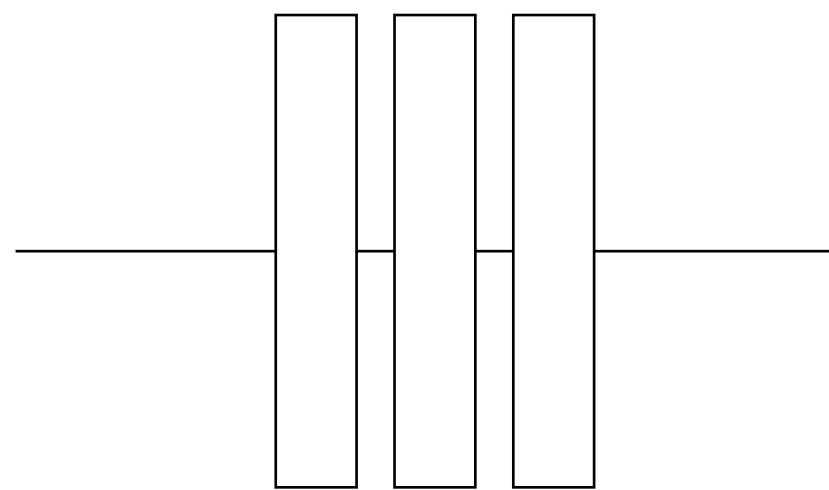
G



$G(x)$



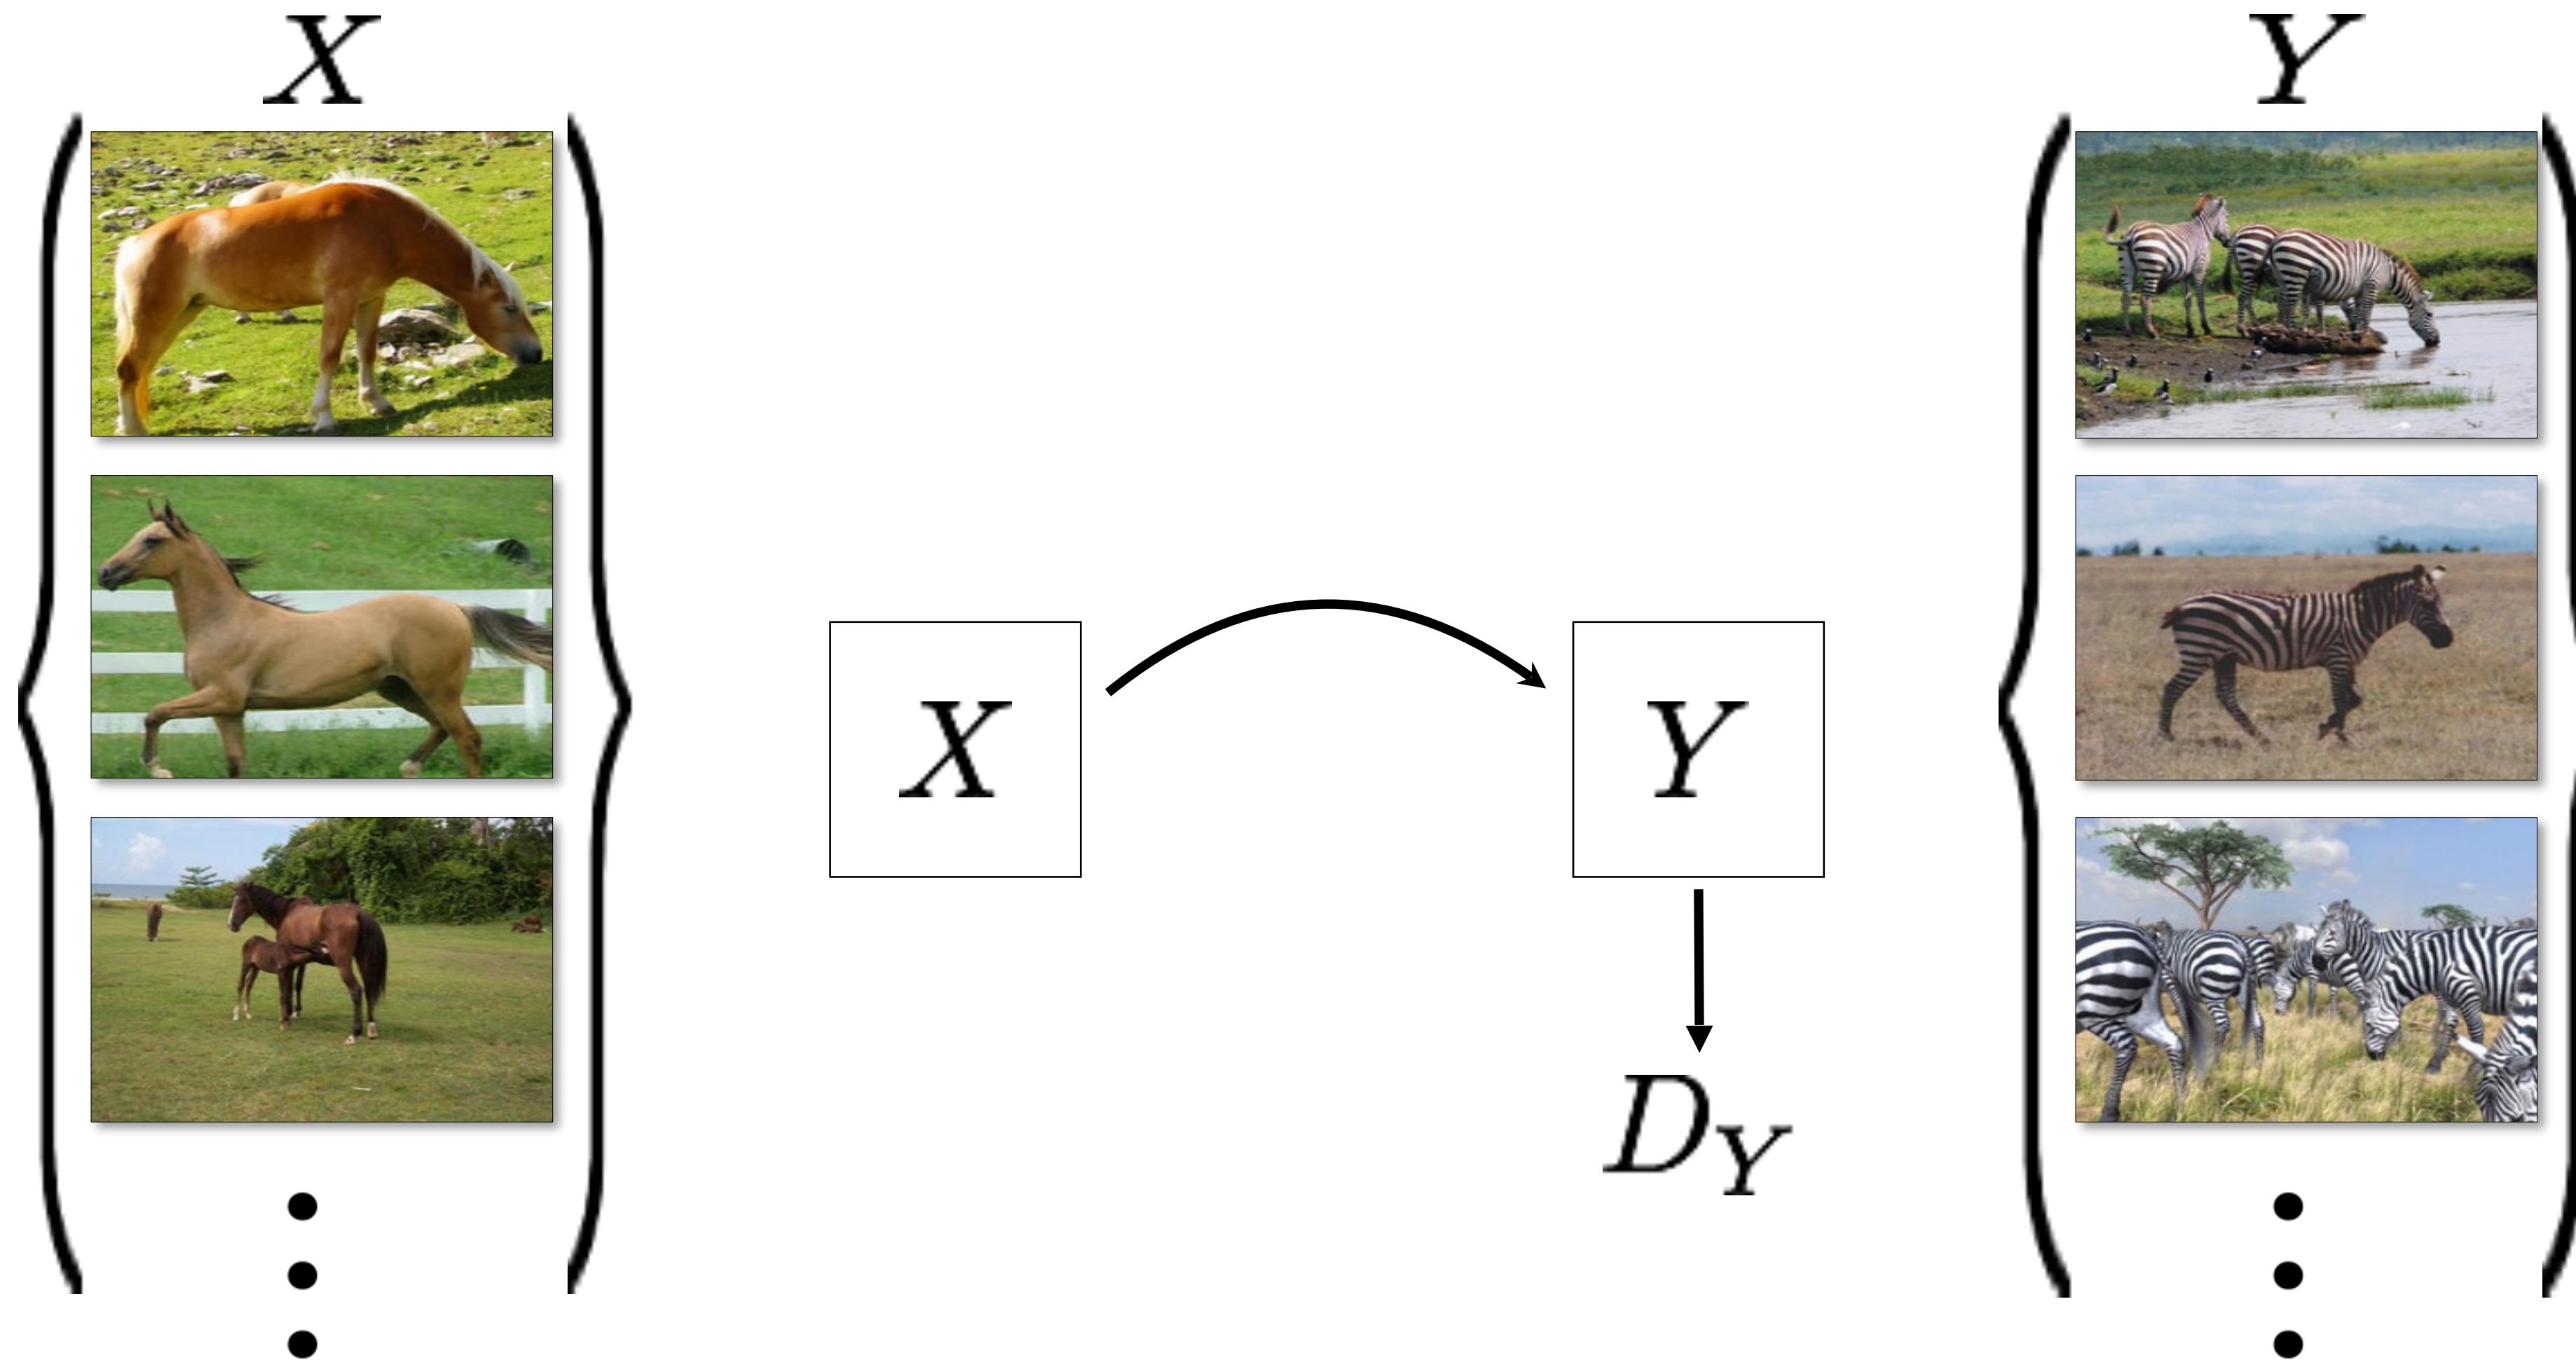
D



Real too!

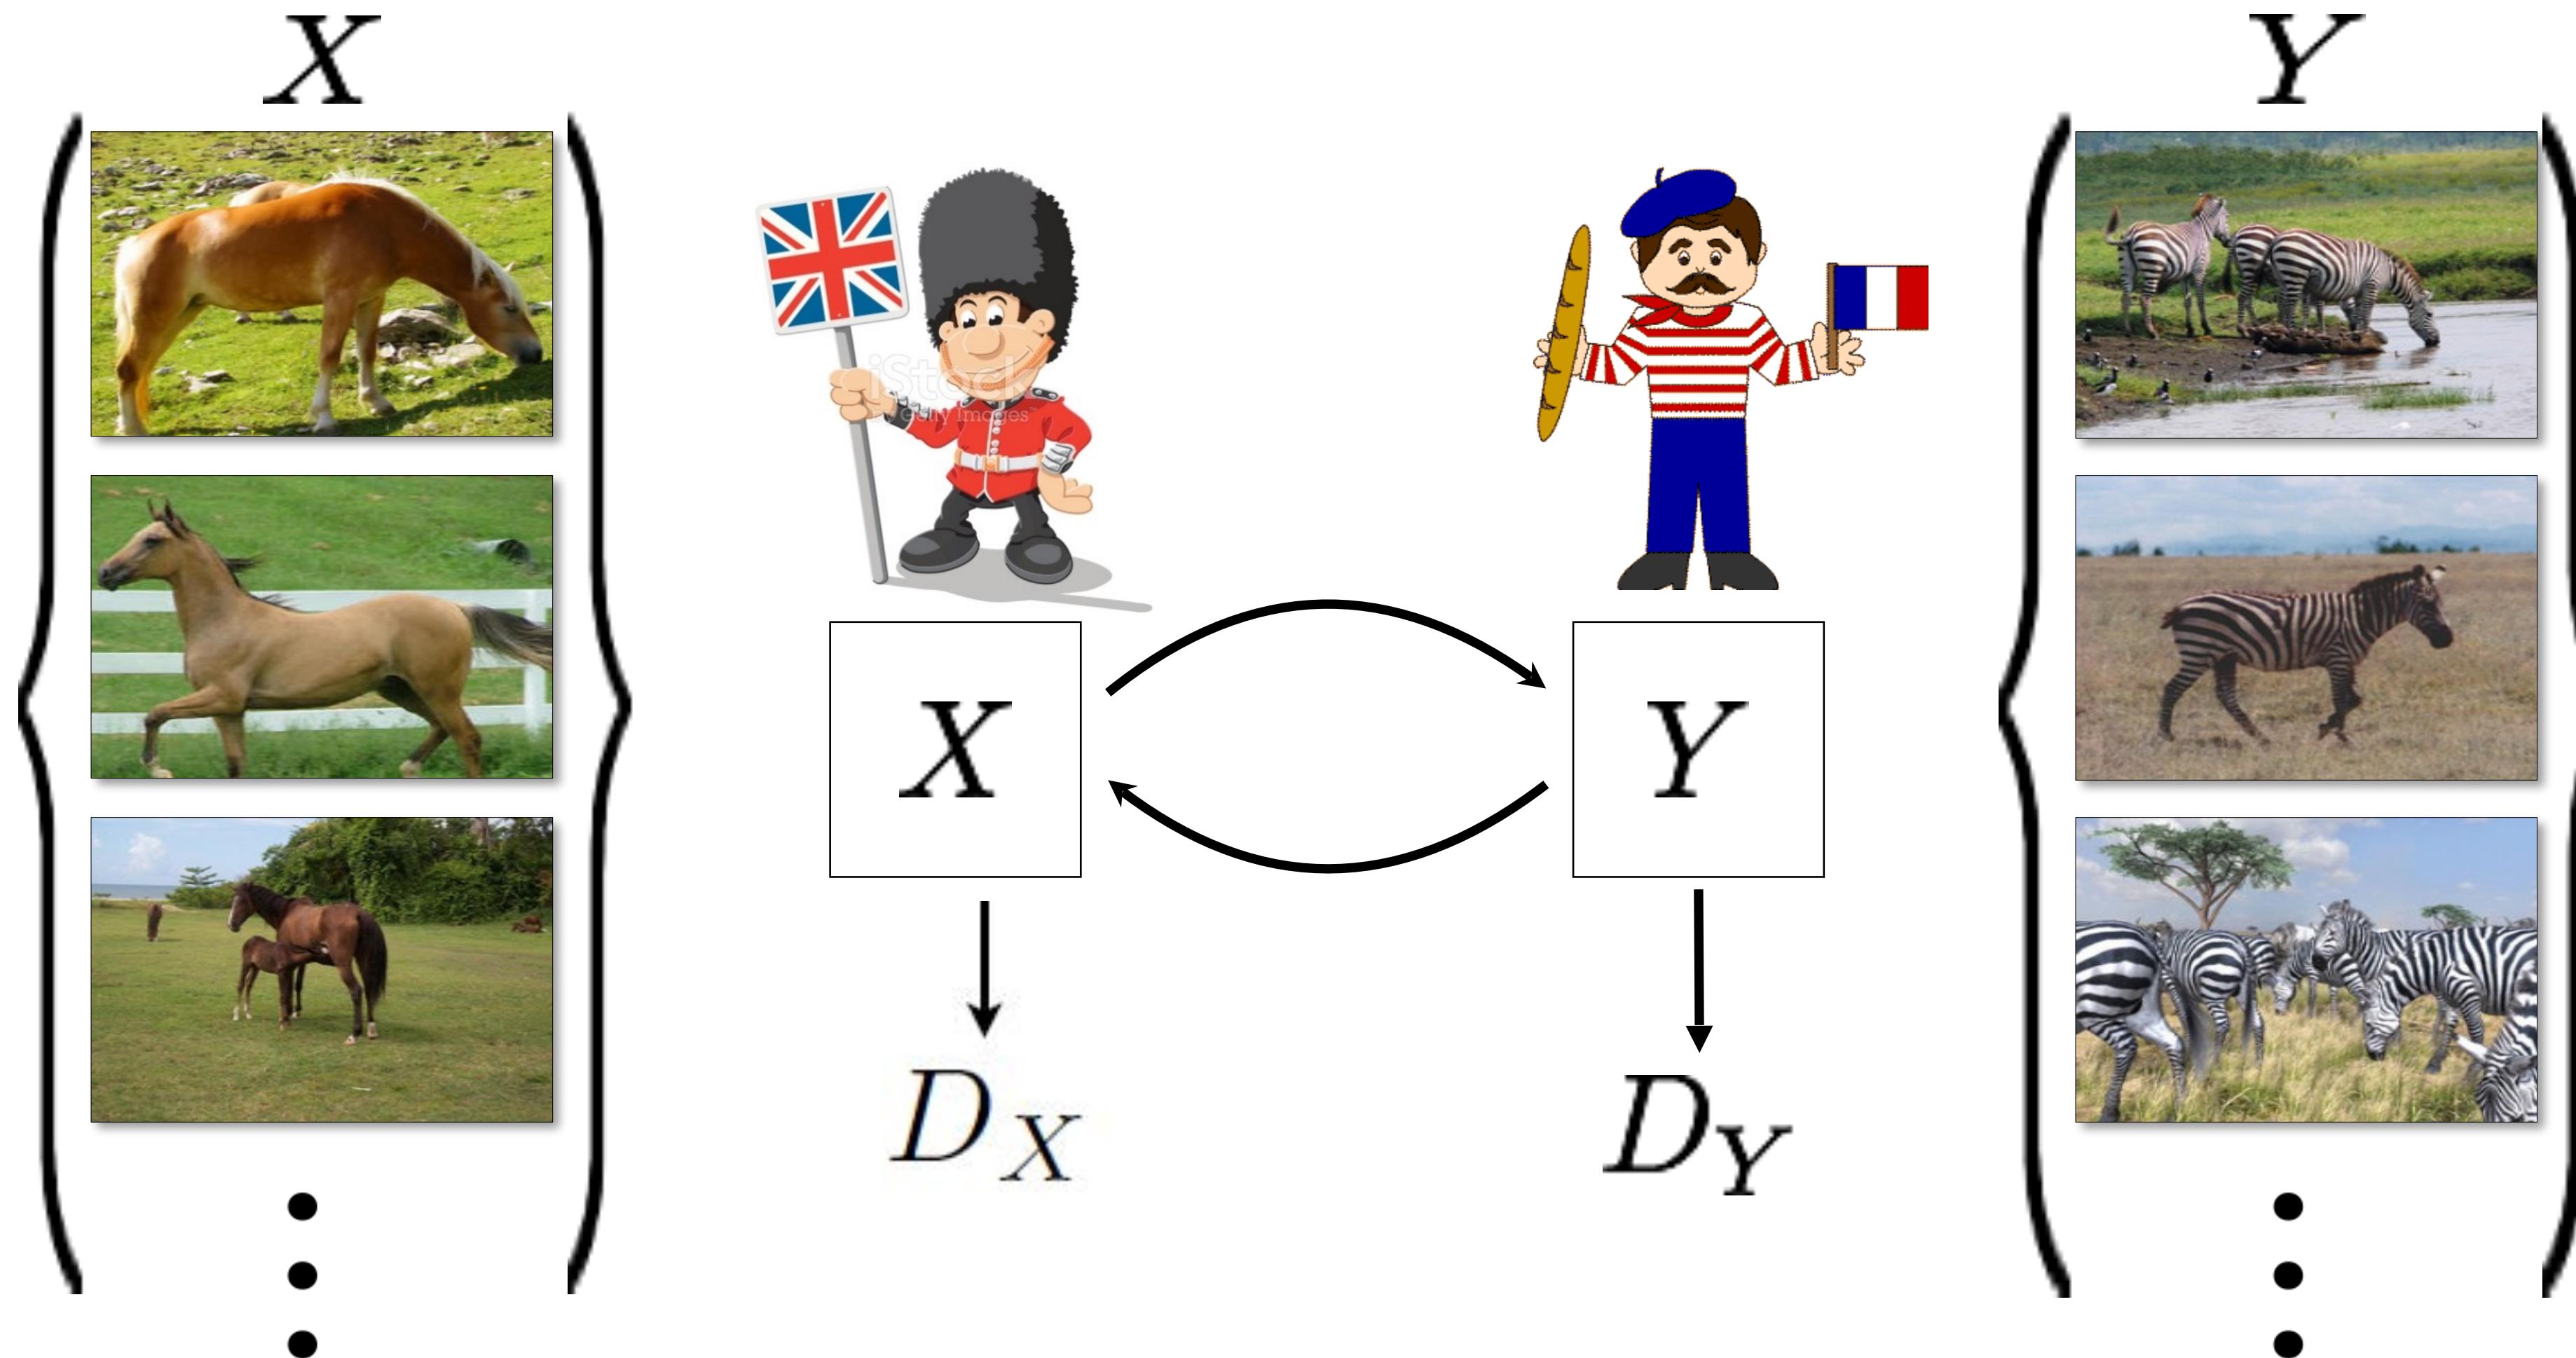
Nothing to force output to correspond to input

Cycle-Consistent Adversarial Networks

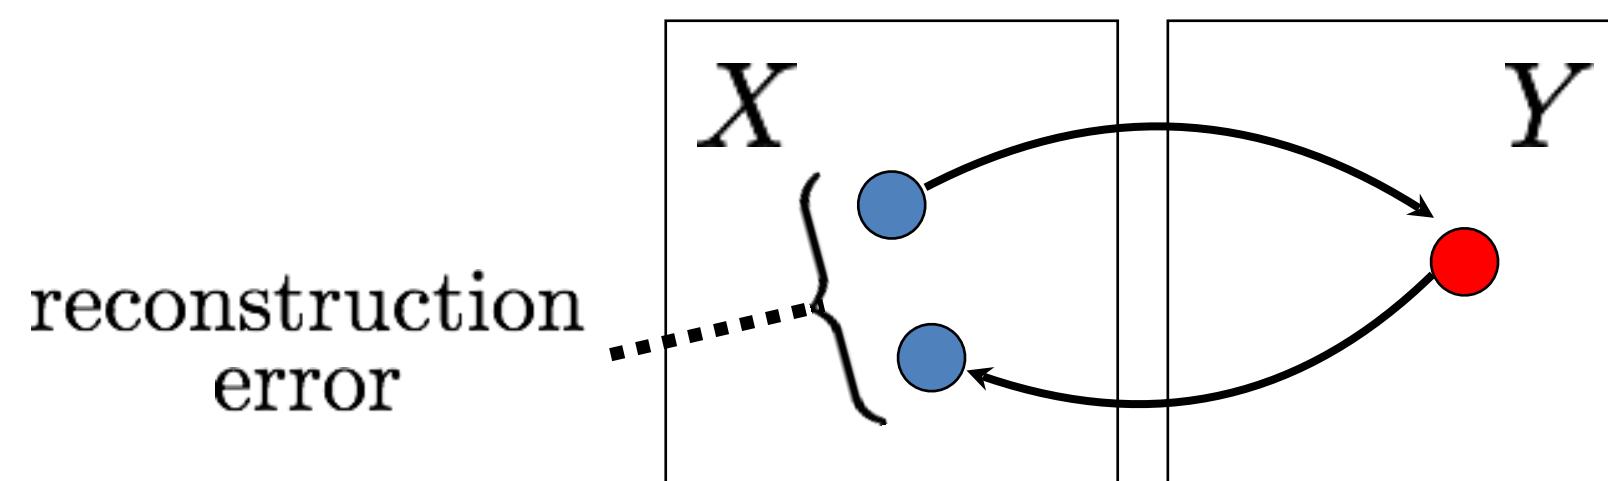
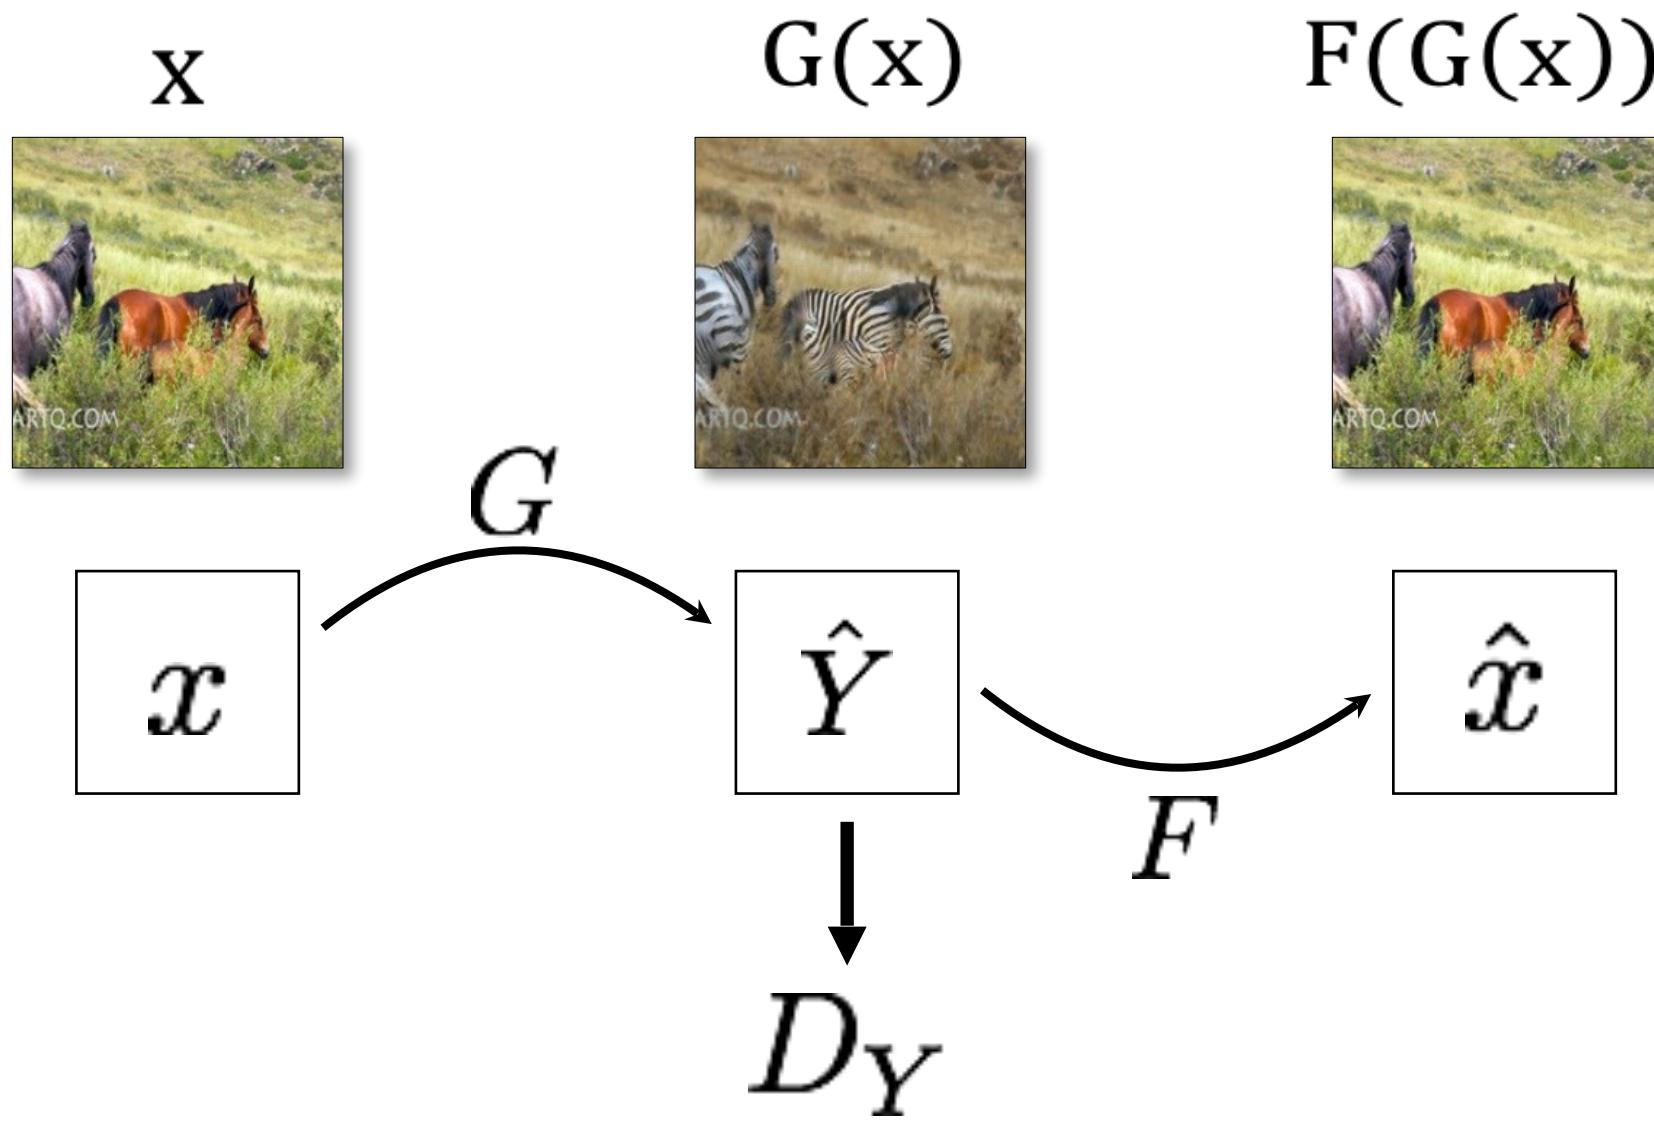


[Zhu et al. 2017], [Yi et al. 2017], [Kim et al. 2017]

Cycle-Consistent Adversarial Networks

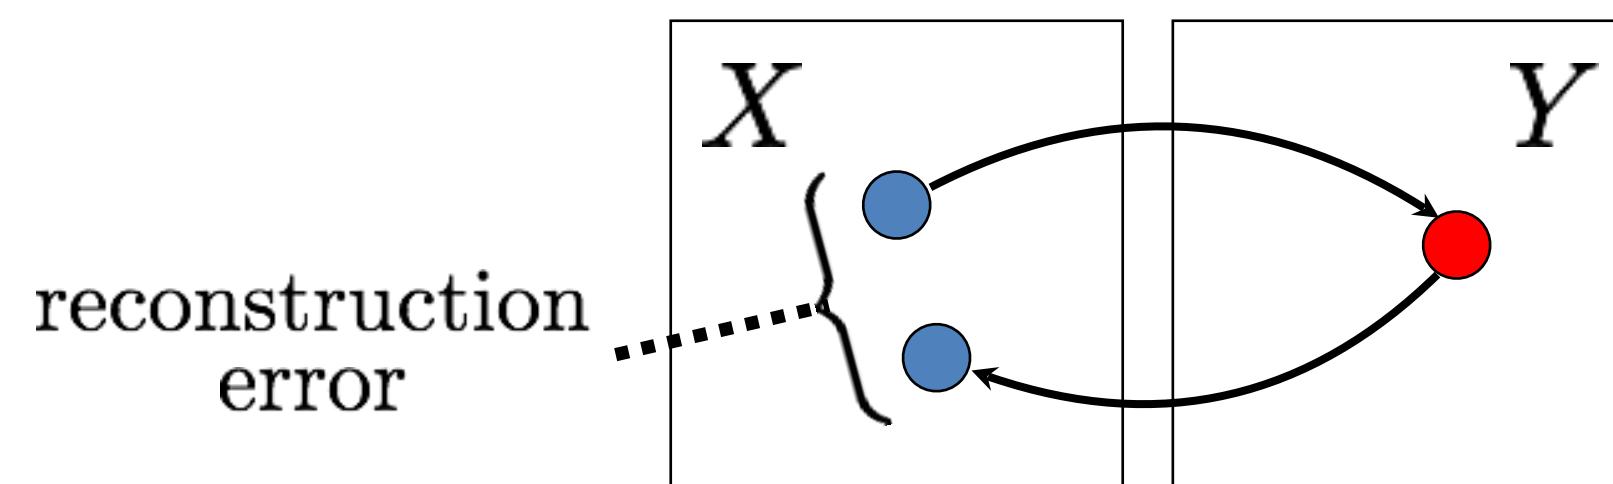
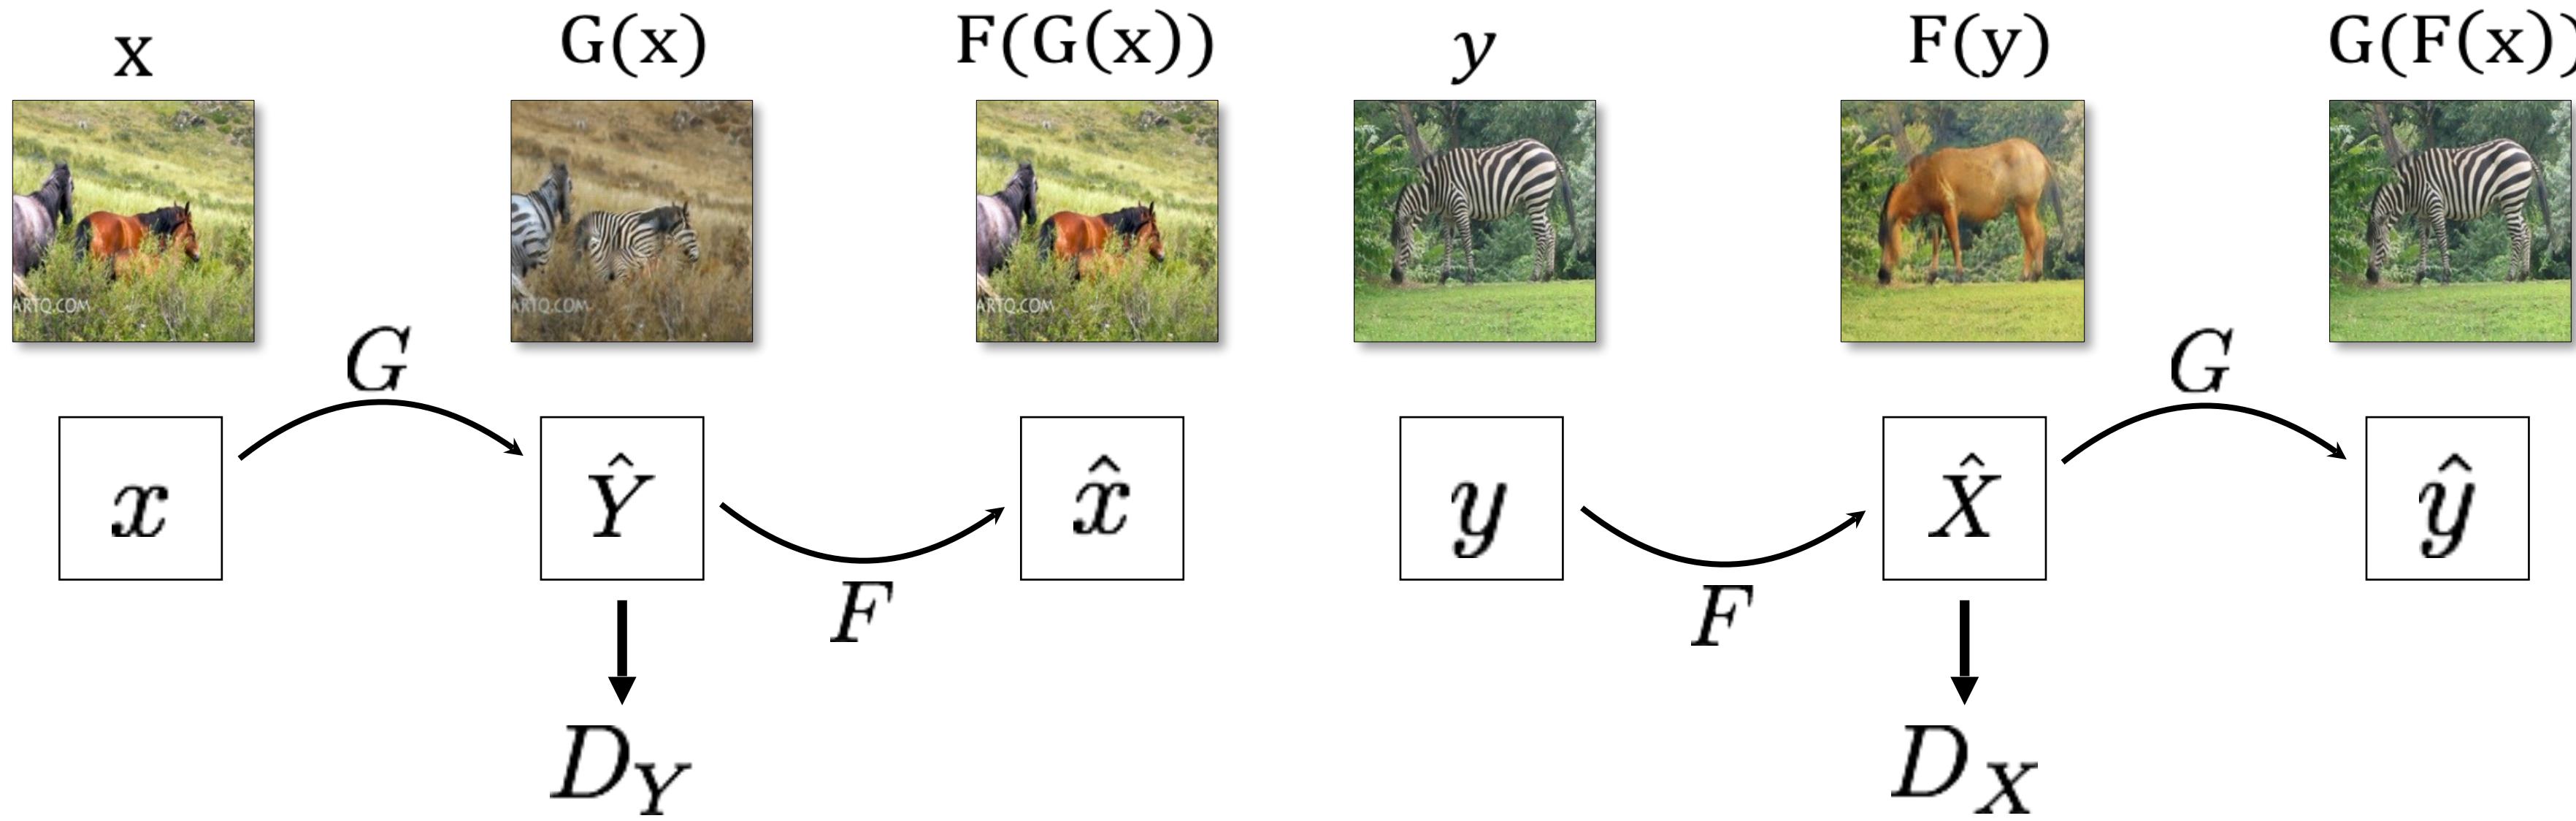


Cycle Consistency Loss

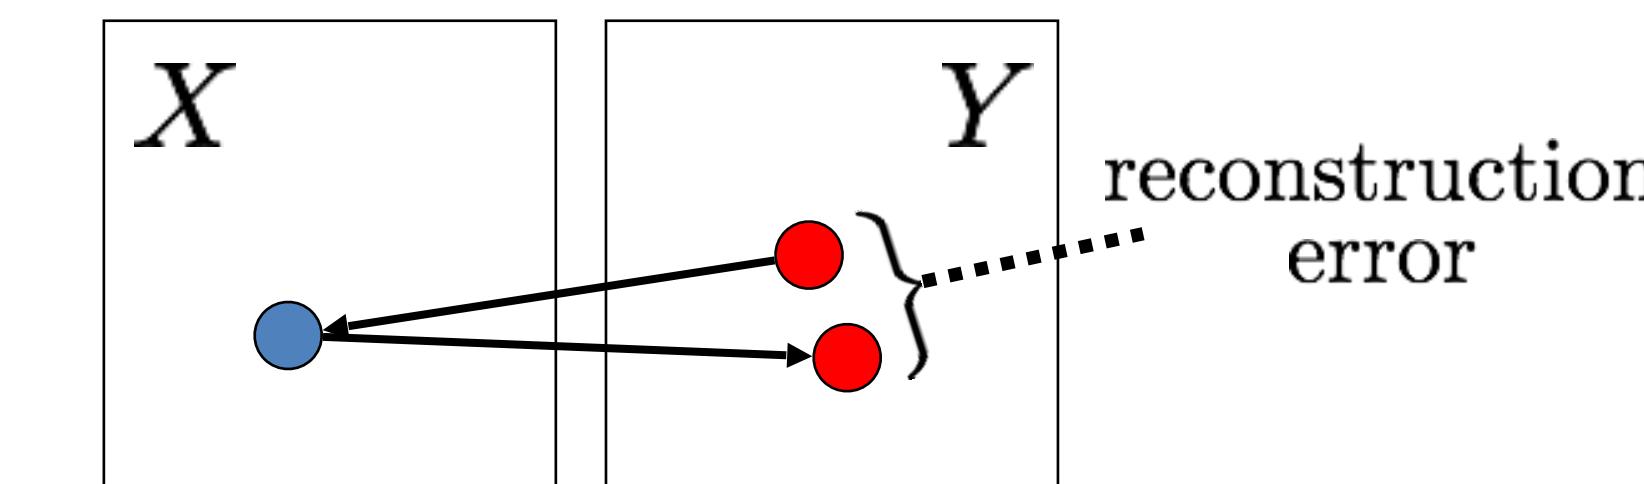


$$\|F(G(x)) - x\|_1$$

Cycle Consistency Loss



$$\|F(G(x)) - x\|_1$$



$$\|G(F(y)) - y\|_1$$





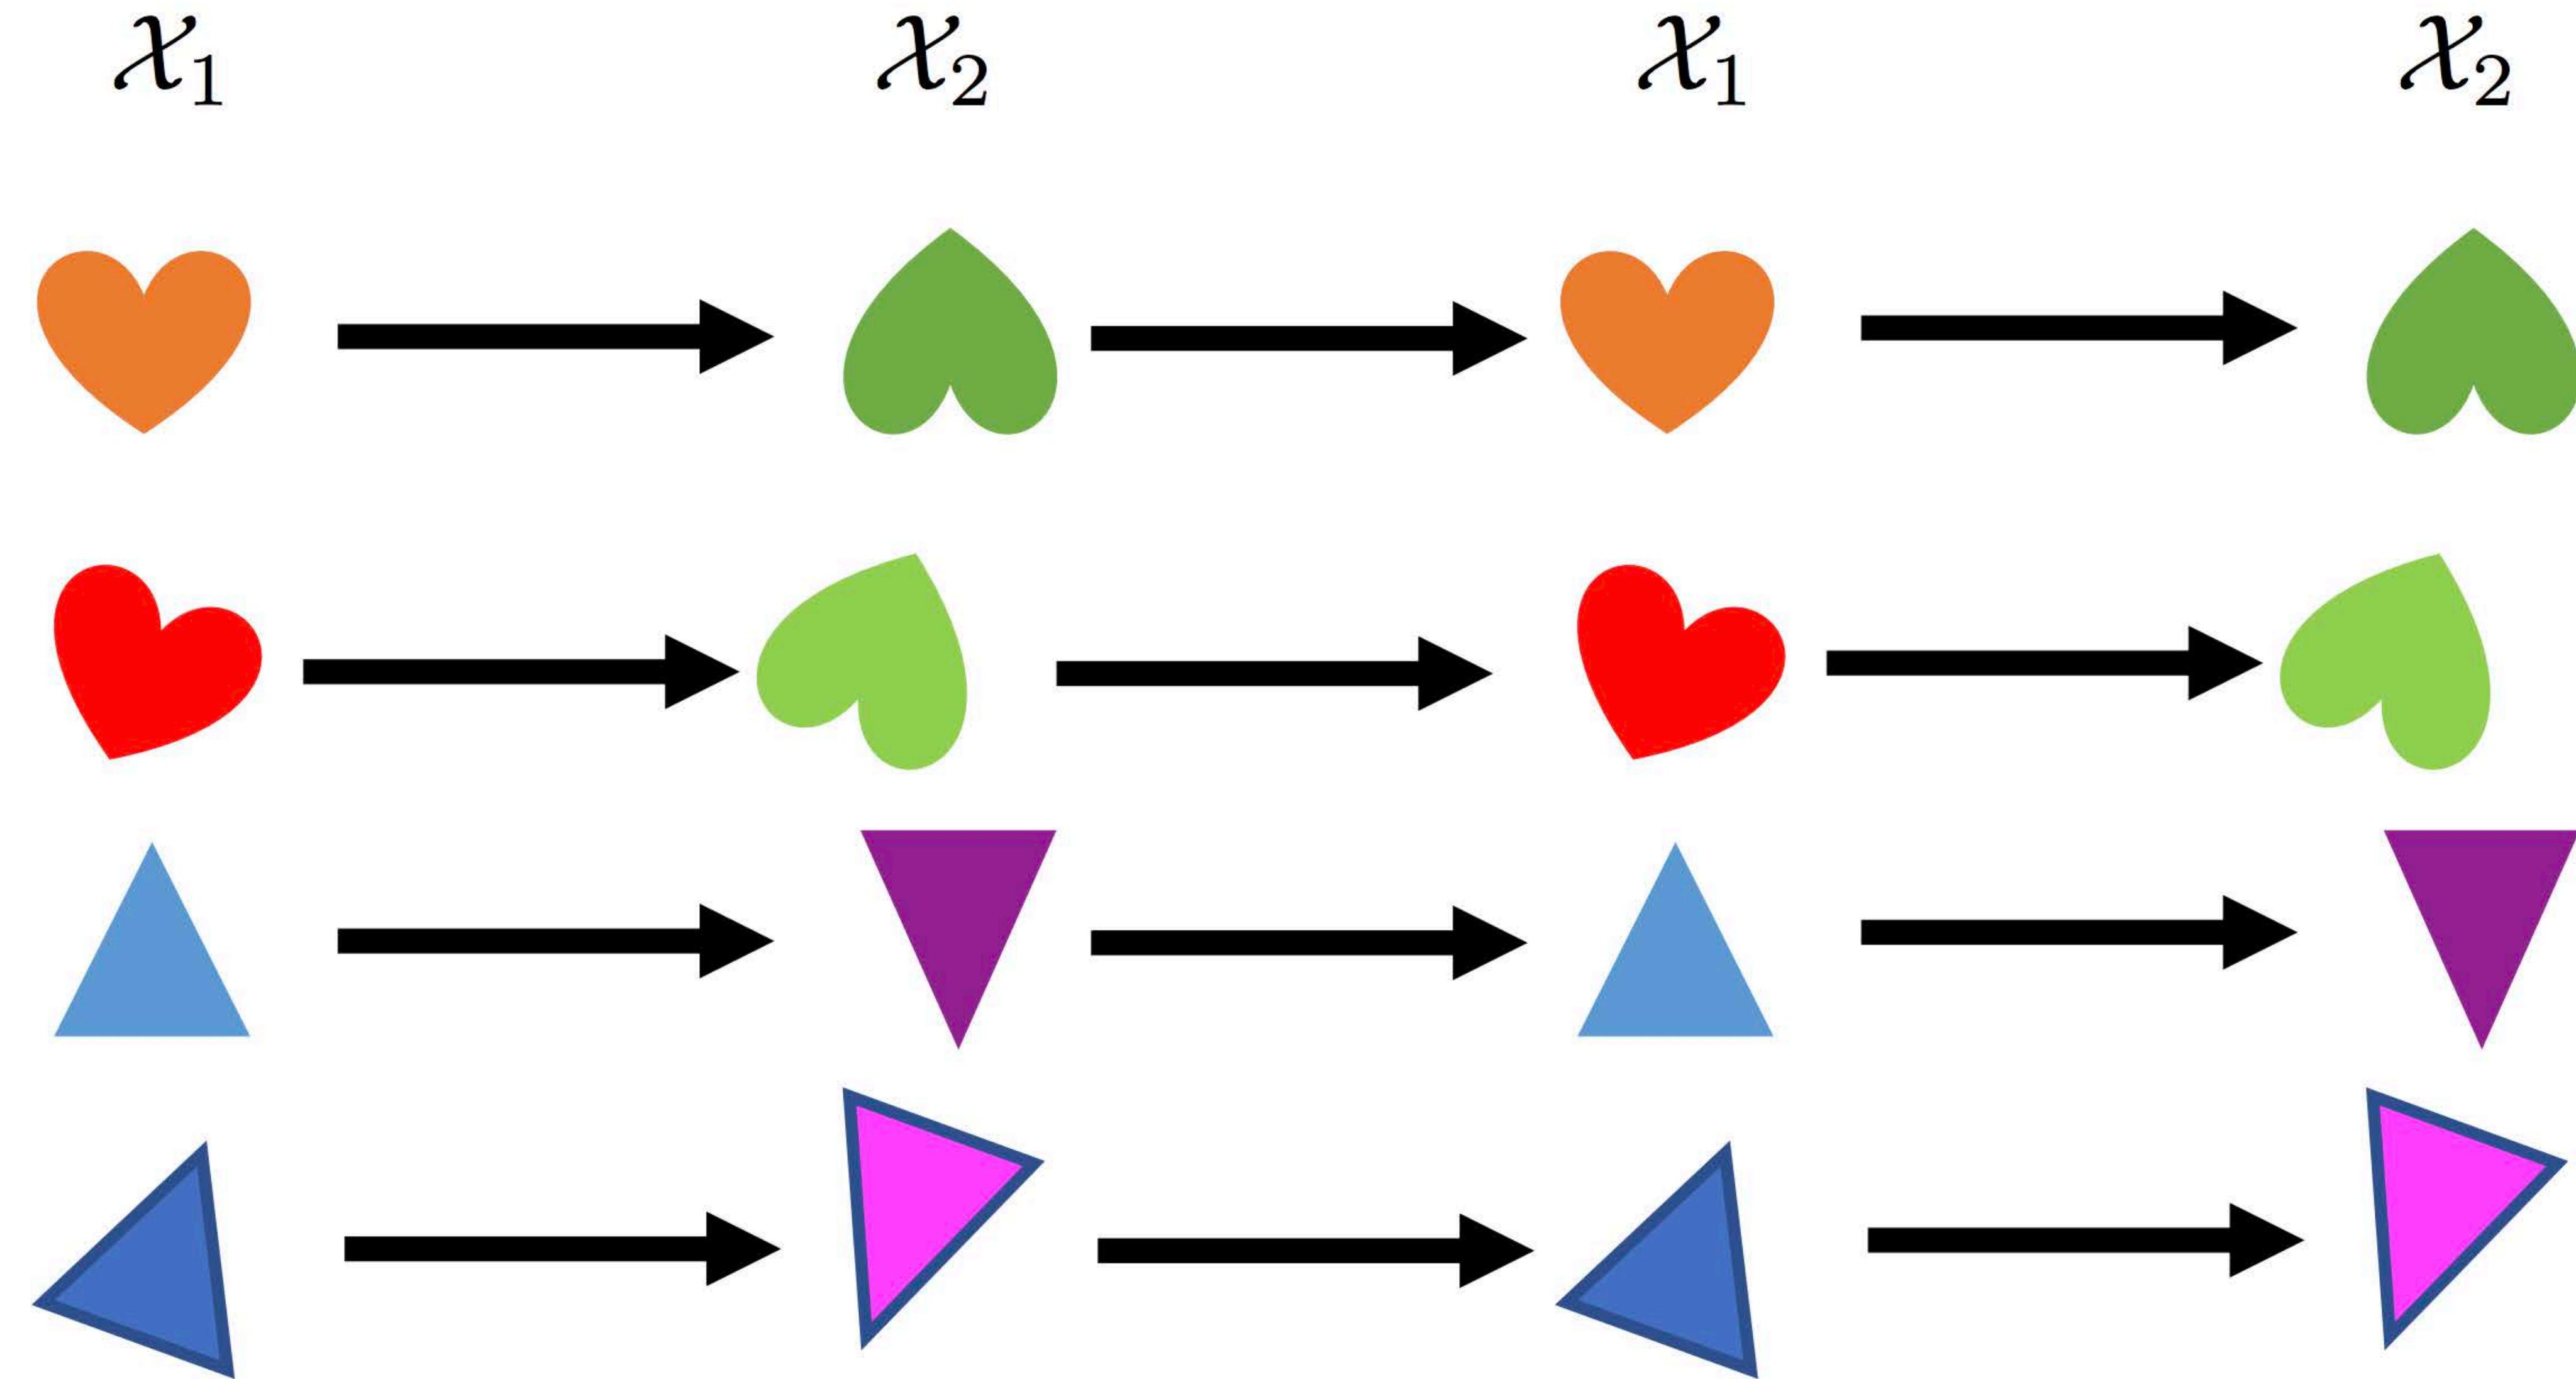
Failure case



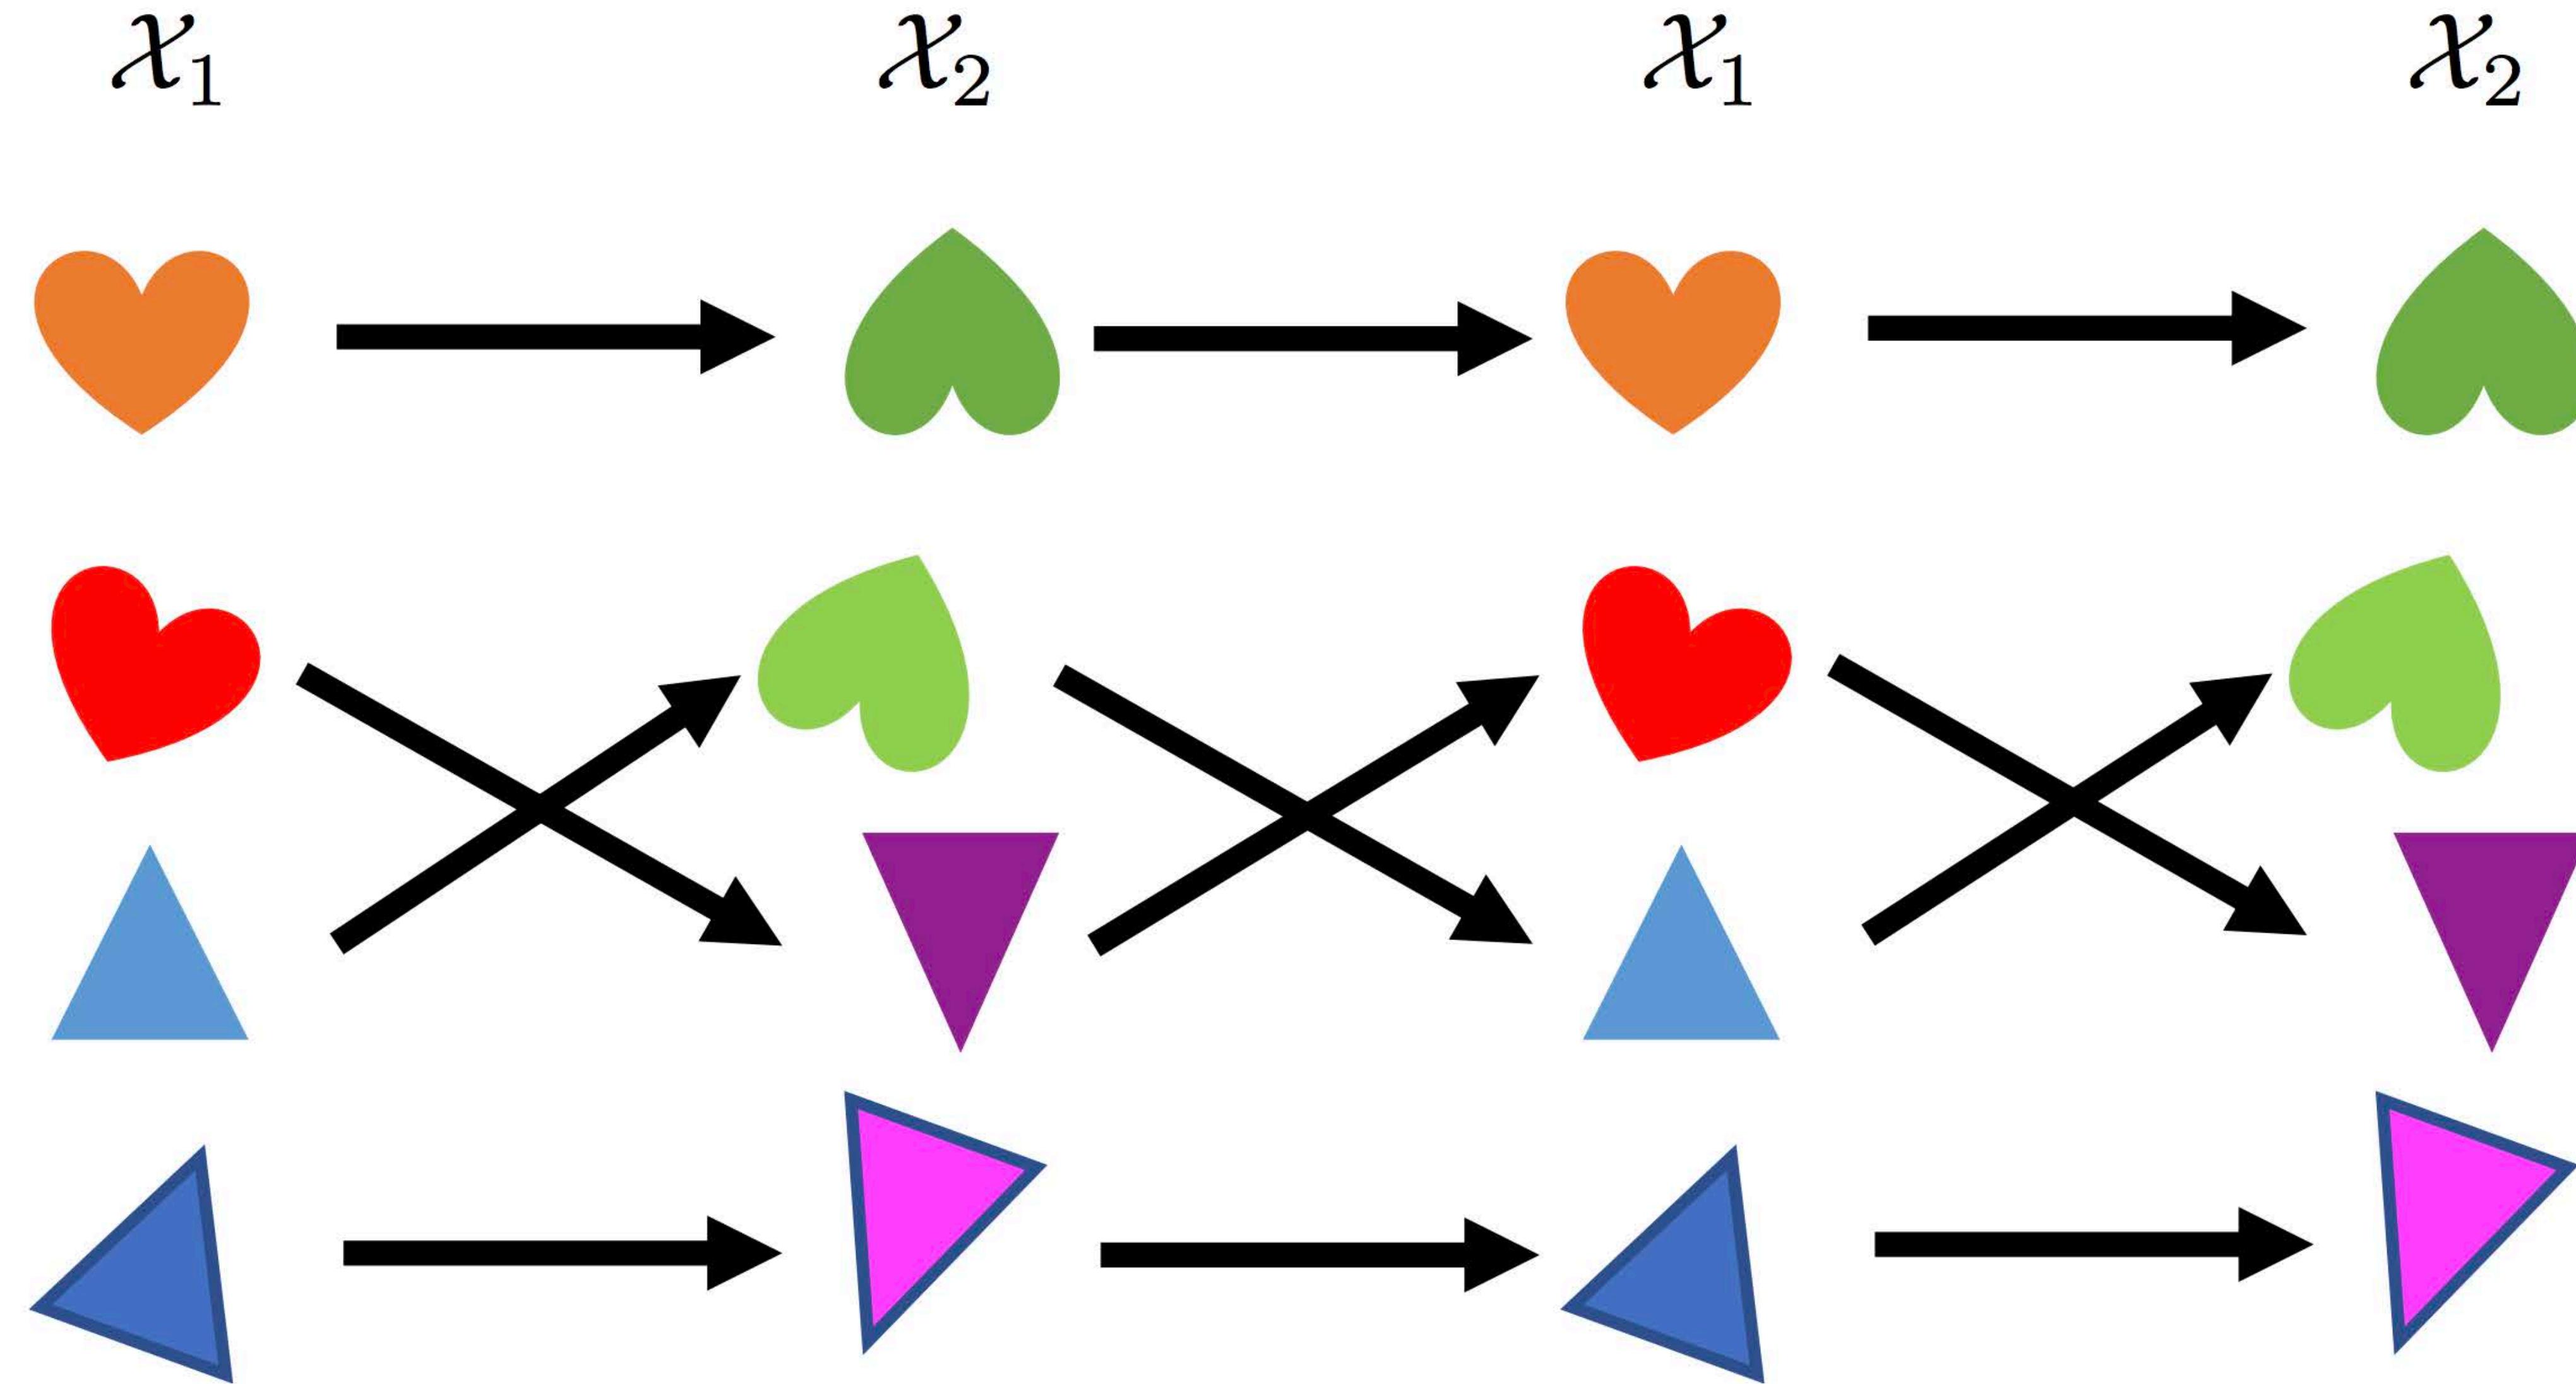
Failure case



Why does CycleGAN work?



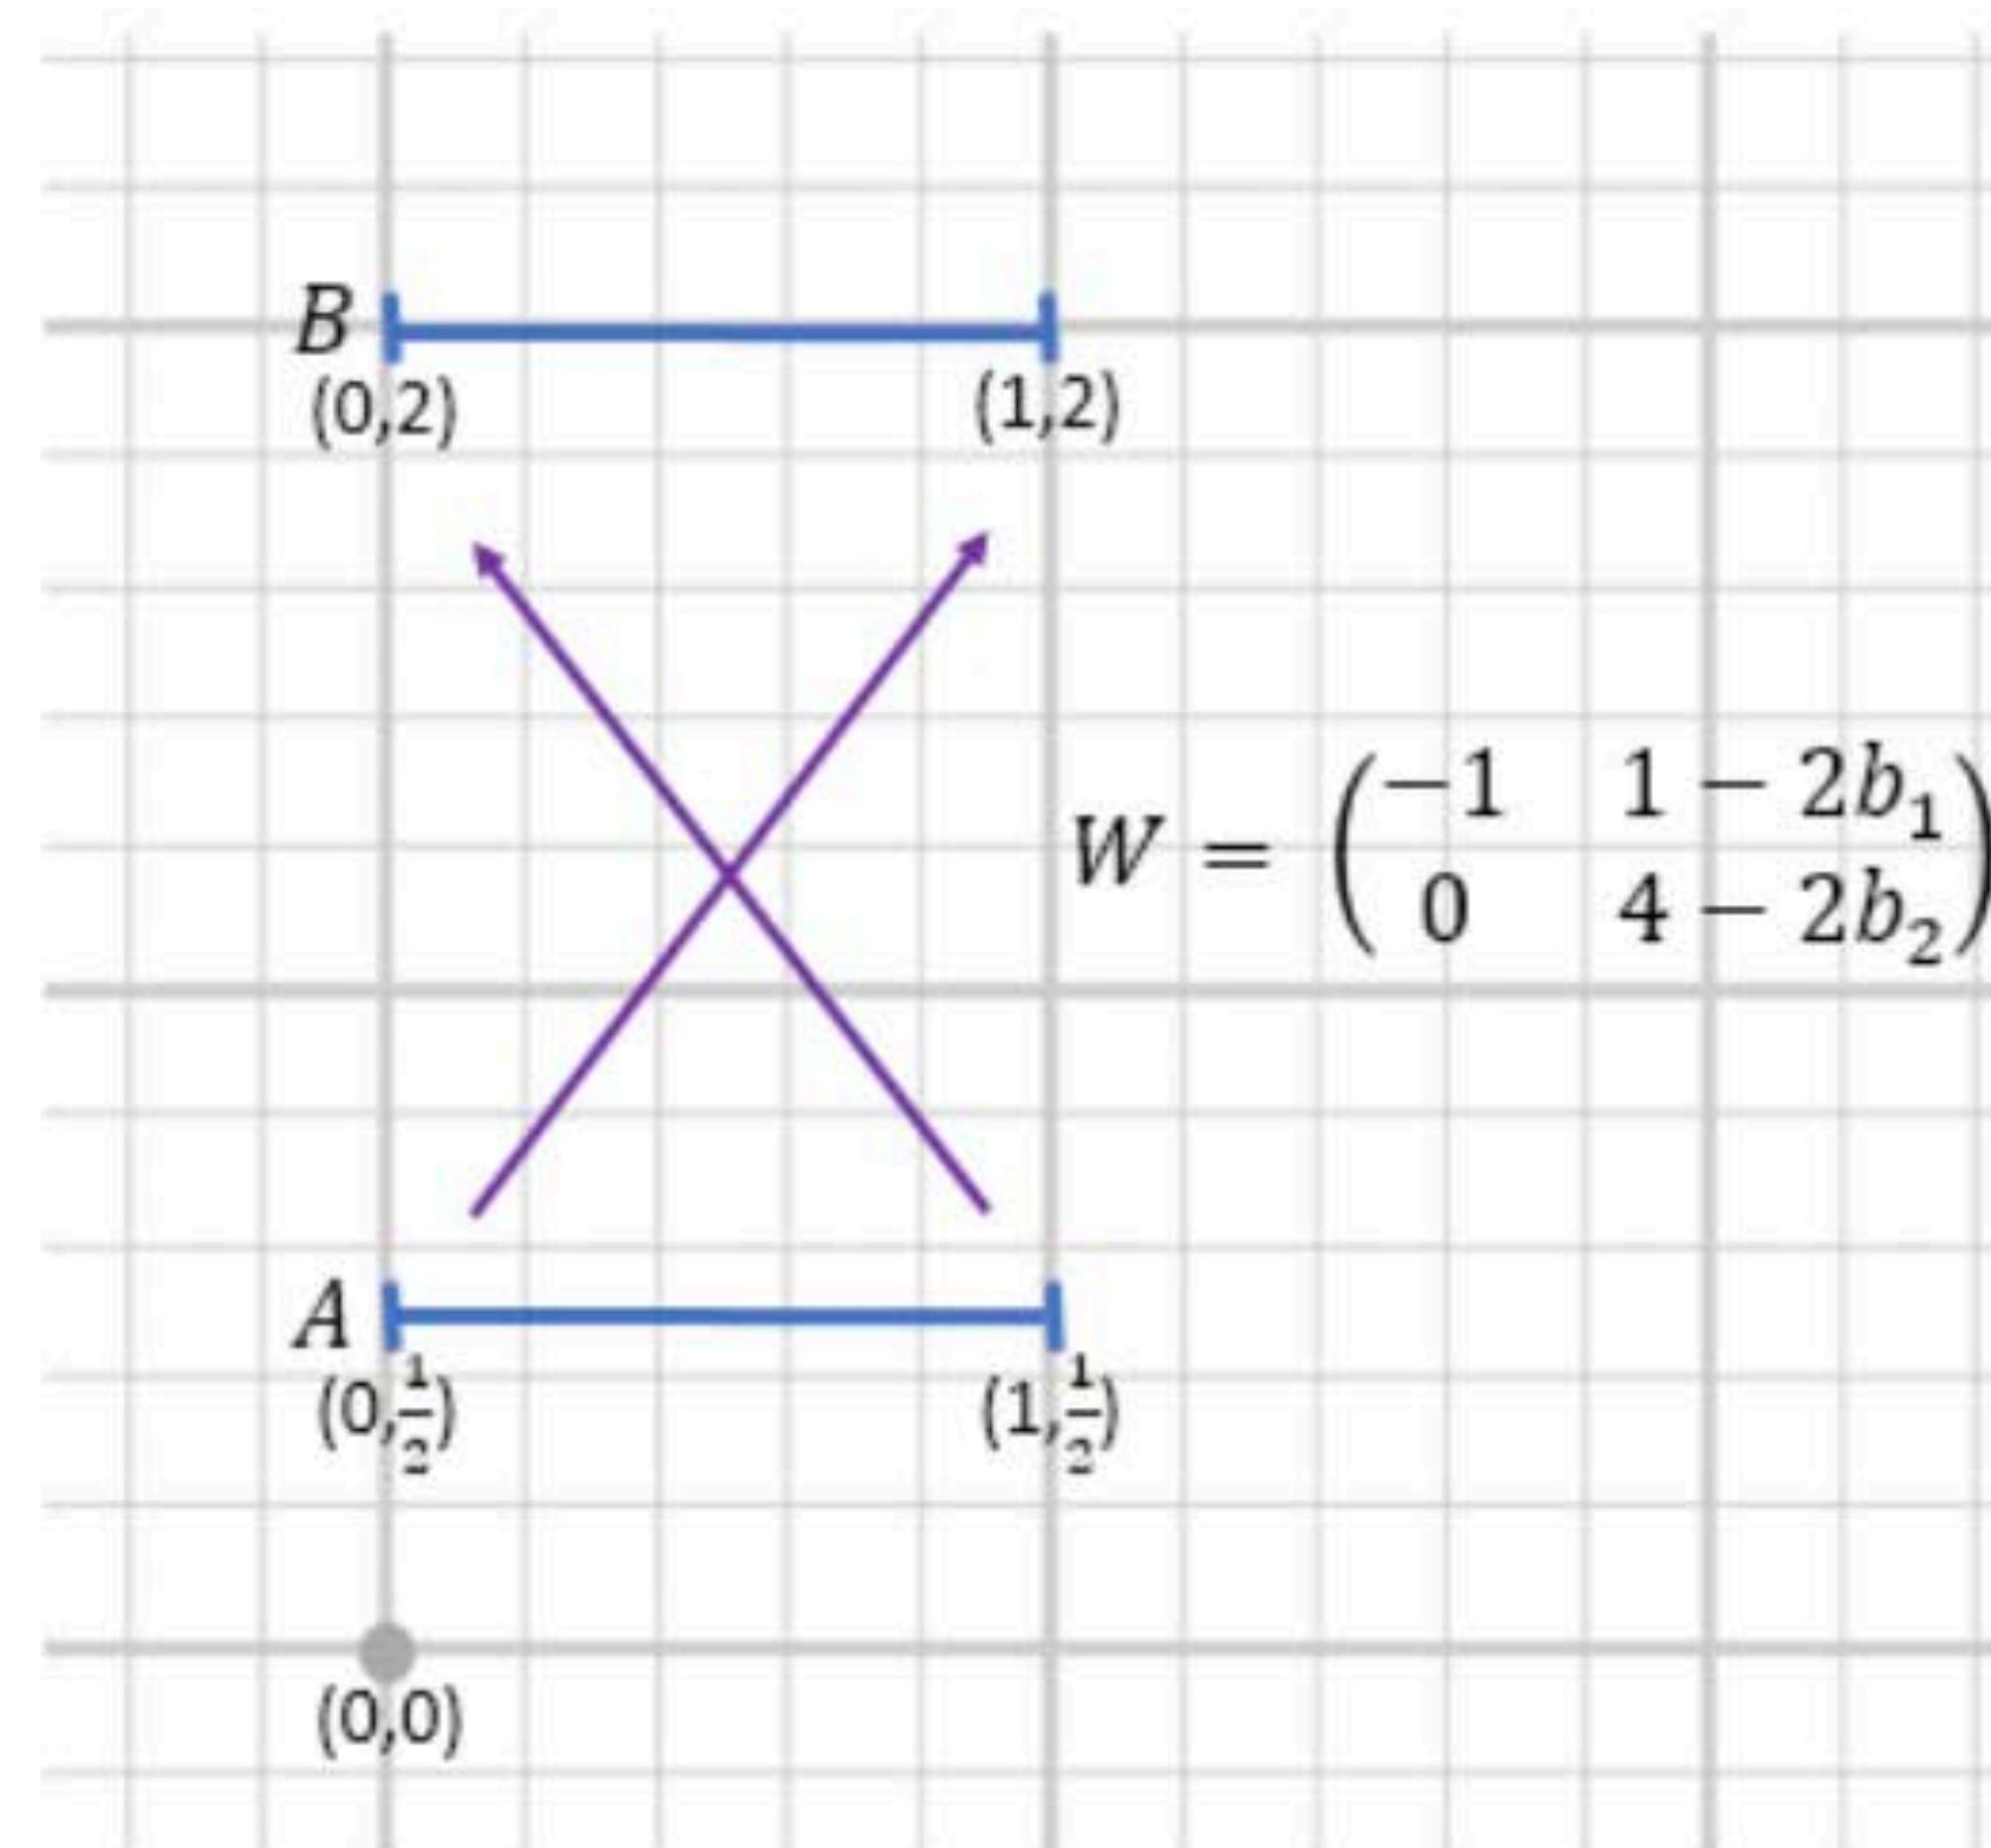
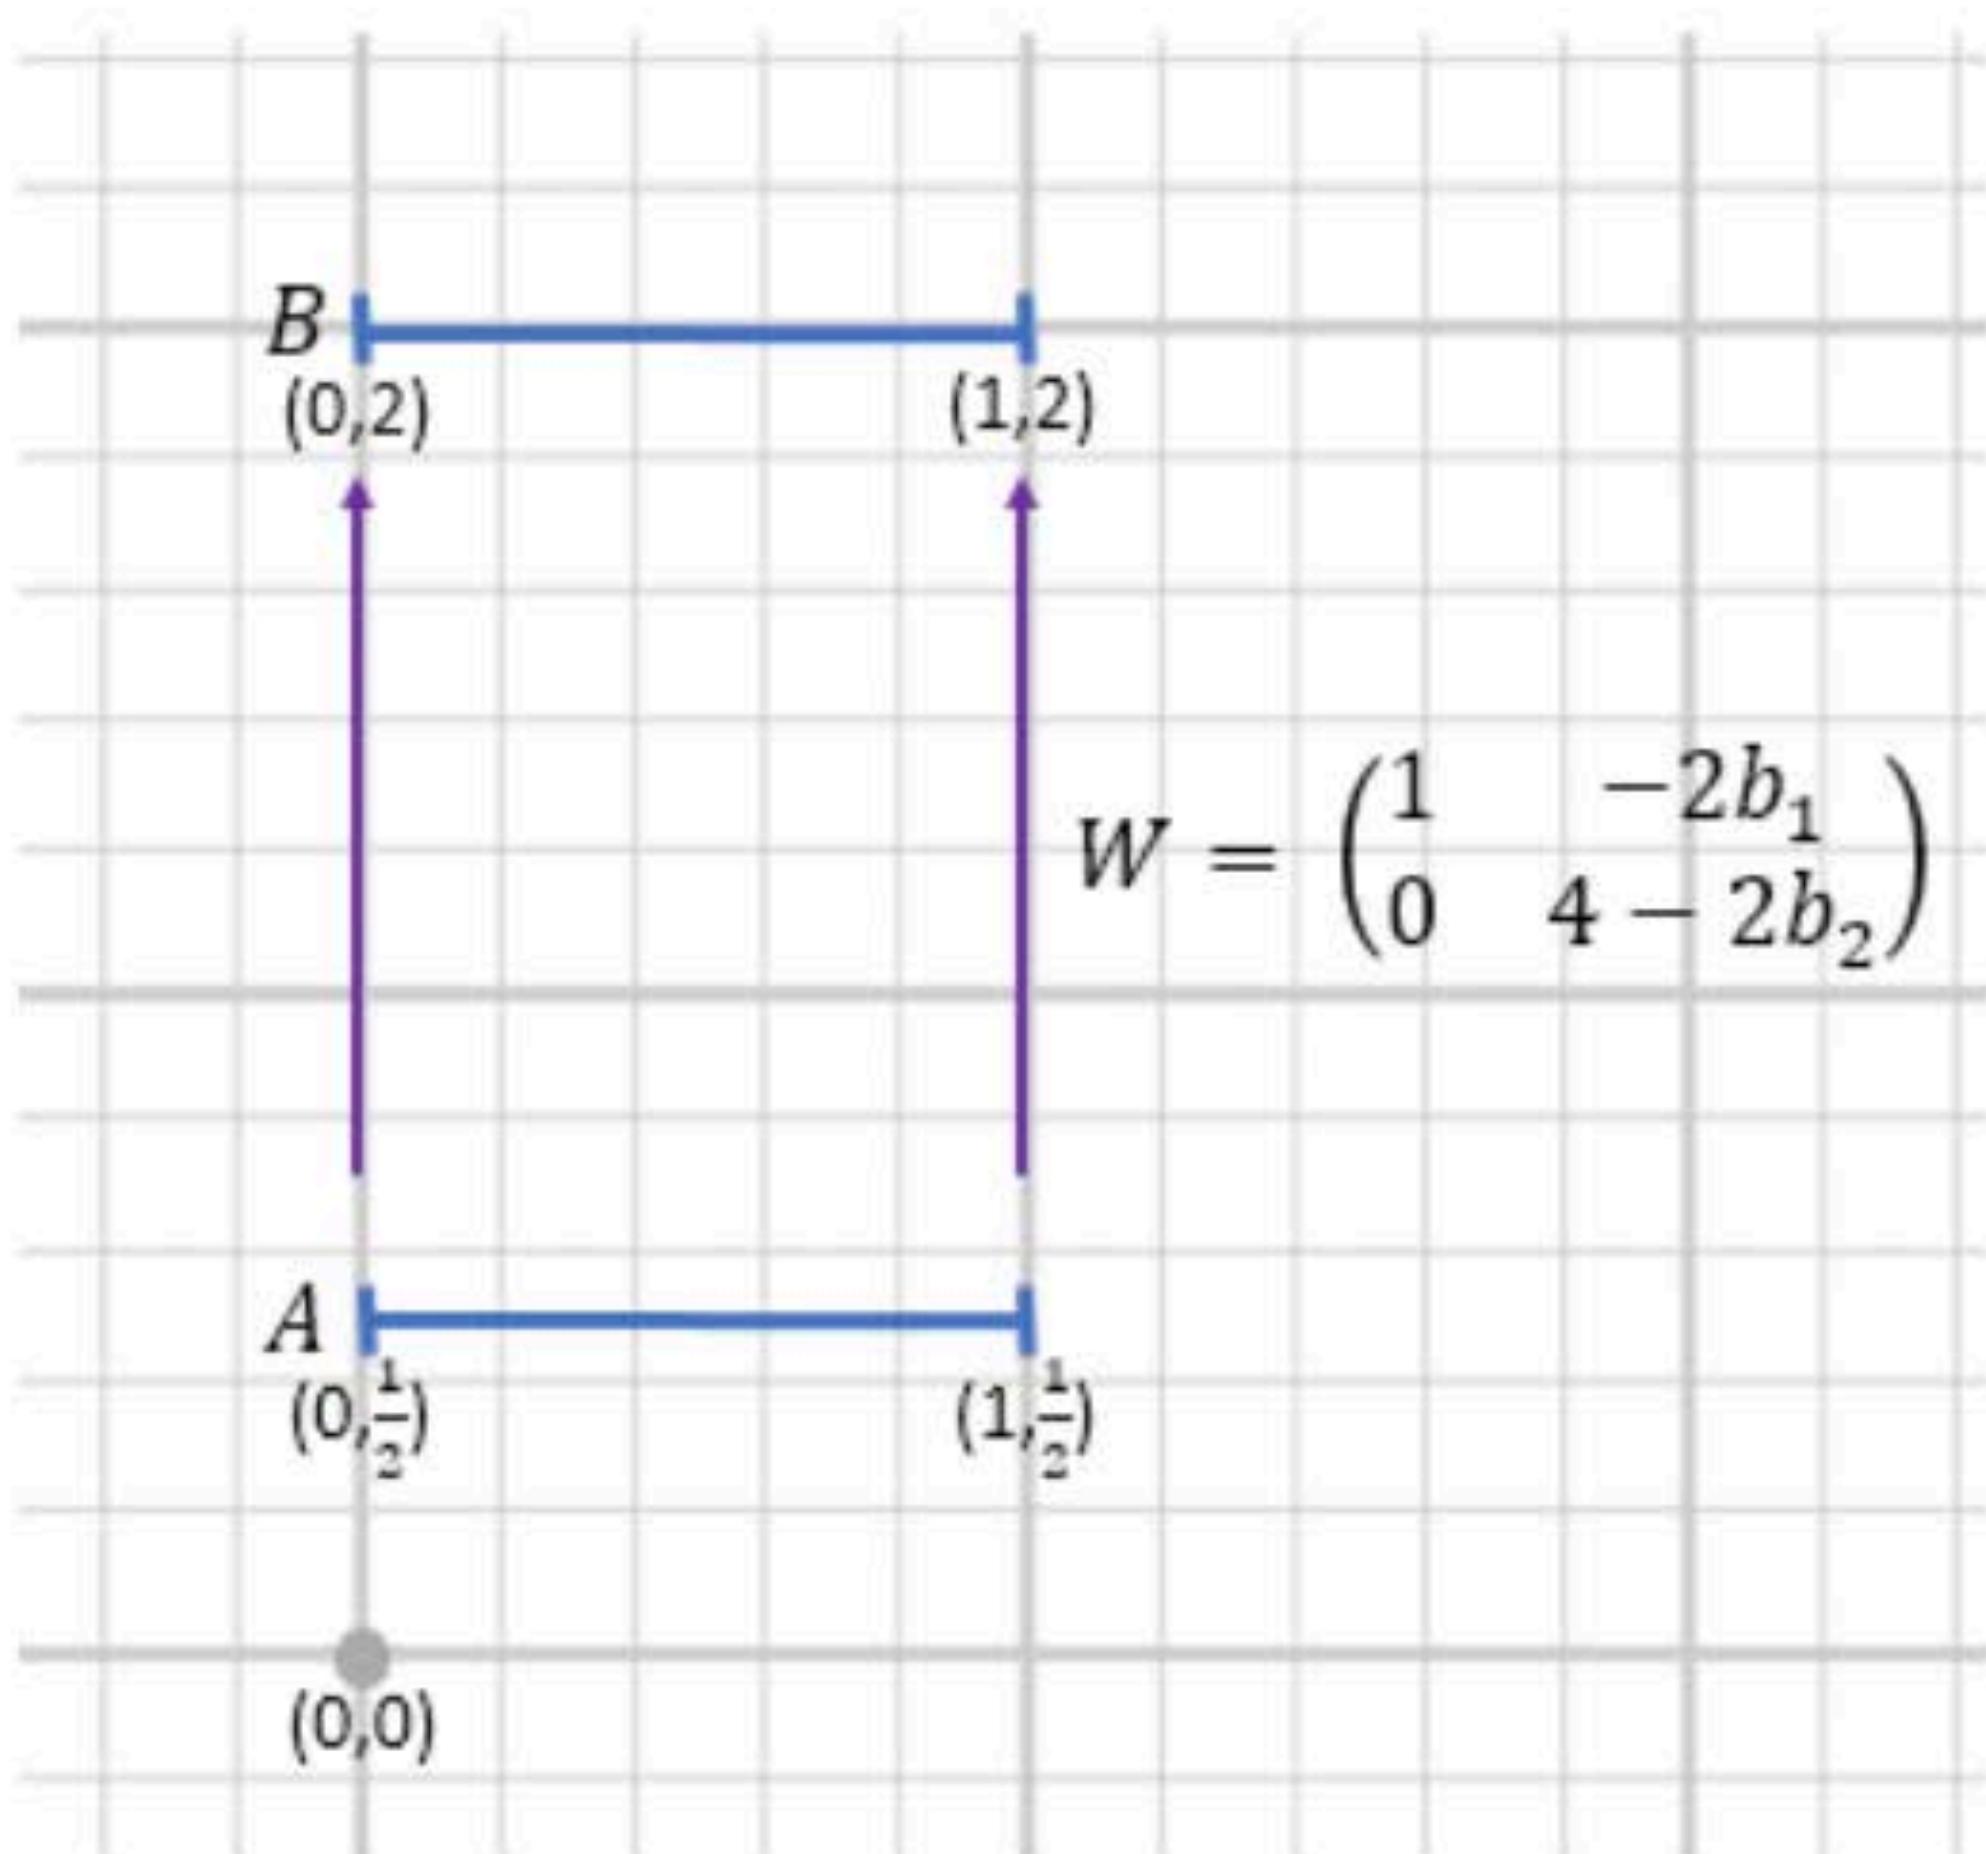
Slide credit: Ming-Yu Liu



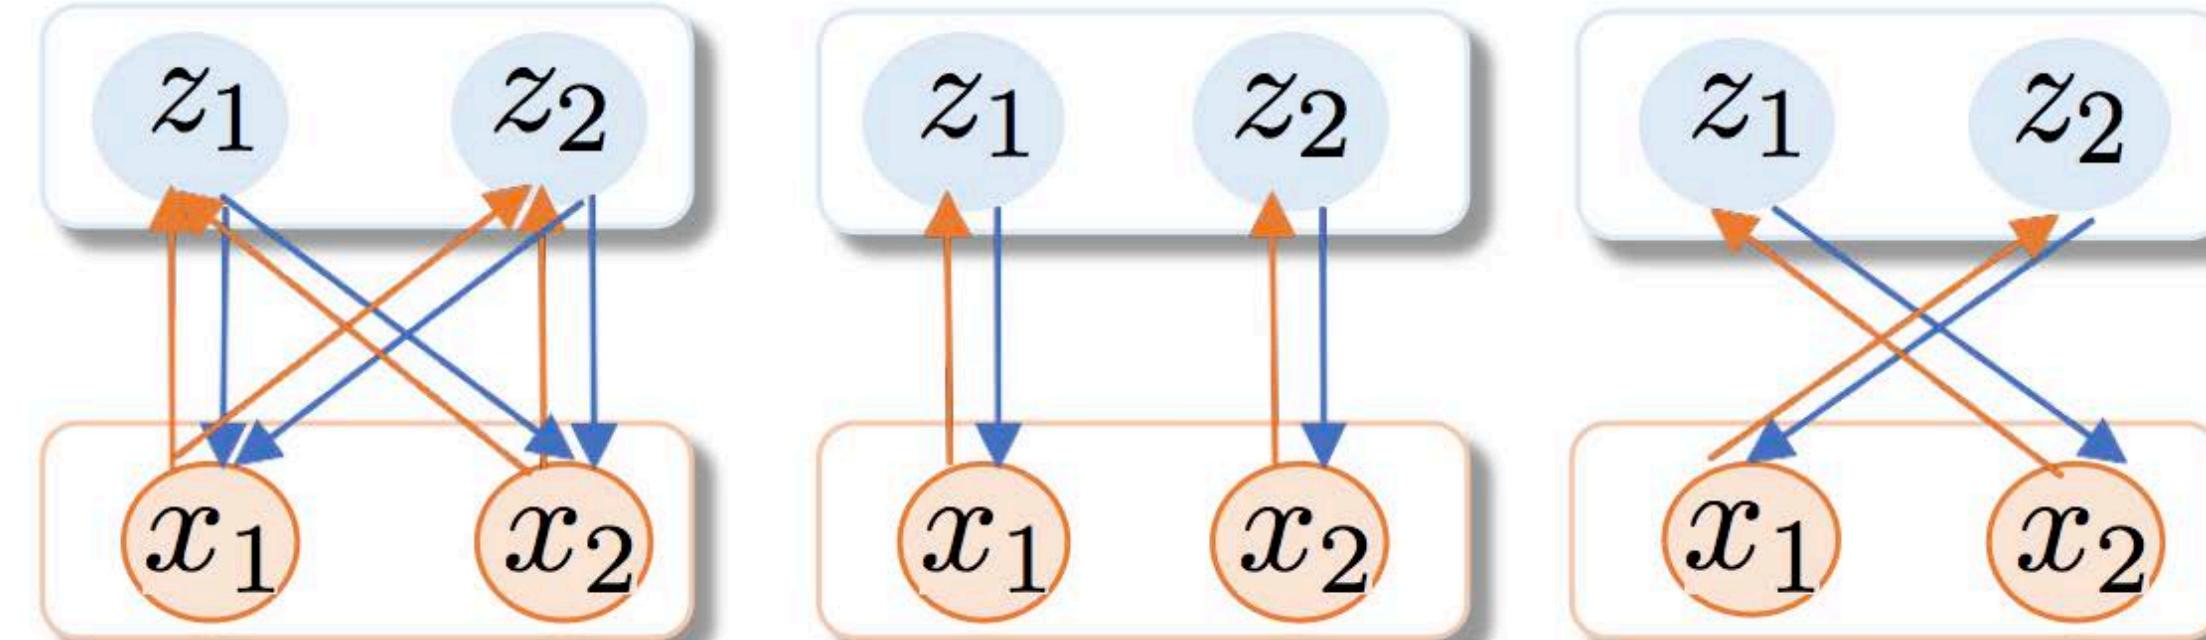
Slide credit: Ming-Yu Liu

Simplicity hypothesis

[Galanti, Wolf, Benaim, 2018]



Cycle Loss upper bounds Conditional Entropy



	z_1	z_2
x_1	$\delta/2$	$(1-\delta)/2$
x_2	$(1-\delta)/2$	$\delta/2$

	z_1	z_2
x_1	$1/2$	0
x_2	0	$1/2$

	z_1	z_2
x_1	0	$1/2$
x_2	$1/2$	0

High
Conditional
Entropy

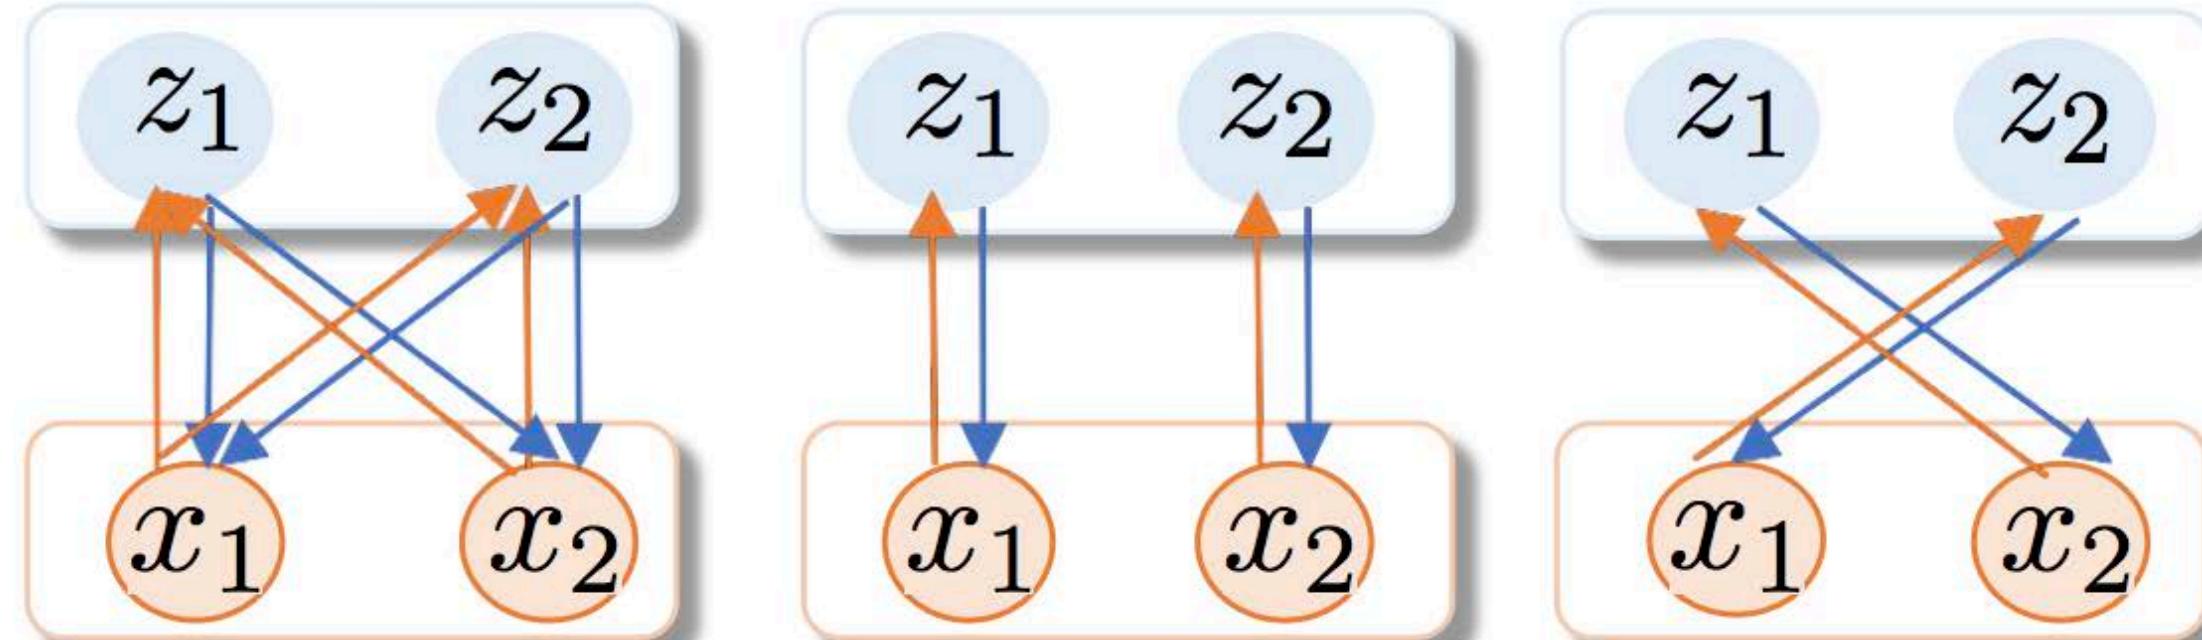
Low
Conditional
Entropy

Conditional Entropy

$$H^\pi(\mathbf{x}|\mathbf{z}) \triangleq -\mathbb{E}_{\pi(\mathbf{x}, \mathbf{z})}[\log \pi(\mathbf{x}|\mathbf{z})]$$

“ALICE: Towards Understanding Adversarial Learning for Joint Distribution Matching” [Li et al. NIPS 2017]. Also see [Tiao et al. 2018] “CycleGAN as Approximate Bayesian Inference”

Cycle Loss upper bounds Conditional Entropy



	z_1	z_2
x_1	$\delta/2$	$(1-\delta)/2$
x_2	$(1-\delta)/2$	$\delta/2$

	z_1	z_2
x_1	$1/2$	0
x_2	0	$1/2$

	z_1	z_2
x_1	0	$1/2$
x_2	$1/2$	0

Conditional Entropy

$$H^\pi(\mathbf{x}|\mathbf{z}) \triangleq -\mathbb{E}_{\pi(\mathbf{x}, \mathbf{z})}[\log \pi(\mathbf{x}|\mathbf{z})]$$

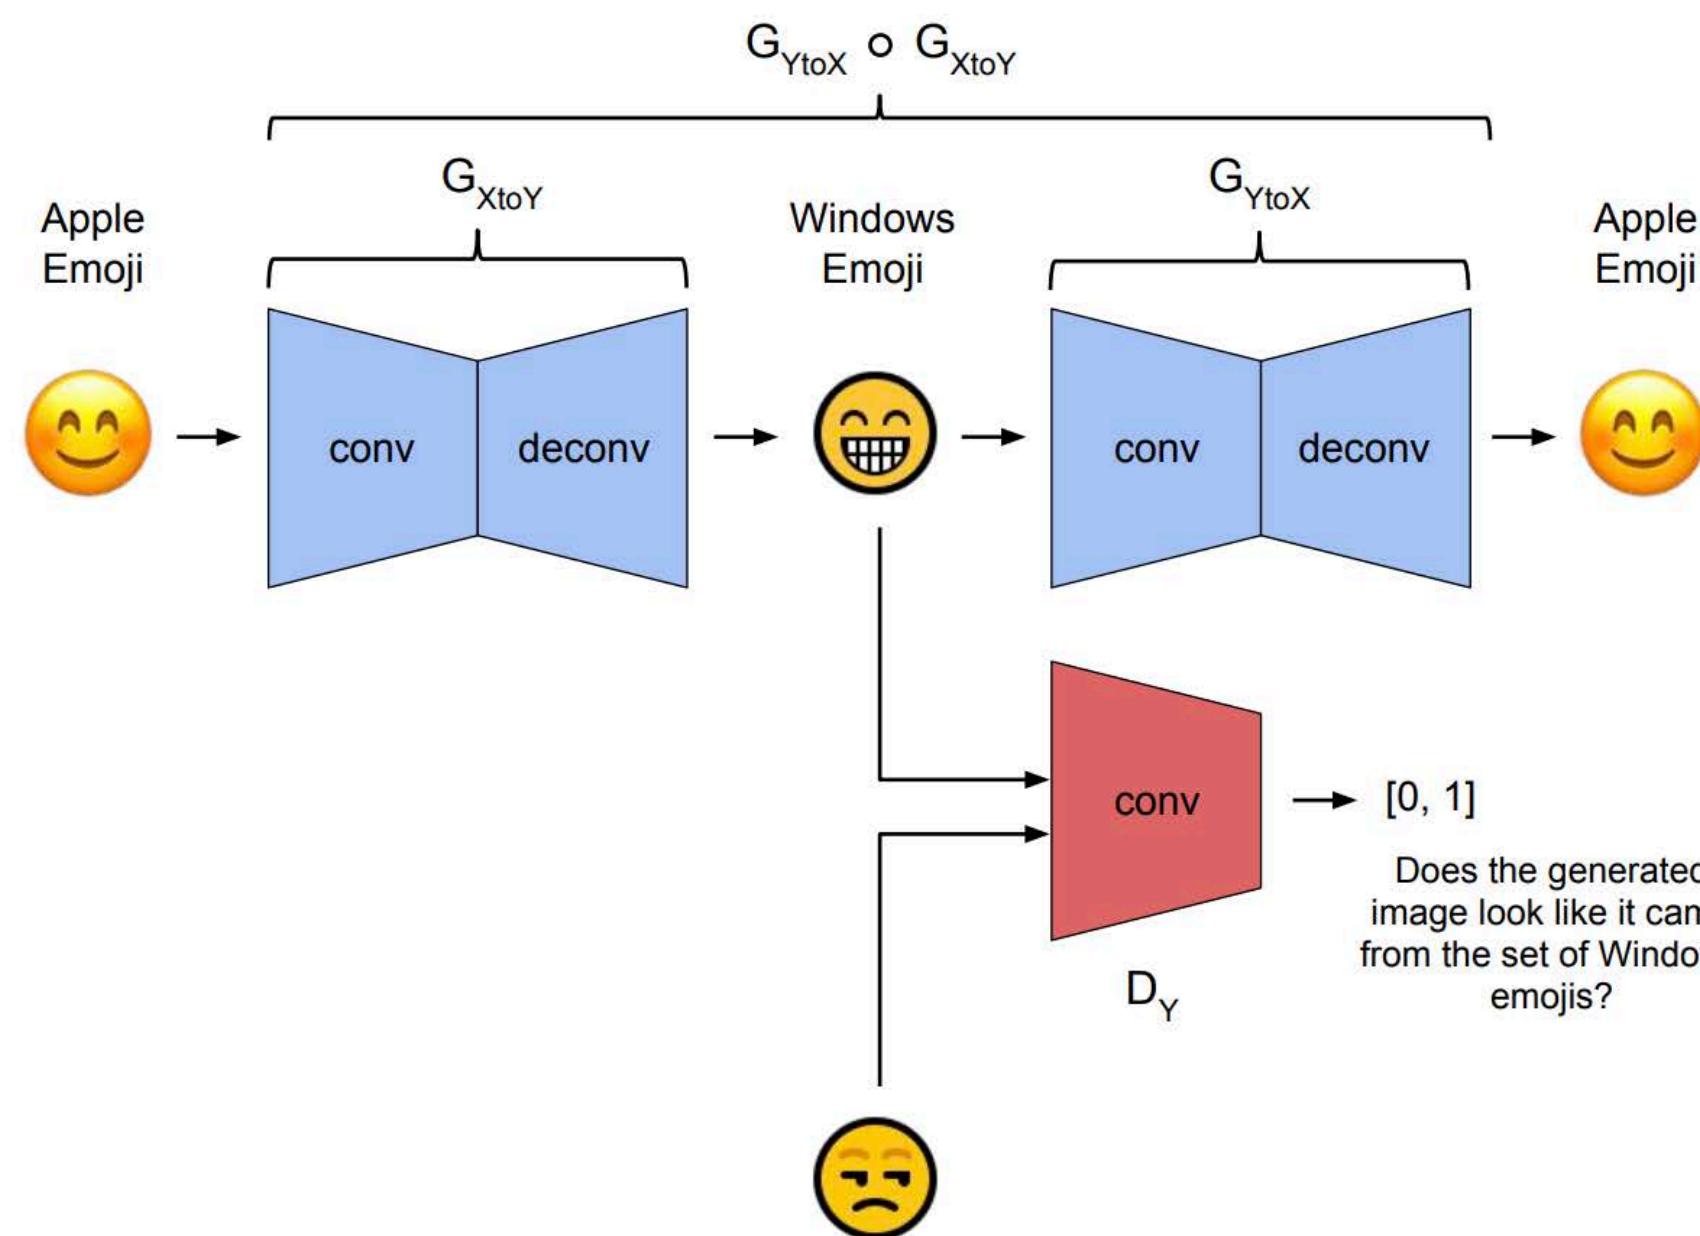
Lemma 3 For joint distributions $p_\theta(\mathbf{x}, \mathbf{z})$ or $q_\phi(\mathbf{x}, \mathbf{z})$, we have

$$\begin{aligned} H^{q_\phi}(\mathbf{x}|\mathbf{z}) &\triangleq -\mathbb{E}_{q_\phi(\mathbf{x}, \mathbf{z})}[\log q_\phi(\mathbf{x}|\mathbf{z})] = -\mathbb{E}_{q_\phi(\mathbf{x}, \mathbf{z})}[\log p_\theta(\mathbf{x}|\mathbf{z})] - \mathbb{E}_{q_\phi(\mathbf{z})}[\text{KL}(q_\phi(\mathbf{x}|\mathbf{z})||p_\theta(\mathbf{x}|\mathbf{z}))] \\ &\leq -\mathbb{E}_{q_\phi(\mathbf{x}, \mathbf{z})}[\log p_\theta(\mathbf{x}|\mathbf{z})] \triangleq \mathcal{L}_{\text{Cycle}}(\boldsymbol{\theta}, \boldsymbol{\phi}). \end{aligned} \quad (6)$$

“ALICE: Towards Understanding Adversarial Learning for Joint Distribution Matching” [Li et al. NIPS 2017]. Also see [Tiao et al. 2018] “CycleGAN as Approximate Bayesian Inference”

CycleGAN at School

- Course assignment [code](#) and [handout](#) designed by Prof. [Roger Grosse](#) for [CSC321](#) “Intro to Neural Networks and Machine Learning” at University of Toronto.



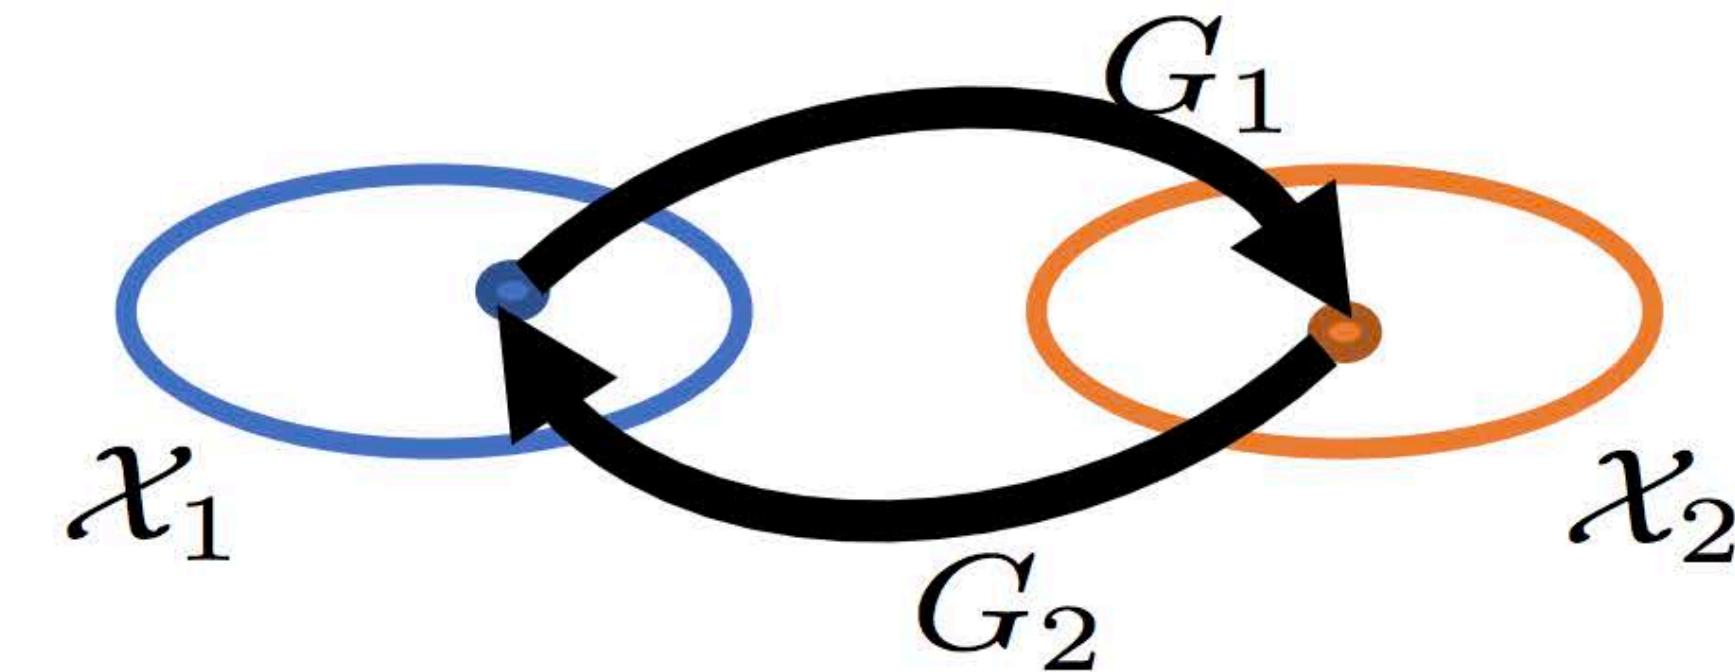
```
## FILL THIS IN: CREATE ARCHITECTURE ##
#####
# 1. Define the encoder part of the generator
# self.conv1 = ...
# self.conv2 = ...

# 2. Define the transformation part of the generator
# self.resnet_block = ...

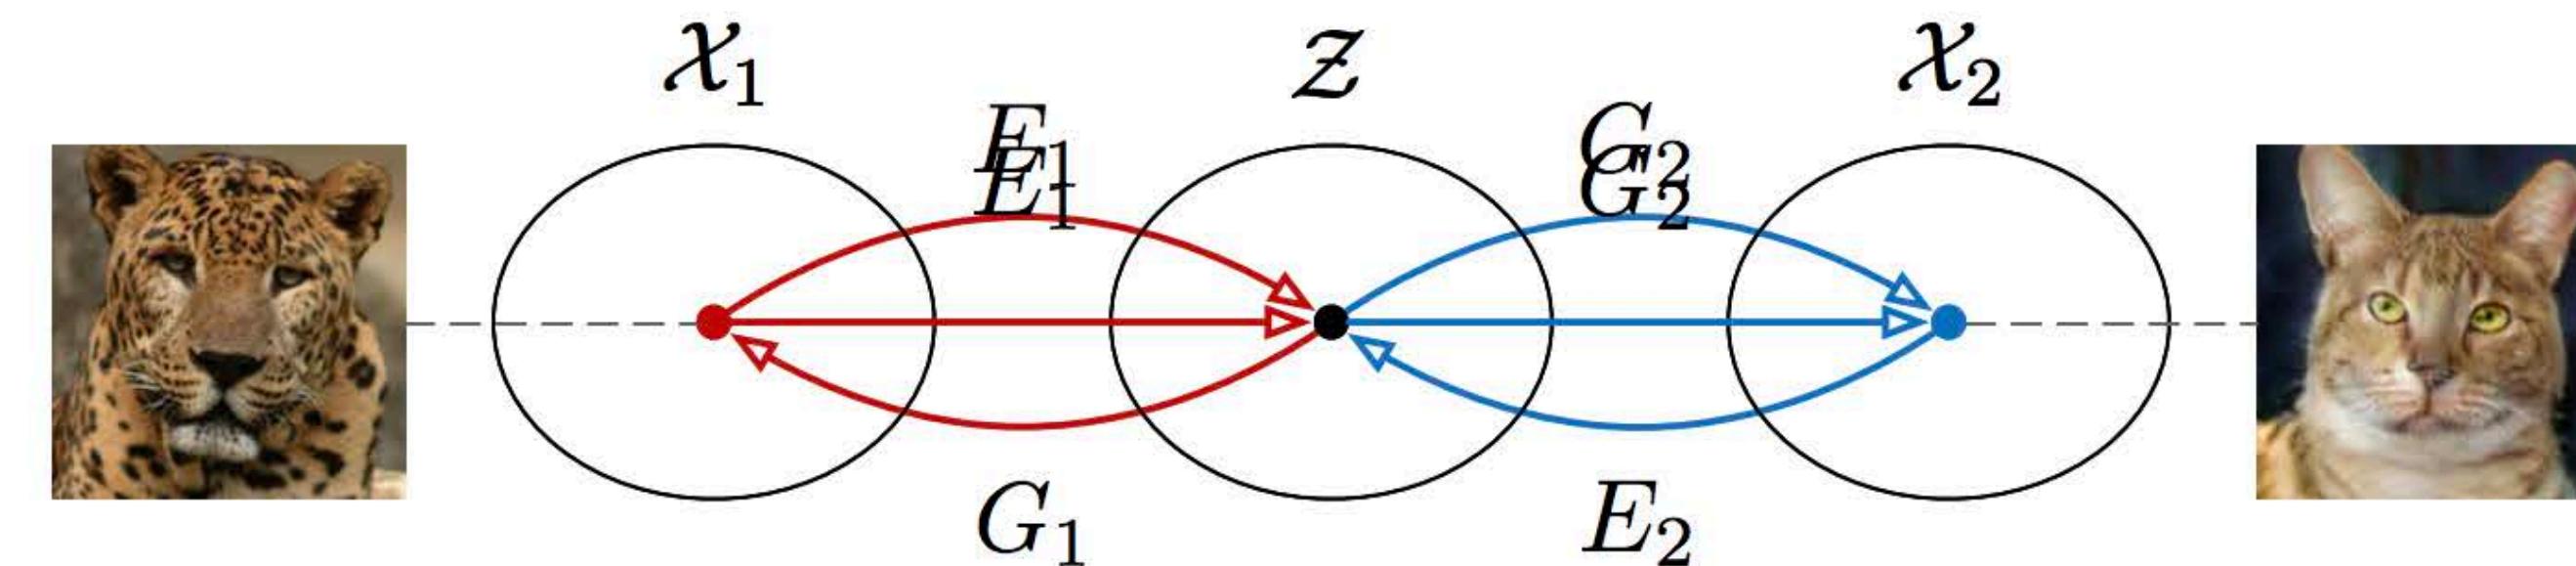
# 3. Define the decoder part of the generator
# self.deconv1 = ...
# self.deconv2 = ...
```

Can other generative models do this?

CycleGAN

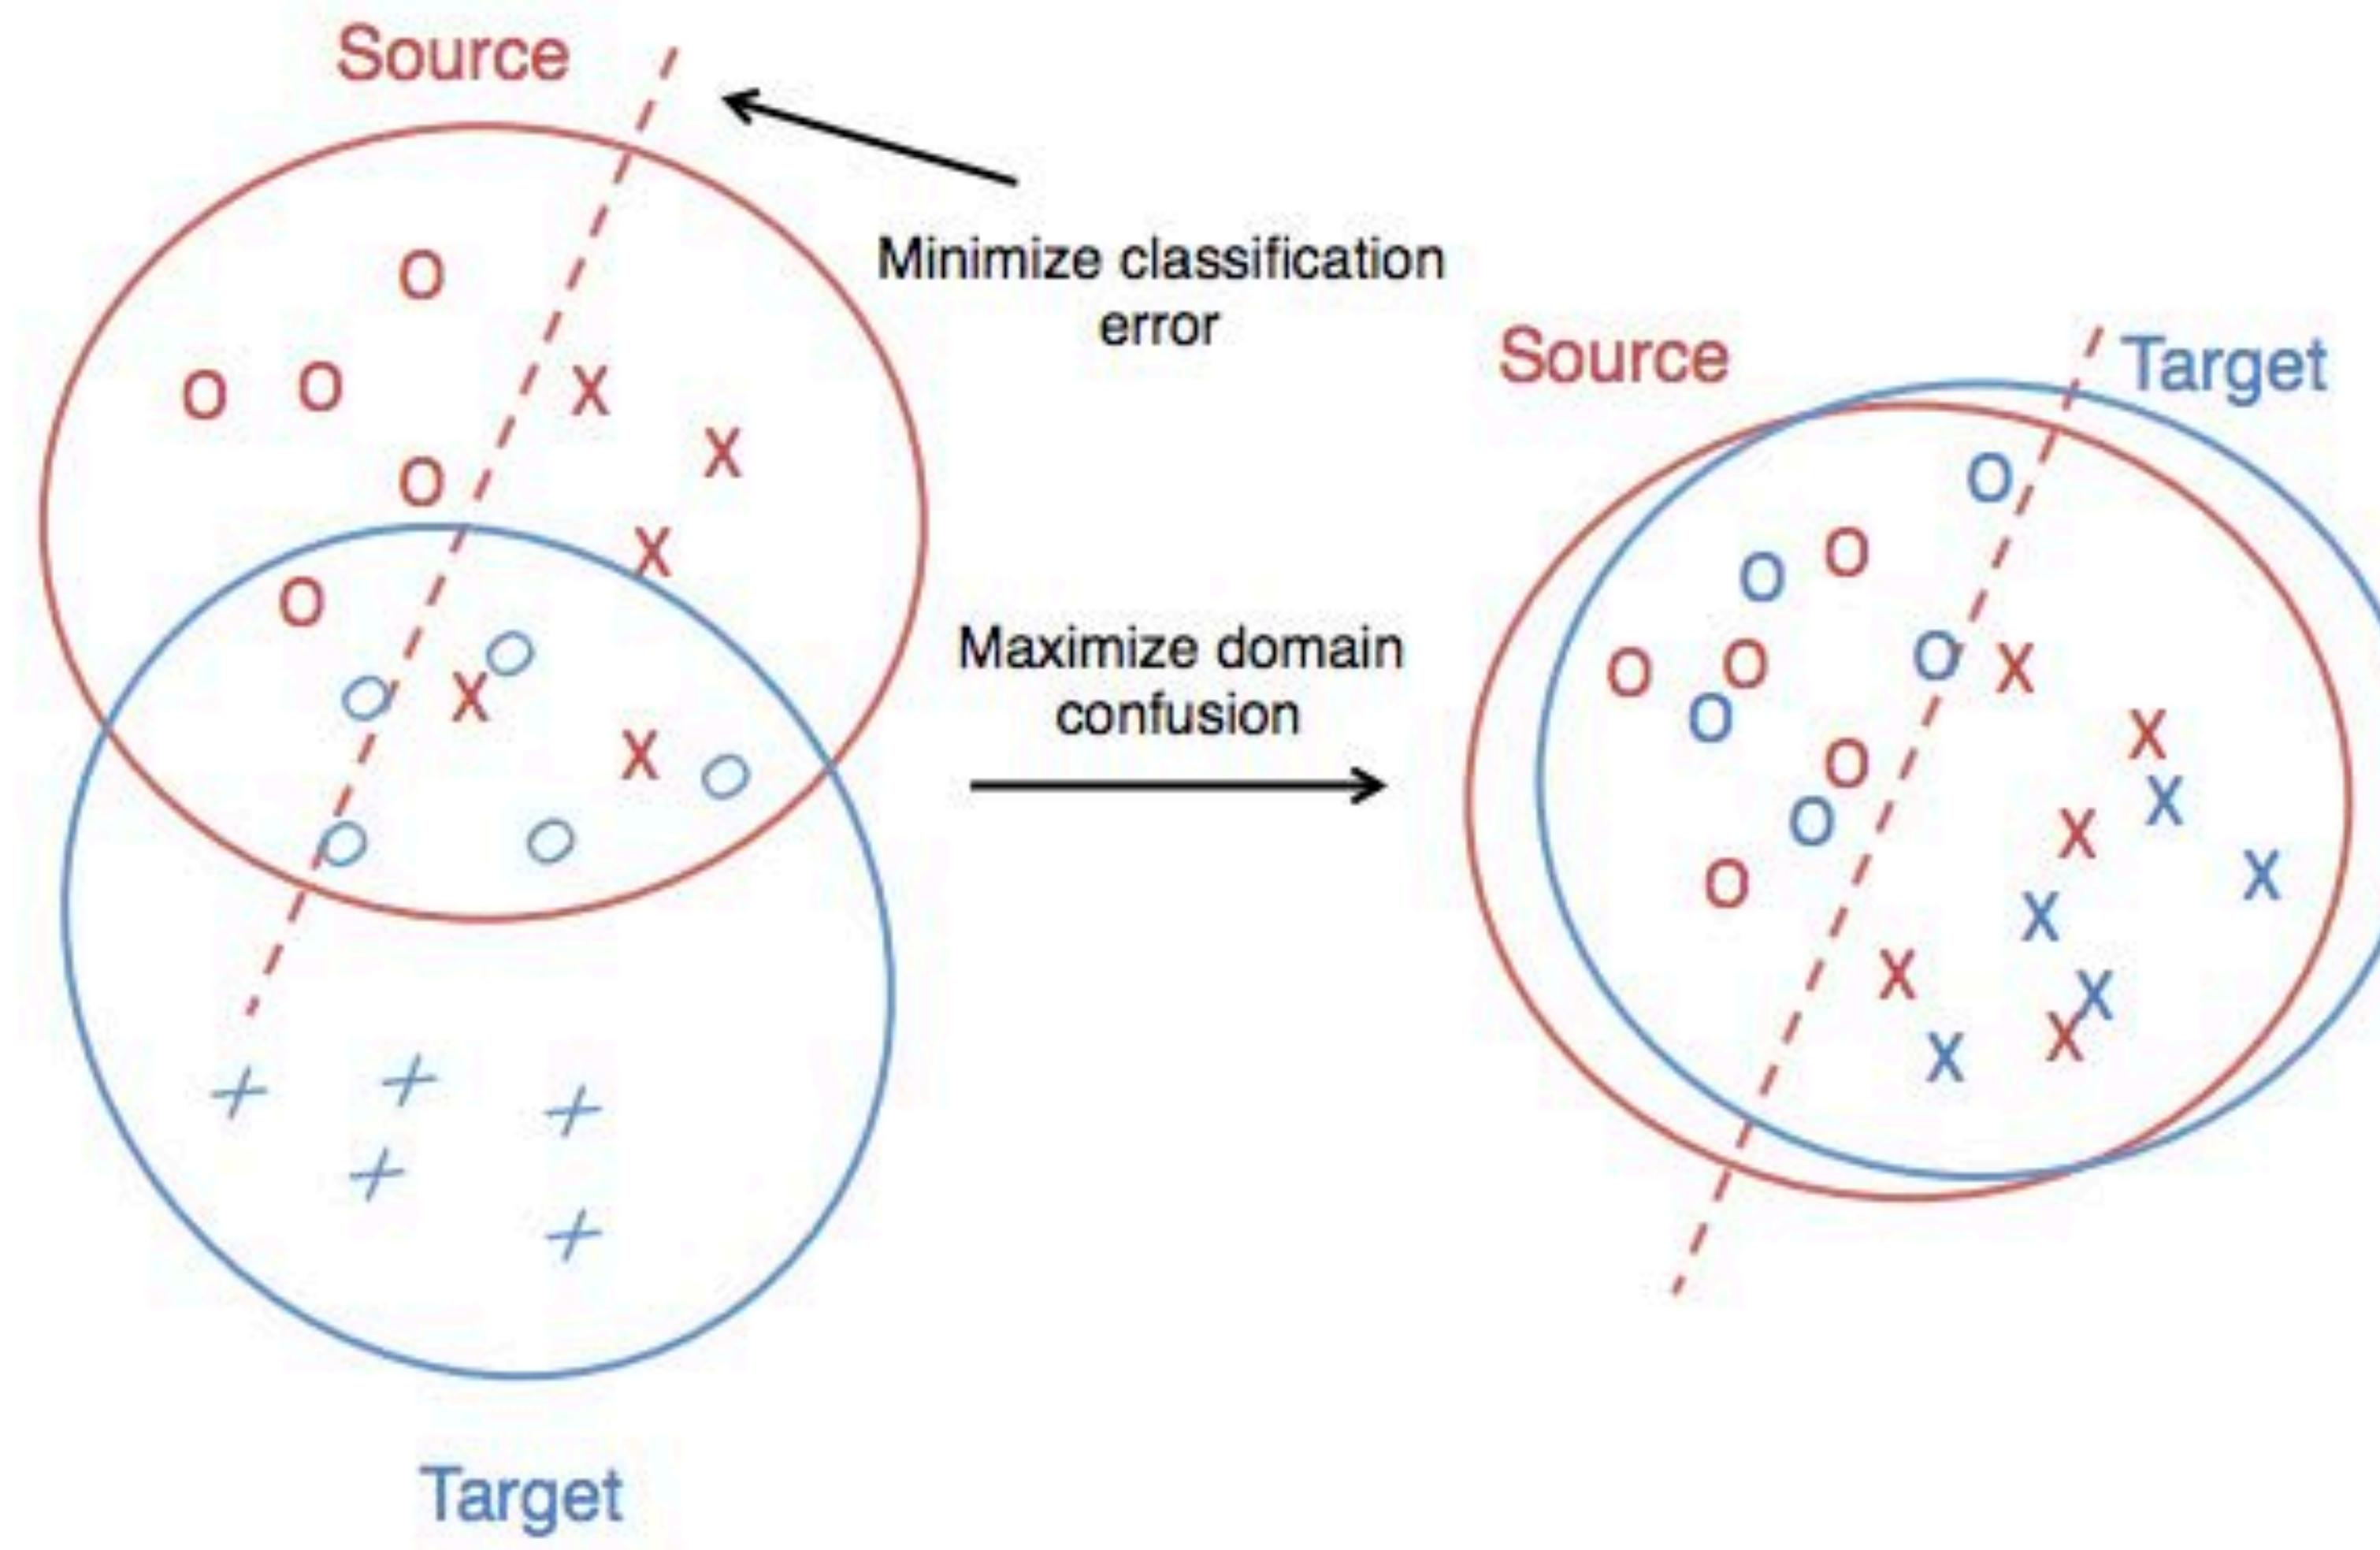


UNIT



Unsupervised image-to-image translation networks
[Liu, Breuel, Kautz, 2017]

Domain Adaptation



[Tzeng et al. 2014]

Sim2real

Simulated data



,



Real data



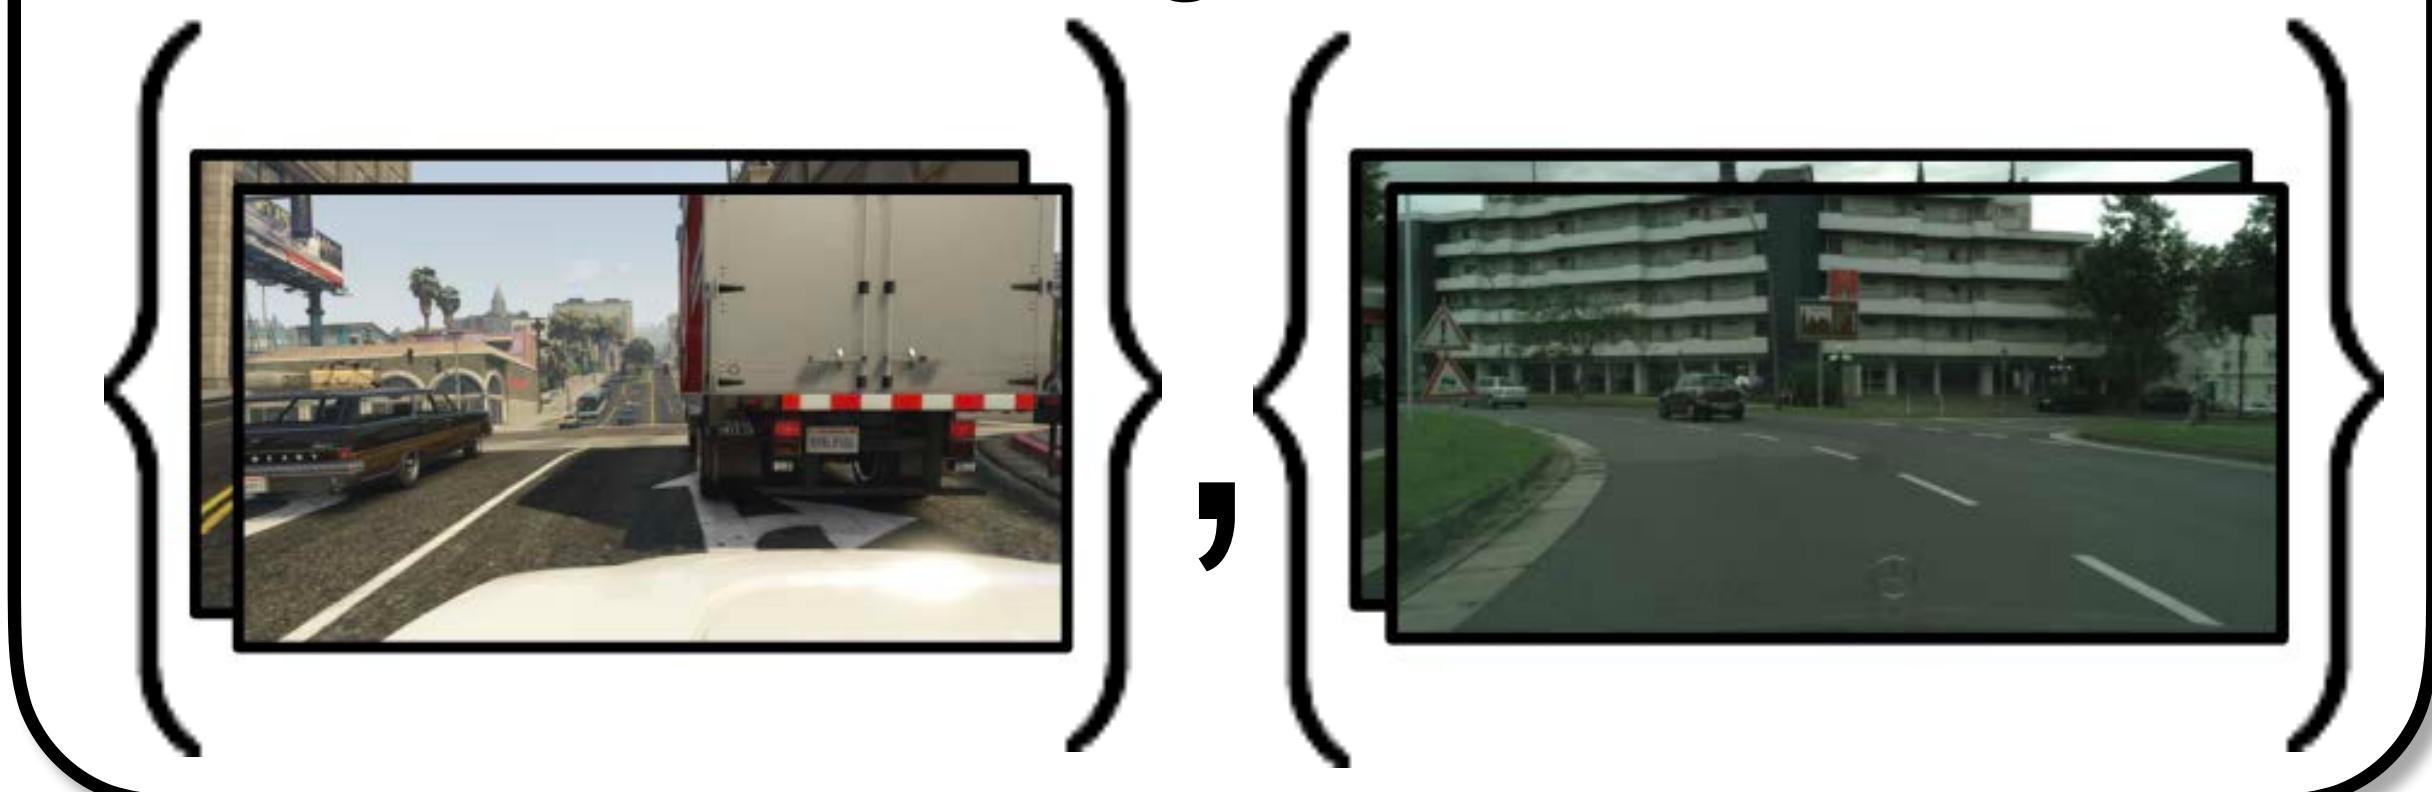
,

?

CycleGAN

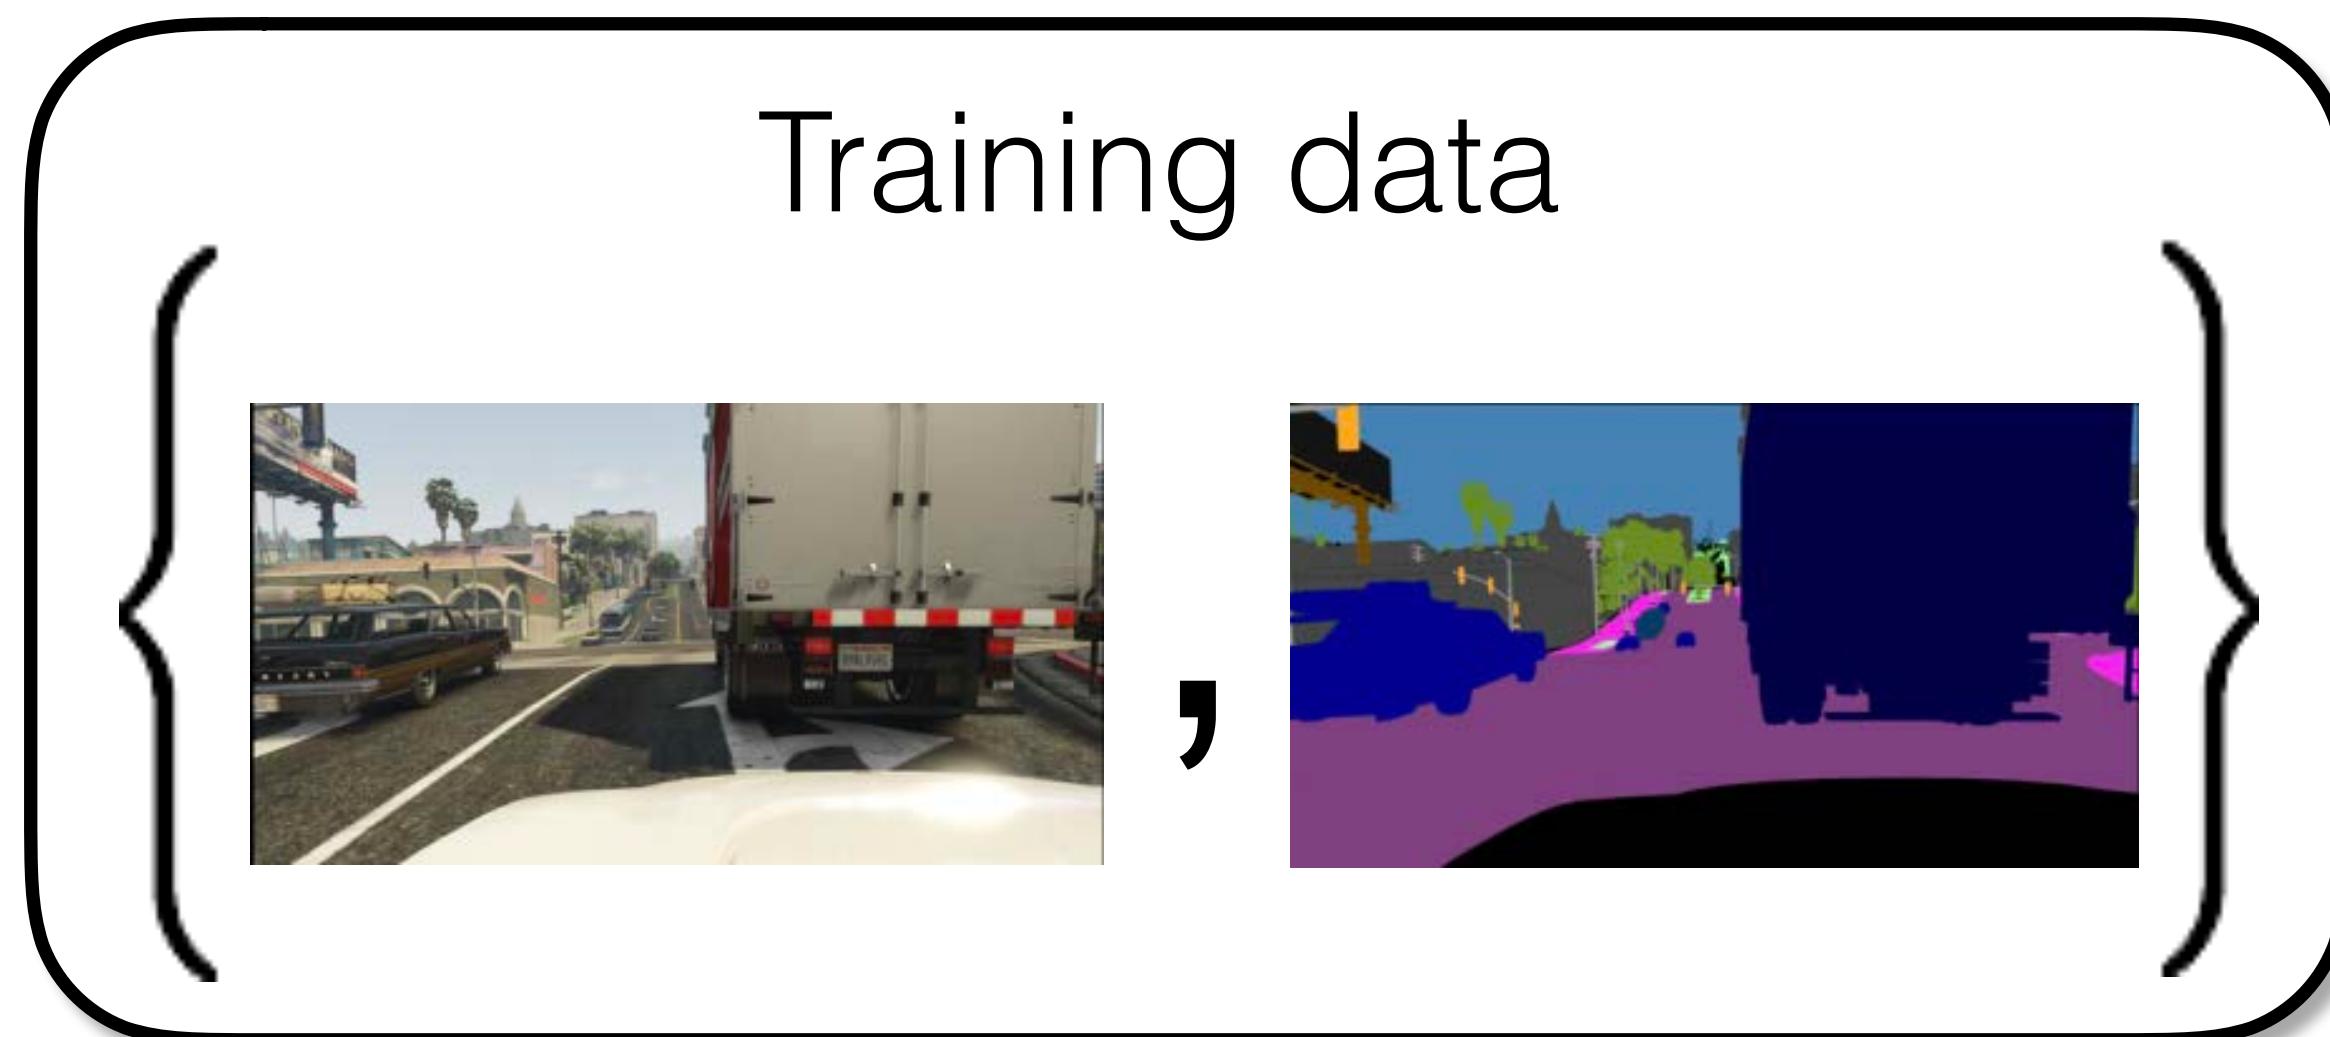


Training data



[Hoffman, Tzeng, Park, Zhu, Isola, Saenko, Darrell, Efros, 2018]

CycleGAN



[Hoffman, Tzeng, Park, Zhu, Isola, Saenko, Darrell, Efros, 2018]

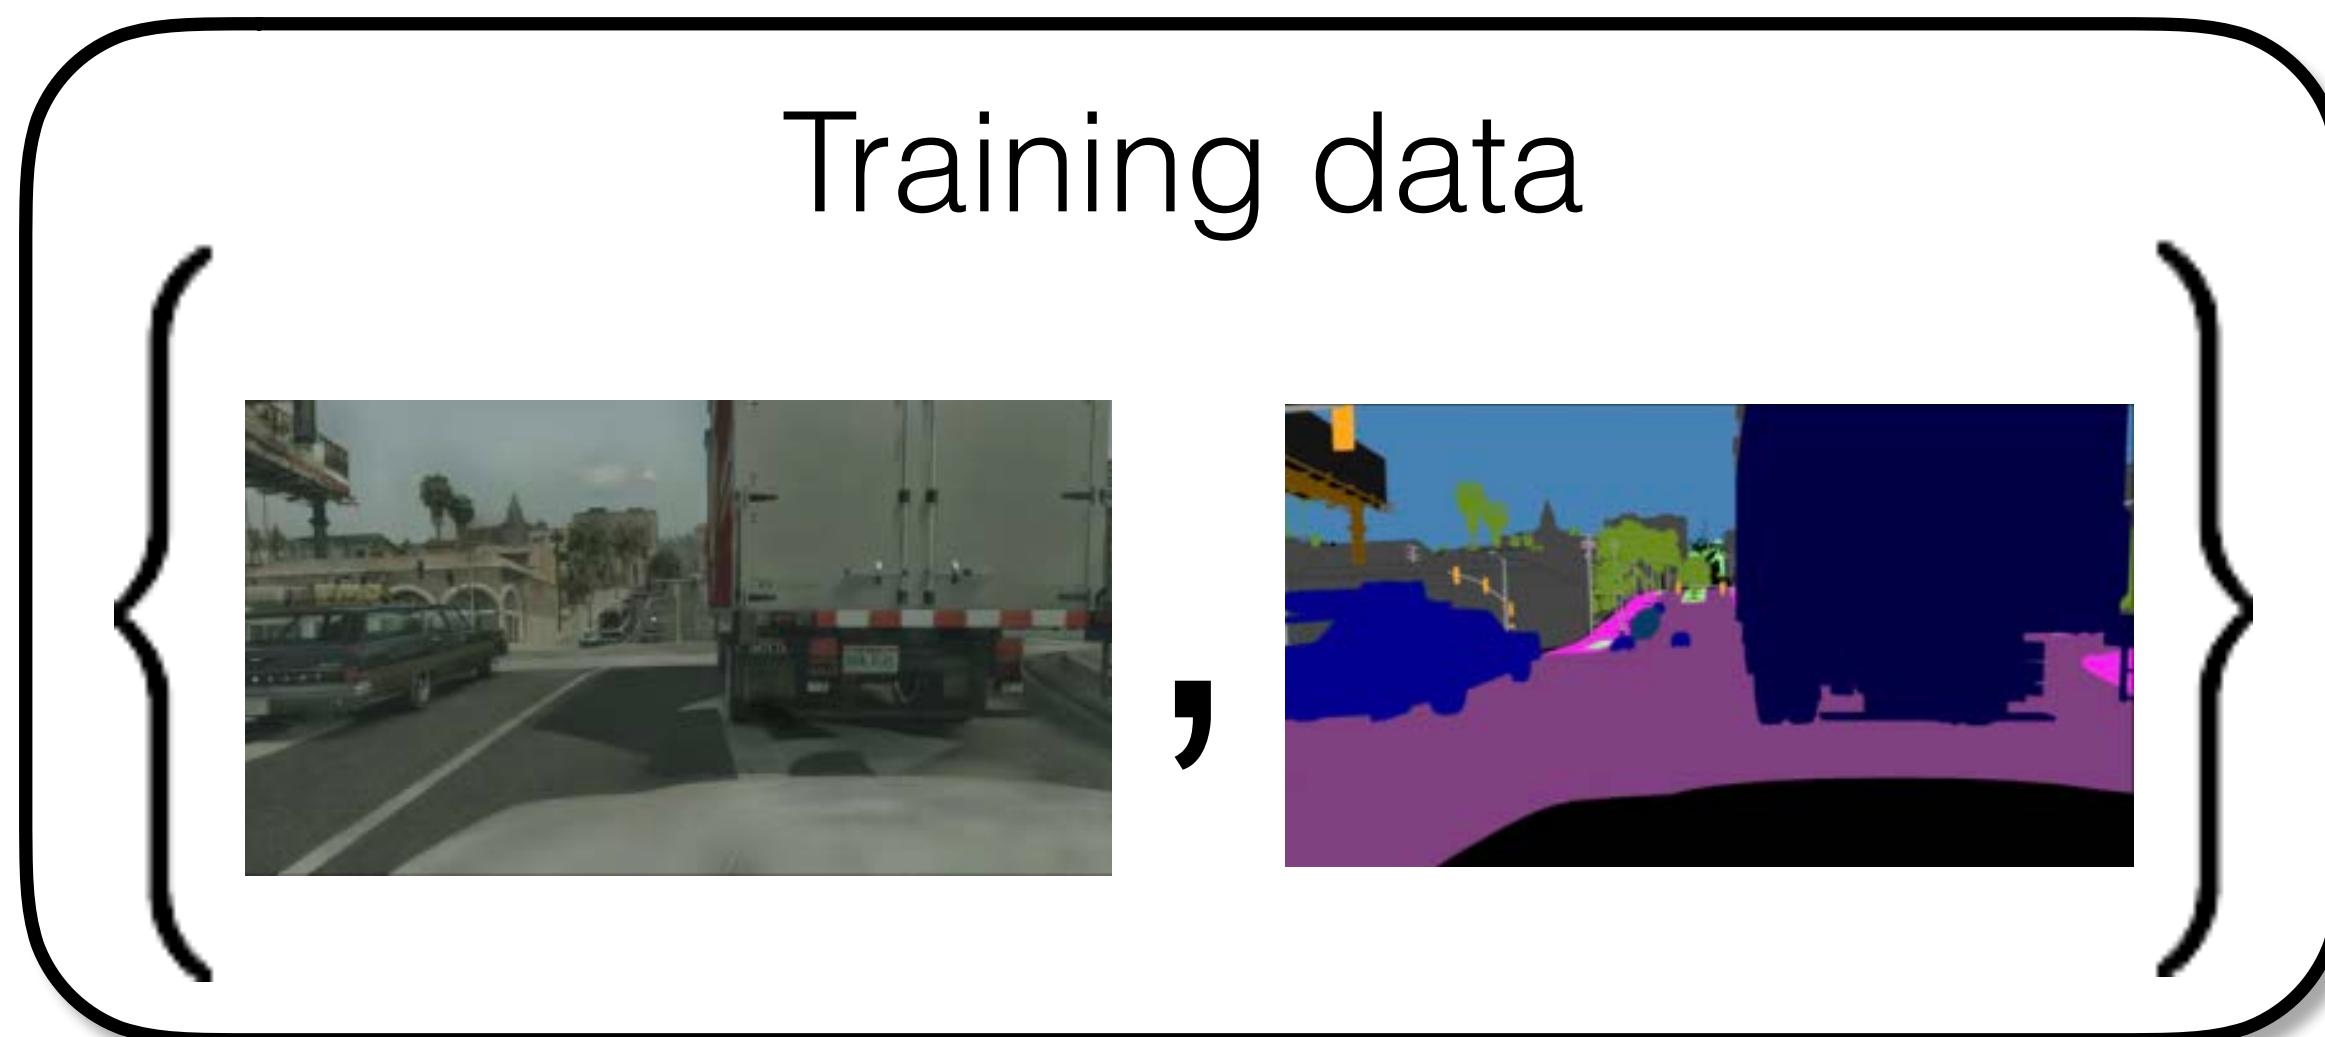
CycleGAN



FCN



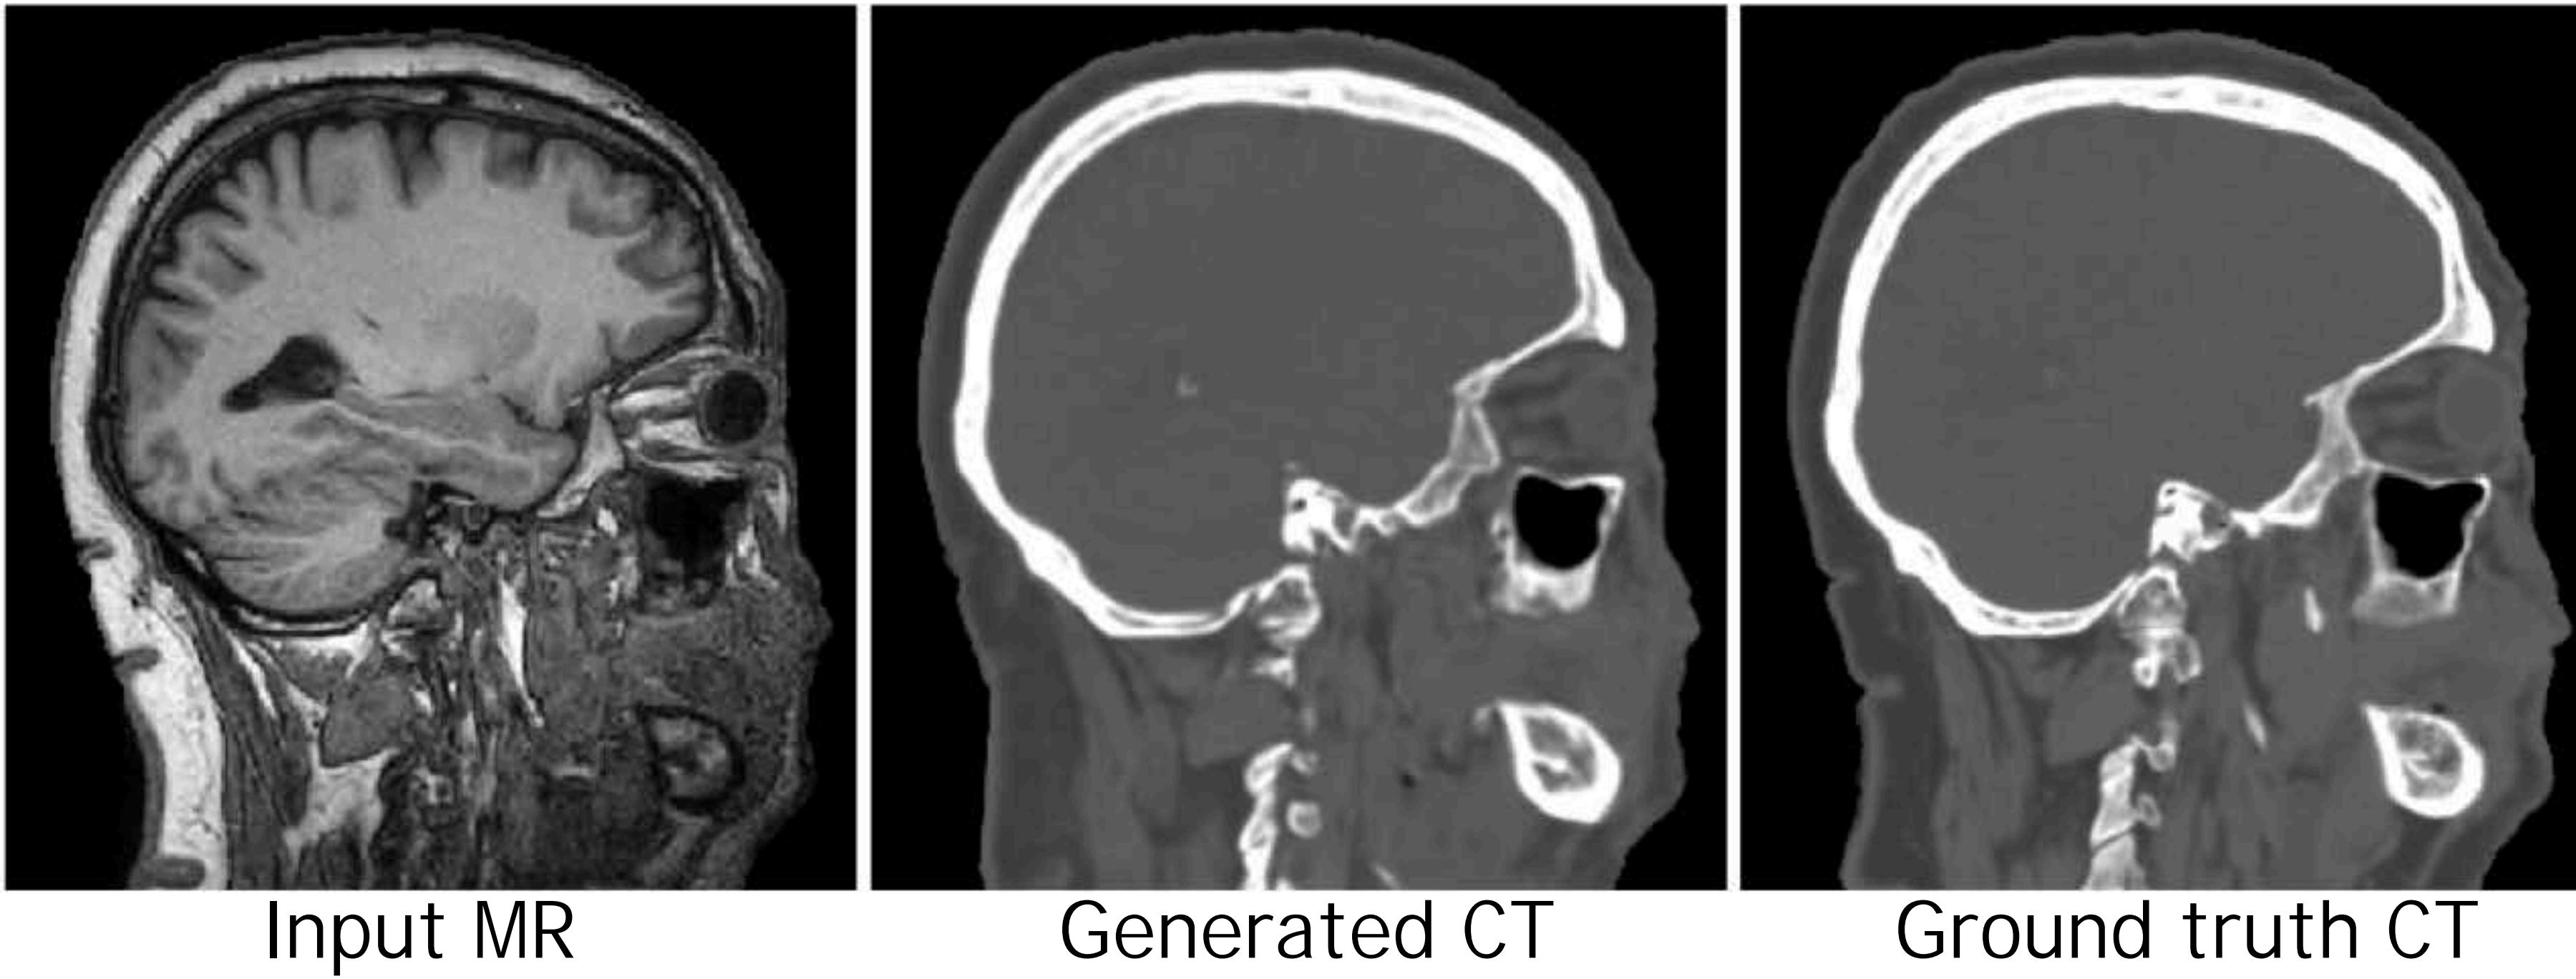
Training data



[Hoffman, Tzeng, Park, Zhu, Isola, Saenko, Darrell, Efros, 2018]

Medical domain adaptation

MR → CT [Wolterink et al] arxiv: 1708.01155



- MRI reconstruction [Quan et al.] arxiv:1709.00753
- Cardiac MR images from CT [Chartsias et al. 2017]

What can you do with generative models?

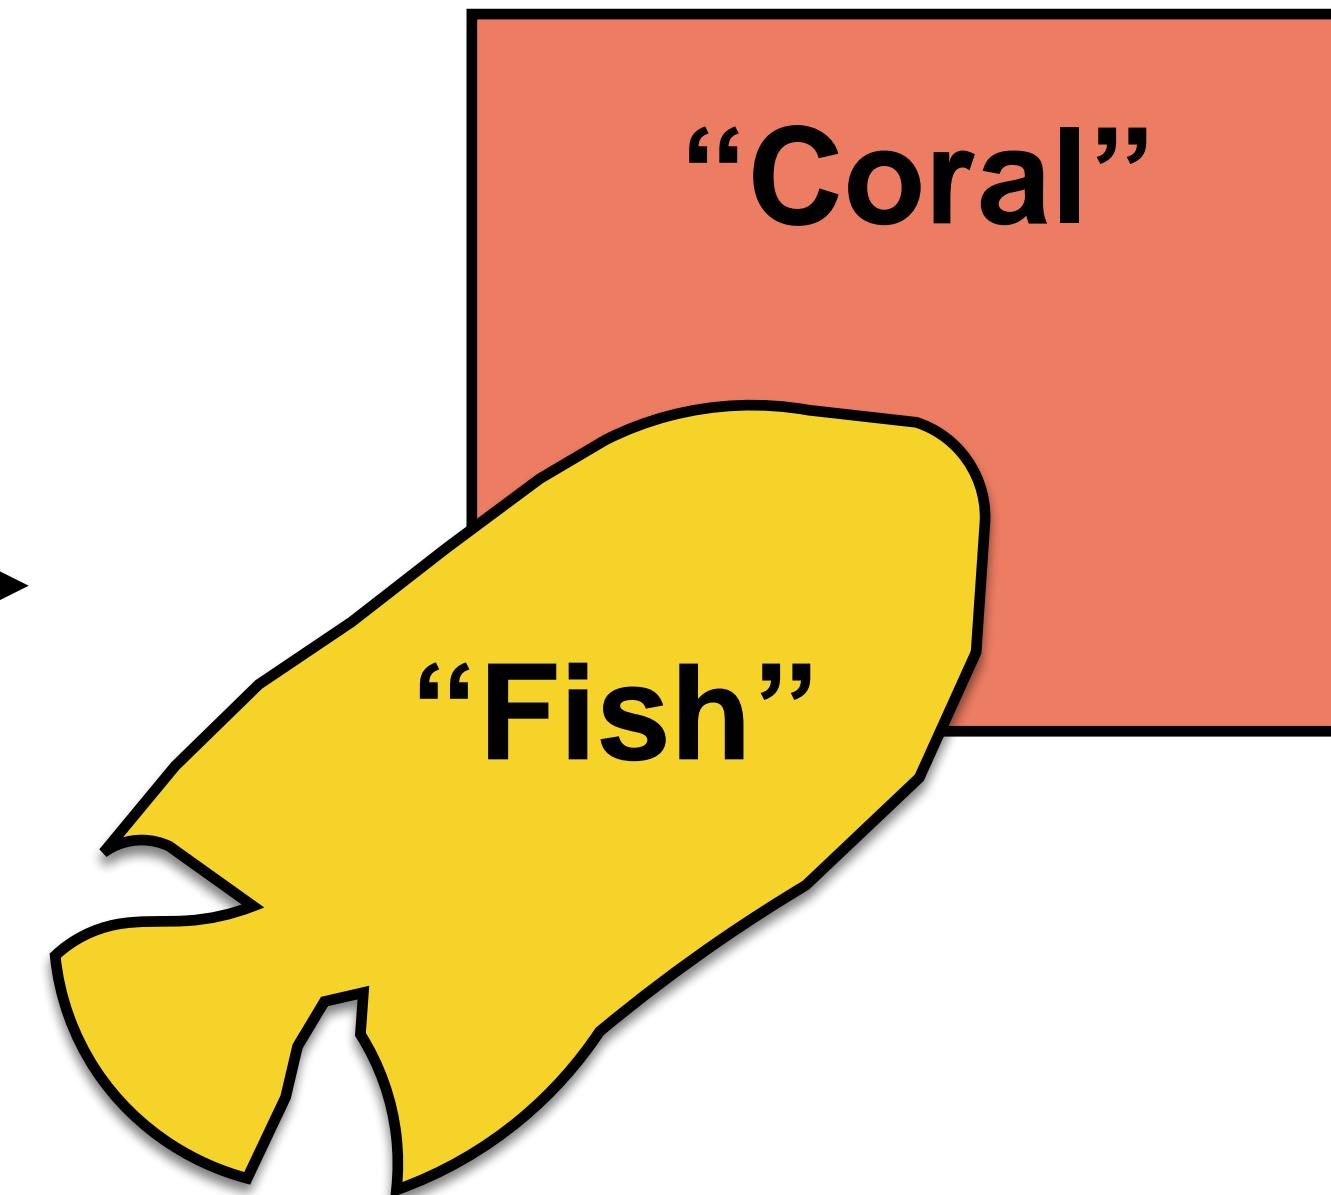
1. Data prediction
- 2. Domain mapping**
3. Representation learning
4. Model-based intelligence

What can you do with generative models?

1. Data prediction
2. Domain mapping
- 3. Representation learning**
4. Model-based intelligence

Representation Learning

X



Image

Compact mental
representation

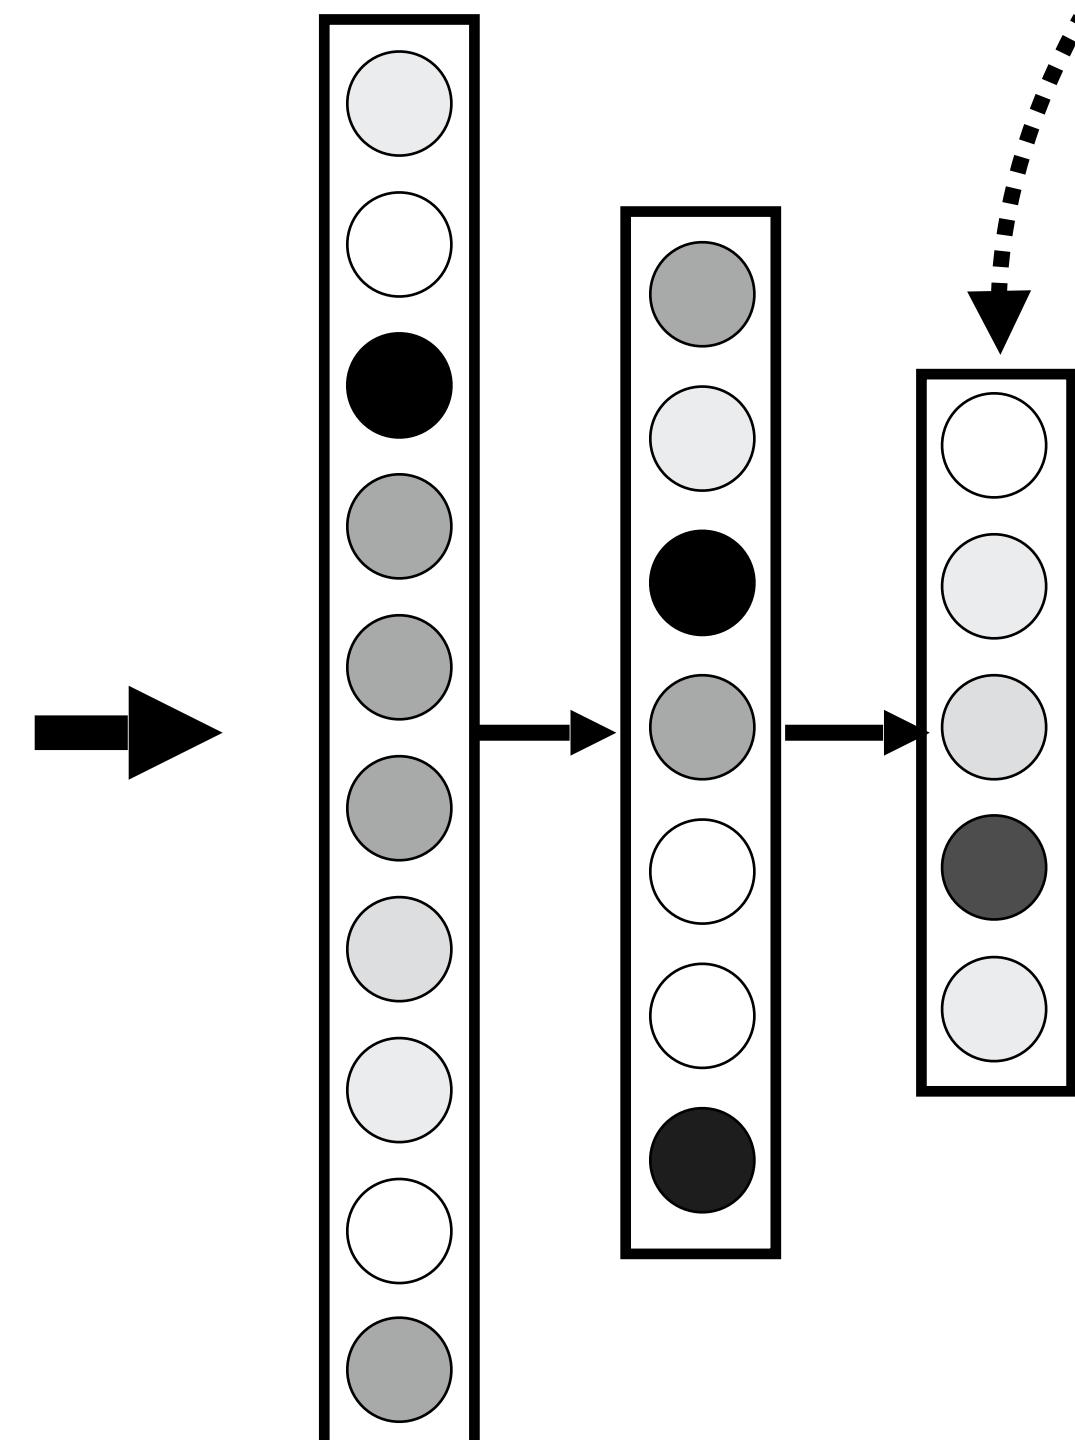
Representation Learning

compressed image code
(vector \mathbf{z})

\mathbf{X}

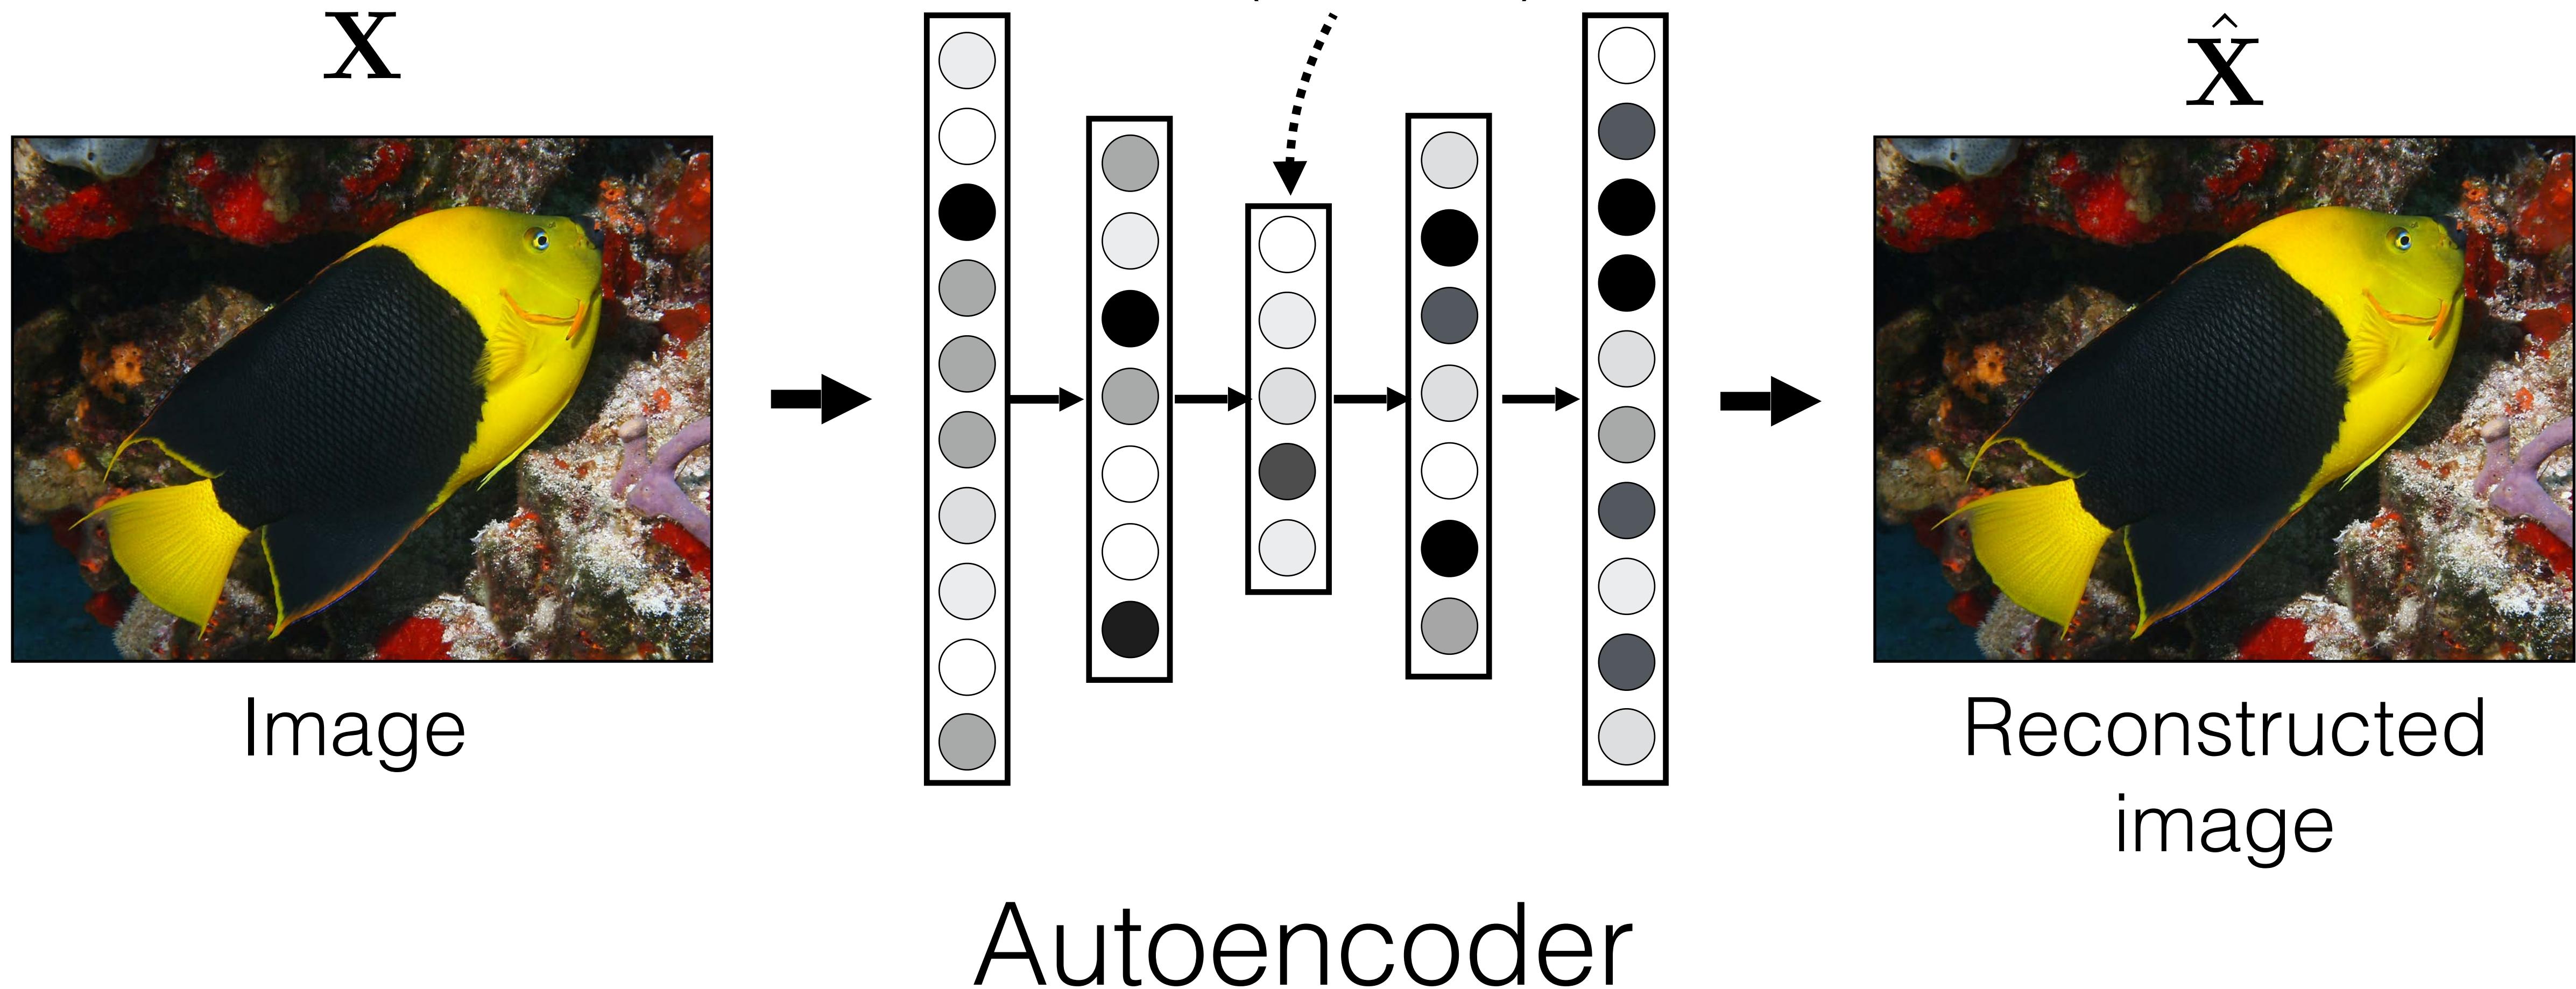


Image

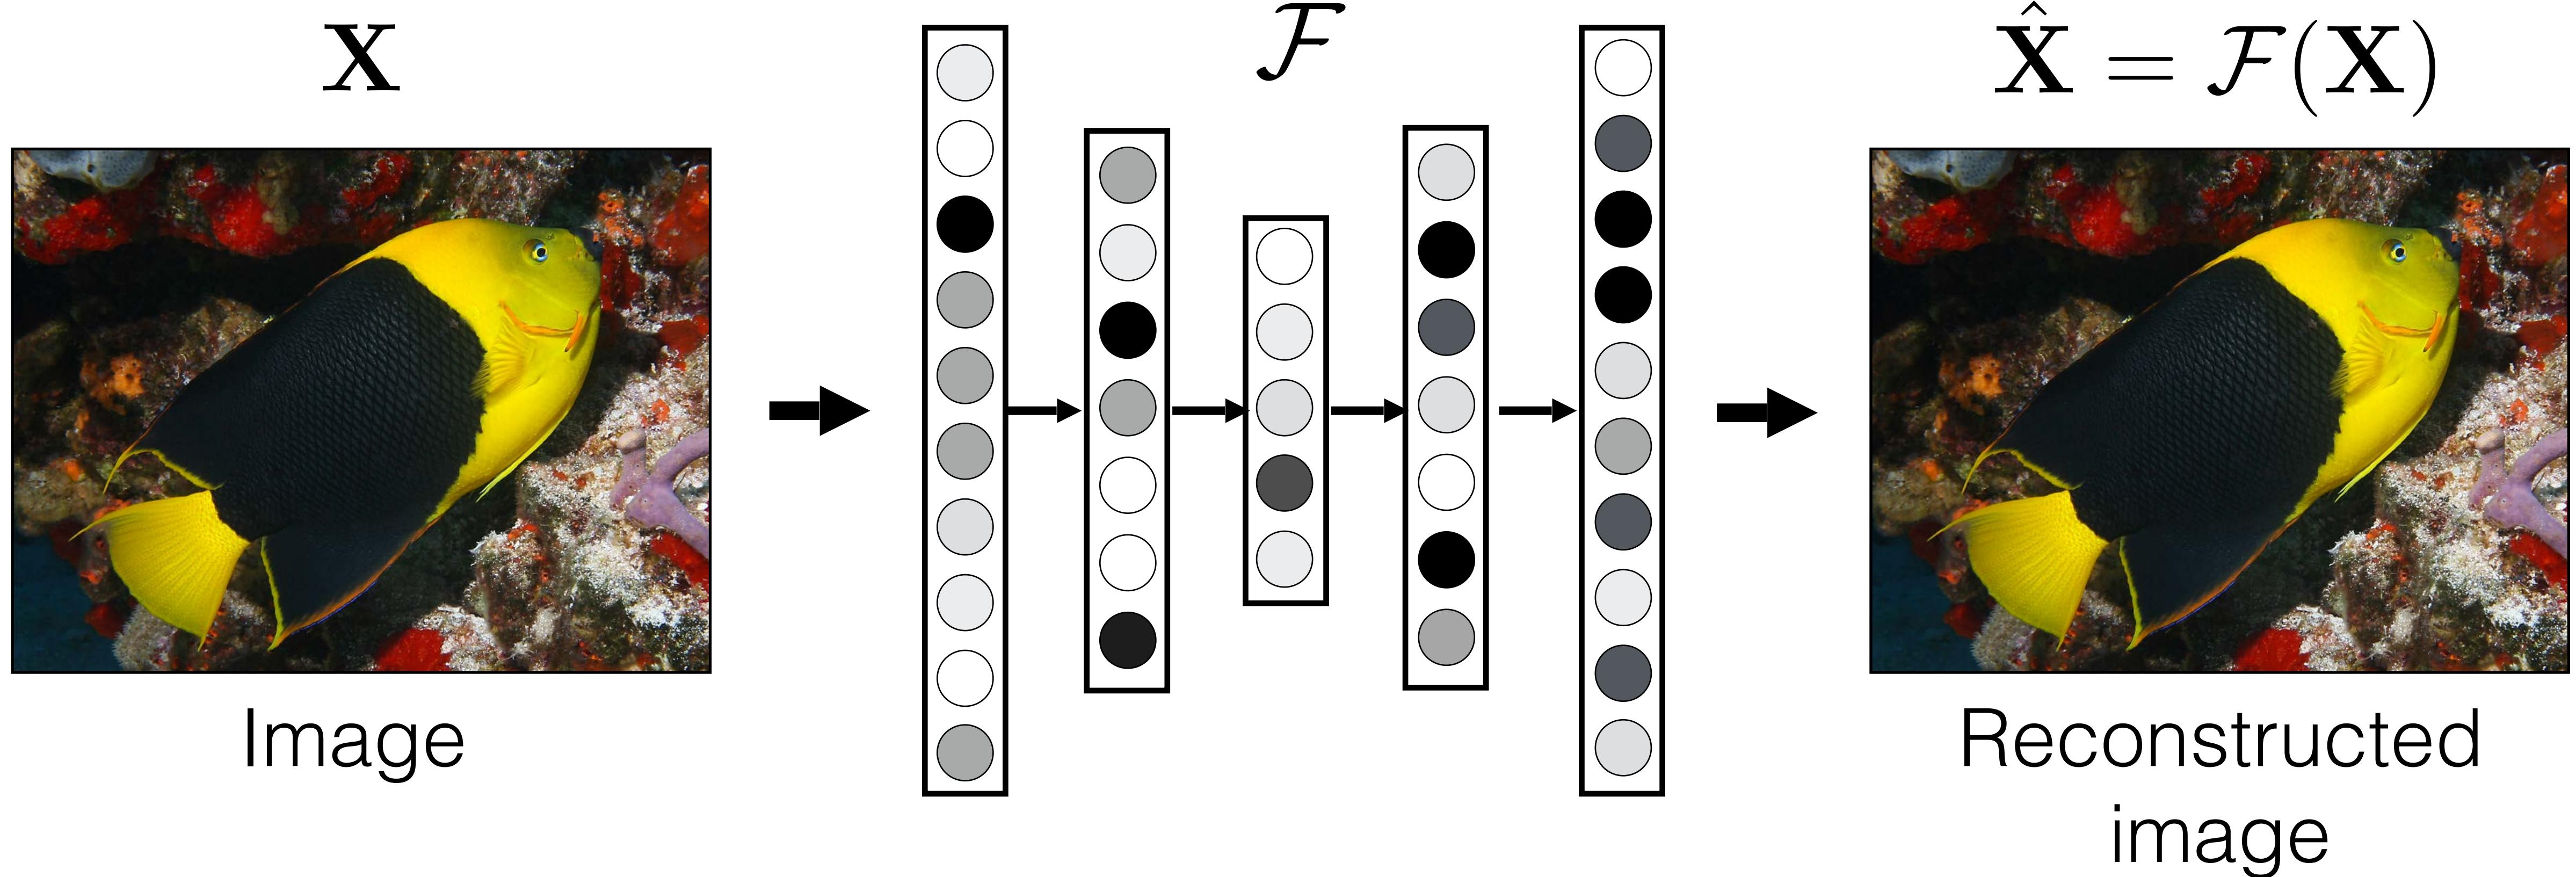


Representation Learning

compressed image code
(vector \mathbf{z})



Autoencoder

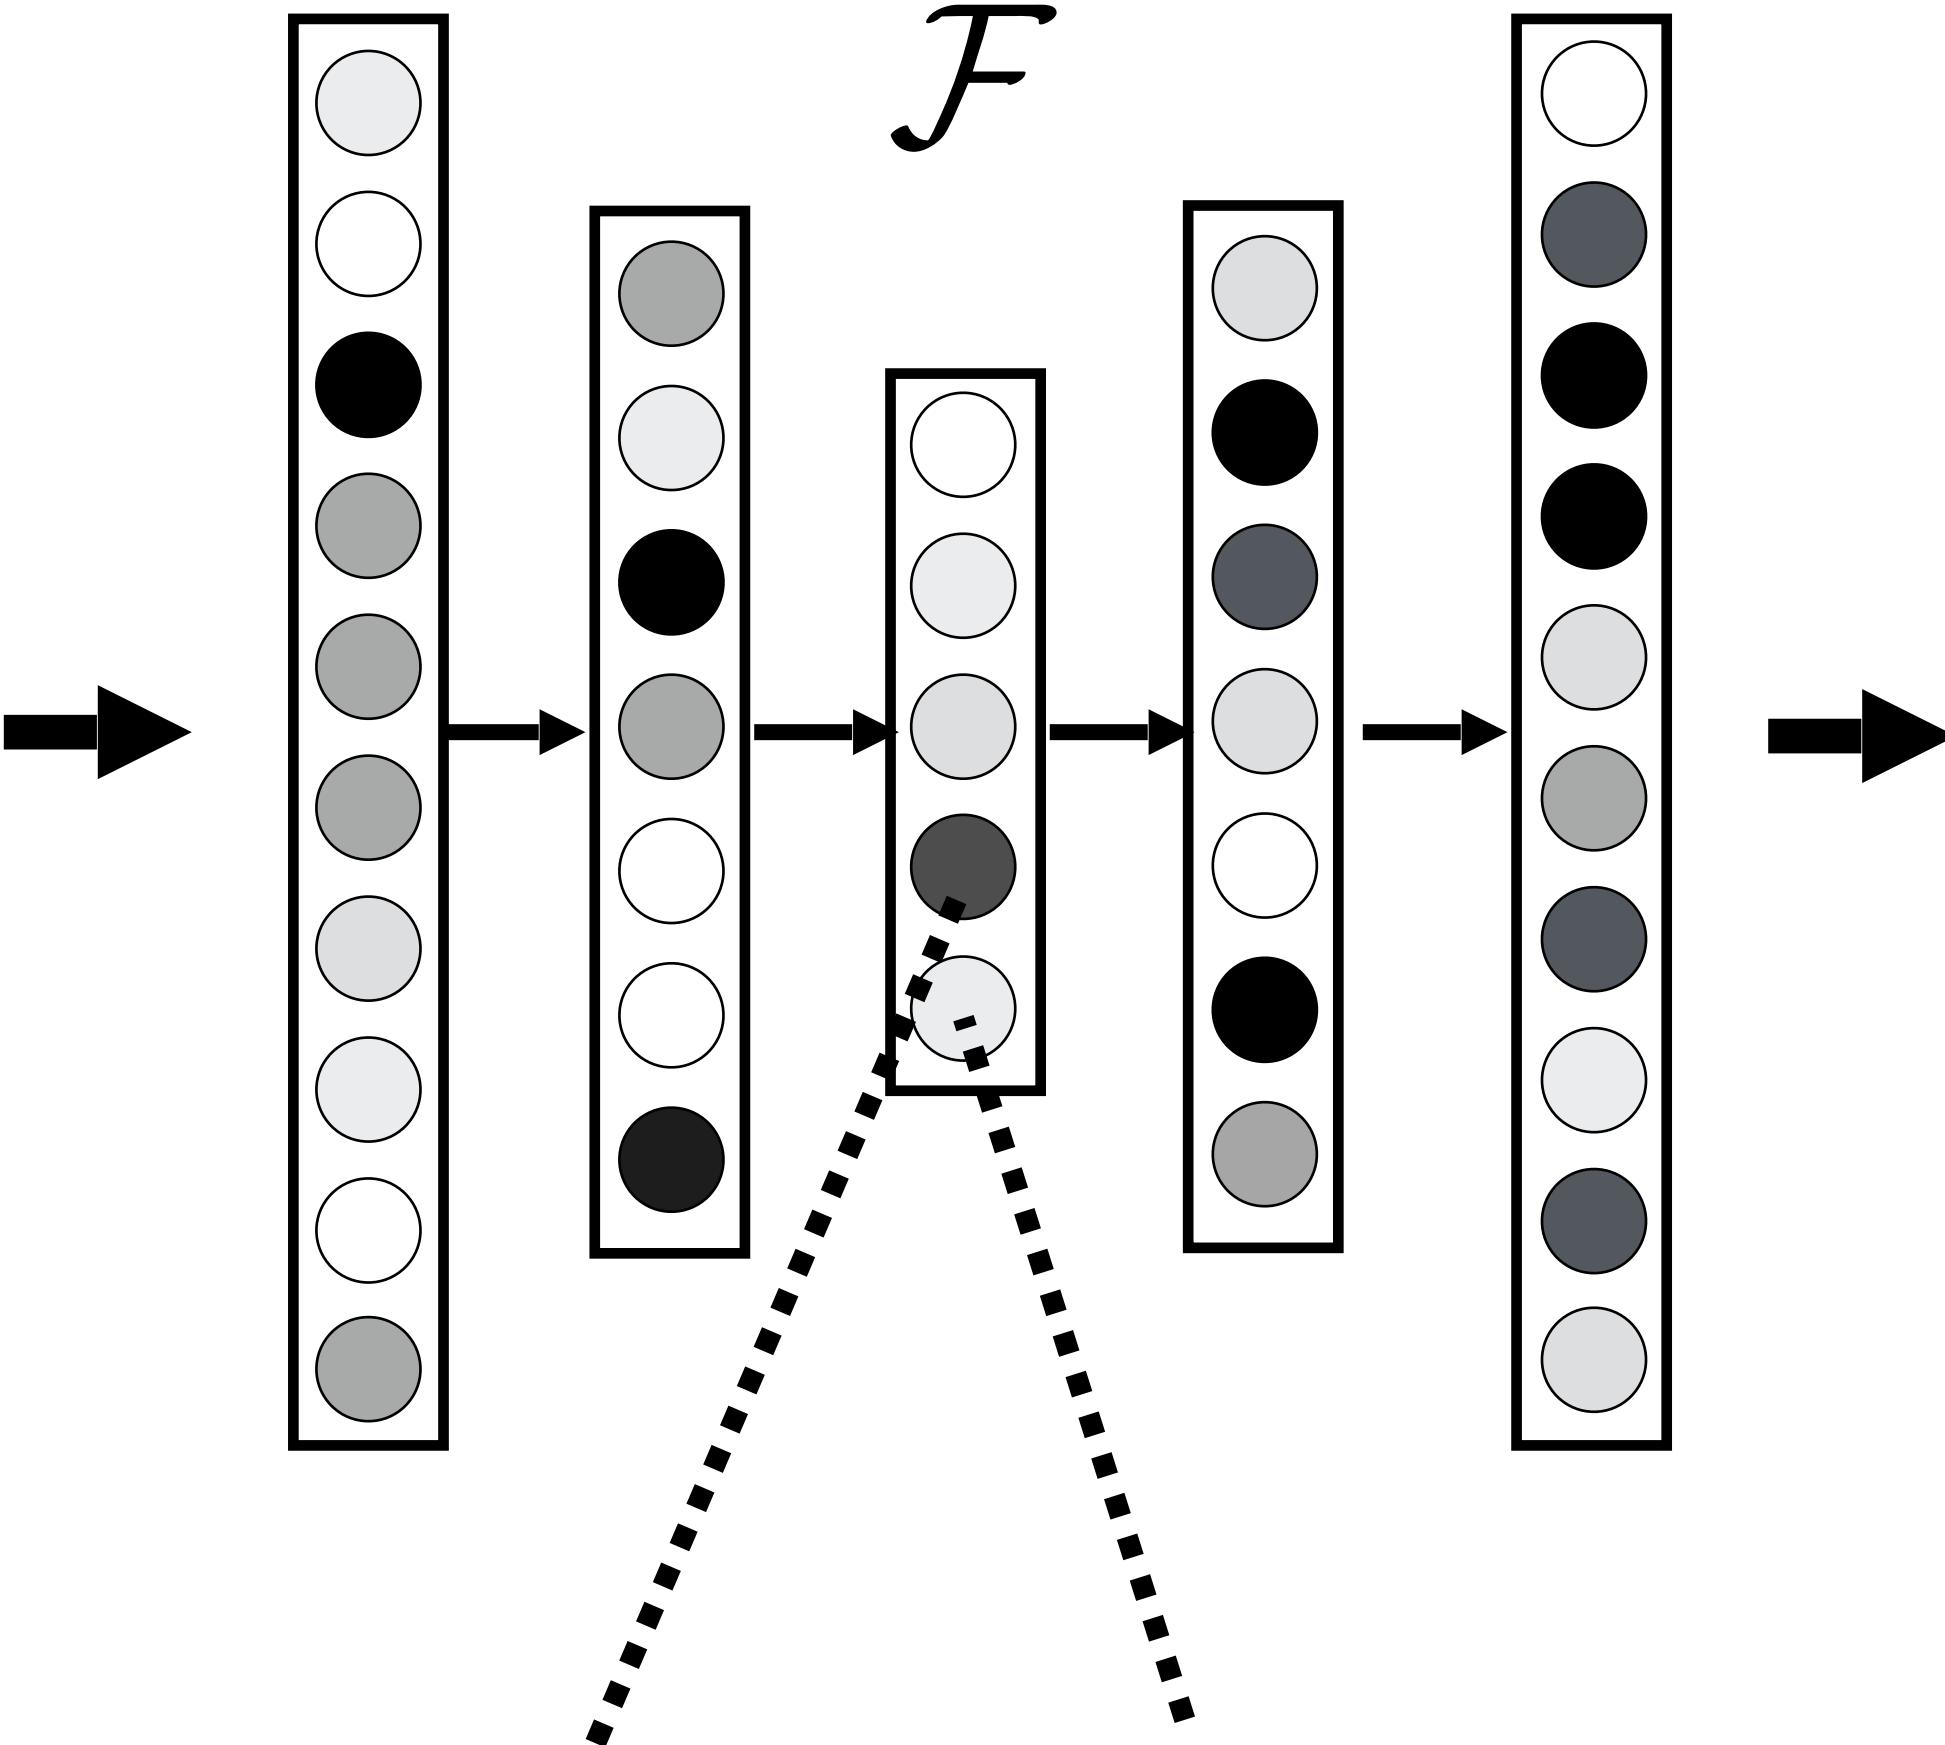


\mathbf{X}



Image

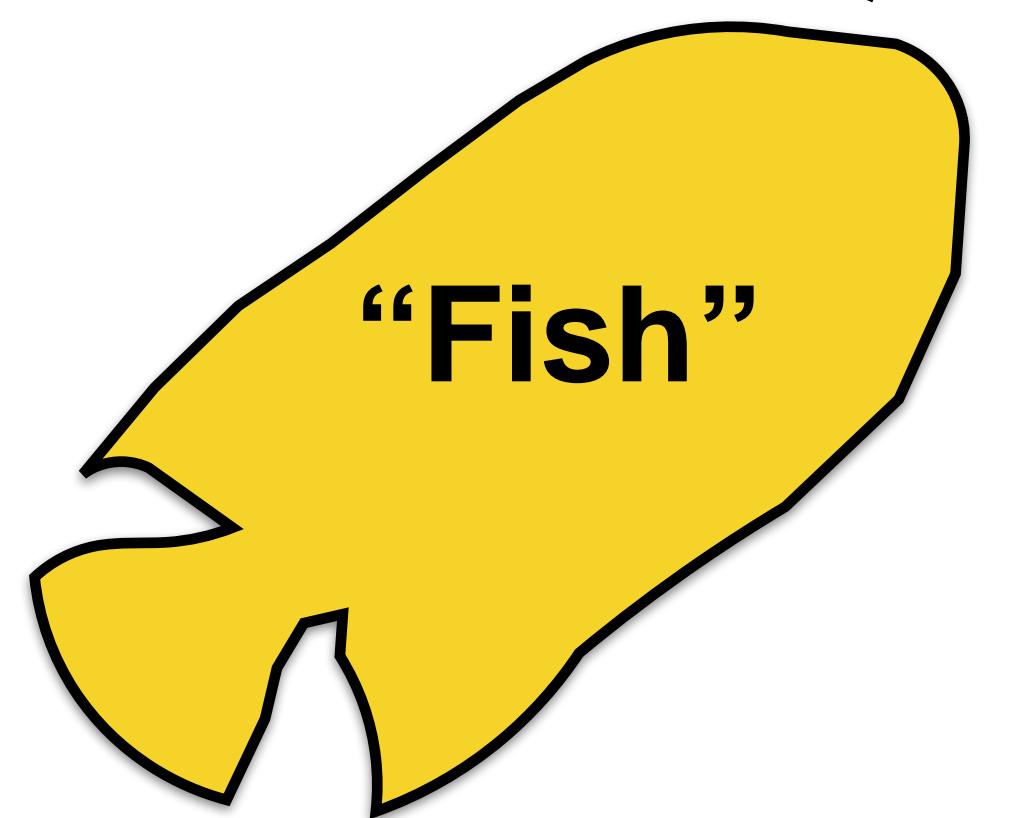
\mathcal{F}

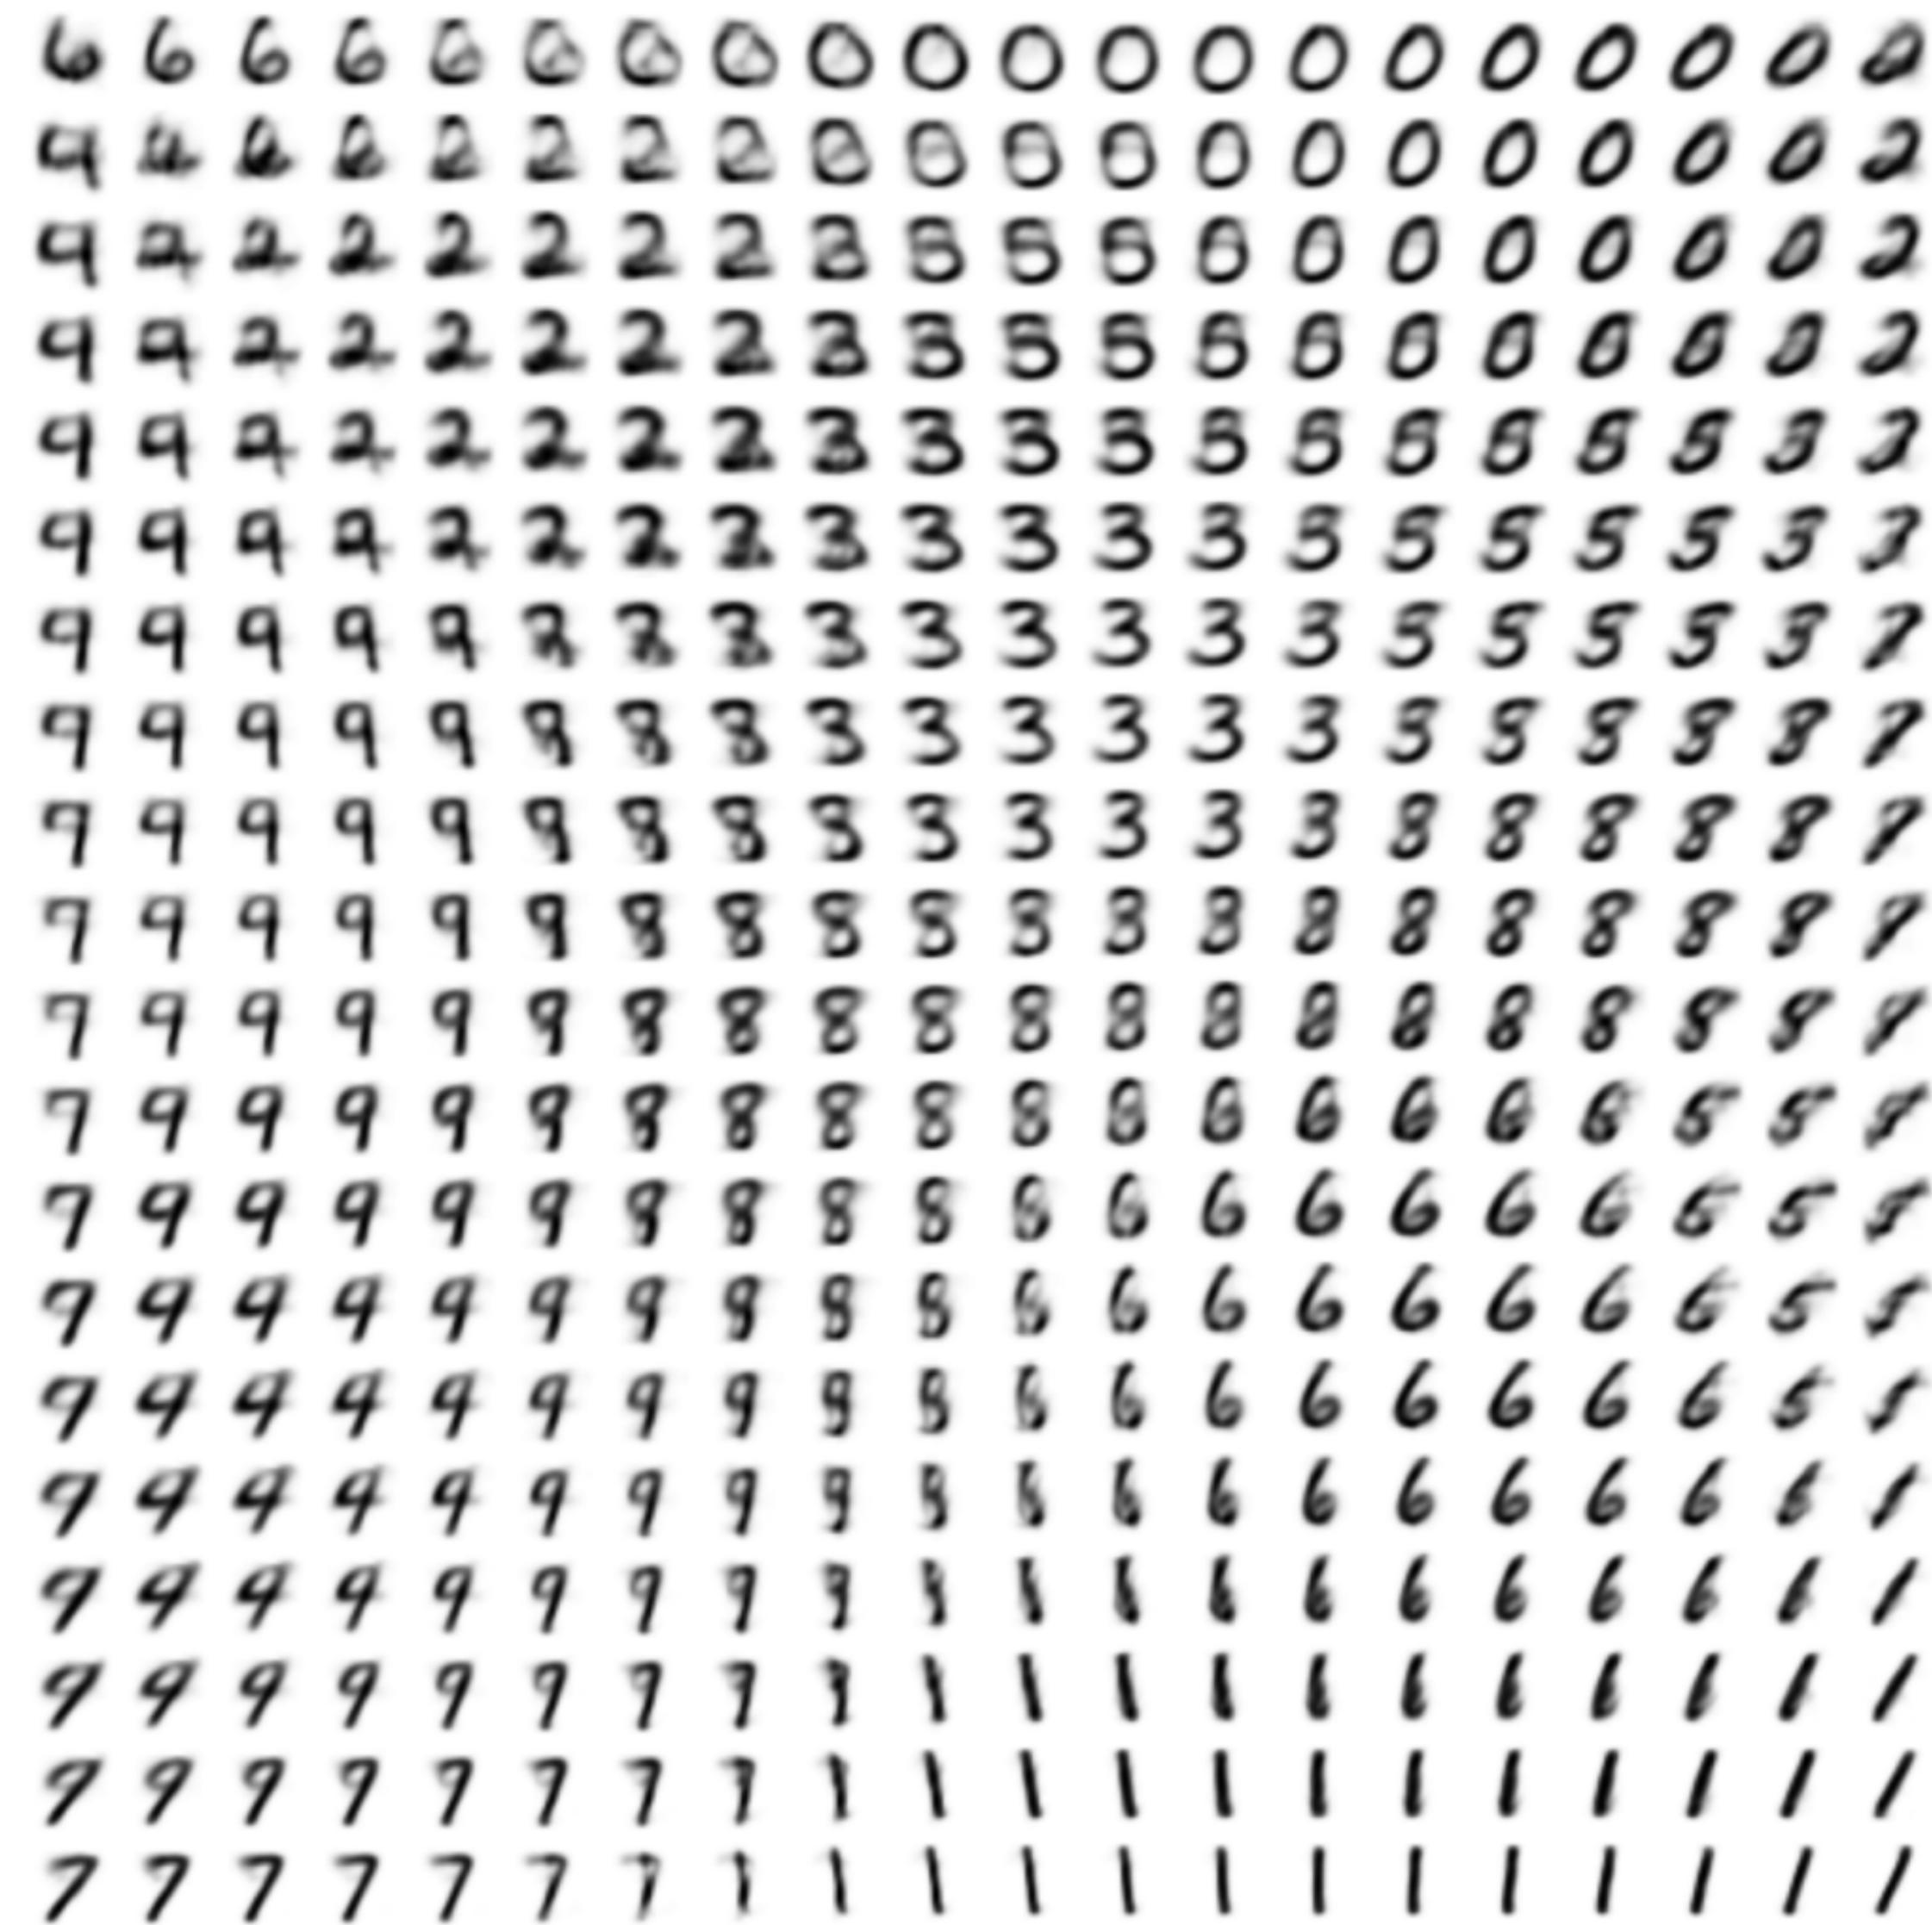


$\hat{\mathbf{X}} = \mathcal{F}(\mathbf{X})$



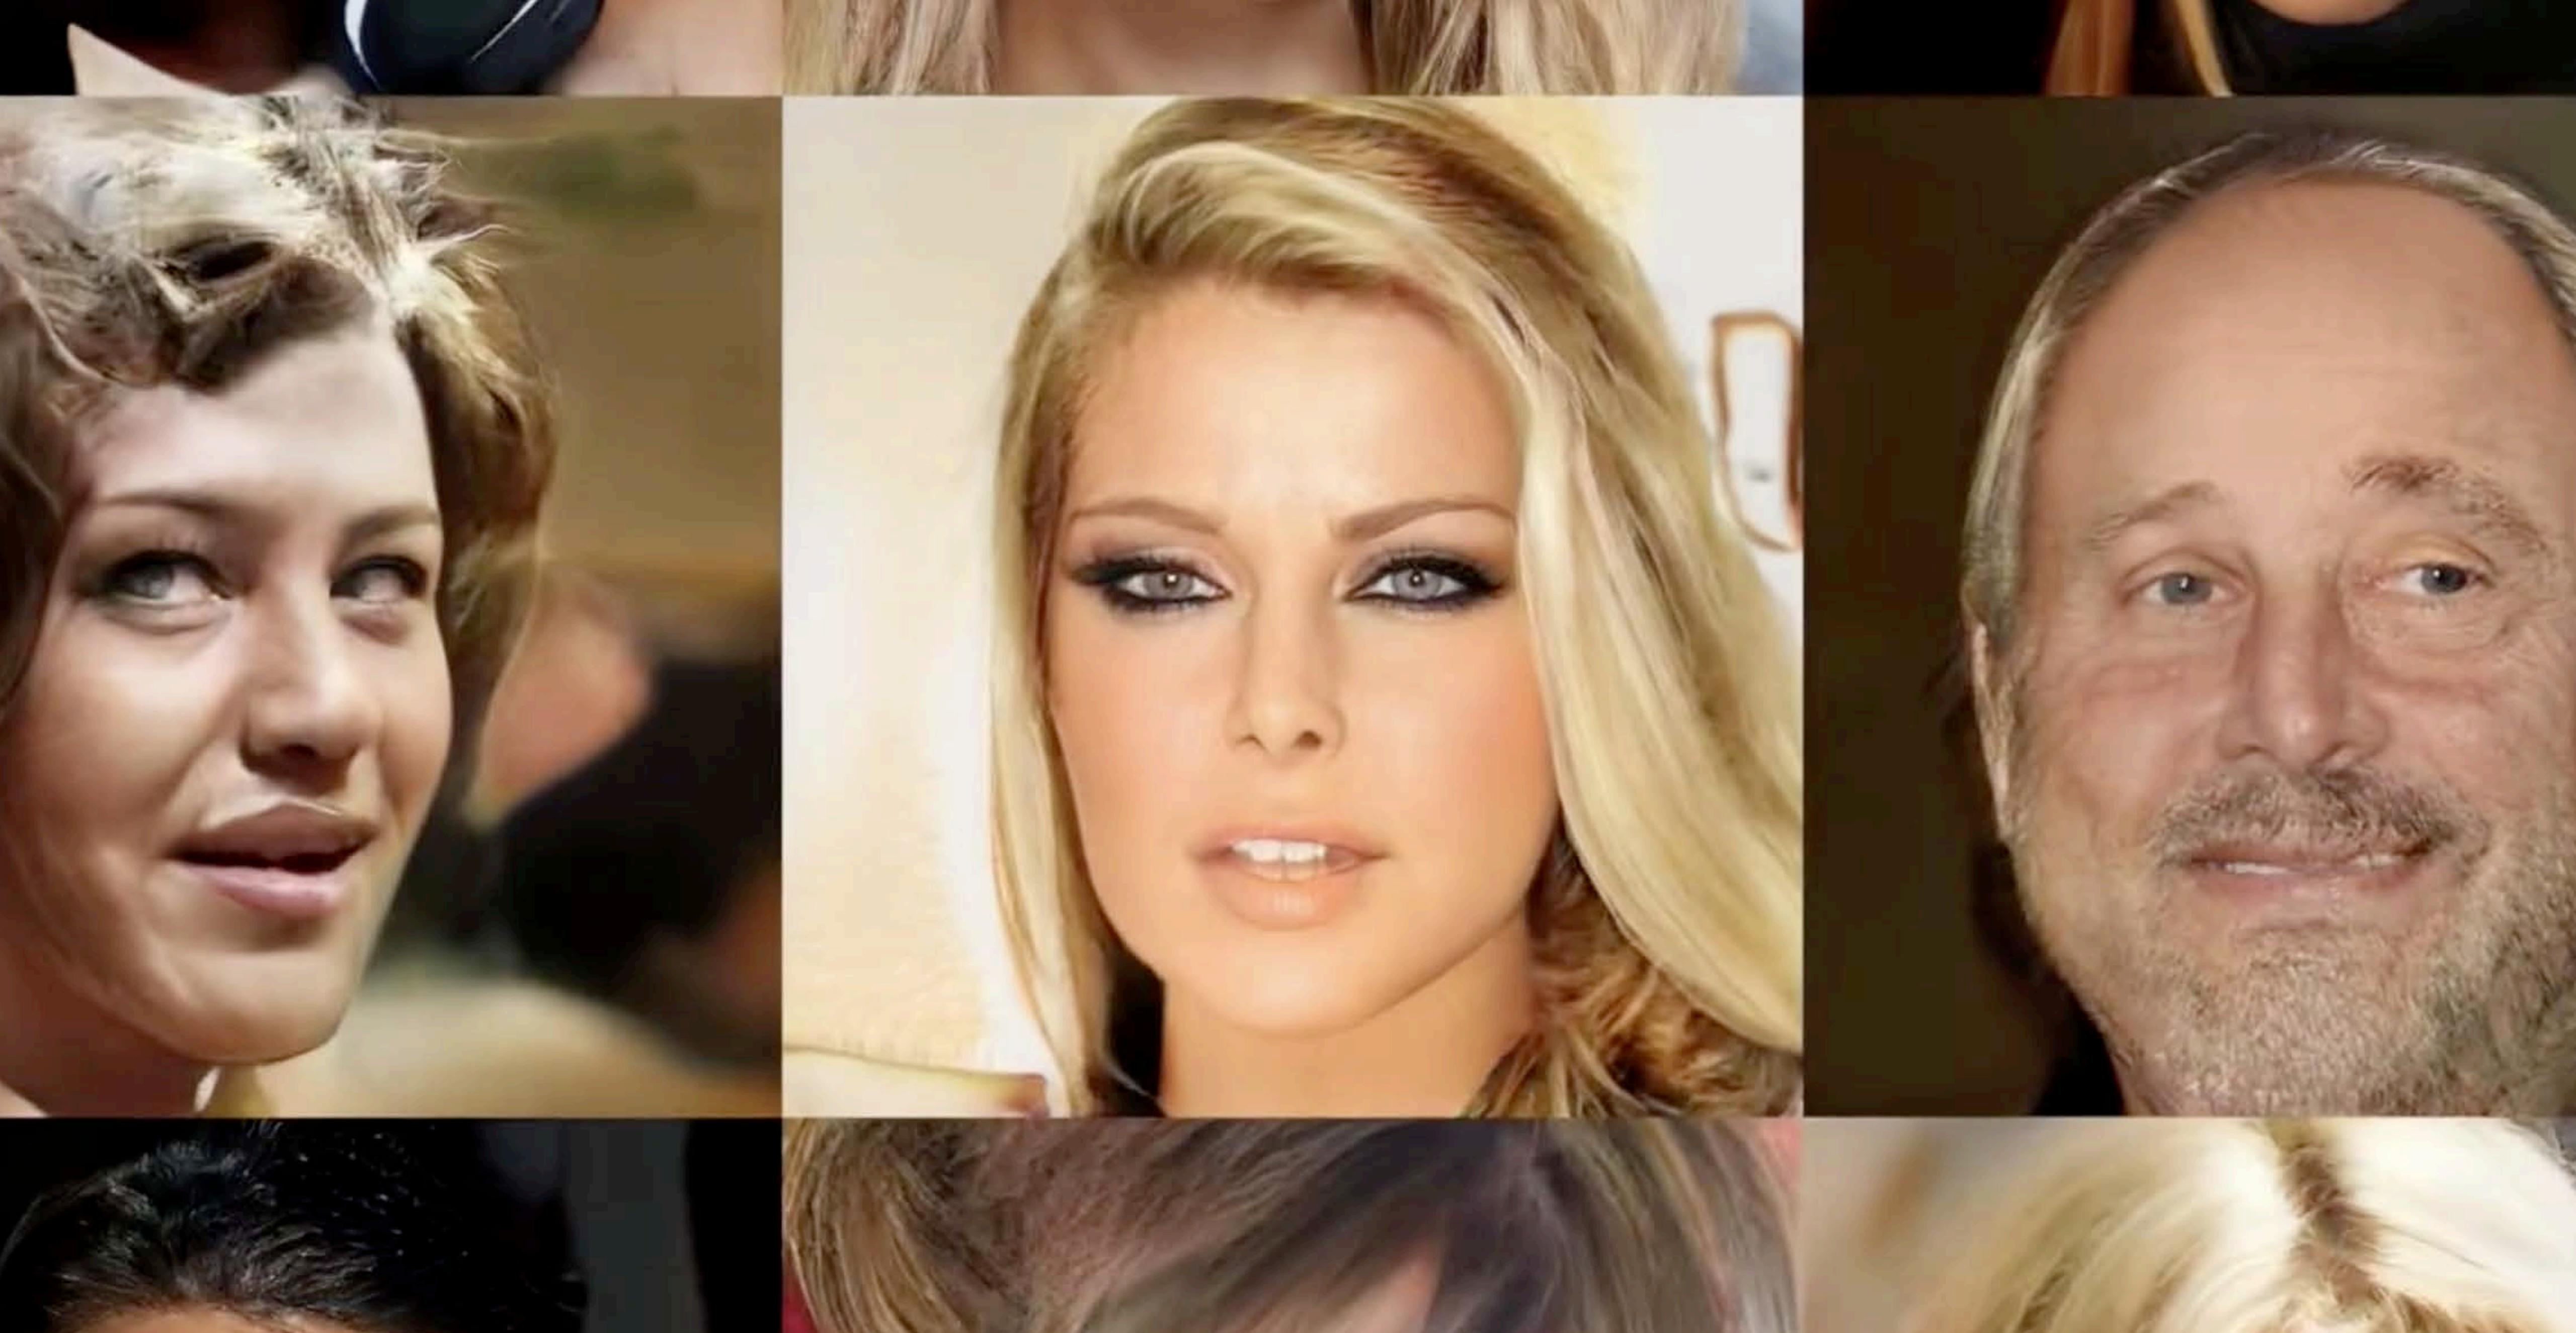
Reconstructed
image



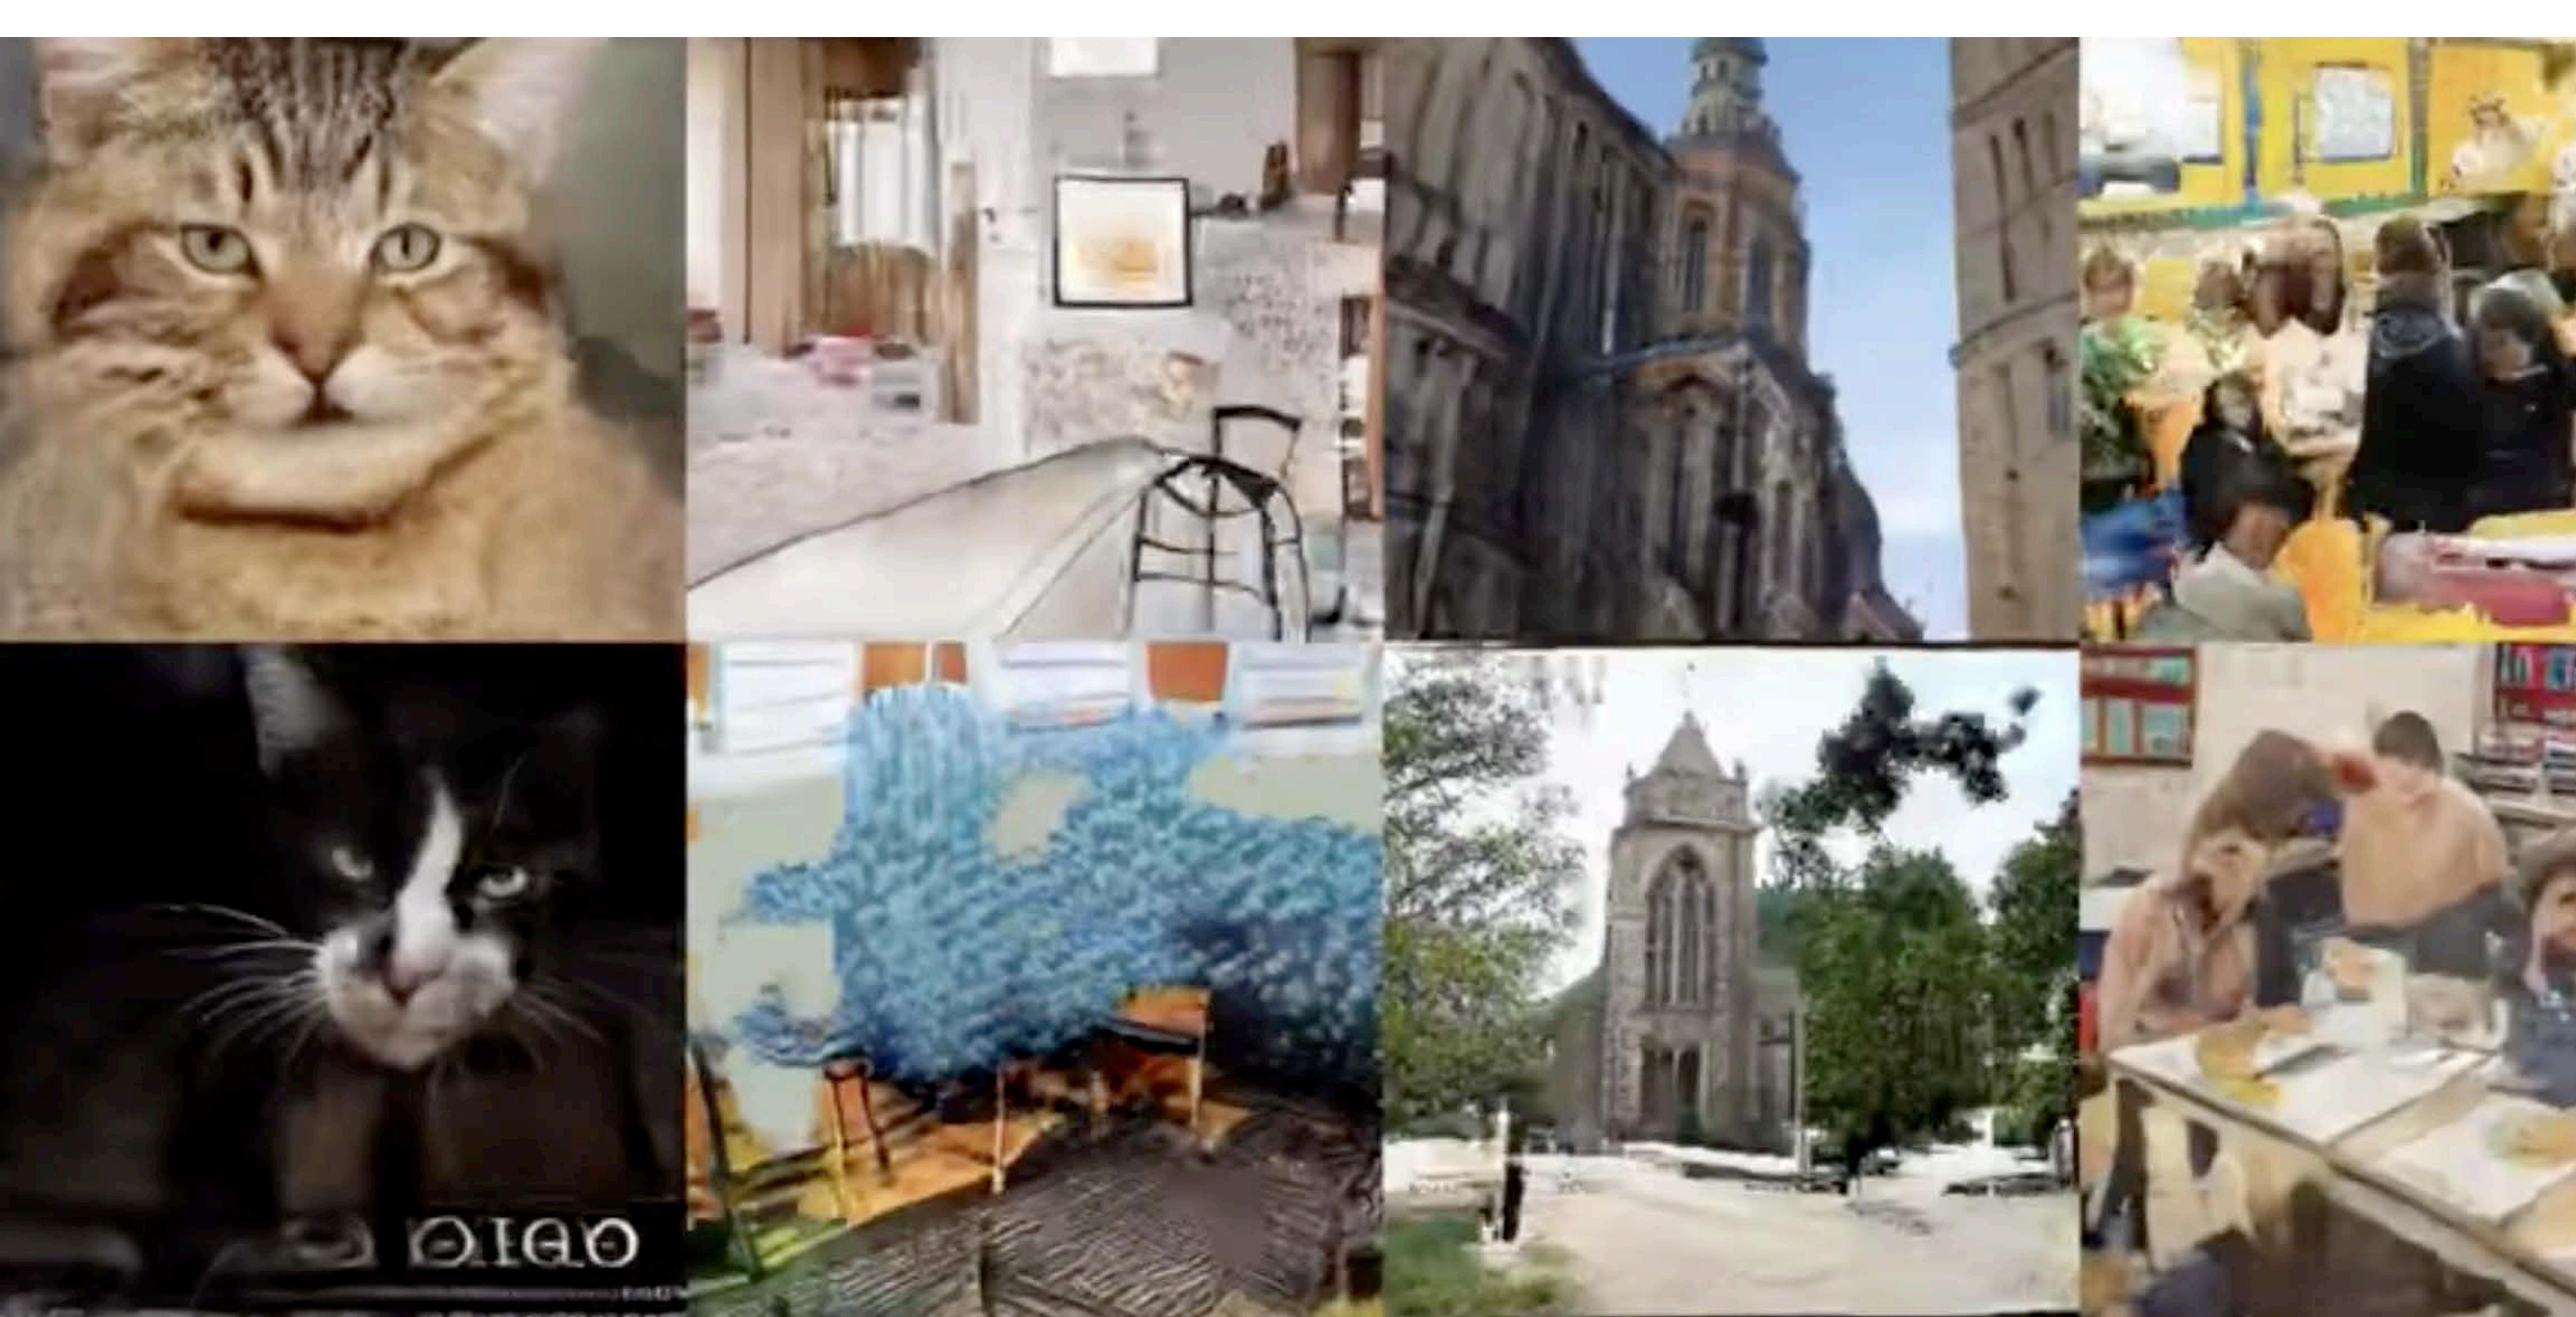


MNIST manifold learned by a VAE

[Kingma & Welling, 2014]

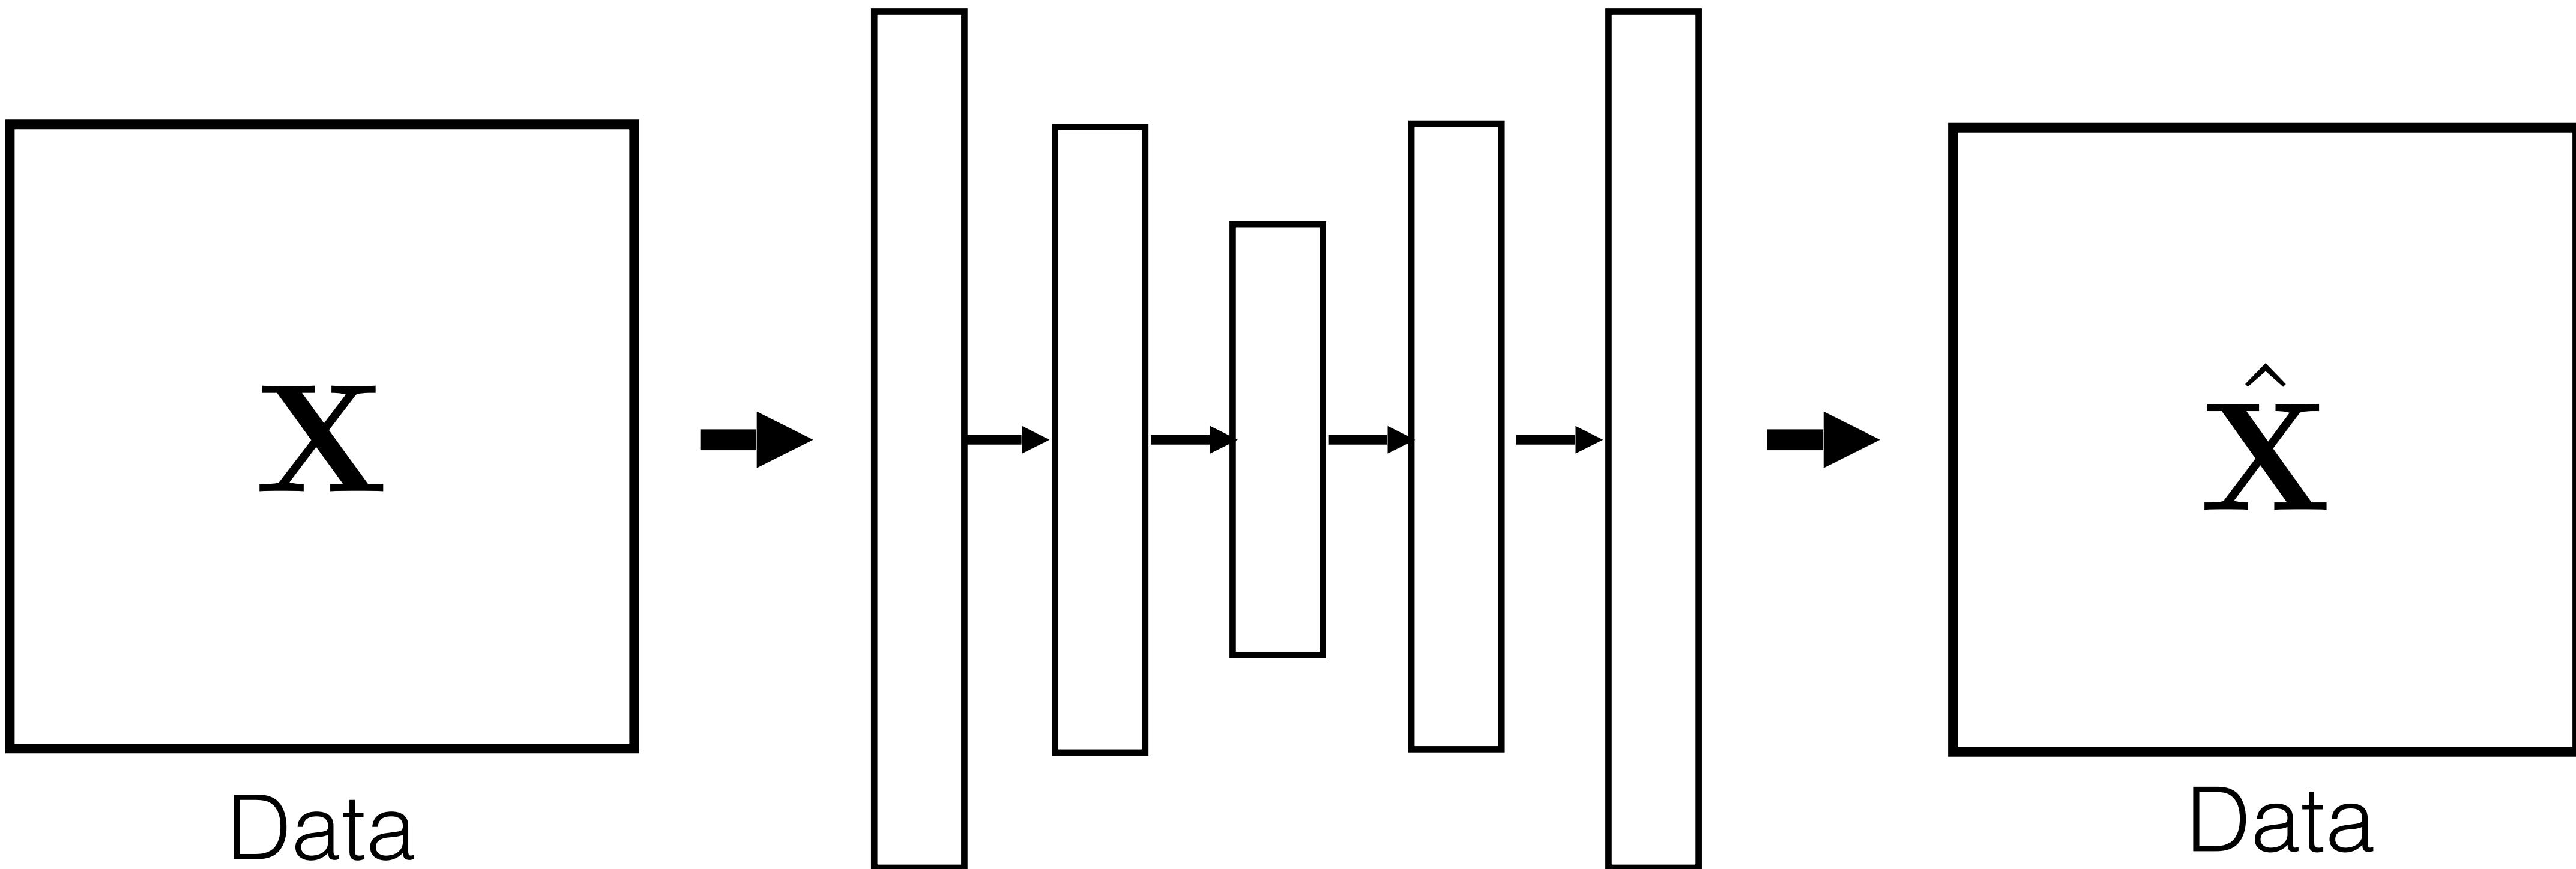


Progressive GAN [Karras et al., 2018]

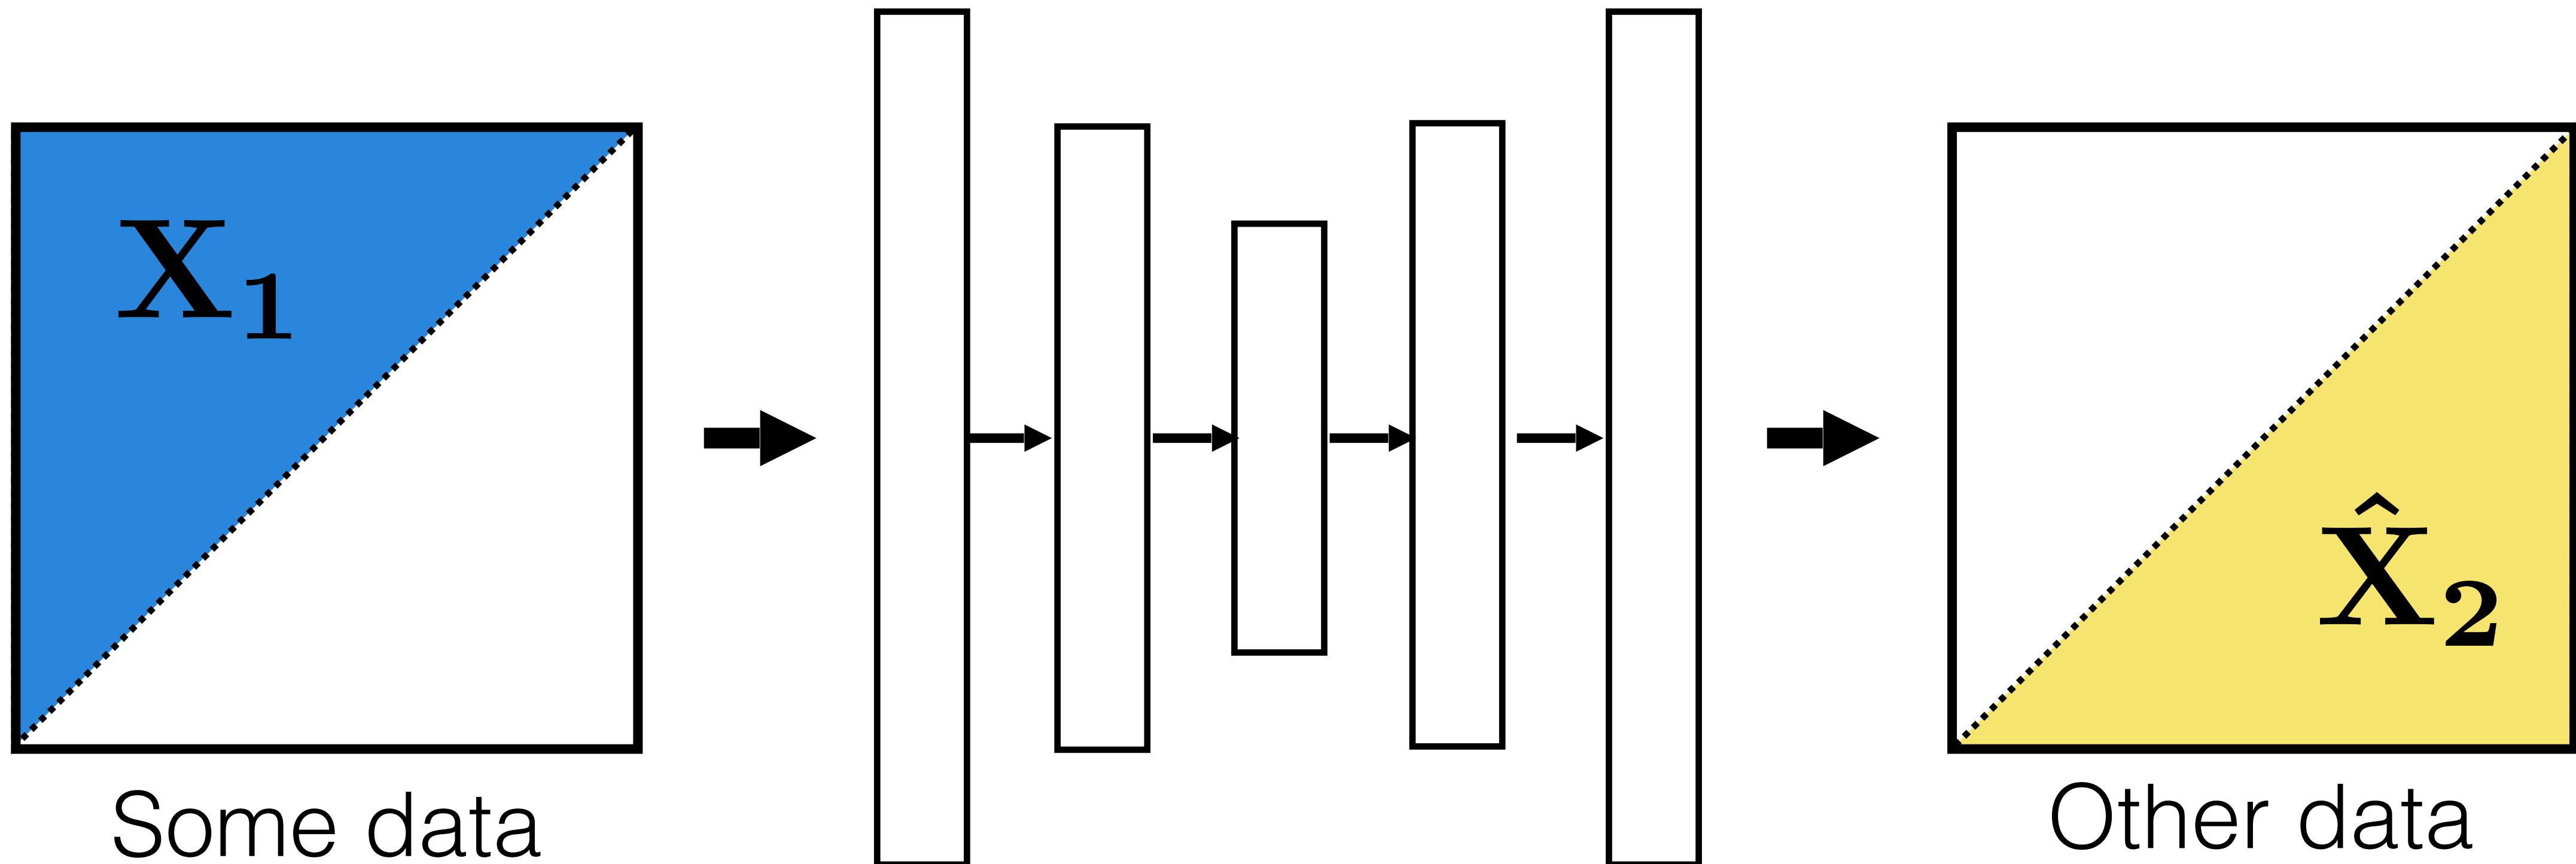


Progressive GAN [Karras et al., 2018]

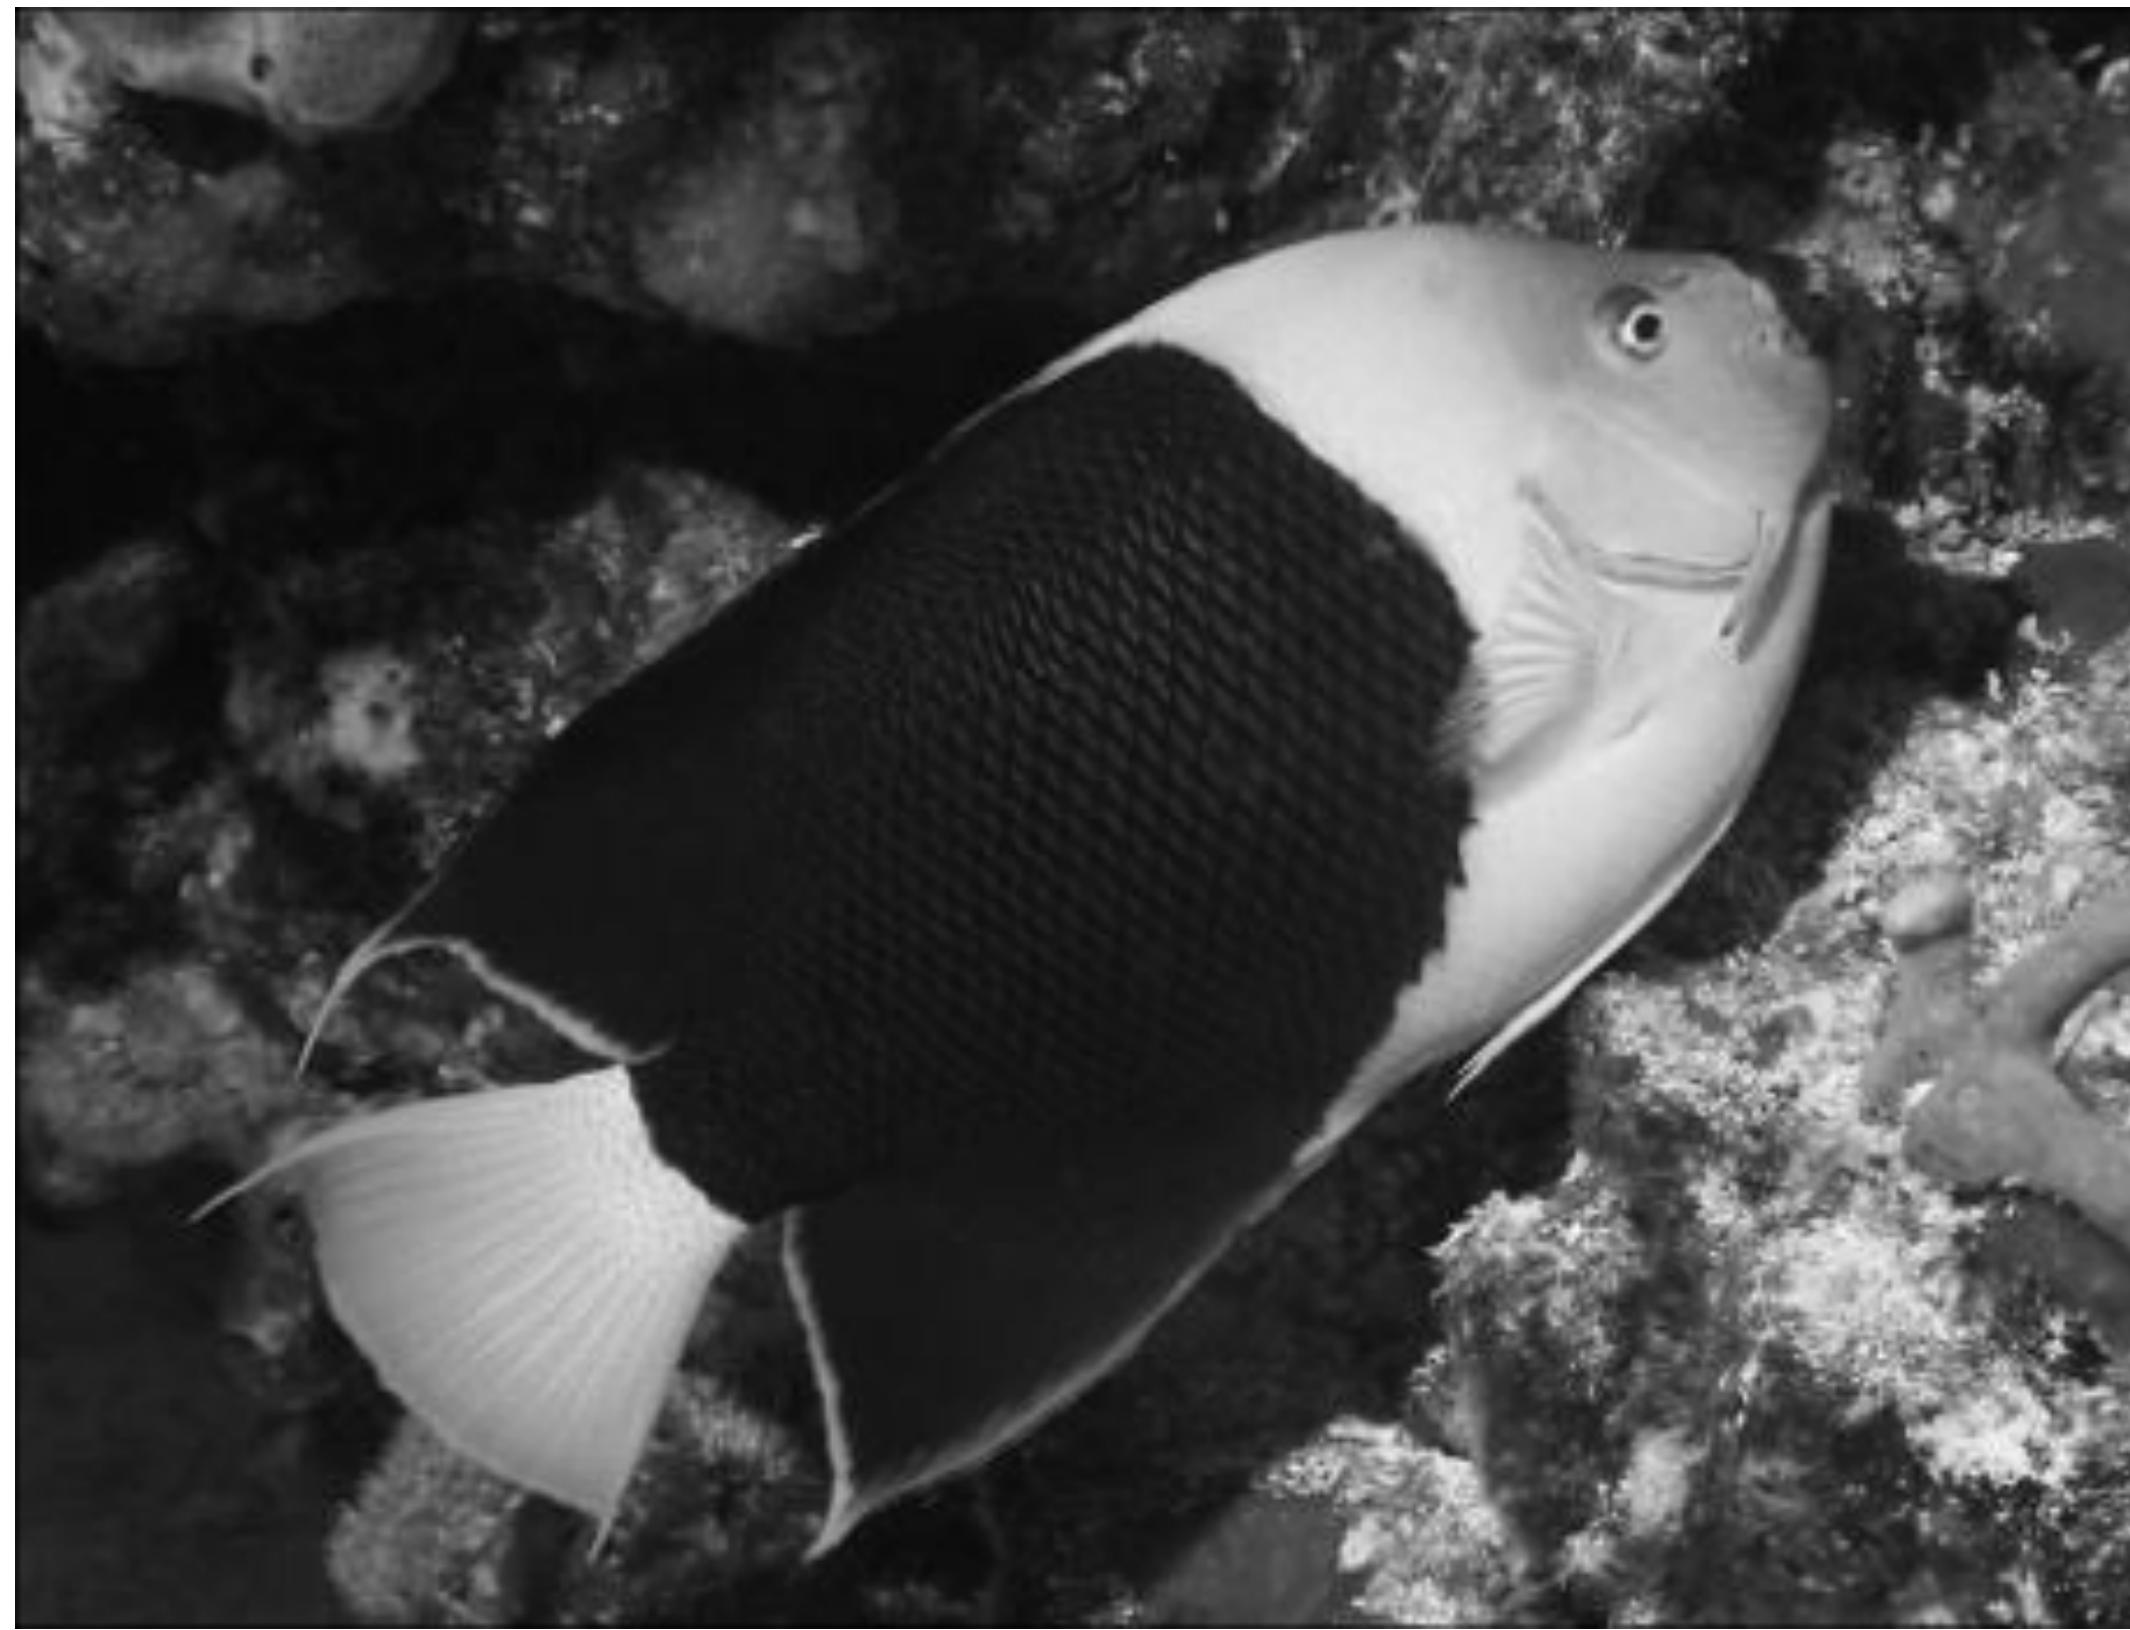
Data compression



Data prediction



see also “Predictive coding”, Denoising autoencoders [Vincent et al., 2008]



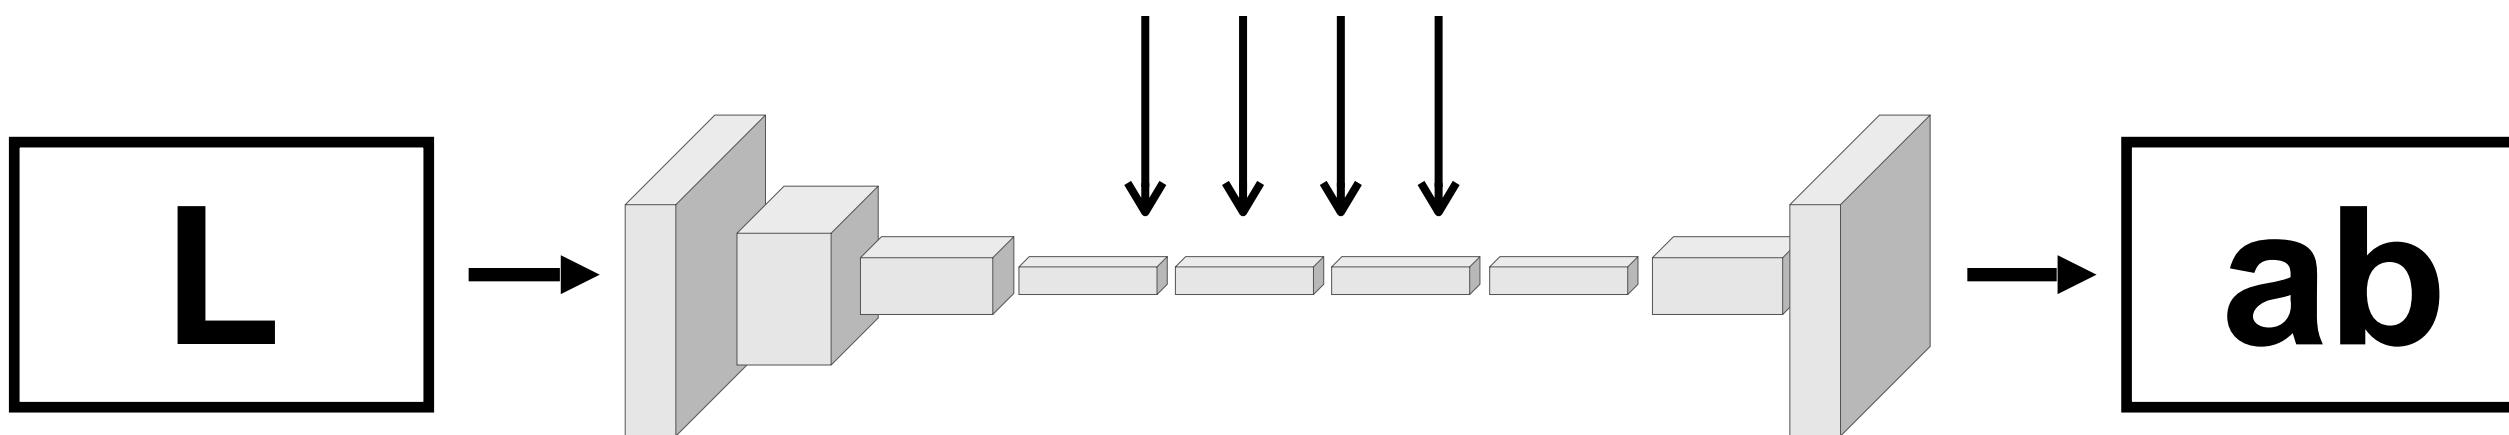
$$\xrightarrow{\mathcal{F}}$$



Grayscale image: L channel
 $\mathbf{X} \in \mathbb{R}^{H \times W \times 1}$

Semantics? Higher-level abstraction?

information: ab channels
 $\hat{\mathbf{Y}} \in \mathbb{R}^{H \times W \times 2}$



[Zhang, Isola, Efros, ECCV 2016]



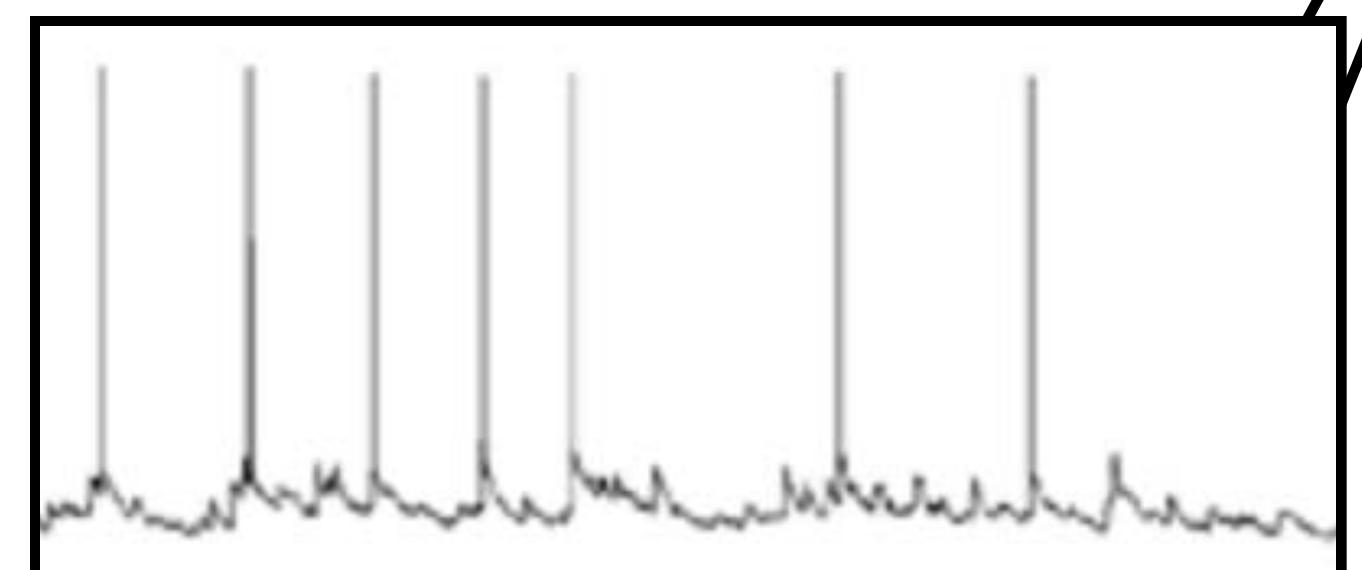
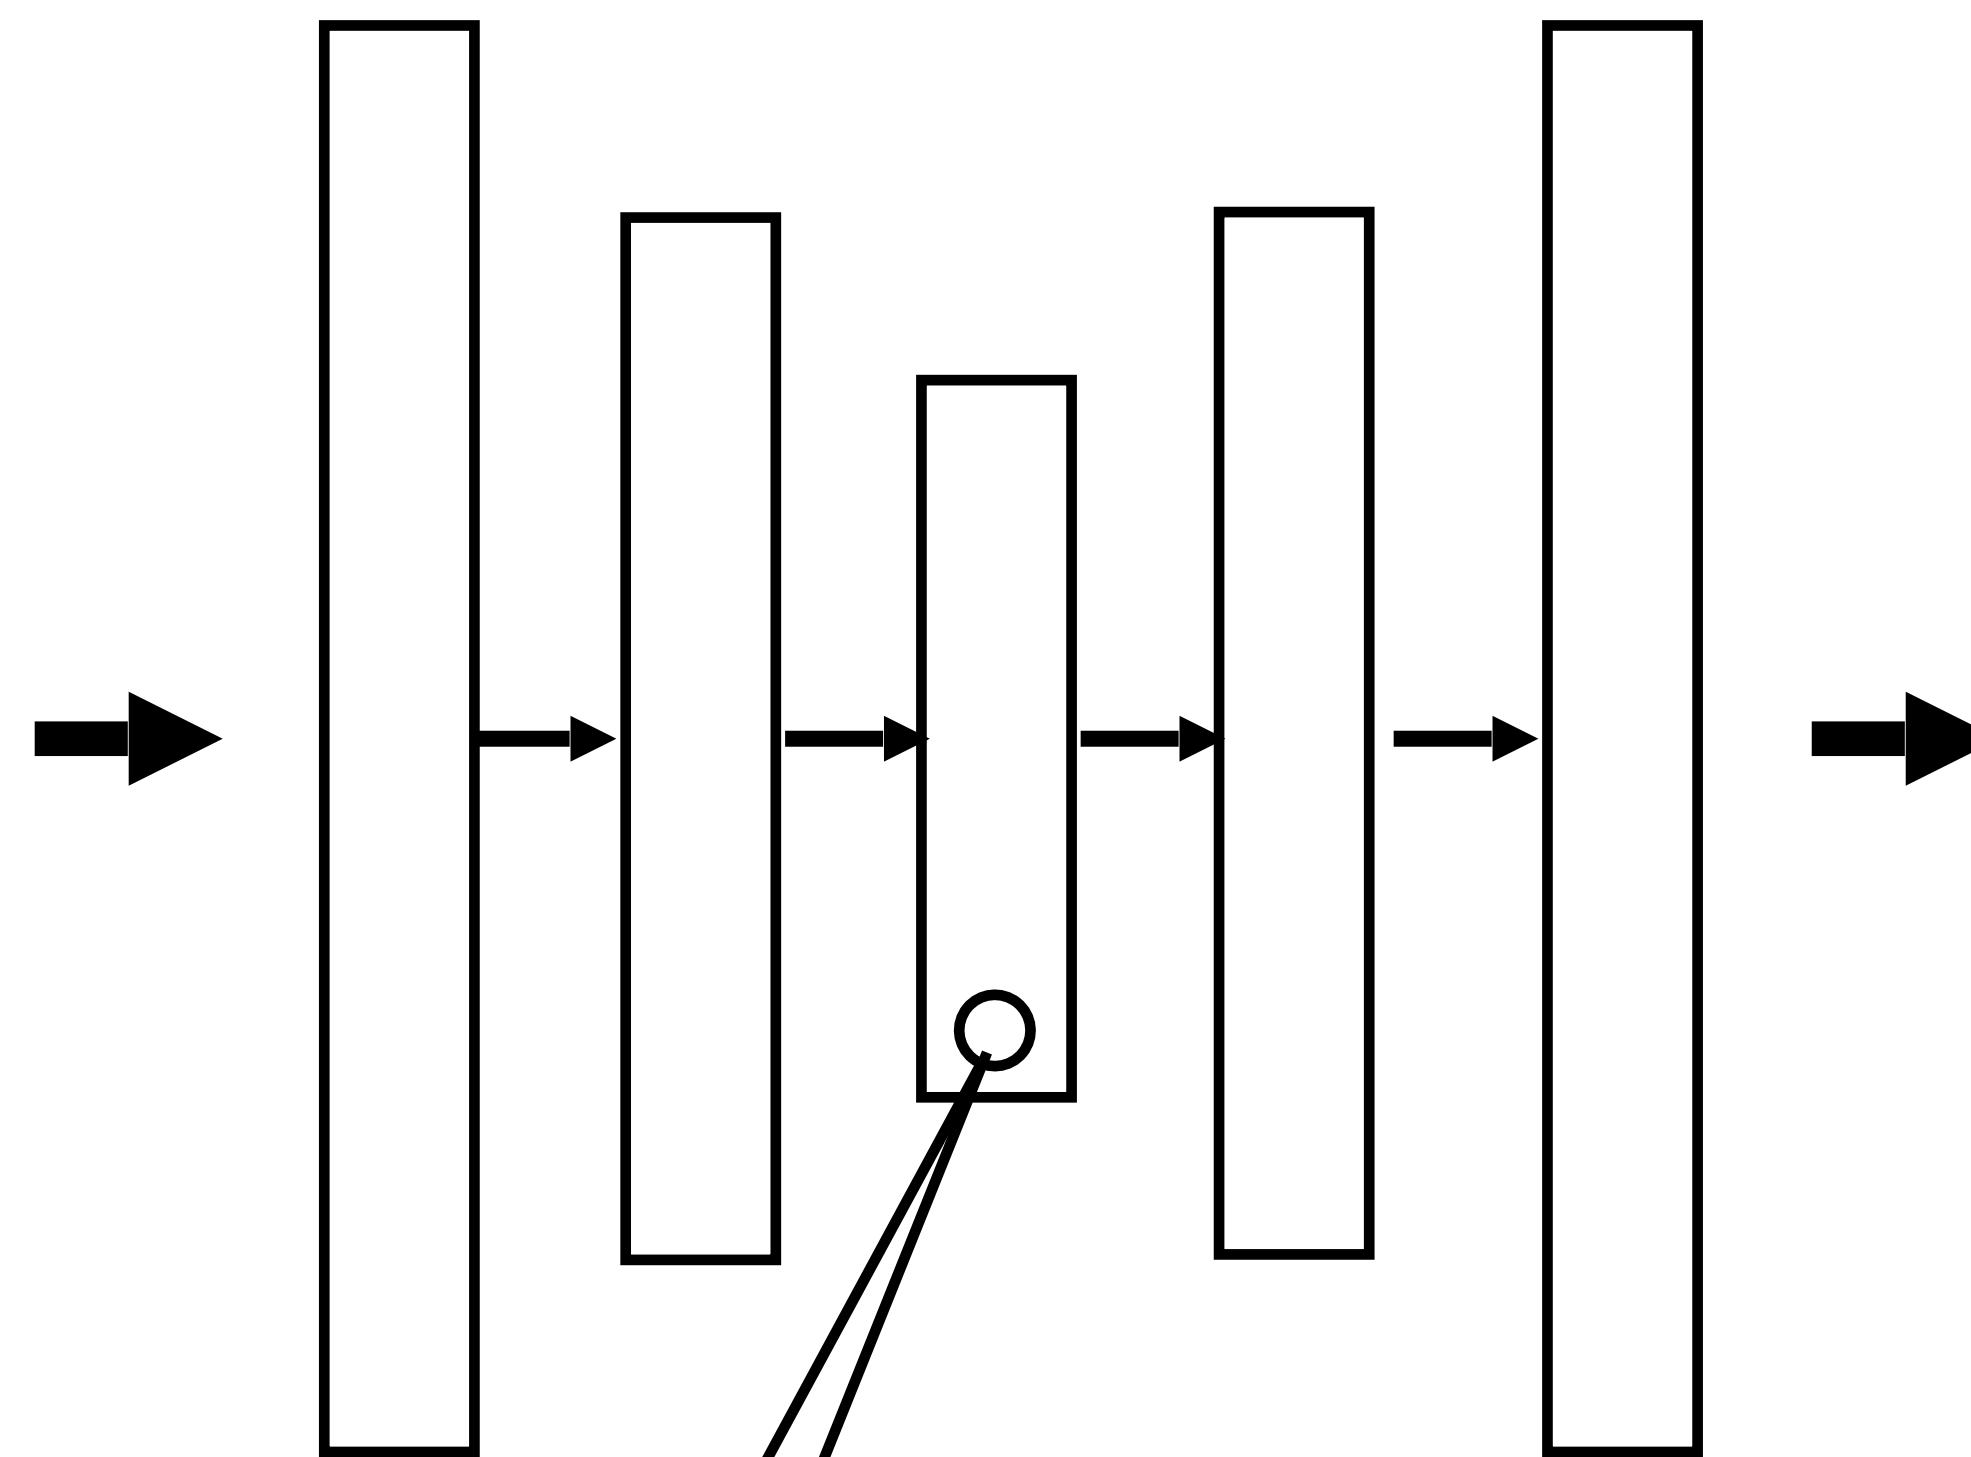
Instructive failure



Instructive failure



Deep Net “Electrophysiology”



[Zeiler & Fergus, ECCV 2014]
[Zhou et al., ICLR 2015]

Stimuli that drive selected neurons (conv5 layer)

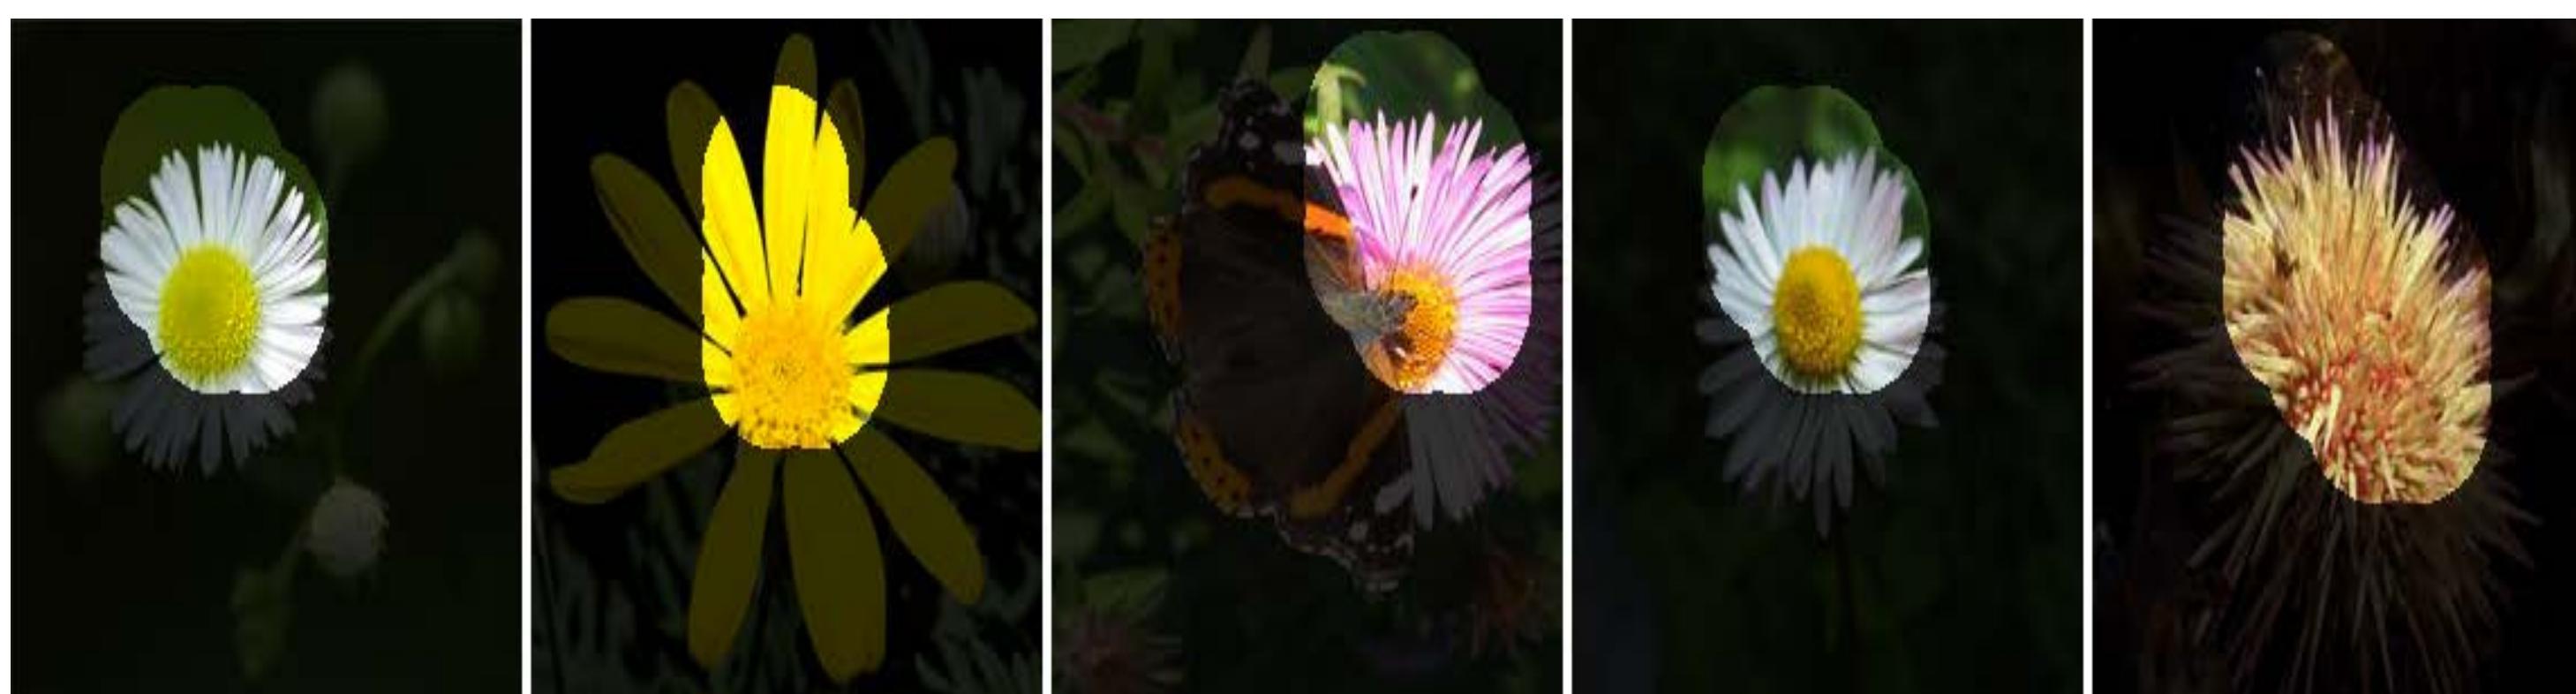
faces



dog
faces

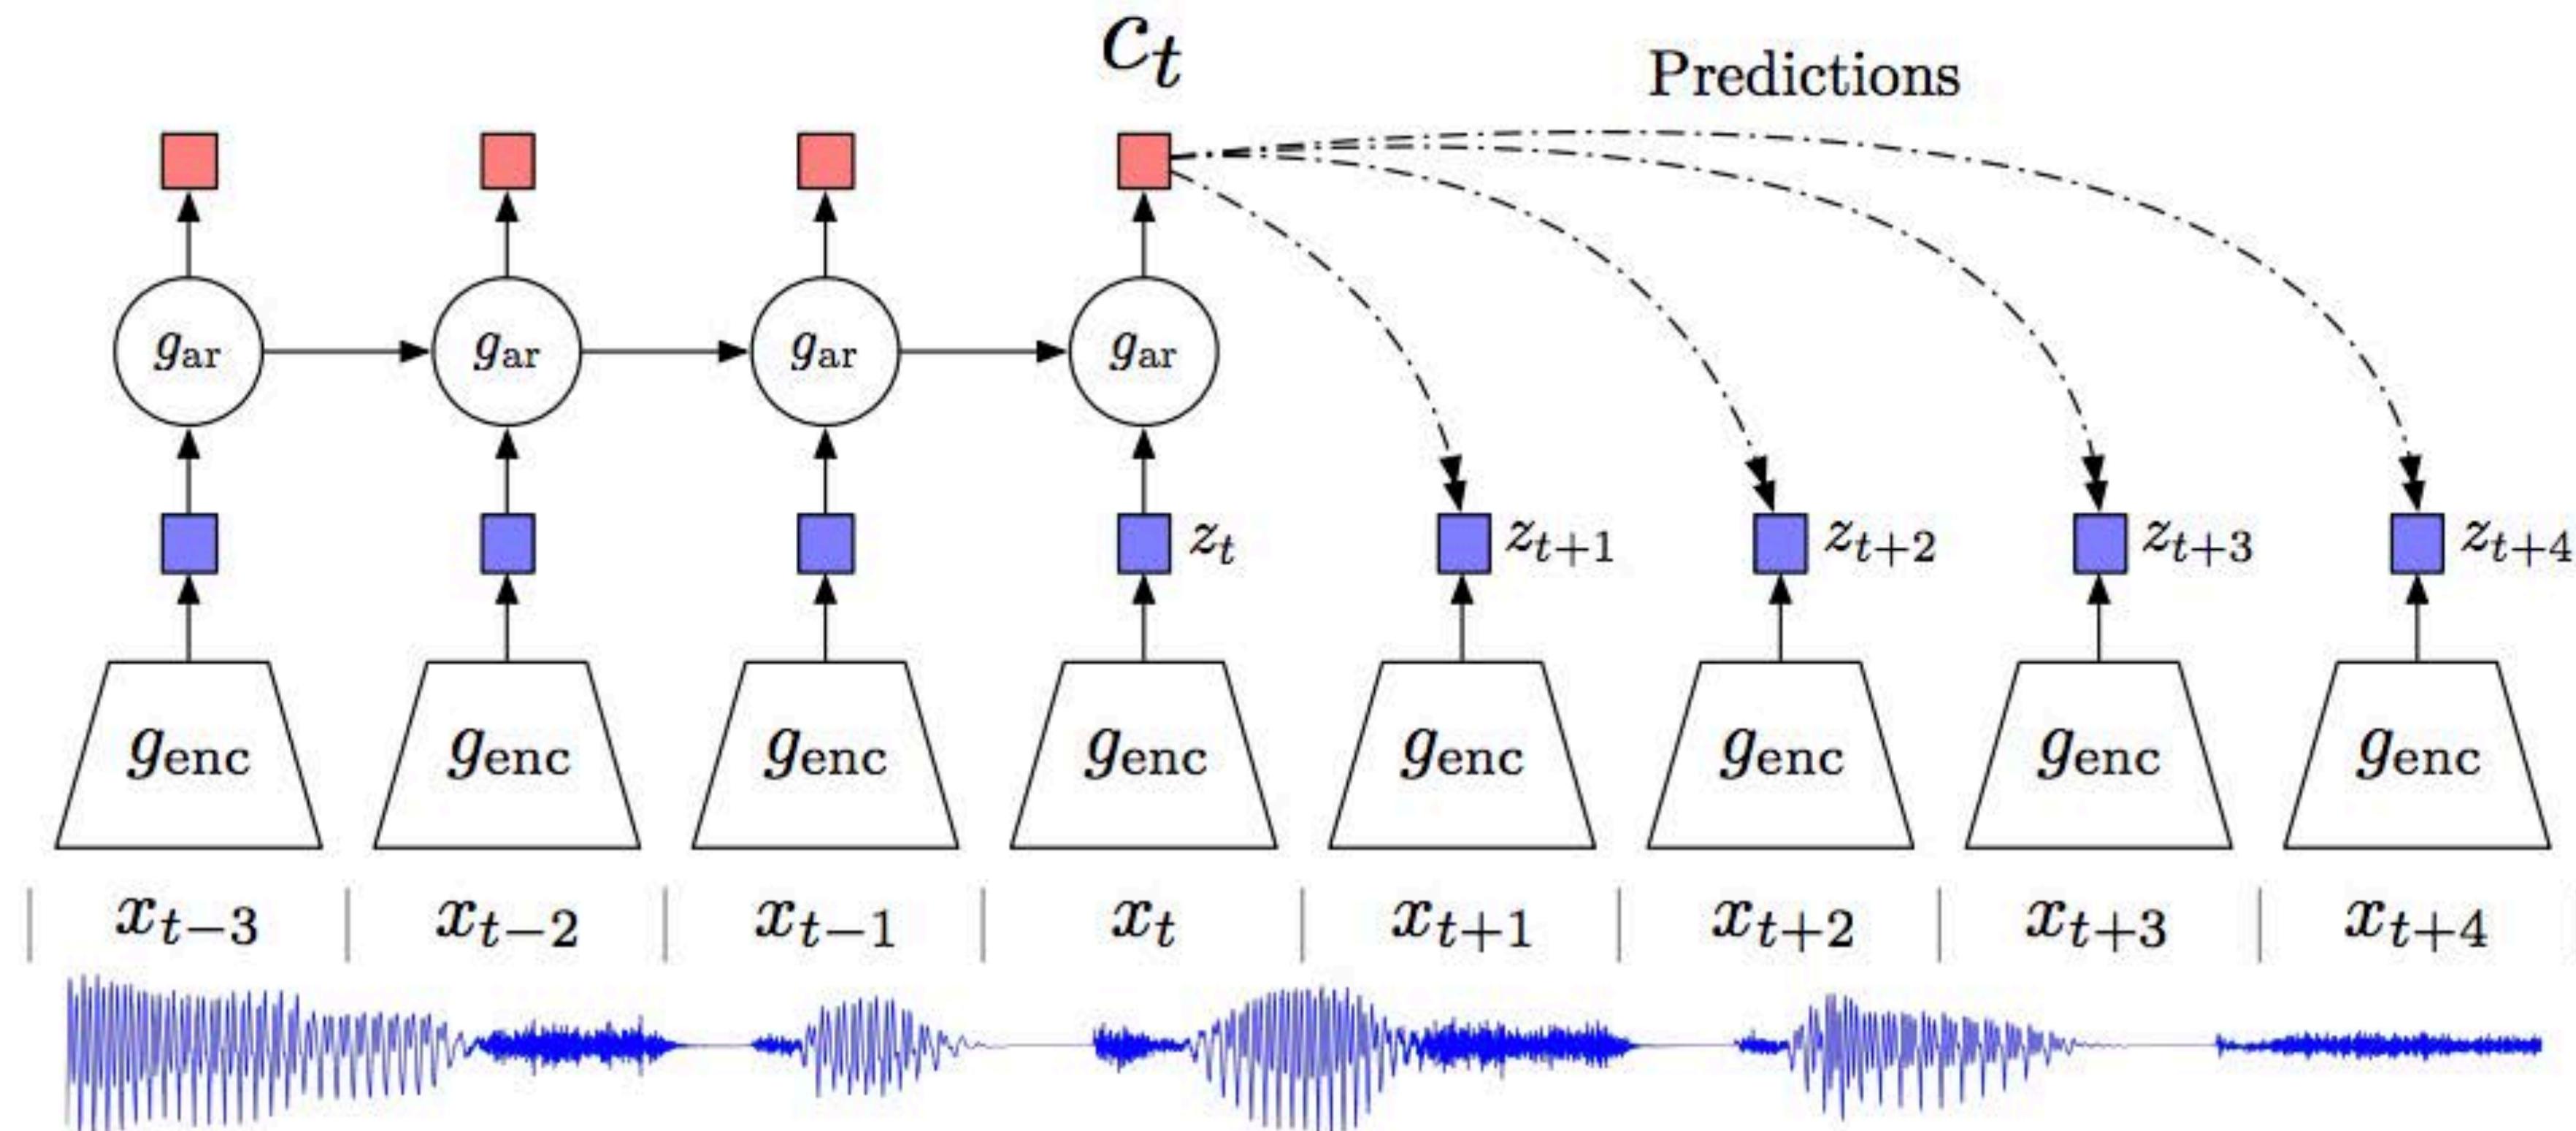


flowers



Representation Learning with Contrastive Predictive Coding

[van den Oord, Li, and Vinyals, 2018]

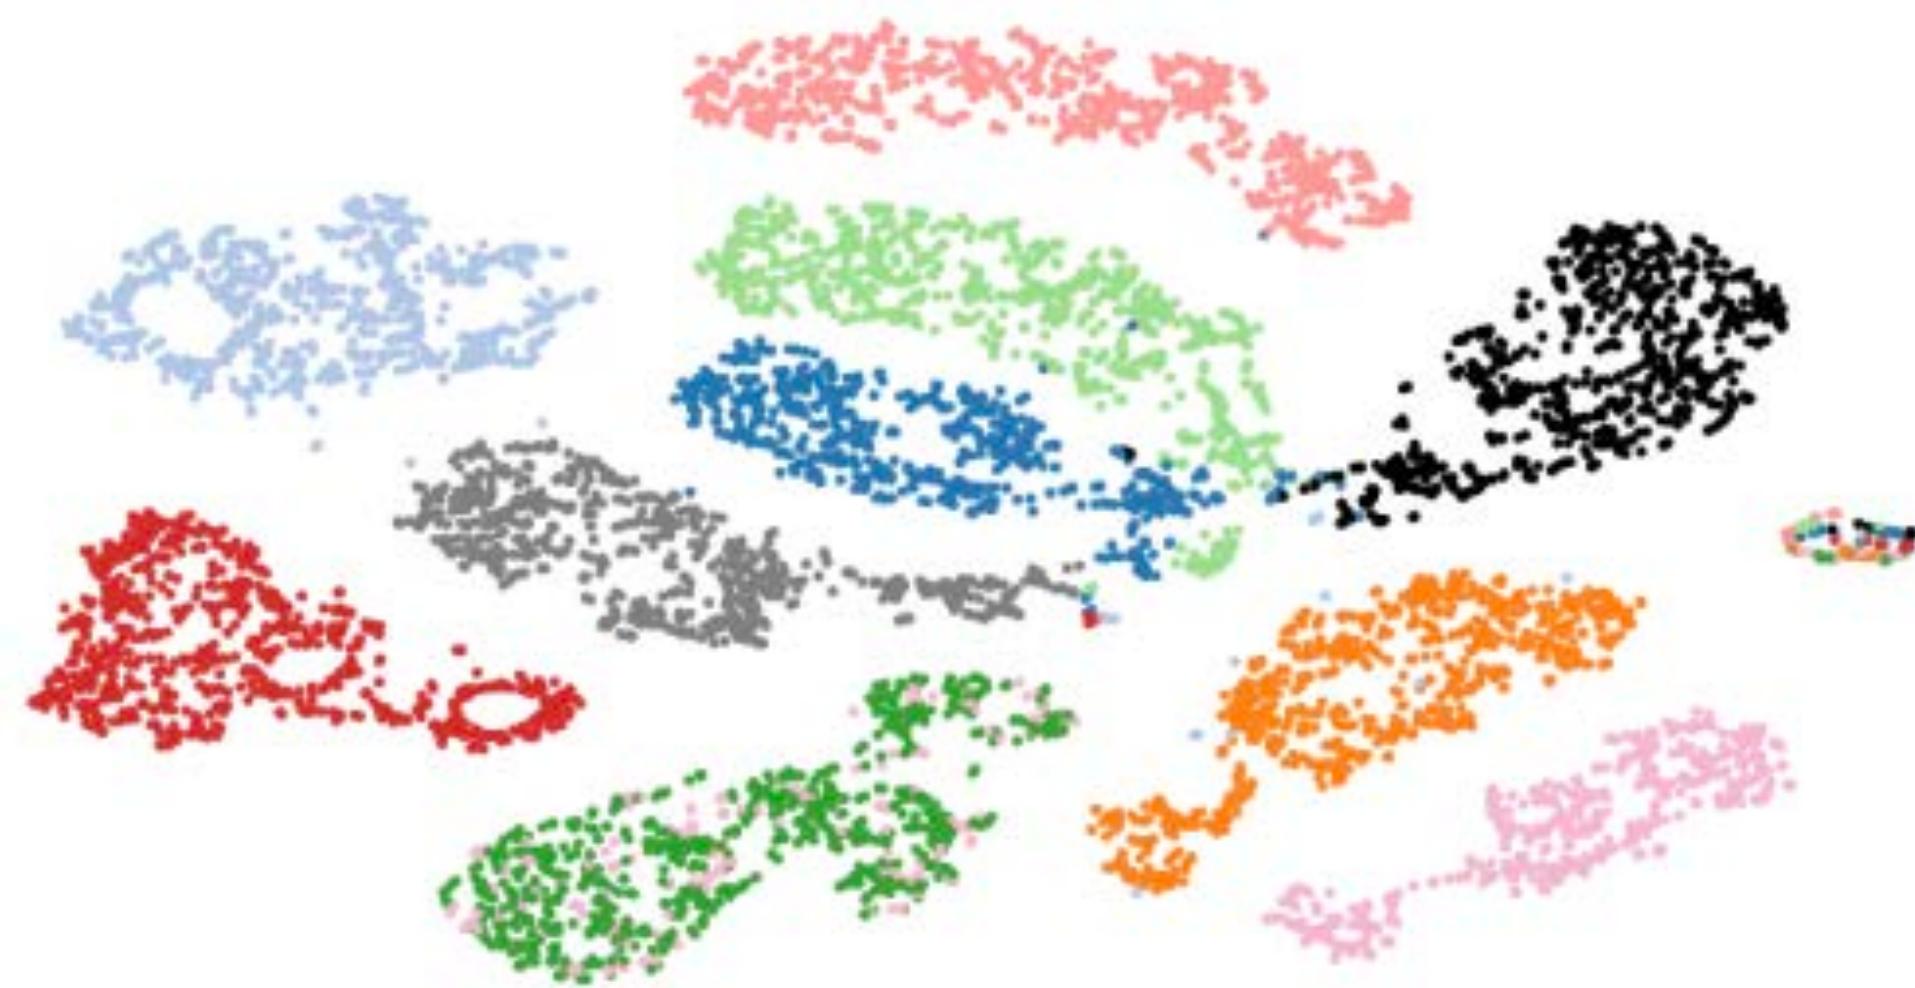


$$I(x; c) = \sum_{x, c} p(x, c) \log \frac{p(x|c)}{p(x)}$$

Representation Learning with Contrastive Predictive Coding

[van den Oord, Li, and Vinyals, 2018]

Sound representation



Each color represents a different speaker

Image representation

Method	Top-1 ACC
Using AlexNet conv5	
Video [27]	29.8
Relative Position [11]	30.4
BiGan [34]	34.8
Colorization [10]	35.2
Jigsaw [28] *	38.1
Using ResNet-V2	
Motion Segmentation [35]	27.6
Exemplar [35]	31.5
Relative Position [35]	36.2
Colorization [35]	39.6
CPC	48.7

What can you do with generative models?

1. Data prediction
2. Domain mapping
- 3. Representation learning**
4. Model-based intelligence

What can you do with generative models?

1. Data prediction
2. Domain mapping
3. Representation learning
- 4. Model-based intelligence**

Model-based intelligence

Predict consequences of my actions.

Then use that predictive model to plan good actions (or learn good policies).

Can be faster, safer, and more abstracted than directly trying actions in real world (or in simulator).

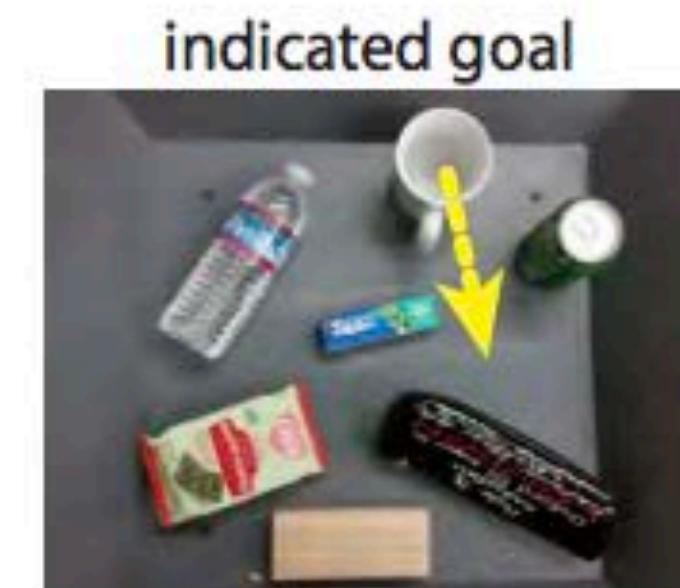
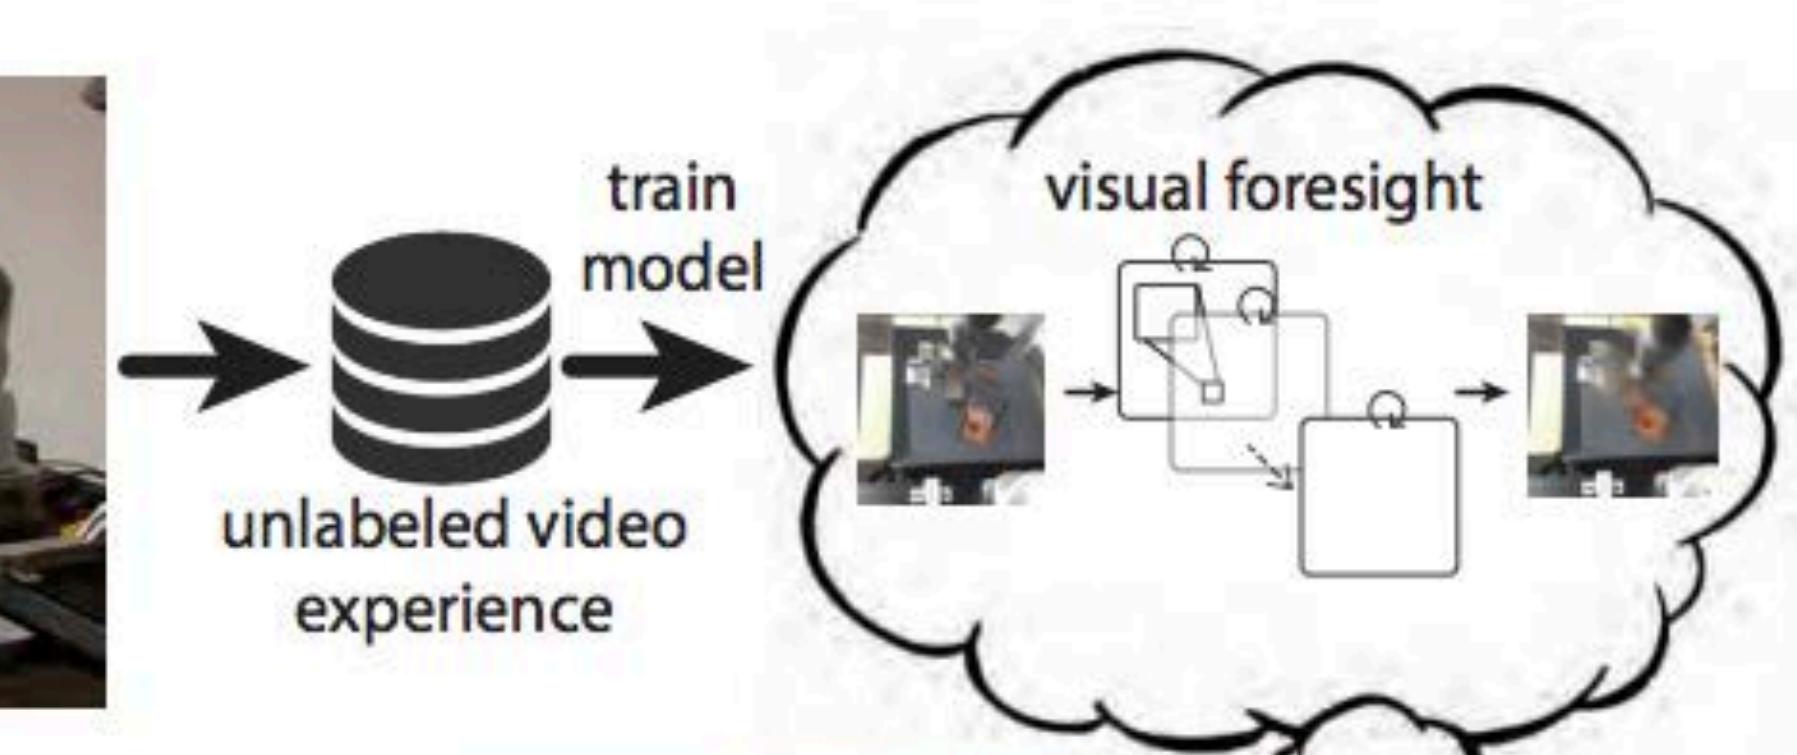


Yann LeCun's cake

Deep Visual Foresight For Robotic Planning

[Finn & Levine 2017]

1. Train model to predict future frame given action
2. Specify target future frame
3. Use prediction model to pick action that maximizes probability of reaching target frame

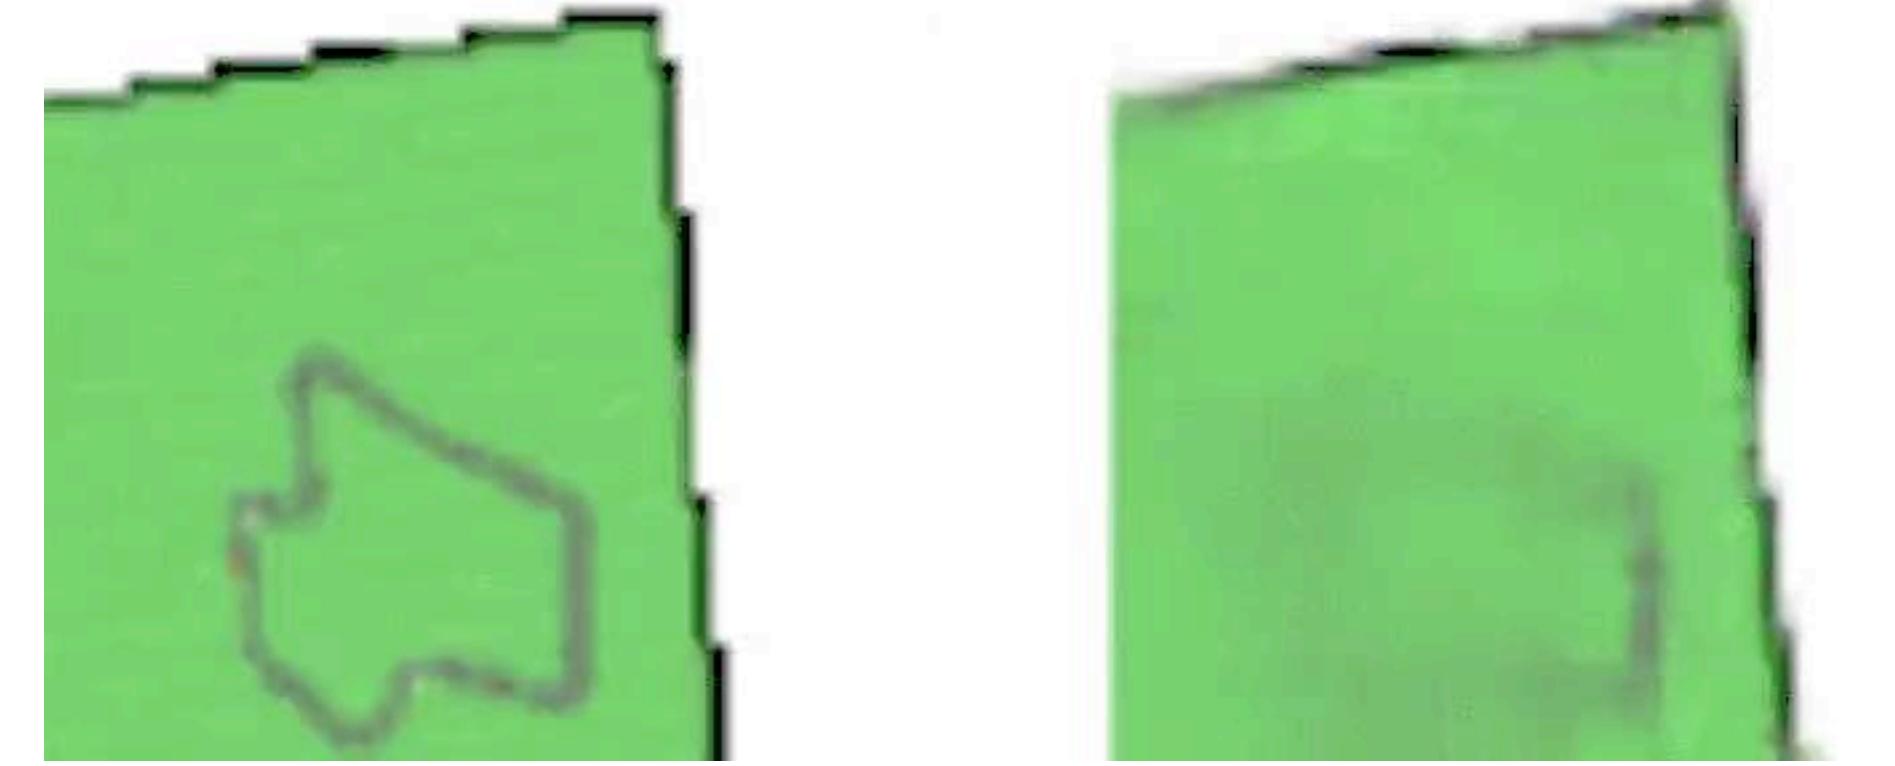


c.f. [Weber et al., 2017], [...]

World Models

[Ha & Schmidhuber 2018]

1. Train an RNN to simulate a video game
2. Do policy optimization using the RNN to simulate trajectories, rather than the actual game (policy reasons on top of latent space)



Model	Parameter Count
VAE	4,348,547
MDN-RNN	422,368
Controller	867

Incentivizing exploration in reinforcement learning with deep predictive models

[Stadie et al., 2015]

1. Train model to predict next state given action and previous state: $p(s_{t+1}|s_t, a_t)$
2. Reward agent for reaching states that surprise it: $r_t = -\log p(s_{t+1}|s_t, a_t)$



[Pathak et al., 2017]

What can you do with generative models?

1. Data prediction
2. Domain mapping
3. Representation learning
4. Model-based intelligence

Thank you!